
Amazon API Gateway

Developer Guide



Amazon API Gateway: Developer Guide

Copyright © 2016 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

What Is Amazon API Gateway?	1
API Gateway Concepts	2
Getting Started	3
Get Ready to Use API Gateway	3
Sign Up for AWS	4
Create an IAM User, Group or Role in Your AWS Account	4
Grant IAM Users Permissions to Access API Gateway Control and Execution Services	4
Next Step	5
Build and Test an API from Example	5
Build and Deploy an API Step by Step	13
Invoke Lambda Functions Synchronously	20
Step 1: Prerequisites	20
Step 2: Create an API	20
Step 3: Create a Resource	20
Step 4: Create Lambda Functions	21
Step 5: Create and Test a GET Method	24
Step 6: Create and Test a POST Method	25
Step 7: Deploy the API	26
Step 8: Test the API	26
Step 9: Clean Up	27
Next Steps	28
Create Lambda Invocation and Execution Roles	28
Map API Request Parameter for HTTP Endpoints	30
Prerequisites	33
Step 1: Create Resources	33
Step 2: Create GET and POST Methods	33
Step 3: Specify Method Settings and Test the Methods	34
Step 4: Deploy the API	36
Step 5: Test the API	36
Next Steps	38
Use Modes and Mapping Templates	39
Prerequisites	40
Step 1: Create Models	41
Step 2: Create Resources	43
Step 3: Create GET Methods	44
Step 4: Create a Lambda Function	44
Step 5: Specify Method Settings and Test the Methods	45
Step 6: Deploy the API	49
Step 7: Test the API	50
Step 8: Clean Up	51
Next Steps	52
Get Started with AWS Proxies	52
Prerequisites	52
Step 1: Create the Resource	53
Step 2: Create the GET Method	53
Step 3: Create the AWS Service Proxy Execution Role	53
Step 4: Specify Method Settings and Test the Method	55
Step 5: Deploy the API	56
Step 6: Test the API	56
Step 7: Clean Up	57
Creating an API	58
Create an API in API Gateway	58
Create an API Using the API Gateway Console	59
Create an API Using the API Gateway Control Service API	59
Create an API Using the AWS SDK for API Gateway	59

Create an API Using the AWS CLI	59
Set up Method and Integration	59
Before Configuring Methods	59
After Setting Up Methods and Integration	59
Configure How a Method Is Integrated with a Back End	60
Configure How a User Calls an API Method	62
Configure How Data Is Mapped between Method and Integration	64
Configure Mock Integration for a Method	65
Set Up Request and Response Payload Mappings	68
Models	69
Mapping Templates	71
Tasks for Models and Mapping Templates	74
Create a Model	74
View a List of Models	74
Delete a Model	75
Photos Example	75
News Article Example	79
Sales Invoice Example	82
Employee Record Example	87
Request and Response Parameter-Mapping Reference	93
Map Data to Integration Request Parameters	93
Map Data to Method Response Headers	94
Transform Request and Response Bodies	95
Request and Response Payload-Mapping Reference	96
Accessing the \$context Variable	96
Accessing the \$input Variable	98
Accessing the \$stageVariables Variable	100
Accessing the \$util Variable	100
Enable CORS for a Resource	101
Prerequisites	102
Enable CORS Using the Console	102
Enable CORS Using Swagger Definition	103
Use an API Key	105
Prerequisites	105
Use an API Key with the API Gateway Console	105
Enable Client-Side SSL Authentication of an API	106
Generate a Client Certificate	106
Configure an API to Use SSL Certificates	107
Test Invoke	107
Configure Back End to Authenticate API	108
Enable Custom authorization	108
Custom authorization Overview	108
Create the Custom Authorizer Lambda Function	109
Input to a Custom Authorizer	111
Output from a Custom Authorizer	111
Configure Custom authorization	112
Call an API with Custom authorization	114
Import and Export API	116
Import an API	117
Export an API	120
API Gateway Extensions to Swagger	122
Create an API as an Amazon S3 Proxy	132
Create an IAM Role and Policy for the API to Access Amazon S3	132
Create API Resources for Amazon S3 Features	134
Expose a GET Method on an API Root as Get Service Action in Amazon S3	135
Expose Methods on an API Folder Resource as Bucket Actions in Amazon S3	138
Expose Methods on an API Item in a Folder as Actions on an Amazon S3 Object in a Bucket	140

Swagger Definitions of a Sample API as an Amazon S3 Proxy	141
Create an API as a Lambda Proxy	150
Set Up an IAM Role and Policy for an API to Invoke Lambda Functions	151
Create a Lambda Function in the Back End	152
Create API Resources for the Lambda Function	153
Create a GET Method with Query Strings to Call the Lambda Function	154
Create a POST Method with a JSON Payload to Call the Lambda Function	156
Create a GET Method with Path Parameters to Call the Lambda Function	158
Swagger Definitions of a Sample API as Lambda Proxy	162
Create an API as an Amazon Kinesis Proxy	165
Create an IAM Role and Policy for the API to Access Amazon Kinesis	166
Start to Create an API as an Amazon Kinesis Proxy	167
List Streams in Amazon Kinesis	168
Create, Describe, and Delete a Stream in Amazon Kinesis	170
Get Records from and Add Records to a Stream in Amazon Kinesis	178
Swagger Definitions of an API as a Kinesis Proxy	187
Maintaining an API	196
View a List of APIs	196
Prerequisites	196
View a List of APIs with the API Gateway Console	196
Delete an API	196
Prerequisites	197
Delete an API with the API Gateway Console	197
Delete a Resource	197
Delete a Resource with the API Gateway Console	197
View a Methods List	197
Prerequisites	197
View a Methods List with the API Gateway Console	198
Delete a Method	198
Delete a Method with the API Gateway Console	198
Deploying an API	199
Deploy an API with the API Gateway Console	199
Prerequisites	199
Deploy an API with the API Gateway Console	199
Update deployment configuration with the API Gateway Console	200
Change a Stage to Use a Different Deployment with the API Gateway Console	200
Manage API Request Throttling	201
Account-Level Throttling	201
Stage-Level and Method-Level Throttling	201
Deploy an API in Stages	201
Create a Stage	201
View a List of Stages	202
Specify a Stage's Settings	202
Delete a Stage	205
Enable API Caching	205
API Caching Overview	205
Enable API Caching	206
Override Stage Caching for Method Caching	207
Use Method/Integration Parameters as Cache Keys	207
Flush the API Stage Cache in API Gateway	208
Invalidate an API Gateway Cache Entry	209
Deploy an API with Stage Variables	210
Use Cases	211
Examples	211
Set Stage Variables	212
Use Stage Variables	214
Stage Variables Reference	220
Generate an SDK for an API	221

Prerequisites	222
Generate an SDK for an API with the API Gateway Console	222
Integrate an API Gateway-Generated Android SDK into Your Android Project	223
Integrate an API Gateway-Generated iOS SDK into Your iOS Project	224
Integrate an API Gateway-Generated JavaScript SDK into Your JavaScript Code	225
Use a Custom Domain Name	227
Prerequisites	227
Specify a Custom Domain Name with the API Gateway Console	228
Specify API Mappings for a Custom Domain Name with the API Gateway Console	229
Call Your API with Custom Domain Names	229
Calling an API	231
Prerequisites	231
Call an API with the API Gateway Console	231
Use the API Dashboard	232
Prerequisites	232
Use the API Dashboard in the API Gateway Console	232
Test a Method	233
Prerequisites	233
Test a Method with the API Gateway Console	233
Log API Calls	234
API Gateway Information in CloudTrail	234
Understanding API Gateway Log File Entries	234
Access Permissions	236
Create and Attach a Policy to an IAM User	237
Action and Resource Syntax for Managing Entities in API Gateway	238
Permissions to call an API in API Gateway	240
Limits and Pricing	243
API Gateway Pricing	243
API Gateway Limits	243
Known Issues	244
Control Service API	245
Document History	246
AWS Glossary	248

What Is Amazon API Gateway?

Amazon API Gateway consists of two services: the API Gateway control service and API Gateway execution service. The control service lets you create a RESTful API to expose selected back-end features. The back end can be another AWS service, such as AWS Lambda or Amazon DynamoDB, or it can be an existing web application. The execution service lets an app call the API to access the exposed back-end features. The app can interact with the API using standard HTTP protocols or using a platform- or language-specific SDK generated by the API creator.

The API you create in API Gateway consists of a set of resources and methods. A resource is a logical entity that can be accessed through a resource path using the API. A resource can have one or more operations that are defined by appropriate HTTP verbs such as GET, POST, and DELETE. A combination of a resource path and an operation identify a method in the API. Each method corresponds to a REST API request submitted by the user of your API and the corresponding response returned to the user. API Gateway integrates the method with a targeted back end by mapping the method request to an integration request acceptable by the back end and then mapping the integration response from the back end to the method response returned to the user. As an API developer, you can configure how methods are mapped to integrations and vice versa by stipulating what parameters to use and specifying mapping templates to transform payloads of given data models.

You can create an API by using the API Gateway management console, described in [Getting Started \(p. 3\)](#), or by using the [API Gateway Control Service API \(p. 245\)](#). In addition, you can integrate API creation with [AWS CloudFormation templates](#) or [API Gateway Extensions to Swagger \(p. 122\)](#). For a list of regions where API Gateway is available, as well as the associated control service endpoints, see [Regions and Endpoints](#).

API Gateway helps developers deliver robust, secure, and scalable mobile and web application back ends. API Gateway allows developers to securely connect mobile and web applications to business logic hosted on AWS Lambda, APIs hosted on Amazon EC2, or other publicly addressable web services hosted inside or outside of AWS. With API Gateway, developers can create and operate APIs for their back-end services without developing and maintaining infrastructure to handle authorization and access control, traffic management, monitoring and analytics, version management, and software development kit (SDK) generation.

API Gateway is designed for web and mobile developers who want to provide secure, reliable access to back-end APIs for access from mobile apps, web apps, and server apps that are built internally or by third-party ecosystem partners. The business logic behind the APIs can either be provided by a publicly accessible endpoint that API Gateway proxies call, or it can be entirely run as a Lambda function.

To better understand the terminology used in this documentation, you may find it useful to peruse the [API Gateway Concepts \(p. 2\)](#) section.

Amazon API Gateway Concepts

API developer or API owner	An AWS account that owns an API Gateway deployment (for example, a service provider who also supports programmatic access.)
App developer or client developer	An app creator who may or may not have an AWS account and interacts with the API deployed by the API developer. An app developer can be represented by an API Key.
App user, end user or client endpoint	An entity that uses the application built by an app developer that interacts with APIs in Amazon API Gateway. An app user can be represented by an Amazon Cognito identity or a bearer token.
Method request	The public interface of an API method in API Gateway that defines the parameters and body that an app developer must send in the requests to access the back end through the API.
Integration request	An API Gateway internal interface that defines how API Gateway maps the parameters and body of a method request into the formats required by the back end.
Integration response	An API Gateway internal interface that defines how API Gateway maps data. The integration response includes the status codes, headers, and payload that are received from the back end into the formats defined for an app developer.
Method response	The public interface of an API that defines the status codes, headers, and body models that an app developer should expect from API Gateway.

Getting Started with Amazon API Gateway

The following walkthroughs include hands-on exercises, using the API Gateway console, to help you learn about API Gateway.

Topics

- [Get Ready to Use Amazon API Gateway \(p. 3\)](#)
- [Build and Test an API Gateway API from an Example \(p. 5\)](#)
- [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#)
- [Make Synchronous Calls to Lambda Functions \(p. 20\)](#)
- [Map API Gateway API Request Parameters for HTTP Endpoints \(p. 30\)](#)
- [Use Models and Mapping Templates for Payload Mappings \(p. 39\)](#)
- [Get Started with Creating API Gateway API as an AWS Service Proxy \(p. 52\)](#)

Get Ready to Use Amazon API Gateway

Before using API Gateway for the first time, you must have an AWS account set up. To create, configure and deploy an API in API Gateway, you must have appropriate IAM policy provisioned with permissible access rights to the API Gateway control service. To permit your API clients to invoke your API in API Gateway, you must set up the right IAM policy to allow the clients to call the API Gateway execution service. To allow API Gateway to invoke an AWS service in the back end, API Gateway must have permissions to assume the roles required to call the back-end AWS service. When an API Gateway API is set up to use AWS IAM roles and policies to control client access, the client must sign API Gateway API requests with [Signature Version 4](#).

Understanding of these topics are important to use API Gateway and to follow the tutorials and instructions presented here. This section provides brief discussions of or quick references to these topics.

Topics

- [Sign Up for AWS \(p. 4\)](#)
- [Create an IAM User, Group or Role in Your AWS Account \(p. 4\)](#)
- [Grant IAM Users Permissions to Access API Gateway Control and Execution Services \(p. 4\)](#)
- [Next Step \(p. 5\)](#)

Sign Up for AWS

Go to <http://aws.amazon.com/>, choose **Create an AWS Account**, and follow the instructions therein.

Create an IAM User, Group or Role in Your AWS Account

For better security practices, you should refrain from using your AWS root account to access API Gateway. Instead, create a new AWS Identity and Access Management (IAM) user or use an existing one in your AWS account, and then access API Gateway with that IAM user credentials.

To manage access for a user, you can create an IAM user, grant the user API Gateway access permissions. To create a new IAM user, see [Creating an IAM User](#).

To manage access for a group of users, you can create an IAM group, grant the group API Gateway access permissions and then add one or more IAM users to the group. To create an IAM group, see [Creating IAM Groups](#).

To delegate access to specific users, apps or service, you can create an IAM role, add the specified users or groups to the role, and grant the users or groups API Gateway access permissions. To create an IAM role, see [Creating IAM Roles](#).

When setting up your API, you need to specify the ARN of an IAM role to control access the API's methods. Make sure that this is ready when creating an API.

Grant IAM Users Permissions to Access API Gateway Control and Execution Services

In AWS, access permissions are stated as [policies](#). A policy created by AWS is a managed policy and one created by a user is an inline policy.

For the API Gateway control service, the managed policy of **AmazonAPIGatewayAdministrator** (`arn:aws:iam::aws:policy/AmazonAPIGatewayAdministrator`) grants the full access to create, configure and deploy an API in API Gateway:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "apigateway:*"  
            ],  
            "Resource": "arn:aws:apigateway:*::/*"  
        }  
    ]  
}
```

To grant the stated permissions to a user, attach the policy to the user, a group containing the user. To attach a policy, see [Attaching Managed Policies](#).

Attaching the preceding policy to an IAM user provides the user with access to all API Gateway control service actions and resources associated with the AWS account. To learn how to restrict IAM users to a limited set of API Gateway control service actions and resources, see [Access Permissions \(p. 236\)](#).

For the API Gateway execution service, the managed policy of **AmazonAPIGatewayInvokeFullAccess** (`arn:aws:iam::aws:policy/AmazonAPIGatewayInvokeFullAccess`) provides full access to invoke an API in API Gateway:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "execute-api:Invoke"  
            ],  
            "Resource": "arn:aws:execute-api:*.*.*"  
        }  
    ]  
}
```

Attaching the preceding policy to an IAM user provides the user with access to all API Gateway execution service actions and resources associated with the AWS account. To learn how to restrict IAM users to a limited set of API Gateway execution service actions and resources, see [Access Permissions \(p. 236\)](#).

To grant the state permissions to a user, attach the policy to the user, a group containing the user. To attach a policy, see [Attaching Managed Policies](#).

In this documentation, we will use managed policies, whenever possible. To create and use inline policies, see [Working with Inline Policies](#).

Note

To complete the steps above, you must have permission to create the IAM policy and attach it to the desired IAM user.

Next Step

You are now ready to start using API Gateway. See [Build and Test an API Gateway API from an Example \(p. 5\)](#).

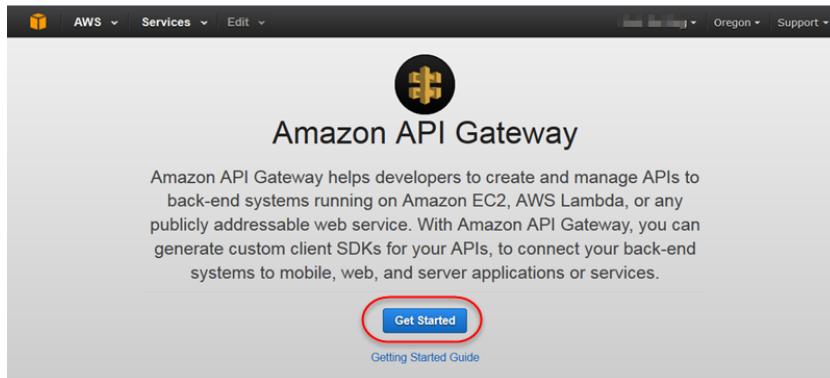
Build and Test an API Gateway API from an Example

The Amazon API Gateway console now provides an option for you to create an API Gateway API by example, with helpful hints provided along the way. If you are new to API Gateway, you may find it useful as a learning tool. The following steps walk you through using this create-by-example option to create and test the example API.

1. Do one of the following:
 - a. For the first API in your account, choose **Get Started** from the API Gateway console welcome page:

Amazon API Gateway Developer Guide

Build and Test an API from Example



Streamline API development
Amazon API Gateway lets you simultaneously run multiple versions and release stages of the same API, allowing you to quickly iterate, test, and release new versions.



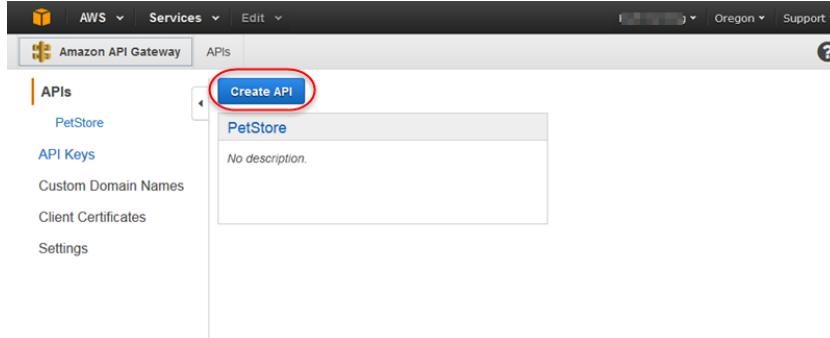
Performance at scale
Amazon API Gateway helps you improve performance by managing traffic to your existing back-end systems, throttling API call spikes, and enabling result caching.



SDK generation
Amazon API Gateway can generate client SDKs for JavaScript, Java, iOS, and Android, which you can use to quickly test new APIs from your applications and distribute SDKs to third-party developers.

If prompted with a modal dialog box containing hints at a stage of the process, choose **OK** to close the modal dialog and continue.

- b. For your next API, choose **Create API** from the API Gateway **APIs** home page:



2. Under **Create new API**, select **Examples API** and then choose **Import** to create the example API. For your first API, the API Gateway console will start with this option as default.

Create new API

In Amazon API Gateway, an API refers to a collection of resources and methods that can be invoked through HTTPS endpoints.

New API Clone from existing API Import from Swagger Example API

Example API

Learn about the service by importing an example API and turning on hints throughout the console.



```

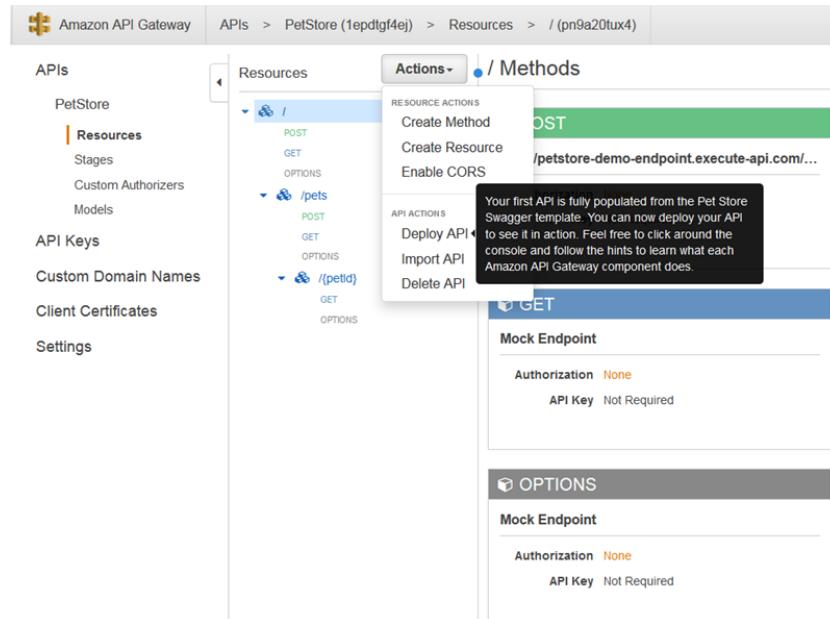
1
2   "swagger": "2.0",
3   "info": {
4     "title": "PetStore"
5   },
6   "schemes": [
7     "https"
8   ],
9   "paths": {
10    "/":
11      "post": {
12        "produces": [
13          "application/json"
14        ],
15        "responses": {
16          "200": {
17            "description": "200 response",
18            "schema": {
19              "$ref": "#/definitions/Empty"
20            }
21          }
22        }
23      }
24    }
25  }
26}

```

Import

You can scroll down the Swagger definition for details of this example API before choosing **Import**.

- The resulting display shows the newly created API:

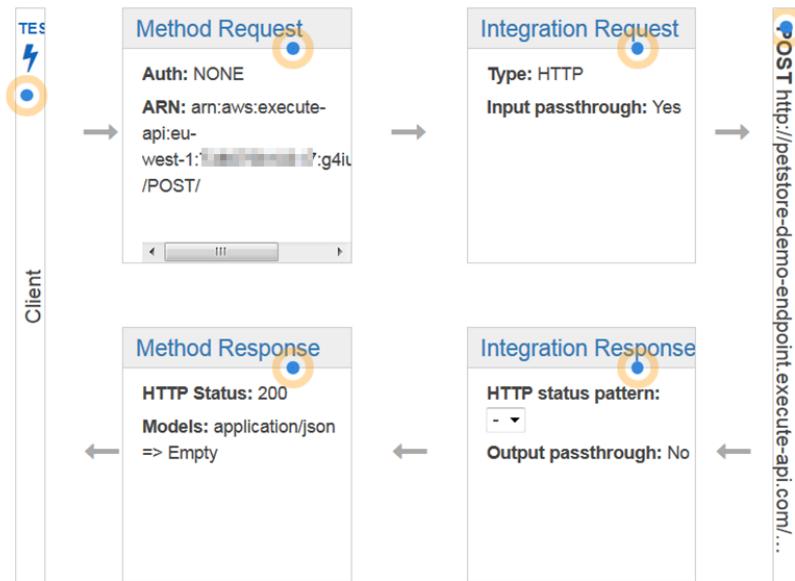


The API Gateway navigation pane on the left shows your available APIs, any API keys, custom domain names and client certificates that you created for your APIs, as well as the settings for logging your APIs' performance metrics. API-specific resources, deployment, custom authorizers and payload-mapping data models are organized under individual APIs.

The **Resources** pane in the middle shows the structure of the selected API as a tree of nodes. API methods defined on each resource are edges of the tree. When a resource is selected, all of its methods are listed in the **Methods** pane on the right. Displayed under each method is a brief summary of the method, including its endpoint URL, authorization type, and API Key requirement.

- To view the details of a method, to modify its set-up, or to test the method invocation, choose the method name from either the method list or the resource tree.

/ - POST - Method Execution



The resulting **Method Execution** pane for the chosen method presents a logical view of the method's structure and behaviors: a client accesses a back-end service by interacting with the API through **Method Request**. API Gateway translates the client request, if necessary, into the form acceptable to the back end before forwarding the request to the back end. The transformed request is known as the integration request and is depicted by **Integration Request** in the display. Similarly, the response from the back end goes through **Integration Response** and then **Request Response** before being received by the client. Again, if necessary, API Gateway maps the response from the form shaped in the back end to a form expected by the client.

For the POST method on this API's root (/) resource, the method's integration request shows that the method is integrated with the endpoint of

`https://http://petstore-demo-endpoint.execute-api.com/petstore/pets` in the back end. The method request payload will be passed through to the integration request without modification. The GET / method request uses the MOCK integration type and is not tied to any endpoint in the back end. When the method is called, the API Gateway simply accepts the request and immediately returns a response, by way of from **Integration Response** to **Method Response**. You can use the mock integration to test an API without requiring a back-end endpoint. You can also use it to serve a local response. In fact, the example API uses it to return a local HTML page as the home page of the API. It uses a mapping template to generate the home page in **Integration Response**.

As an API developer, you control the behaviors of your API's front-end interactions by configuring the method request and a method response. You control the behaviors of your API's back-end interactions by setting up the integration request and integration response. They involve data mappings between a method and its corresponding integration. We will cover the method setup in [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#). For now, we focus on testing the API to provide an end-to-end user experience.

5. Choose **Test** shown on **Client** (as shown in the previous image) to start testing. Enter the following `{"type": "dog", "price": 249.99}` payload into the **Request Body** before choosing the **Test** button.

[Method Execution](#) / - POST - Method Test

Make a test call to your method with the provided input

Path

No path parameters exist for this resource. You can define path parameters by using the syntax **{myPathParam}** in a resource path.

Query Strings

No query string parameters exist for this method. You can add them via Method Request.

Headers

No header parameters exist for this method. You can add them via Method Request.

Stage Variables

No **stage variables** exist for this method.

Client Certificate

No client certificates have been generated.

Request Body

1 `{"type": "dog", "price": 249.99}`

 Test

The input specifies the attributes of the pet that we wish to add to the list of pets on the PetStore website.

6. The results display as follows:

```
Request: /
Status: 200
Latency: 1445 ms
Response Body

{
  "pet": {
    "type": "dog",
    "price": 249.99
  },
  "message": "success"
}

Response Headers

{"Access-Control-Allow-Origin":"*","Content-Type":"application/json"}

Logs

Execution log for request test-request
Mon Apr 04 04:59:05 UTC 2016 : Starting execution for request: test-invoke-request
Mon Apr 04 04:59:05 UTC 2016 : HTTP Method: POST, Resource Path: /
Mon Apr 04 04:59:05 UTC 2016 : Method request path: {}
Mon Apr 04 04:59:05 UTC 2016 : Method request query string: {}
Mon Apr 04 04:59:05 UTC 2016 : Method request headers: {}
Mon Apr 04 04:59:05 UTC 2016 : Method request body before transformations: {"type": "dog","price": 249.99}
Mon Apr 04 04:59:05 UTC 2016 : Endpoint request URI: http://petstore-demo-endpoint.execute-api.com/petstore/pets
Mon Apr 04 04:59:05 UTC 2016 : Endpoint request headers: {x-amzn-apigateway-api-id=g4iukm23bf, Accept=application/json, User-Agent=AmazonAPIGateway_g4iukm23bf, Content-Type=application/json}
Mon Apr 04 04:59:05 UTC 2016 : Endpoint request body after transformations: {"type": "dog","price": 249.99}
Mon Apr 04 04:59:06 UTC 2016 : Endpoint response body before transformations: {
  "pet": {
    "type": "dog",
    "price": 249.99
  },
  "message": "success"
}
Mon Apr 04 04:59:06 UTC 2016 : Endpoint response headers: {date=Mon, 04 Apr 2016 04:59:06 GMT, content-length=81, x-powered-by=Express, content-type=application/json; charset=utf-8, connection=keep-alive}
Mon Apr 04 04:59:06 UTC 2016 : Method response body after transformations: {
  "pet": {
    "type": "dog",
    "price": 249.99
  },
  "message": "success"
}
Mon Apr 04 04:59:06 UTC 2016 : Method response headers: {Access-Control-Allow-Origin=*, Content-Type=application/json}
Mon Apr 04 04:59:06 UTC 2016 : Successfully completed execution
Mon Apr 04 04:59:06 UTC 2016 : Method completed with status: 200
```

The **Logs** entry of the output shows the state changes from the method request to the integration request and from the integration response to the method response. This can be useful for troubleshooting any mapping errors that cause the request to fail. In this example, no mapping is applied: the method request is identical to the integration request and the integration response is the same as the method response.

To test the API using a client other than the API Gateway test-invoke-request feature, you must first deploy the API to a stage.

7. To deploy the sample API, select the **PetStore** API and the root / resource, and then choose **Deploy API** from the **Actions** menu.

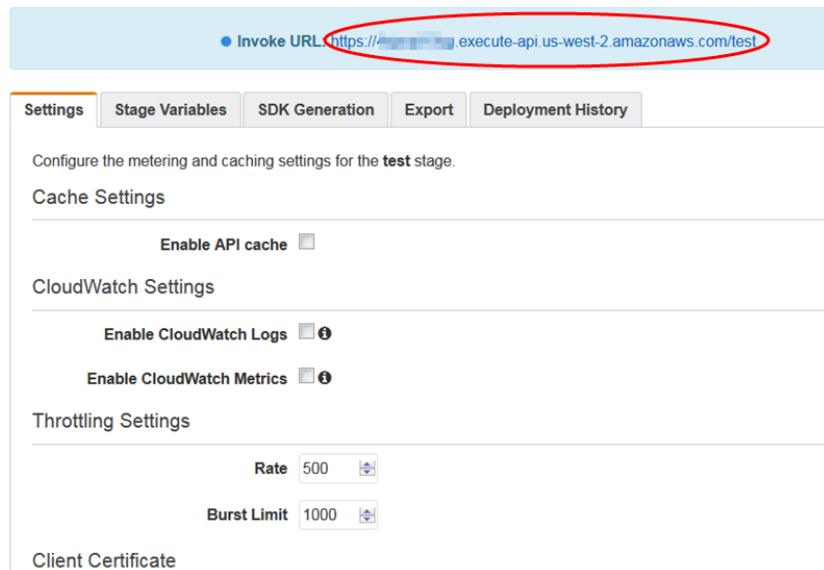
The screenshot shows the AWS Lambda console interface. On the left, there's a sidebar with 'APIs' and 'PetStore' selected. Under 'PetStore', 'Resources' is expanded, showing three resources: '/', '/pets', and '/{petId}'. Each resource has associated 'POST', 'GET', and 'OPTIONS' methods. On the right, the main area shows the 'Actions' dropdown menu open. The 'API ACTIONS' section contains 'Create Method', 'Create Resource', 'Enable CORS', and 'Deploy API'. The 'Deploy API' option is circled in red. Below it are 'Import API' and 'Delete API'. To the right of the actions, there are three sections: 'POST' (with endpoint information), 'GET' (labeled 'Mock Endpoint'), and 'OPTIONS' (also labeled 'Mock Endpoint').

In **Deploy API**, for **Deployment stage**, choose **[New Stage]** because this is the first deployment of the API. Type a name (e.g., `test`) in **Stage name** and, optionally, type descriptions in **Stage description** and **Deployment description**. Choose **Deploy**.

The screenshot shows the 'Deploy API' dialog box. It has fields for 'Deployment stage' (set to '[New Stage]'), 'Stage name*' (set to 'test'), 'Stage description' (set to 'test stage'), and 'Deployment description' (set to 'sample API first deployment'). At the bottom right are 'Cancel' and 'Deploy' buttons, with 'Deploy' circled in red.

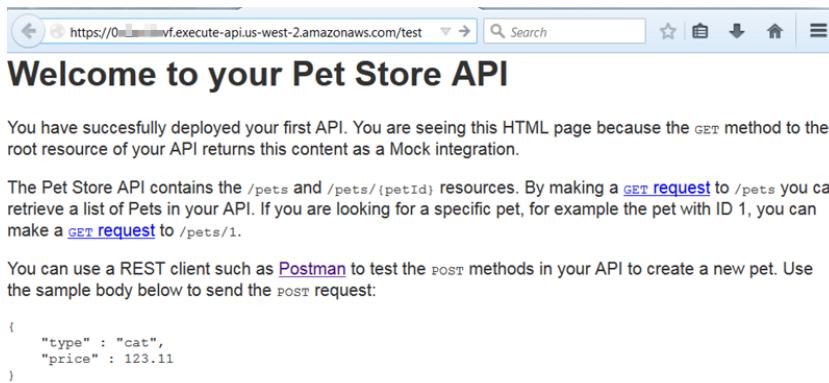
In the resulting **Stage Editor** pane, **Invoke URL** displays the URL to invoke the API's `GET /` method request.

test Stage Editor



The screenshot shows the 'test' Stage Editor interface. At the top, there's a blue header bar with a red circle highlighting the 'Invoke URL' link, which contains the URL [https://\[REDACTED\].execute-api.us-west-2.amazonaws.com/test](https://[REDACTED].execute-api.us-west-2.amazonaws.com/test). Below the header are tabs for 'Settings', 'Stage Variables', 'SDK Generation', 'Export', and 'Deployment History'. The 'Settings' tab is selected. Under 'Settings', there are sections for 'Cache Settings' (with 'Enable API cache' checked), 'CloudWatch Settings' (with 'Enable CloudWatch Logs' and 'Enable CloudWatch Metrics' checked), and 'Throttling Settings' (with Rate set to 500 and Burst Limit set to 1000). There's also a 'Client Certificate' section.

- On **Stage Editor**, follow the **Invoke URL** link to submit the `GET /` method request in a browser. The result, generated from the mapping template in the integration response, is shown as follows:



The screenshot shows a web browser displaying the 'Welcome to your Pet Store API' page. The URL in the address bar is [https://\[REDACTED\].execute-api.us-west-2.amazonaws.com/test](https://[REDACTED].execute-api.us-west-2.amazonaws.com/test). The page content includes a success message: 'You have successfully deployed your first API. You are seeing this HTML page because the `GET` method to the root resource of your API returns this content as a Mock integration.' It also provides information about the Pet Store API resources (`/pets` and `/pets/{petId}`) and how to make `GET` requests to them. Below this, there's a note about using Postman to test the `POST` methods and a sample JSON body for a `POST` request.

```
{
  "type": "cat",
  "price": 123.11
}
```

- In the **Stages** navigation pane, expand the **test** stage, select **GET** on `/pets/{petId}`, and then copy the **Invoke URL** value of [{petId}](https://[api-id].execute-api.[region].amazonaws.com/test/pets/{petId}). `{petId}` stands for a path variable.

Paste the **Invoke URL** value (obtained in the previous step) into the address bar of a browser, replacing `{petId}` by, for example, 1, and press Enter to submit the request. A 200 OK response should return with the following JSON payload:

```
{
  "id": 1,
  "type": "dog",
  "price": 249.99
}
```

Invoking the API method as shown is possible because its **Authorization** type is set to **NONE**. If the **AWS_IAM** authorization were used, you would sign the request using the Signature Version 4 protocols. For an example of such a request, see [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).

Build and Deploy an API Gateway API Step by Step

You can create an API in the Amazon API Gateway console from the ground up. In essence, you use the console as an API design studio to scope the API features, to experiment with its behaviors, to build the API, and to deploy your API in stages.

This section walks you through the steps to create resources, expose methods on a resource, configure a method to achieve the desired API behaviors, and to test and deploy the API.

1. From **Create new API**, select **New API**, type a name in **API Name**, optionally add a description in **Description**, and then choose **Create API**.

Create new API

In Amazon API Gateway, an API refers to a collection of resources and methods that can be invoked through HTTPS endpoints.

New API Clone from existing API Import from Swagger Example API

Name and description

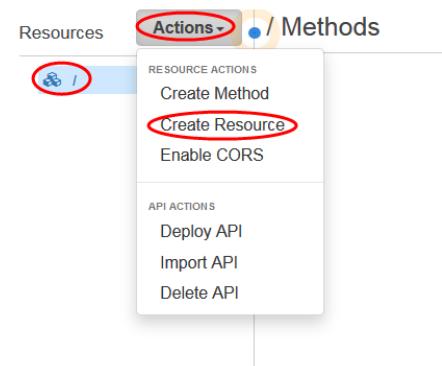
Choose a friendly name and description for your API.

API name* Description

* Required

As a result, an empty API is created. The **Resources** tree shows the root resource (/) without any methods. In this exercise, we will build the API as an HTTP proxy of the PetStore demo website (<http://petstore-demo-endpoint.execute-api.com/>) For illustration purposes, we will create a /pets resource as a child of the root and expose a GET method on this resource for a client to retrieve a list of available Pets items from the PetStore website.

2. To create the /pets resource, select the root, choose **Actions** and then choose **Create Resource**.



Type **Pets** in **Resource Name**, leave the **Resource Path** value as given, and choose **Create Resource**.

New Child Resource

Use this page to create a new child resource for your resource.

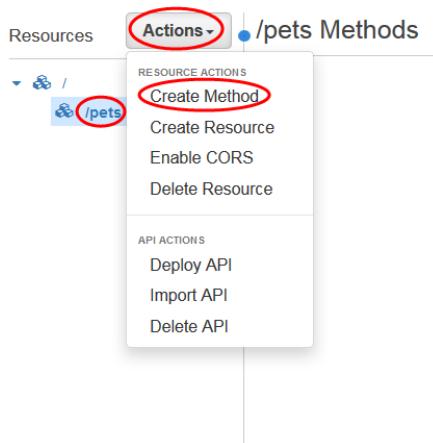
Resource Name*

Resource Path*

You can add path parameters using brackets. For example, the resource path `{username}` represents a path parameter called 'username'.

* Required [Cancel](#) Create Resource

3. To expose a GET method on the `/pets` resource, choose **Actions** and then **Create Method**.



Choose **GET** from the list under the `/pets` resource node and choose the checkmark icon to finish creating the method.



The method created is not yet integrated with the back end. The next step sets this up.

4. In the method's **Setup** pane, select **HTTP Proxy** for **Integration type**, select **GET** from the **HTTP method** drop-down list, type `http://petstore-demo-endpoint.execute-api.com/petstore/pets` as the **Endpoint URL** value, and then choose **Save**.

/pets - GET - Setup

Choose the integration point for your new method. i

Integration type	<input checked="" type="radio"/> Lambda Function <input checked="" type="radio"/> HTTP Proxy <input type="radio"/> Mock Integration
Show advanced	
HTTP method	<input checked="" type="radio"/> GET
Endpoint URL	http://petstore-demo-endpoint.execute-api.com/petstore/pets
<input style="border: 1px solid blue; border-radius: 5px; padding: 2px 10px;" type="button" value="Save"/>	

When the method setup finishes, you are presented with the **Method Execution** pane, where you can further configure the method request to add query string or custom header parameters. You can also update the integration request to map input data from the method request to the format required by the back end.

The PetStore website allows you to retrieve a list of Pet items by the pet type (e.g., "Dog" or "Cat") on a given page. It uses the `type` and `page` query string parameters to accept such input. As such, we must add the query string parameters to the method request and map them into the corresponding query strings of the integration request.

5. In the GET method's **Method Execution** pane, choose **Method Request**, select `AWS_IAM` for **Authorization**, expand the **URL Query String Parameters** section, and choose **Add query string** to create two query string parameters named `type` and `page`. Choose the checkmark icon to save the newly added query string parameters.

[Method Execution](#) /pets - GET - Method Request

Provide information about this method's authorization settings and the parameters it can receive.

Authorization Settings

Authorization **AWS_IAM** i

API Key Required false

URL Query String Parameters

Name	Caching	
type	<input type="checkbox"/>	
page	<input type="checkbox"/>	

[Add query string](#)

HTTP Request Headers

Request Models [Create a Model](#)

The client can now supply a pet type and a page number as query string parameters when submitting a request. These input parameters must be mapped into the integration's query string parameters to forward the input values to our PetStore website in the back end. Because the method uses AWS_IAM, you must sign the request to invoke the method.

6. From the method's **Integration Request** page, expand the **URL Query String Parameters** section. By default, the method request query string parameters are mapped to the like-named integration request query string parameters. This default mapping works for our demo API. We will leave them as given. To map a different method request parameter to the corresponding integration request parameter, choose the pencil icon for the parameter to edit the mapping expression, shown in the **Mapped from** column. To map a method request parameter to a different integration request parameter, first choose the delete icon to remove the existing integration request parameter, choose **Add query string** to specify a new name and the desired method request parameter mapping expression.

[◀ Method Execution /pets - GET](#) **- Integration Request**

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration

[Show advanced](#)

HTTP method **GET**

Endpoint URL <http://petstore-demo-endpoint.execute-api.com/petstore/pets>

▶ URL Path Parameters

▼ URL Query String Parameters

Name	Mapped from	Caching
type	method.request.querystring.type	
page	method.request.querystring.page	

[Add query string](#)

▶ HTTP Headers

▶ Body Mapping Templates

This completes building the simple demo API. It's time to test the API.

7. To test the API using the API Gateway console, choose **Test** from the GET-on-Pets method's **Method Execution** pane. In the **Method Test** pane, enter **Dog** and **2** for the **type** and **page** query strings, respectively, and then choose **Test**.

[Method Execution](#) /pets - GET - **Method Test**

Make a test call to your method with the provided input

Path

No path parameters exist for this resource. You can define path parameters by using the syntax `{myPathParam}` in a resource path.

Query Strings

type

Dog

page

2

Headers

No header parameters exist for this method. You can add them via Method Request.

Stage Variables

No stage variables exist for this method.

Client Certificate

No client certificates have been generated.

Request Body

Request Body is not supported for GET methods.



The result is shown as follows. (You may need to scroll down to see the test result.)

Request: /pets?type=Dog&page=2

Status: 200

Latency: 1036 ms

Response Body

```
[  
  {  
    "id": 4,  
    "type": "Dog",  
    "price": 999.99  
  },  
  {  
    "id": 5,  
    "type": "Dog",  
    "price": 249.99  
  },  
  {  
    "id": 6,  
    "type": "Dog",  
    "price": 49.97  
  }  
]
```

Response Headers

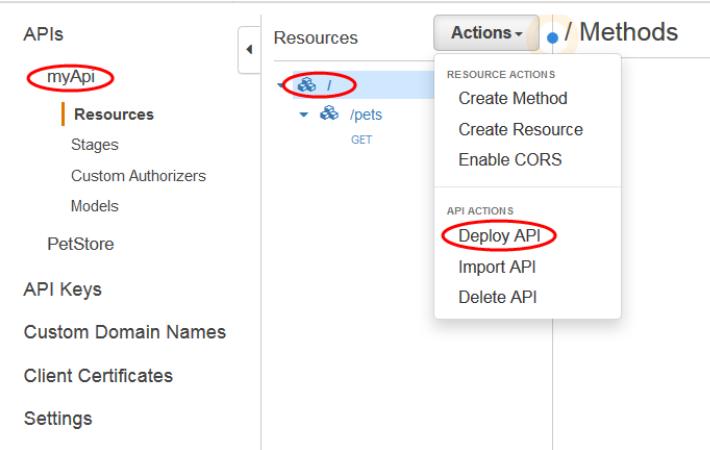
```
{"Content-Type": "application/json"}
```

Logs

```
Execution log for request test-request  
Mon Apr 04 05:48:01 UTC 2016 : Starting execution for request: test-invoke-request  
Mon Apr 04 05:48:01 UTC 2016 : HTTP Method: GET, Resource Path: /pets  
Mon Apr 04 05:48:01 UTC 2016 : Method request path: {}  
Mon Apr 04 05:48:01 UTC 2016 : Method request query string: {page=2, type=Dog}  
Mon Apr 04 05:48:01 UTC 2016 : Method request headers: {}  
Mon Apr 04 05:48:01 UTC 2016 : Method request body before transformations: null
```

Now that the test is successful, we can deploy the API to make it publicly available.

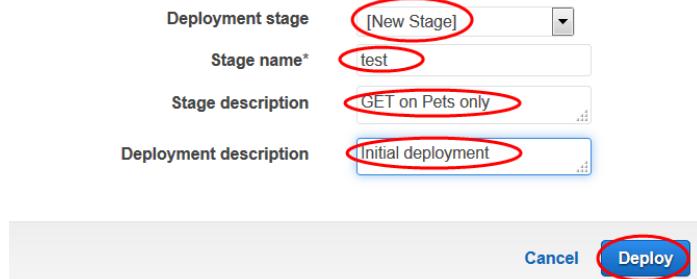
8. To deploy the API, select the API and then choose **Deploy API** from the **Actions** drop-down menu.



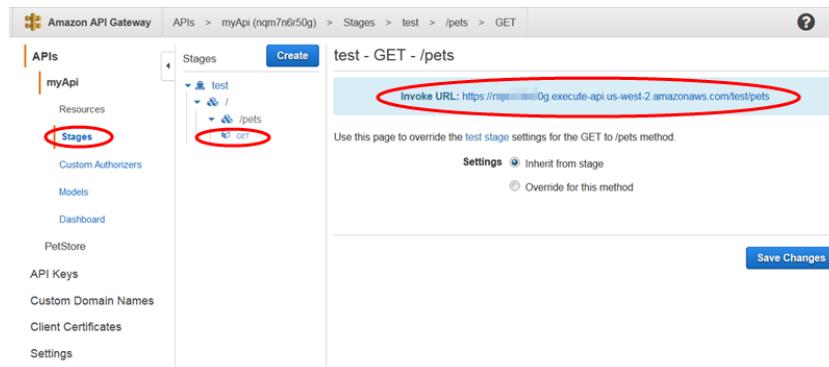
In the **Deploy API** dialog, choose a stage (or [New Stage] for the API's first deployment); enter a name (e.g., "test", "prod", "dev", etc.) in the **Stage name** input field; optionally, provide a description in **Stage description** and/or **Deployment description**; and then choose **Deploy**.



Choose a stage where your API will be deployed. For example, a test version of your API could be deployed to a stage named beta.



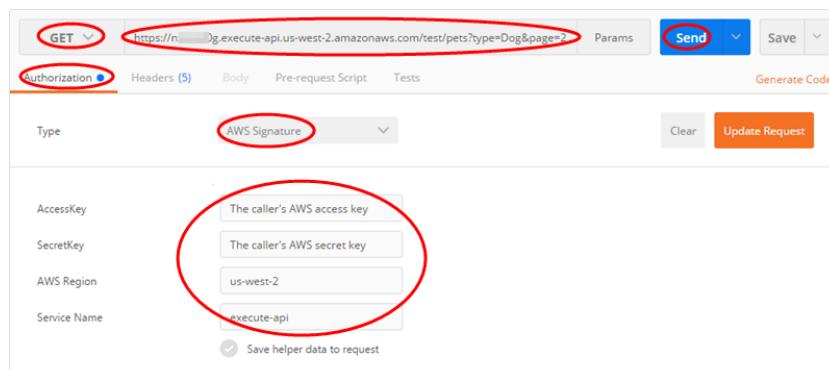
Once deployed, you can obtain the invocation URLs (**Invoke URL**) of the API's endpoints. For example, the GET on Pets method's invocation URL is as follows:



To invoke this API method in a client (e.g., a [Postman](#) browser), append the query string parameters to the stage-specific method invocation URL (as shown in the previous image) to create the complete method request URL:

```
https://api-id.execute-api.region.amazonaws.com/test/pets?type=Dog&page=2
```

Specify this URL in the address bar of the browser. Choose **GET** as the HTTP verb. Select **AWS Signature** for the **Authorization** type and then specify the required properties (as shown), following the [Signature Version 4](#) protocols. Finally, send the request.



If you use an SDK to create a client, you can call the methods exposed by the SDK to sign the request. For implementation details, see the [AWS SDK](#) of your choosing.

Note

When changes are made to your API, you must redeploy the API to make the new or updated features available before invoking the request URL again.

Make Synchronous Calls to Lambda Functions

AWS Lambda provides an easy way to build back ends without managing servers. API Gateway and Lambda together can be powerful to create and deploy serverless Web applications. In this walkthrough, you learn how to create Lambda functions and build an API Gateway API to enable a Web client to call the Lambda functions synchronously. For more information about Lambda, see the [AWS Lambda Developer Guide](#). For information about asynchronous invocation of Lambda functions, see [Create an API as a Lambda Proxy](#) (p. 150).

Topics

- [Step 1: Prerequisites](#) (p. 20)
- [Step 2: Create an API](#) (p. 20)
- [Step 3: Create a Resource](#) (p. 20)
- [Step 4: Create Lambda Functions](#) (p. 21)
- [Step 5: Create and Test a GET Method](#) (p. 24)
- [Step 6: Create and Test a POST Method](#) (p. 25)
- [Step 7: Deploy the API](#) (p. 26)
- [Step 8: Test the API](#) (p. 26)
- [Step 9: Clean Up](#) (p. 27)
- [Next Steps](#) (p. 28)
- [Create Lambda Invocation and Execution Roles](#) (p. 28)

Step 1: Prerequisites

You must grant API Gateway access permission to the IAM user who will perform the tasks discussed here. The IAM user must have full access to work with Lambda. For this, you can customize the managed policy of **AWSLambdaFullAccess** (`arn:aws:iam::aws:policy/AWSLambdaFullAccess`) and attach it to the IAM user. For more information, see [Get Ready to Use API Gateway](#) (p. 3). The IAM user must also be allowed to create policies and roles in IAM. For this you can customize the managed policy of **IAMFullAccess** (`arn:aws:iam::aws:policy/IAMFullAccess`) and attach it to the user.

Step 2: Create an API

In this step, you will create a new API named `MyDemoAPI`. To create the new API, follow the steps in [Build and Deploy an API Gateway API Step by Step](#) (p. 13).

Step 3: Create a Resource

In this step, you will create a new resource named `MyDemoResource`. To create this resource, follow the steps in [Build and Deploy an API Gateway API Step by Step](#) (p. 13).

Step 4: Create Lambda Functions

Note

Creating Lambda functions may result in charges to your AWS account.

In this step, you will create two new Lambda functions. The first Lambda function, `GetHelloWorld`, will log the call to Amazon CloudWatch and return the JSON object `{"Hello": "World"}`. For more information about JSON, see [Introducing JSON](#).

The second Lambda function, `GetHelloWithName`, will take an input ("name"), log the call to CloudWatch, and return the JSON object `{"Hello": user-supplied-input-value}`. If no input value is provided, the value will be "No-Name".

You will use the Lambda console to create the Lambda functions and set up the required execution role/policy. You will then use the API Gateway console to create an API to integrate API methods with the Lambda functions; the API Gateway console will set up the required Lambda invocation role/policy. If you set up the API without using the API Gateway console, such as when [importing an API from Swagger](#), you must explicitly create, if necessary, and set up an invocation role/policy for API Gateway to invoke the Lambda functions. For more information on how to set up Lambda invocation and execution roles, see [Create Lambda Invocation and Execution Roles \(p. 28\)](#). For more information about Lambda see [AWS Lambda Developer Guide](#).

To create the `GetHelloWorld` Lambda function

1. Open the AWS Lambda console at <https://console.aws.amazon.com/lambda/>.
2. Do one of the following:
 - If the welcome page appears, choose **Get Started Now**.
 - If the **Lambda: Function list** page appears, choose **Create a Lambda function**.
3. From **Select blueprint**, select the **hello-world** blueprint for node.js. You may need to type Hello as the search filter to bring the blueprint in focus.
4. For **Name**, type `GetHelloWorld`.
5. For **Description**, type `Returns {"Hello": "World"}`.
6. For **Runtime**, choose or leave as-is: **Node.js**.
7. Under **Lambda function code**, replace the default code statements in the inline code editor with the following:

```
'use strict';
console.log('Loading event');

exports.handler = function(event, context) {
    console.log('Hello:"World"');
    context.done(null, {"Hello":"World"}); // SUCCESS with message
};
```

Tip

The preceding code is written in Node.js. The `console.log` method writes information to an Amazon CloudWatch Log. The `event` parameter contains the event's data. The `context` parameter contains callback context. Lambda uses `context.done` to perform follow-up actions. For more information about how to write Lambda function code, see the "Programming Model" section in [AWS Lambda: How it Works](#) and the sample walkthroughs in the [AWS Lambda Developer Guide](#).

8. Under **Lambda function handler and role**, leave the default of `index.handler` for **Handler**.
9. For **Role**, choose * Basic execution role under **Create new role**.
 - a. Leave the default selection of `lambda_basic_execution` for **IAM Role**.
 - b. Leave the default selection of `Create a new Role Policy` for **Policy Name**.
 - c. Choose **Allow**.
10. For **Advanced settings** leave the default setting as is.
11. Choose **Next**
12. Choose **Create function**.
13. For the newly created **GetHelloWorld** function, note the AWS region where you created this function. You will need it later.
14. To test the newly created function, as a good practice, choose **Actions** and then select **Configure test event**.
15. For **Input test event**, replace any default code statements with the following, and then choose **Save and test**.

```
{  
}  
}
```

Tip

This function does not use any input. Therefore, we provide an empty JSON object as the input.

16. Choose **Test** to invoke the function. The **Execution result** section shows `{"Hello": "World"}`. The output is also written to CloudWatch Logs.
17. Go to the **Functions** list to create the next Lambda function.

In addition to the Lambda function, an IAM role (`lambda_basic_execution`) is also created as the result of this procedure. You can view this in the IAM console. Attached to this IAM role is the following inline policy that grants users of your AWS account permission to call the CloudWatch `CreateLogGroup`, `CreateLogStreams`, and `PutLogEvents` actions on any of the CloudWatch resources.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "logs:CreateLogGroup",  
                "logs:CreateLogStream",  
                "logs:PutLogEvents"  
            ],  
            "Resource": "arn:aws:logs:*:*:  
        }  
    ]  
}
```

A trusted entity of this IAM role is `lambda.amazonaws.com`, which has the following trust relationship:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "lambda.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

The combination of this trust relationship and the inline policy makes it possible for the user to invoke the Lambda function and for Lambda to call the supported CloudWatch actions on the user's behalf.

To create the **GetHelloWithName** Lambda function

1. Choose **Create a Lambda function**.
2. From **Select blueprint**, select the **hello-world** blueprint for node.js.
3. Type **GetHelloWithName** for **Name**.
4. For **Description**, type **Returns {"Hello": "a user-provided string, and "}**.
5. For **Runtime**, choose **Node.js**.
6. In the code editor under **Lambda function code** replace the default code statements with the following:

```
'use strict';
console.log('Loading event');

exports.handler = function(event, context) {
    var name = (event.name === undefined ? 'No-Name' : event.name);
    console.log('Hello:' + name + '');
    context.done(null, {"Hello":name}); // SUCCESS with message
};
```

7. Under **Lambda function handler and role**, leave the default of `index.handler` for **Handler**.
8. For **Role**, choose `lambda_basic_execution` under **Use existing role**, assuming you have created the `lambda_basic_execution` role in the previous procedure.
9. Leave the default values for **Advanced settings**. Then choose **Next**.
10. Choose **Create function**.
11. For the newly created **GetHelloWorldName** function, note the AWS region where you created this function. You will need it in later steps.
12. To test this newly created function, choose **Actions** and then **Configure test event**.
13. In **Input test event**, replace any default code statements with the following, and then choose **Save and test**.

```
{
    "name": "User"
}
```

Tip

The function calls `context.name` to read the input name. We expect it to return `{"Hello": "User"}`, given the above input.

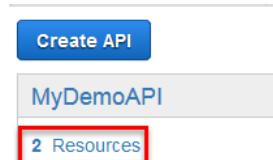
You can experiment with this function by removing `"name": "User"` from the **Invoke** function input and choosing **Save and test** again. You should see the output of `{"Hello": "No-Name"}` under **Execution result** in the Lambda console, as well as in CloudWatch Logs.

Step 5: Create and Test a GET Method

In this step, you will create a new GET method, connect it to your `GetHelloWorld` function in Lambda, and then test it. You use a GET method primarily to retrieve or read a representation of a resource. If successful, the GET method typically returns to the caller the resource representation in JSON. In this walkthrough, the GET method will simply return a JSON-formatted object.

To create and test the GET method

1. Return to the API Gateway console.
2. If the **MyDemoAPI** box is displayed, choose **2 Resources**, as shown in the following screenshot.



3. In the **Resources** pane, choose `/mydemoresource`, and then choose **Create Method**.

Note

If `/mydemoresource` is not displayed, in the secondary navigation bar, choose **APIs** in the first list, **MyDemoAPI** in the second list, and **Resources** in the third list. Then expand the resource root to display `/mydemoresource`.

4. For the HTTP method, choose **GET**, and then choose the check mark button to save your choice.

Note

Other options here include:

- **POST**, which is primarily used to create child resources.
- **PUT**, which is primarily used to update existing resources (and, although not recommended, can be used to create child resources).
- **DELETE**, which is used to delete resources.
- **PATCH**, which can be used to update resources.
- **HEAD**, which is the same as GET but does not return the resource representation. HEAD is used primarily in some testing scenarios.
- **OPTIONS**, which can be used by callers to get information about available communication options for the target service.

5. In the **Setup** pane, for **Integration type**, choose **Lambda Function**.
6. For **Lambda Region**, choose the region identifier that corresponds to the region name in which you created the `GetHelloWorld` Lambda function. For example, if you created this Lambda function in the US East (N. Virginia) region, you would choose `us-east-1`. For a list of region names and identifiers, see [AWS Lambda](#) in the *Amazon Web Services General Reference*.
7. For **Lambda Function**, type `GetHelloWorld`, and then choose **Save**.

8. When you are prompted to give API Gateway permission to invoke your Lambda function, choose **OK**.
9. In the **Method Execution** pane, in the **Client** box, choose **TEST**, and then choose **Test**. If successful, **Response Body** will display the following:

```
{  
    "Hello": "World"  
}
```

By default, API Gateway will pass through the request from the API caller. For the GET method call you just created, as well as for HEAD method calls, a Lambda function will receive an empty JSON response by default and then return the response from the Lambda function without modifications. In the next step, you will create a POST method call. For POST and PUT method calls, you can pass in a request body in JSON format, which the Lambda function will receive as its input event. Optionally, you can transform the input to the Lambda function by using mapping templates in API Gateway.

Step 6: Create and Test a POST Method

In this step, you will create a new POST method, connect it to your `GetHelloWithName` function in Lambda, and then test it. If successful, the POST method typically returns to the caller the URI of the newly created resource. In this walkthrough, the POST method will simply return a JSON-formatted object.

To create and test the POST method

1. In the **Resources** pane, choose `/mydemoresource`, and then choose **Create Method**.
2. For the HTTP method, choose **POST**, and then choose the check mark button to save your choice.
3. In the **Setup** pane, for **Integration Type**, choose **Lambda Function**.
4. For **Lambda Region**, choose the region identifier that corresponds to the region name in which you created the `GetHelloWithName` Lambda function.
5. For **Lambda Function**, type `GetHelloWithName`, and then choose **Save**.
6. When you are prompted to give API Gateway permission to invoke your Lambda function, choose **OK**.
7. In the **Method Execution** pane, in the **Client** box, and then choose **TEST**. Expand **Request Body**, and type the following:

```
{  
    "name": "User"  
}
```

8. Choose **Test**. If successful, **Response Body** will display the following:

```
{  
    "Hello": "User"  
}
```

9. Change **Request Body** by removing `"name": "User"` so that only a set of curly braces (`{ }`) remain, and then choose **Test** again. If successful, **Response Body** will display the following:

```
{  
    "Hello": "No-Name"  
}
```

Step 7: Deploy the API

You are now ready to deploy your API so that you can call it outside of the API Gateway console. In this step, you will create a stage. In API Gateway, a stage defines the path through which an API deployment is accessible. For example, you can define a test stage and deploy your API to it, so that a resource named `MyDemoAPI` is accessible through a URI that ends in `.../test/MyDemoAPI`.

To deploy the API

1. In the **Resources** pane, choose **Deploy API**.
2. For **Deployment stage**, choose **New Stage**.
3. For **Stage name**, type `test`.
4. For **Stage description**, type `This is a test`.
5. For **Deployment description**, type `Calling Lambda functions walkthrough`.
6. Choose **Deploy**.

Step 8: Test the API

In this step, you will go outside of the API Gateway console to call the GET and POST methods in the API you just deployed.

To test the MyDemoGetMethod method

1. In the **Stage Editor** pane, next to **Invoke URL**, copy the URL to the clipboard. It should look something like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test
```

2. In a separate web browser tab or window, paste the URL into the address box. In the URL, append the path to your resource (`/mydemoresource`). The URL should look something like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/mydemoresource
```

3. Browse to this URL. If the GET method is successfully called, the web page will display:

```
{"Hello": "World"}
```

To test the MyDemoPOSTMethod method

1. You will not be able to test a POST method request with your web browser's address bar. Instead, use a web debugging proxy tool or the cURL command-line tool.
2. Send a POST method request to the URL from the previous procedure. The URL should look something like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/mydemoresource
```

Be sure to append to the request headers the following header:

```
Content-Type: application/json
```

Also, be sure to add the following code to the request body:

```
{  
    "name": "User"  
}
```

For example, if you use the cURL command-line tool, run a command similar to the following:

```
curl -H "Content-Type: application/json" -X POST -d "{\"name\": \"User\"}"  
https://my-api-id.execute-api.region-id.amazonaws.com/test/mydemoresource
```

If the POST method is successfully called, the response should contain:

```
{"Hello": "User"}
```

Step 9: Clean Up

If you no longer need the Lambda functions you created for this walkthrough, you can delete them now. You can also delete the accompanying IAM resources.

Caution

If you plan to complete the other walkthroughs in this series, do not delete the Lambda execution role or the Lambda invocation role. If you delete a Lambda function your APIs rely on, those APIs will no longer work. Deleting a Lambda function cannot be undone. If you want to use the Lambda function again, you must re-create the function.

If you delete an IAM resource that a Lambda function relies on, that Lambda function will no longer work, and any APIs that rely on that function will no longer work. Deleting an IAM resource cannot be undone. If you want to use the IAM resource again, you must re-create the resource.

To delete the Lambda functions

1. Sign in to the AWS Management Console and open the AWS Lambda console at <https://console.aws.amazon.com/lambda/>.
2. On the **Lambda: Function List** page, in the list of functions, choose the option button next to **GetHelloWorld**, and then choose **Actions, Delete**. When prompted, choose **Delete** again.
3. Choose the option button next to **GetHelloWithName**, and then choose **Actions, Delete**. When prompted, choose **Delete** again.

To delete the associated IAM resources

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the **Details** area, choose **Roles**.
3. Select **APIGatewayLambdaExecRole**, and then choose **Role Actions, Delete Role**. When prompted, choose **Yes, Delete**.

4. In the **Details** area, choose **Policies**.
5. Select **APIGatewayLambdaExecPolicy**, and then choose **Policy Actions**, **Delete**. When prompted, choose **Delete**.

You have now reached the end of this walkthrough.

Next Steps

You may want to proceed to the next walkthrough, which shows you how to use API Gateway to connect your API to an HTTP endpoint through a REST proxy.

For more information about API Gateway, see [What Is Amazon API Gateway? \(p. 1\)](#). For more information about REST, see [RESTful Web Services: A Tutorial](#).

Create Lambda Invocation and Execution Roles

Before you create these Lambda functions, you must assign appropriate permissions for the functions to execute the specified CloudWatch action (namely, writing to the CloudWatch log) and for API Gateway to invoke the Lambda functions. You set up the permissions using IAM roles and policies for API Gateway to invoke your code and for Lambda to execute your code. For more information about invocation and execution roles/policies in Lambda see [Permission Model](#) in the *AWS Lambda Developer Guide*.

To create the Lambda invocation role and its policy

1. Open the IAM console: <https://console.aws.amazon.com/iam/>.
If using the IAM-managed **AWSLambdaRole** policy, skip Steps 2-7 (below) and proceed from Step 8 to create an invocation role.
2. In the **Details** area, choose **Policies**.
3. Do one of the following:
 - If a list of policies appears, choose **Create Policy**.
 - If the **Welcome to Managed Policies** page appears, choose **Get Started**, and then choose **Create Policy**.
4. Next to **Create Your Own Policy**, choose **Select**.
5. For **Policy Name**, type a name for the policy, for example **APIGatewayLambdaInvokePolicy**.
6. For **Description**, type **Enables API Gateway to call Lambda functions**.
7. For **Policy Document**, type the following, and then choose **Create Policy**.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Resource": ["*"],  
            "Action": ["lambda:InvokeFunction"]  
        }  
    ]  
}
```

```
    ]  
}
```

8. In the **Details** area, choose **Roles**.
9. Choose **Create New Role**.
10. For **Role Name**, type a name for the invocation role, for example `APIGatewayLambdaInvokeRole`, and then choose **Next Step**.
11. Under **Select Role Type**, with the option button next to **AWS Service Roles** already chosen, next to **Amazon API Gateway**, choose **Select**.
12. In **Attach Policy**, if the listed policies contains the one you want, choose it before choosing **Next Step**. Otherwise, simply choose **Next Step** to proceed.
13. For **Role ARN**, make a note of the invocation role's Amazon Resource Name (ARN). You will need this ARN in later steps when you must specify the invocation role explicitly. The ARN should look similar to this: `arn:aws:iam::123456789012:role/APIGatewayLambdaInvokeRole`, where `123456789012` is your AWS account ID.
14. Choose **Create Role**.

The newly created IAM role will have the following policy document created for its Trusted Relationship.

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Sid": "",  
      "Effect": "Allow",  
      "Principal": {  
        "Service": "apigateway.amazonaws.com"  
      },  
      "Action": "sts:AssumeRole"  
    }  
  ]  
}
```

The preceding policy document enables API Gateway to assume roles taken up by and, hence, take actions on behalf of your AWS account.

To create the Lambda execution role and its policies

1. Open the IAM console: <https://console.aws.amazon.com/iam/>.
2. In the **Details** area, choose **Policies**.
3. Choose **Create Policy**.
4. Next to **Create Your Own Policy**, choose **Select**.
5. For **Policy Name**, type a name for the policy (for example, `APIGatewayLambdaExecPolicy`).
6. For **Description**, type `Enables Lambda to execute code`.
7. For **Policy Document**, type the following, and then choose **Create Policy**.

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Action": [
```

```
        "logs:/*"
    ],
    "Effect": "Allow",
    "Resource": "arn:aws:logs:*:*:*"
}
]
```

Note

The preceding policy document permits all log actions on Amazon CloudWatch Logs. Typically, you would also add other permissions required by your Lambda function to interact with AWS services, such as uploading an object to an Amazon S3 bucket. In this walkthrough, the Lambda functions you will create are very simple; they do not interact with AWS services.

8. In the **Details** area, choose **Roles**.
9. Choose **Create New Role**.
10. In the **Role Name** box, type a name for the execution role (for example, **APIGatewayLambdaExecRole**), and then choose **Next Step**.
11. Next to **AWS Lambda**, choose **Select**.

Note

IAM will attach the following resource-policy document in **Trust Relationships**:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "",
            "Effect": "Allow",
            "Principal": {
                "Service": "lambda.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

This policy document enables Lambda to assume roles taken up by and, hence, to take actions on behalf of your AWS account.

Map API Gateway API Request Parameters for HTTP Endpoints

In this walkthrough, you will learn how to use API Gateway to connect a custom API to an HTTP endpoint through a REST proxy. This walkthrough builds on the instructions and concepts in the [Invoke Lambda Functions Synchronously \(p. 20\)](#), which shows you how to use API Gateway to create a custom API, connect it to a set of AWS Lambda functions, and then call the Lambda functions from your API. If you have not yet completed that walkthrough, we suggest you do it first.

In this walkthrough, you will use API Gateway to interact with some example data through a publicly accessible HTTP endpoint:

```
http://petstore-demo-endpoint.execute-api.com/petstore/pets
```

If you use a web browser to browse to this endpoint, the following JSON-formatted information will be displayed:

```
[  
  {  
    "id": 1,  
    "type": "dog",  
    "price": 249.99  
  },  
  {  
    "id": 2,  
    "type": "cat",  
    "price": 124.99  
  },  
  {  
    "id": 3,  
    "type": "fish",  
    "price": 0.99  
  }  
]
```

This endpoint supports two parameters: `type` and `page`. For example, if you browse to the following:

```
http://petstore-demo-endpoint.execute-api.com/petstore/pets?type=cat&page=2
```

The following JSON-formatted information about cats on page 2 is displayed:

```
[  
  {  
    "id": 4,  
    "type": "cat",  
    "price": 999.99  
  },  
  {  
    "id": 5,  
    "type": "cat",  
    "price": 249.99  
  },  
  {  
    "id": 6,  
    "type": "cat",  
    "price": 49.97  
  }  
]
```

This endpoint supports the use of an item ID. For example, if you browse to the following:

```
http://petstore-demo-endpoint.execute-api.com/petstore/pets/1
```

The following JSON-formatted information about the item with an ID of 1 is displayed:

```
{  
  "id": 1,  
  "type": "dog",  
  "price": 249.99  
}
```

You can send a payload to this endpoint. For example, if you use a web debugging proxy tool or the cURL command-line tool to send a POST method request to the following:

```
http://petstore-demo-endpoint.execute-api.com/petstore/pets
```

Making sure to include the header Content-type: application/json and the following request body:

```
{  
  "type": "dog",  
  "price": 249.99  
}
```

Then the following information will be returned in the response body:

```
{  
  "pet": {  
    "type": "dog",  
    "price": 249.99  
  },  
  "message": "success"  
}
```

This walkthrough shows you how to use API Gateway to interact with this HTTP endpoint as follows:

- Use a REST proxy with the format
`https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/pets` to interact with the HTTP endpoint of
`http://petstore-demo-endpoint.execute-api.com/petstore/pets`.
- Enable the REST proxy to specify two parameters, `petType` and `petsPage`, and send them to the HTTP endpoint's `type` and `page` parameters, respectively.
- Enable the REST proxy to specify an item ID to the HTTP endpoint.
- Enable the REST proxy to send a payload to the HTTP endpoint.

Topics

- [Prerequisites \(p. 33\)](#)
- [Step 1: Create Resources \(p. 33\)](#)
- [Step 2: Create GET and POST Methods \(p. 33\)](#)
- [Step 3: Specify Method Settings and Test the Methods \(p. 34\)](#)
- [Step 4: Deploy the API \(p. 36\)](#)
- [Step 5: Test the API \(p. 36\)](#)
- [Next Steps \(p. 38\)](#)

Prerequisites

Before you begin this walkthrough, you should do the following:

1. Complete the steps in [Get Ready to Use API Gateway \(p. 3\)](#), including assigning API Gateway access permission to the IAM user.
2. At a minimum, open the API Gateway console and create a new API named `MyDemoAPI`.

Step 1: Create Resources

In this step, you will create three resources that will enable the REST proxy to interact with the HTTP endpoint.

To create the first resource

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. If `MyDemoAPI` is displayed, choose **Resources**.
3. In the **Resources** pane, choose the resource root, represented by a single forward slash (/), and then choose **Create Resource**.
4. For **Resource Name**, type `petstorewalkthrough`.
This maps to `petstore` in the HTTP endpoint.
5. For **Resource Path**, accept the default of `/petstorewalkthrough`, and then choose **Create Resource**.
This maps to `/petstore` in the HTTP endpoint.

To create the second resource

1. In the **Resources** pane, choose `/petstorewalkthrough`, and then choose **Create Resource**.
2. For **Resource Name**, type `pets`.
This maps to `pets` in the HTTP endpoint.
3. For **Resource Path**, accept the default of `/petstorewalkthrough/pets`, and then choose **Create Resource**.
This maps to `/petstore/pets` in the HTTP endpoint.

To create the third resource

1. In the **Resources** pane, choose `/petstorewalkthrough/pets`, and then choose **Create Resource**.
2. For **Resource Name**, type `petId`. This maps to the item ID in the HTTP endpoint.
3. For **Resource Path**, overwrite `petid` with `{petId}`. Be sure to use curly braces ({}) around `petId` so that `/petstorewalkthrough/pets/{petId}` is displayed, and then choose **Create Resource**.
This maps to `/petstore/pets/my-item-id` in the HTTP endpoint.

Step 2: Create GET and POST Methods

In this step, you will create two GET methods and a POST method that will enable the REST proxy to interact with the HTTP endpoint.

To create the first GET method

1. In the **Resources** pane, choose `/petstorewalkthrough/pets`, and then choose **Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

To create the second GET method

1. In the **Resources** pane, choose `/petstorewalkthrough/pets/{petId}`, and then choose **Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

To create the POST method

1. In the **Resources** pane, choose `/petstorewalkthrough/pets` again, and then choose **Create Method**.
2. For the HTTP method, choose **POST**, and then save your choice.

Step 3: Specify Method Settings and Test the Methods

In this step, you will map parts of the REST proxy to their counterparts in the HTTP endpoint. You will then test the methods.

To specify settings for the first GET method and then test it

1. In the **Resources** pane, in `/petstorewalkthrough/pets`, choose **GET**.
2. In the **Setup** pane, for **HTTP method**, choose **GET**.
3. For **Endpoint URL**, type
`http://petstore-demo-endpoint.execute-api.com/petstore/pets`.
4. Choose **Save**.
5. In the **Method Execution** pane, choose **Method Request**, and then choose the arrow next to **URL Query String Parameters**.
6. Choose **Add query string**.
7. For **Name**, type `petType`.

This specifies `petType` in the REST proxy.

8. Choose **Create a new query string** (the check mark icon).
9. Choose **Add query string** again.
10. For **Name**, type `petsPage`.

This specifies `petsPage` in the REST proxy.

11. Choose **Create a new query string**.
12. Choose **Method Execution**, choose **Integration Request**, and then choose the arrow next to **URL Query String Parameters**.
13. Choose **Add query string**.
14. For **Name**, type `type`.
15. For **Mapped from**, type `method.request.querystring.petType`.

This maps `type` in the HTTP endpoint to `petType` in the REST proxy.

16. Choose **Create** (the check mark icon).
17. Choose **Add query string** again.

18. For **Name**, type `page`.
19. For **Mapped from**, type `method.request.querystring.petsPage`.
This maps `page` in the HTTP endpoint to `petsPage` in the REST proxy.
20. Choose **Create**.
21. Choose **Method Execution**, and in the **Client** box, choose **TEST**. In the **Query Strings** area, for `petType`, type `cat`. For `petsPage`, type `2`.
22. Choose **Test**. If successful, **Response Body** will display the following:

```
[  
  {  
    "id": 4,  
    "type": "cat",  
    "price": 999.99  
  },  
  {  
    "id": 5,  
    "type": "cat",  
    "price": 249.99  
  },  
  {  
    "id": 6,  
    "type": "cat",  
    "price": 49.97  
  }  
]
```

To specify settings for the second GET method and then test it

1. In the **Resources** list, in `/petstorewalkthrough/pets/{petId}`, choose **GET**.
2. In the **Setup** pane, for **HTTP method**, choose **GET**.
3. For **Endpoint URL**, type
`http://petstore-demo-endpoint.execute-api.com/petstore/pets/{id}`.
4. Choose **Save**.
5. In the **Method Execution** pane, choose **Integration Request**, and then choose the arrow next to **URL Path Parameters**.
6. Choose **Add path**.
7. For **Name**, type `id`.
8. For **Mapped from**, type `method.request.path.petId`.
9. Choose **Create**.
10. Choose **Method Execution**, and in the **Client** box, choose **TEST**. In the **Path** area, for `petId`, type `1`.
11. Choose **Test**. If successful, **Response Body** will display the following:

```
{  
  "id": 1,  
  "type": "dog",  
  "price": 249.99  
}
```

To specify settings for the POST method and then test it

1. In the **Resources** pane, in `/petstorewalkthrough/pets`, choose **POST**.
2. In the **Setup** pane, for **HTTP method**, choose **POST**.
3. For **Endpoint URL**, type
`http://petstore-demo-endpoint.execute-api.com/petstore/pets`.
4. Choose **Save**.
5. In the **Method Execution** pane, in the **Client** box, choose **TEST**. Expand **Request Body**, and then type the following:

```
{  
    "type": "dog",  
    "price": 249.99  
}
```

6. Choose **Test**. If successful, **Response Body** will display the following:

```
{  
    "pet": {  
        "type": "dog",  
        "price": 249.99  
    },  
    "message": "success"  
}
```

Step 4: Deploy the API

In this step, you will deploy the API so that you can begin calling it outside of the API Gateway console.

To deploy the API

1. In the **Resources** pane, choose **Deploy API**.
2. For **Deployment stage**, choose `test`.
3. For **Deployment description**, type `Calling HTTP endpoint walkthrough`.
4. Choose **Deploy**.

Step 5: Test the API

In this step, you will go outside of the API Gateway console and use your API's REST proxy to interact with the HTTP endpoint.

1. In the **Stage Editor** pane, next to **Invoke URL**, copy the URL to the clipboard. It should look something like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test
```

2. Paste this URL in the address box of a new browser tab.
3. Append `/petstorewalkthrough/pets` so that it looks like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/pets
```

Browse to the URL. The following information should be displayed:

```
[  
  {  
    "id": 1,  
    "type": "dog",  
    "price": 249.99  
  },  
  {  
    "id": 2,  
    "type": "cat",  
    "price": 124.99  
  },  
  {  
    "id": 3,  
    "type": "fish",  
    "price": 0.99  
  }  
]
```

4. After petstorewalkthrough/pets, type ?petType=cat&petsPage=2 so that it looks like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/pets?petType=cat&petsPage=2
```

5. Browse to the URL. The following information should be displayed:

```
[  
  {  
    "id": 4,  
    "type": "cat",  
    "price": 999.99  
  },  
  {  
    "id": 5,  
    "type": "cat",  
    "price": 249.99  
  },  
  {  
    "id": 6,  
    "type": "cat",  
    "price": 49.97  
  }  
]
```

6. After petstorewalkthrough/pets, replace ?petType=cat&petsPage=2 with /1 so that it looks like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/pets/1
```

7. Browse to the URL. The following information should be displayed:

```
{  
    "id": 1,  
    "type": "dog",  
    "price": 249.99  
}
```

8. Using a web debugging proxy tool or the cURL command-line tool, send a POST method request to the URL from the previous procedure. Be sure to append /petstorewalkthrough/pets so that it looks like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/pets
```

Also, be sure to append the following header:

```
Content-Type: application/json
```

And be sure to add the following code to the request body:

```
{  
    "type": "dog",  
    "price": 249.99  
}
```

For example, if you use the cURL command-line tool, run a command similar to the following:

```
curl -H "Content-Type: application/json" -X POST -d "{\"type\": \"dog\", \"price\": 249.99}" https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/pets
```

The following information should be returned in the response body:

```
{  
    "pet": {  
        "type": "dog",  
        "price": 249.99  
    },  
    "message": "success"  
}
```

You have reached the end of this walkthrough.

Next Steps

You may want to begin the next walkthrough, which shows you how to use models and mappings in API Gateway to transform the output of an API call from one data format to another. See [Use Modes and Mapping Templates \(p. 39\)](#).

Use Models and Mapping Templates for Payload Mappings

In this walkthrough, you will learn how to use models and mapping templates in API Gateway to transform the output of an API call from one data schema to another. This walkthrough builds on the instructions and concepts in the [Invoke Lambda Functions Synchronously \(p. 20\)](#) and the [Map API Request Parameter for HTTP Endpoints \(p. 30\)](#). If you have not yet completed those walkthroughs, we suggest you do them first.

In this walkthrough, you will use API Gateway to get example data from a publicly-accessible HTTP endpoint and from an AWS Lambda function you will create. Both the HTTP endpoint and the Lambda function return the same example data:

```
[  
  {  
    "id": 1,  
    "type": "dog",  
    "price": 249.99  
  },  
  {  
    "id": 2,  
    "type": "cat",  
    "price": 124.99  
  },  
  {  
    "id": 3,  
    "type": "fish",  
    "price": 0.99  
  }  
]
```

You will use models and mapping templates to transform this data to one or more output formats. In API Gateway, a model defines the format, also known as the schema or shape, of some data. In API Gateway, a mapping template is used to transform some data from one format to another. For more information, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).

The first model and mapping template is used to rename `id` to `number`, `type` to `class`, and `price` to `salesPrice`, as follows:

```
[  
  {  
    "number": 1,  
    "class": "dog",  
    "salesPrice": 249.99  
  },  
  {  
    "number": 2,  
    "class": "cat",  
    "salesPrice": 124.99  
  },  
  {  
    "number": 3,  
    "class": "fish",  
    "salesPrice": 0.99  
  }
```

```
}
```

The second model and mapping template is used to combine `id` and `type` into `description`, and to rename `price` to `askingPrice`, as follows:

```
[  
 {  
     "description": "Item 1 is a dog.",  
     "askingPrice": 249.99  
 },  
 {  
     "description": "Item 2 is a cat.",  
     "askingPrice": 124.99  
 },  
 {  
     "description": "Item 3 is a fish.",  
     "askingPrice": 0.99  
 }  
 ]
```

The third model and mapping template is used to combine `id`, `type`, and `price` into a set of `listings`, as follows:

```
{  
     "listings": [  
         "Item 1 is a dog. The asking price is 249.99.",  
         "Item 2 is a cat. The asking price is 124.99.",  
         "Item 3 is a fish. The asking price is 0.99."  
     ]  
 }
```

Topics

- [Prerequisites \(p. 40\)](#)
- [Step 1: Create Models \(p. 41\)](#)
- [Step 2: Create Resources \(p. 43\)](#)
- [Step 3: Create GET Methods \(p. 44\)](#)
- [Step 4: Create a Lambda Function \(p. 44\)](#)
- [Step 5: Specify Method Settings and Test the Methods \(p. 45\)](#)
- [Step 6: Deploy the API \(p. 49\)](#)
- [Step 7: Test the API \(p. 50\)](#)
- [Step 8: Clean Up \(p. 51\)](#)
- [Next Steps \(p. 52\)](#)

Prerequisites

Before you begin this walkthrough, you should have already done the following:

1. Complete the steps in [Get Ready to Use API Gateway \(p. 3\)](#), including assigning API Gateway access permission to an IAM user.

2. Open the API Gateway console and create a new API named `MyDemoAPI`. For more information, see [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).
3. Create two resources: `petstorewalkthrough` and `pets`. For more information, see [Create Resources \(p. 33\)](#) in the [Map API Request Parameter for HTTP Endpoints \(p. 30\)](#).
4. To use the Lambda portions of this walkthrough, make sure the IAM user has full access to work with Lambda. You can use IAM to attach the `AWSLambdaFullAccess` AWS managed policy to the IAM user.
5. Make sure the IAM user has access to create policies and roles in IAM. If you have not done so already, create a Lambda execution role named `APIGatewayLambdaExecRole` in IAM. For more information, see [Create Lambda Functions \(p. 21\)](#) in the [Invoke Lambda Functions Synchronously \(p. 20\)](#).

Step 1: Create Models

In this step, you will create four models. The first three models represent the data output formats for use with the HTTP endpoint and the Lambda function. The last model represents the data input schema for use with the Lambda function.

To create the first output model

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. If `MyDemoAPI` is displayed, choose **Models**.
3. Choose **Create**.
4. For **Model name**, type `PetsModelNoFlatten`.
5. For **Content type**, type `application/json`.
6. For **Model description**, type `Changes id to number, type to class, and price to salesPrice`.
7. For **Model schema**, type the following JSON Schema-compatible definition:

```
{  
    "$schema": "http://json-schema.org/draft-04/schema#",  
    "title": "PetsModelNoFlatten",  
    "type": "array",  
    "items": {  
        "type": "object",  
        "properties": {  
            "number": { "type": "integer" },  
            "class": { "type": "string" },  
            "salesPrice": { "type": "number" }  
        }  
    }  
}
```

8. Choose **Create model**.

To create the second output model

1. Choose **Create**.
2. For **Model name**, type `PetsModelFlattenSome`.
3. For **Content type**, type `application/json`.
4. For **Model description**, type `Combines id and type into description, and changes price to askingPrice`.

5. For **Model schema**, type the following:

```
{  
  "$schema": "http://json-schema.org/draft-04/schema#",  
  "title": "PetsModelFlattenSome",  
  "type": "array",  
  "items": {  
    "type": "object",  
    "properties": {  
      "description": { "type": "string" },  
      "askingPrice": { "type": "number" }  
    }  
  }  
}
```

6. Choose **Create model**.

To create the third output model

1. Choose **Create**.
2. For **Model name**, type `PetsModelFlattenAll`.
3. For **Content type**, type `application/json`.
4. For **Model description**, type `Combines id, type, and price into a set of listings.`
5. For **Model schema**, type the following:

```
{  
  "$schema": "http://json-schema.org/draft-04/schema#",  
  "title": "PetsModelFlattenAll",  
  "type": "object",  
  "properties": {  
    "listings": {  
      "type": "array",  
      "items": {  
        "type": "string"  
      }  
    }  
  }  
}
```

6. Choose **Create model**.

To create the input model

1. Choose **Create**.
2. For **Model name**, type `PetsLambdaModel`.
3. For **Content type**, type `application/json`.
4. For **Model description**, type `GetPetsInfo model`.
5. For **Model schema**, type the following:

```
{  
  "$schema": "http://json-schema.org/draft-04/schema#",  
  "title": "PetsLambdaModel",  
  "type": "object",  
  "properties": {  
    "method": { "type": "string" },  
    "path": { "type": "string" },  
    "httpMethod": { "type": "string" },  
    "requestParameters": {  
      "type": "object",  
      "properties": {  
        "method": { "type": "string" },  
        "path": { "type": "string" }  
      }  
    },  
    "responses": {  
      "type": "object",  
      "properties": {  
        "method": { "type": "string" },  
        "path": { "type": "string" },  
        "responseCode": { "type": "string" },  
        "responseDescription": { "type": "string" },  
        "schema": { "type": "string" }  
      }  
    }  
  }  
}
```

```
"type": "array",
"items": {
    "type": "object",
    "properties": {
        "id": { "type": "integer" },
        "type": { "type": "string" },
        "price": { "type": "number" }
    }
}
```

6. Choose **Create model**.

Step 2: Create Resources

In this step, you will create four resources. The first three resources will enable you to get the example data from the HTTP endpoint in the three output formats. The last resource will enable you to get the example data from the Lambda function in the output schema that combines `id` and `type` into `description` and renames `price` to `askingPrice`.

To create the first resource

1. In the links list, choose **Resources**.
2. In the **Resources** pane, choose `/petstorewalkthrough`, and then choose **Create Resource**.
3. For **Resource Name**, type `NoFlatten`.
4. For **Resource Path**, accept the default of `/petstorewalkthrough/noflatten`, and then choose **Create Resource**.

To create the second resource

1. In the **Resources** pane, choose `/petstorewalkthrough` again, and then choose **Create Resource**.
2. For **Resource Name**, type `FlattenSome`.
3. For **Resource Path**, accept the default of `/petstorewalkthrough/flattensome`, and then choose **Create Resource**.

To create the third resource

1. In the **Resources** pane, choose `/petstorewalkthrough` again, and then choose **Create Resource**.
2. For **Resource Name**, type `FlattenAll`.
3. For **Resource Path**, accept the default of `/petstorewalkthrough/flattenall`, and then choose **Create Resource**.

To create the fourth resource

1. In the **Resources** pane, choose `/petstorewalkthrough` again, and then choose **Create Resource**.
2. For **Resource Name**, type `LambdaFlattenSome`.
3. For **Resource Path**, accept the default of `/petstorewalkthrough/lambdaflattensome`, and then choose **Create Resource**.

Step 3: Create GET Methods

In this step, you will create a GET method for each of the resources you created in the previous step.

To create the first GET method

1. In the **Resources** list, choose **/petstorewalkthrough/flatternall**, and then choose **Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

To create the second GET method

1. In the **Resources** list, choose **/petstorewalkthrough/lambdaflattensome**, and then choose **Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

To create the third GET method

1. In the **Resources** list, choose **/petstorewalkthrough/flattensome**, and then choose **Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

To create the fourth GET method

1. In the **Resources** list, choose **/petstorewalkthrough/noflatten**, and then choose **Actions, Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

Step 4: Create a Lambda Function

In this step, you will create a Lambda function that returns the sample data.

To create the Lambda function

1. Open the AWS Lambda console at <https://console.aws.amazon.com/lambda/>.
2. Do one of the following:
 - If a welcome page appears, choose **Get Started Now**.
 - If the **Lambda: Function list** page appears, choose **Create a Lambda function**.
3. For **Name**, type **GetPetsInfo**.
4. For **Description**, type **Gets information about pets**.
5. For **Code template**, choose **None**.
6. Type the following code:

```
console.log('Loading event');

exports.handler = function(event, context) {
  context.done(null,
    [ {"id": 1, "type": "dog", "price": 249.99},
      {"id": 2, "type": "cat", "price": 124.99},
```

```
{ "id": 3, "type": "fish", "price": 0.99}]); // SUCCESS with message  
};
```

Tip

In the preceding code, written in Node.js, `console.log` writes information to an Amazon CloudWatch log. `event` contains the event's data. `context` contains callback context. Lambda uses `context.done` to perform follow-up actions. For more information about how to write Lambda function code, see the "Programming Model" section in [AWS Lambda: How it Works](#) and the sample walkthroughs in the [AWS Lambda Developer Guide](#).

7. For **Handler name**, leave the default of `index.handler`.
8. For **Role**, choose the Lambda execution role, **APIGatewayLambdaExecRole**, you created in the [Invoke Lambda Functions Synchronously \(p. 20\)](#).
9. Choose **Create Lambda function**.
10. In the list of functions, choose **GetPetsInfo** to show the function's details.
11. Make a note of the AWS region where you created this function. You will need it later.
12. In the pop-up list, choose **Edit or test function**.
13. For **Sample event**, replace any code that appears with the following:

```
{  
}  
}
```

Tip

The empty curly braces mean there are no input values for this Lambda function. This function simply returns the JSON object containing the pets information, so those key/value pairs are not required here.

14. Choose **Invoke**. **Execution result** shows
`[{"id":1,"type":"dog","price":249.99}, {"id":2,"type":"cat","price":124.99}, {"id":3,"type":"fish","price":0.99}]`, which is also written to the CloudWatch logs.
15. Choose **Go to function list**.

Step 5: Specify Method Settings and Test the Methods

In this step, you will specify the URL and data output schema for the three GET methods associated with the HTTP endpoint, testing each method as you proceed. You will also specify the Lambda function name, data input schema, and data output schema for the GET method associated with the Lambda function. You will then test this method.

To specify settings for the first GET method and then test it

1. In the **Resources** pane, in `/petstorewalkthrough/flattenall`, choose **GET**.
2. For **HTTP Method**, choose **GET**.
3. In the **Setup** pane, for **Endpoint URL**, type
`http://petstore-demo-endpoint.execute-api.com/petstore/pets`.
4. Choose **Save**.
5. In the **Method Execution** pane, choose **Method Response**, and then choose the arrow next to **200**.
6. In the **Response Models for 200** area, for **application/json**, choose **Edit**. For **Models**, choose **PetsModelFlattenAll**, and then choose **Save**.

7. Choose **Method Execution**, choose **Integration Response**, and then choose the arrow next to **200**.
8. In the **Template Mappings** area, for **Content type**, choose **application/json**, and then choose **Edit**. Clear **Output passthrough**. For **Generate template from model**, choose **PetsModelFlattenAll**. This displays the **PetsModelFlattenAll** model as a starting point.
9. Modify the code as follows:

```
#set($inputRoot = $input.path('$'))
{
    "listings" : [
#foreach($elem in $inputRoot)
        "Item number $elem.id is a $elem.type. The asking price is
$elem.price."#if($foreach.hasNext),#end

#end
    ]
}
```

10. Choose **Update**.
11. Choose **Method Execution**, and in the **Client** box, choose **TEST**, and then choose **Test**. If successful, **Response Body** will display the following:

```
{
    "listings" : [
        "Item number 1 is a dog. The asking price is 249.99.",
        "Item number 2 is a cat. The asking price is 124.99.",
        "Item number 3 is a fish. The asking price is 0.99."
    ]
}
```

To specify settings for the second GET method and then test it

1. In the **Resources** pane, in **/petstorewalkthrough/lambdaflattensome**, choose **GET**.
2. In the **Setup** pane, for **Execution Type**, choose **Lambda Function**.
3. For **Lambda Region**, choose the region identifier that corresponds to the region in which you created the **GetPetsInfo** Lambda function. For example, if you created this Lambda in the US East (N. Virginia) region, you would choose **us-east-1**. For a list of region names and identifiers, see [AWS Lambda](#) in the *Amazon Web Services General Reference*.
4. For **Lambda Function**, type **GetPetsInfo**, and then choose **Save**.
5. When you are prompted to give API Gateway permission to invoke your Lambda function, choose **Ok**.
6. In the **Method Execution** pane, choose **Integration Request**.
7. Next to **Templates**, choose **Add**.
8. For **Content-Type**, type **application/json**.
9. Leave **Input passthrough** cleared. For **Generate template from model**, choose **PetsLambdaModel**. This displays the **PetsLambdaModel** model as a starting point.
10. In the **Input mapping** area, modify the code as follows, and then choose **Update**:

```
#set($inputRoot = $input.path('$'))
[
#foreach($elem in $inputRoot)
{
```

```

    "id" : $elem.id,
    "type" : "$elem.type",
    "price" : $elem.price
}#if($foreach.hasNext),#end

#endif
]

```

11. Choose **Method Execution**, choose **Method Response**, and then choose the arrow next to **200**.
12. In the **Response Models** for **200** area, for **application/json**, choose **Edit**. For **Models**, choose **PetsModelFlattenSome**, and then choose **Save**.
13. Choose **Method Execution**, choose **Integration Response**, and then choose the arrow next to **200**.
14. In the **Template Mappings** area, for **Content type**, choose **application/json**, and then choose **Edit**. Clear **Output passthrough**. For **Generate template from model**, choose **PetsModelFlattenSome**. This displays the **PetsModelFlattenSome** model as a starting point.
15. Modify the code as follows, and then choose **Update**:

```

#set($inputRoot = $input.path('$.'))
[
#foreach($elem in $inputRoot)
{
    "description" : "Item $elem.id is a $elem.type.",
    "askingPrice" : $elem.price
}#if($foreach.hasNext),#end

#endif
]

```

16. Choose **Method Execution**, and in the **Client** box, choose **TEST**, and then choose **Test**. If successful, **Response Body** will display the following:

```

[
{
    "description" : "Item 1 is a dog.",
    "askingPrice" : 249.99
},
{
    "description" : "Item 2 is a cat.",
    "askingPrice" : 124.99
},
{
    "description" : "Item 3 is a fish.",
    "askingPrice" : 0.99
}
]

```

To specify settings for the third GET method and then test it

1. In the **Resources** pane, in **/petstorewalkthrough/flattensome**, choose **GET**.
2. In the **Setup** pane, for **HTTP method**, choose **GET**.
3. For **Endpoint URL**, type
http://petstore-demo-endpoint.execute-api.com/petstore/pets.

4. Choose **Save**.
5. In the **Method Execution** pane, choose **Method Response**, and then choose the arrow next to **200**.
6. In the **Response Models for 200** area, for **application/json**, choose **Edit**. For **Models**, choose **PetsModelFlattenSome**, and then choose **Save**.
7. Choose **Method Execution**, choose **Integration Response**, and then choose the arrow next to **200**.
8. In the **Template Mappings** area, for **Content type**, choose **application/json**, and then choose **Edit**. Clear **Output passthrough**. For **Generate template from model**, choose **PetsModelFlattenSome**. This displays the **PetsModelFlattenSome** model as a starting point.
9. Modify the code as follows:

```
#set($inputRoot = $input.path('$'))
[
#foreach($elem in $inputRoot)
{
    "description": "Item $elem.id is a $elem.type.",
    "askingPrice": $elem.price
}#if($foreach.hasNext),#end

#end
]
```

10. Choose **Update**.
11. Choose **Method Execution**, and in the **Client** box, choose **TEST**, and then choose **Test**. If successful, **Response Body** will display the following:

```
[
{
    "description": "Item 1 is a dog.",
    "askingPrice": 249.99
},
{
    "description": "Item 2 is a cat.",
    "askingPrice": 124.99
},
{
    "description": "Item 3 is a fish.",
    "askingPrice": 0.99
}
]
```

To specify settings for the fourth GET method and then test it

1. Return to the API Gateway console.
2. If **MyDemoAPI** is displayed, choose **Resources**.
3. In the **Resources** pane, in **/petstorewalkthrough/noflatten**, choose **GET**.
4. In the **Setup** pane, for **HTTP method**, choose **GET**.
5. For **Endpoint URL**, type
http://petstore-demo-endpoint.execute-api.com/petstore/pets.
6. Choose **Save**.
7. In the **Method Execution** pane, choose **Method Response**, and then expand **200**.

8. In the **Response Models for 200** area, for **application/json**, choose **Edit**. For **Models**, choose **PetsModelNoFlatten**, and then choose **Save**.
9. Choose **Method Execution**, choose **Integration Response**, and then choose the arrow next to **200**.
10. In the **Template Mappings** area, for **Content type**, choose **application/json**, and then choose **Edit**. Clear **Output passthrough**. For **Generate template from model**, choose **PetsModelNoFlatten**. This displays the **PetsModelNoFlatten** model as a starting point.
11. Modify the code as follows:

```
#set($inputRoot = $input.path('$'))  
[  
#foreach($elem in $inputRoot)  
{  
    "number": $elem.id,  
    "class": "$elem.type",  
    "salesPrice": $elem.price  
}#if($foreach.hasNext),#end  
  
#end  
]
```

12. Choose **Update**.
13. Choose **Method Execution**, and in the **Client** box, choose **TEST**, and then choose **Test**. If successful, **Response Body** will display the following:

```
[  
{  
    "number": 1,  
    "class": "dog",  
    "salesPrice": 249.99  
,  
{  
    "number": 2,  
    "class": "cat",  
    "salesPrice": 124.99  
,  
{  
    "number": 3,  
    "class": "fish",  
    "salesPrice": 0.99  
}  
]
```

Step 6: Deploy the API

In this step, you will deploy the API so that you can begin calling it outside of the API Gateway console.

To deploy the API

1. In the **Resources** pane, choose **Deploy API**.
2. For **Deployment stage**, choose **test**.
3. For **Deployment description**, type **using models and mapping templates walkthrough**.
4. Choose **Deploy**.

Step 7: Test the API

In this step, you will go outside of the API Gateway console to interact with both the HTTP endpoint and the Lambda function.

1. In the **Stage Editor** pane, next to **Invoke URL**, copy the URL to the clipboard. It should look something like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test
```

2. Paste this URL in the address box of a new browser tab.
3. Append /petstorewalkthrough/noflatten so that it looks like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/petstorewalkthrough/noflatten
```

Browse to the URL. The following information should be displayed:

```
[  
 {  
     "number": 1,  
     "class": "dog",  
     "salesPrice": 249.99  
 },  
 {  
     "number": 2,  
     "class": "cat",  
     "salesPrice": 124.99  
 },  
 {  
     "number": 3,  
     "class": "fish",  
     "salesPrice": 0.99  
 }]
```

4. After petstorewalkthrough/, replace noflatten with flattensome.
5. Browse to the URL. The following information should be displayed:

```
[  
 {  
     "description": "Item 1 is a dog.",  
     "askingPrice": 249.99  
 },  
 {  
     "description": "Item 2 is a cat.",  
     "askingPrice": 124.99  
 },  
 {  
     "description": "Item 3 is a fish.",  
     "askingPrice": 0.99  
 }]
```

6. After `petstorewalkthrough/`, replace `flattensome` with `flattenall`.
7. Browse to the URL. The following information should be displayed:

```
{  
  "listings" : [  
    "Item number 1 is a dog. The asking price is 249.99.",  
    "Item number 2 is a cat. The asking price is 124.99.",  
    "Item number 3 is a fish. The asking price is 0.99."  
  ]  
}
```

8. After `petstorewalkthrough/`, replace `flattenall` with `lambdaflattensome`.
9. Browse to the URL. The following information should be displayed:

```
[  
  {  
    "description" : "Item 1 is a dog.",  
    "askingPrice" : 249.99  
  },  
  {  
    "description" : "Item 2 is a cat.",  
    "askingPrice" : 124.99  
  },  
  {  
    "description" : "Item 3 is a fish.",  
    "askingPrice" : 0.99  
  }  
]
```

Step 8: Clean Up

If you no longer need the Lambda function you created for this walkthrough, you can delete it now. You can also delete the accompanying IAM resources.

Caution

If you delete a Lambda function your APIs rely on, those APIs will no longer work. Deleting a Lambda function cannot be undone. If you want to use the Lambda function again, you must re-create the function.

If you delete an IAM resource a Lambda function relies on, the Lambda function and any APIs that rely on it will no longer work. Deleting an IAM resource cannot be undone. If you want to use the IAM resource again, you must re-create the resource. If you plan to continue experimenting with the resources you created for this and the other walkthroughs, do not delete the Lambda invocation role or the Lambda execution role.

API Gateway does not currently support the deactivation or deletion of APIs that no longer work.

To delete the Lambda function

1. Sign in to the AWS Management Console and open the AWS Lambda console at <https://console.aws.amazon.com/lambda/>.
2. On the **Lambda: Function list** page, in the list of functions, choose the button next to **GetPetsInfo**, and then choose **Actions, Delete**. When prompted, choose **Delete** again.

To delete the associated IAM resources

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the **Details** area, choose **Roles**.
3. Select **APIGatewayLambdaExecRole**, and then choose **Role Actions**, **Delete Role**. When prompted, choose **Yes, Delete**.
4. In the **Details** area, choose **Policies**.
5. Select **APIGatewayLambdaExecPolicy**, and then choose **Policy Actions**, **Delete**. When prompted, choose **Delete**.

You have now reached the end of this walkthrough.

Next Steps

You may want to begin the next walkthrough, which shows you how to integrate API Gateway with an AWS service. See [Get Started with AWS Proxies \(p. 52\)](#).

Get Started with Creating API Gateway API as an AWS Service Proxy

In this walkthrough, you will learn how to use API Gateway to connect a custom API to an AWS service through what we call an AWS service proxy. This enables you to call an AWS service directly instead of through an AWS Lambda function. An AWS service proxy can call only one action in an AWS service, and that action typically does not change. If you want more flexibility, you should call a Lambda function instead.

This walkthrough builds on the instructions and concepts in the [Invoke Lambda Functions Synchronously \(p. 20\)](#), which shows you how to use API Gateway to create a custom API, connect it to a set of AWS Lambda functions, and then call the Lambda functions from your API. If you have not yet completed that walkthrough, we suggest that you do it first.

Topics

- [Prerequisites \(p. 52\)](#)
- [Step 1: Create the Resource \(p. 53\)](#)
- [Step 2: Create the GET Method \(p. 53\)](#)
- [Step 3: Create the AWS Service Proxy Execution Role \(p. 53\)](#)
- [Step 4: Specify Method Settings and Test the Method \(p. 55\)](#)
- [Step 5: Deploy the API \(p. 56\)](#)
- [Step 6: Test the API \(p. 56\)](#)
- [Step 7: Clean Up \(p. 57\)](#)

Prerequisites

Before you begin this walkthrough, you should have already done the following:

1. Complete the steps in [Get Ready to Use API Gateway \(p. 3\)](#).
2. Make sure the IAM user has access to create policies and roles in IAM. You will need to create an IAM policy and role in this walkthrough.

3. At a minimum, open the API Gateway console and create a new API named `MyDemoAPI`. For more information, see [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).
4. Deploy the API at least once to a stage named `test`. For more information, see [Deploy the API \(p. 26\)](#) in the [Invoke Lambda Functions Synchronously \(p. 20\)](#).
5. Complete the rest of the steps in the [Invoke Lambda Functions Synchronously \(p. 20\)](#).
6. Create at least one topic in Amazon Simple Notification Service (Amazon SNS). You will use the deployed API to get a list of topics in Amazon SNS that are associated with your AWS account. To learn how to create a topic in Amazon SNS, see [Create a Topic](#). (You do not need to copy the topic ARN mentioned in step 5.)

Step 1: Create the Resource

In this step, you will create a resource that will enable the AWS service proxy to interact with the AWS service.

To create the resource

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. If **MyDemoAPI** is displayed, choose **Resources**.
3. In the **Resources** pane, choose the resource root, represented by a single forward slash (/), and then choose **Create Resource**.
4. For **Resource Name**, type `MyDemoAWSProxy`, and then choose **Create Resource**.

Step 2: Create the GET Method

In this step, you will create a GET method that will enable the AWS service proxy to interact with the AWS service.

To create the GET method

1. In the **Resources** pane, choose `/mydemoawsproxy`, and then choose **Create Method**.
2. For the HTTP method, choose **GET**, and then save your choice.

Step 3: Create the AWS Service Proxy Execution Role

In this step, you will create an IAM role that your AWS service proxy will use to interact with the AWS service. We call this IAM role an AWS service proxy execution role. Without this role, API Gateway cannot interact with the AWS service. In later steps, you will specify this role in the settings for the GET method you just created.

To create the AWS service proxy execution role and its policy

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Policies**.
3. Do one of the following:
 - If the **Welcome to Managed Policies** page appears, choose **Get Started**, and then choose **Create Policy**.

- If a list of policies appears, choose **Create Policy**.
4. Next to **Create Your Own Policy**, choose **Select**.
5. For **Policy Name**, type a name for the policy (for example, `APIGatewayAWSProxyExecPolicy`).
6. For **Description**, type `Enables API Gateway to call AWS services`.
7. For **Policy Document**, type the following, and then choose **Create Policy**.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Resource": ["*"],  
            "Action": ["sns>ListTopics"]  
        }  
    ]  
}
```

Note

This policy document allows the caller to get a list of the Amazon SNS topics for the AWS account.

8. Choose **Roles**.
9. Choose **Create New Role**.
10. For **Role Name**, type a name for the execution role (for example, `APIGatewayAWSProxyExecRole`), and then choose **Next Step**.
11. Next to **Amazon EC2**, choose **Select**.

Note

You choose **Select** here because you need to choose a standard AWS service role statement before you can continue. There is currently no option to choose a standard API Gateway service role statement. Later in this step, you will modify the standard Amazon EC2 service role statement for use with API Gateway.

12. In the list of policies, select **APIGatewayAWSProxyExecPolicy**, and then choose **Next Step**.
13. For **Role ARN**, make a note of the Amazon Resource Name (ARN) for the execution role. You will need it later. The ARN should look similar to:
`arn:aws:iam::123456789012:role/APIGatewayAWSProxyExecRole`, where 123456789012 is your AWS account ID.
14. Choose **Create Role**.

The invocation role IAM just created enables Amazon EC2 to get a list of the Amazon SNS topics for the AWS account. You will change this role to enable API Gateway to get a list of the Amazon SNS topics for the AWS account instead.

15. In the list of roles, select **APIGatewayAWSProxyExecRole**.
16. In the **Trust Relationships** area, choose **Edit Trust Relationship**.
17. For **Policy Document**, replace `ec2.amazonaws.com` with `apigateway.amazonaws.com` so that the access control policy document now looks as follows:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "",
            "Effect": "Allow",
            "Principal": {
                "Service": "apigateway.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

This policy document enables API Gateway to take actions on behalf of your AWS account.

18. Choose **Update Trust Policy**.

Step 4: Specify Method Settings and Test the Method

In this step, you will specify the settings for the GET method so that it can interact with an AWS service through an AWS service proxy. You will then test the method.

To specify settings for the GET method and then test it

1. In the API Gateway console, in the **Resources** pane for the API named **MyDemoAPI**, in **/mydemoawsproxy**, choose **GET**.
2. In the **Setup** pane, for **Integration type**, choose **Show advanced**, and then choose **AWS Service Proxy**.
3. For **AWS Region**, choose the name of the AWS region where you want to get the Amazon SNS topics.
4. For **AWS Service**, choose **SNS**.
5. For **HTTP method**, choose **GET**.
6. For **Action**, type **ListTopics**.
7. For **Execution Role**, type the ARN for the execution role.
8. Leave **Path Override** blank.
9. Choose **Save**.
10. In the **Method Execution** pane, in the **Client** box, choose **TEST**, and then choose **Test**. If successful, **Response Body** will display a response similar to the following:

```
{
    "ListTopicsResponse": {
        "ListTopicsResult": {
            "NextToken": null,
            "Topics": [
                {
                    "TopicArn": "arn:aws:sns:us-east-1:80398EXAMPLE:MySNSTopic-1"
                },
                {
                    "TopicArn": "arn:aws:sns:us-east-1:80398EXAMPLE:MySNSTopic-2"
                }
            ]
        }
    }
}
```

```
        ...
    },
    "TopicArn": "arn:aws:sns:us-east-1:80398EXAMPLE:MySNSTopic-N"
}
],
{
    "ResponseMetadata": {
        "RequestId": "abc1de23-45fa-6789-b0c1-d2e345fa6b78"
    }
}
}
```

Step 5: Deploy the API

In this step, you will deploy the API so that you can begin calling it from outside of the API Gateway console.

To deploy the API

1. In the **Resources** pane, choose **Deploy API**.
2. For **Deployment stage**, choose **test**.
3. For **Deployment description**, type **Calling AWS service proxy walkthrough**.
4. Choose **Deploy**.

Step 6: Test the API

In this step, you will go outside of the API Gateway console and use your AWS service proxy to interact with the Amazon SNS service.

1. In the **Stage Editor** pane, next to **Invoke URL**, copy the URL to the clipboard. It should look like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test
```

2. Paste the URL into the address box of a new browser tab.
3. Append /mydemoawsproxy so that it looks like this:

```
https://my-api-id.execute-api.region-id.amazonaws.com/test/mydemoawsproxy
```

Browse to the URL. Information similar to the following should be displayed:

```
{"ListTopicsResponse": {"ListTopicsResult": {"NextToken": null, "Topics": [{"TopicArn": "arn:aws:sns:us-east-1:80398EXAMPLE:MySNSTopic-1"}, {"TopicArn": "arn:aws:sns:us-east-1:80398EXAMPLE:MySNSTopic-2"}, {"TopicArn": "arn:aws:sns:us-east-1:80398EXAMPLE:MySNSTopic-N"}]}, "ResponseMetadata": {"RequestId": "abc1de23-45fa-6789-b0c1-d2e345fa6b78"}}}
```

Step 7: Clean Up

You can delete the IAM resources the AWS service proxy needs to work.

Caution

If you delete an IAM resource an AWS service proxy relies on, that AWS service proxy and any APIs that rely on it will no longer work. Deleting an IAM resource cannot be undone. If you want to use the IAM resource again, you must re-create it.

To delete the associated IAM resources

1. Open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the **Details** area, click **Roles**.
3. Select **APIGatewayAWSProxyExecRole**, and then choose **Role Actions, Delete Role**. When prompted, choose **Yes, Delete**.
4. In the **Details** area, choose **Policies**.
5. Select **APIGatewayAWSProxyExecPolicy**, and then choose **Policy Actions, Delete**. When prompted, choose **Delete**.

You have reached the end of this walkthrough. For more detailed discussions about creating API as an AWS service proxy, see [Create an API as an Amazon S3 Proxy \(p. 132\)](#), [Create an API as a Lambda Proxy \(p. 150\)](#) or [Create an API as an Amazon Kinesis Proxy \(p. 165\)](#).

Creating an API in Amazon API Gateway

Topics

- [Create an API in API Gateway \(p. 58\)](#)
- [Set up API Gateway API Method and Integration \(p. 59\)](#)
- [Set Up Amazon API Gateway API Request and Response Payload Mappings \(p. 68\)](#)
- [Amazon API Gateway API Request and Response Parameter-Mapping Reference \(p. 93\)](#)
- [API Gateway API Request and Response Payload-Mapping Template Reference \(p. 96\)](#)
- [Enable CORS for a Resource through a Method in API Gateway \(p. 101\)](#)
- [Use an API Key in API Gateway \(p. 105\)](#)
- [Enable Client-Side SSL Authentication of an API with the API Gateway Console \(p. 106\)](#)
- [Enable Amazon API Gateway Custom Authorization \(p. 108\)](#)
- [Import and Export API Gateway API with Swagger Definition Files \(p. 116\)](#)
- [Create an API as an Amazon S3 Proxy \(p. 132\)](#)
- [Create an API Gateway API as an AWS Lambda Proxy \(p. 150\)](#)
- [Create an API Gateway API as an Amazon Kinesis Proxy \(p. 165\)](#)

Create an API in API Gateway

In Amazon API Gateway you can create an API using the API Gateway console, AWS CLI, the API Gateway control service REST API, and platform-specific or language-specific SDKs.

Topics

- [Create an API Using the API Gateway Console \(p. 59\)](#)
- [Create an API Using the API Gateway Control Service API \(p. 59\)](#)
- [Create an API Using the AWS SDK for API Gateway \(p. 59\)](#)
- [Create an API Using the AWS CLI \(p. 59\)](#)

Create an API Using the API Gateway Console

To create an API Gateway API using the API Gateway console, see [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).

You can learn how to create an API by following an example. For more information, see [Build and Test an API Gateway API from an Example \(p. 5\)](#).

Alternatively, you can create an API by using the API Gateway [Import API \(p. 117\)](#) feature to upload an external API definition, such as one expressed in the [Swagger 2.0 with the API Gateway Extensions to Swagger \(p. 122\)](#). The example provided in [Build and Test an API Gateway API from an Example \(p. 5\)](#) uses the Import API feature.

Create an API Using the API Gateway Control Service API

For more information about the API Gateway Control Service API, see [Amazon API Gateway REST API Reference](#).

Create an API Using the AWS SDK for API Gateway

For more information using a AWS SDK, see [AWS SDKs](#).

Create an API Using the AWS CLI

For an example of creating an API Gateway API Using AWS CLI, see [Create an API Gateway API for Lambda tutorial](#).

Set up API Gateway API Method and Integration Before Configuring Methods

- You must have the method available in API Gateway. Follow the instructions in [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).
- If you want the method to communicate with a Lambda function, you must have already created the Lambda invocation role and Lambda execution role in IAM and created the Lambda function with which your method will communicate in AWS Lambda. To create the roles and function, use the instructions in [Step 4: Create Lambda Functions \(p. 21\)](#) of the [Invoke Lambda Functions Synchronously \(p. 20\)](#).
- If you want the method to communicate with an HTTP proxy, you must have already created and have access to the HTTP endpoint URL with which your method will communicate.

After Setting Up Methods and Integration

The next step is to deploy the API to make it open for access. For instructions, see [Deploying an API \(p. 199\)](#).

Topics

- [Configure How API Gateway Integrates the Method with a Back End \(p. 60\)](#)
- [Configure How an API User Calls an API Method in Amazon API Gateway \(p. 62\)](#)

- [Configure How Data Is Mapped between a Method and its Integration in Amazon API Gateway \(p. 64\)](#)
- [Configure Mock Integration for a Method in API Gateway \(p. 65\)](#)

Configure How API Gateway Integrates the Method with a Back End

The settings of an API method defines the method and describes its behaviors. To create a method, you must specify a resource, including the root (""/"), on which the method is exposed, a method type (GET, POST, etc.), and how it will be integrated with the targeted backend. The method request and response specify the contract with the calling app, stipulating which parameters the API can receive and what the response looks like. The integration request and response specifies how API Gateway interacts with their backend. The following topics describe how to use the API Gateway console to specify a method settings.

1. In the **Resources** pane, choose the method.
2. In the **Method Execution** pane, choose **Integration Request**. For **Integration type**, choose one of the following:
 - Choose **Lambda Function** if your API will be communicating with a Lambda function.
 - Choose **HTTP Proxy** if your API will be communicating with an HTTP endpoint.
 - Choose **Show Advanced, AWS Service Proxy** if your API will be communicating directly with an AWS service.
 - Choose **Mock Integration** if your API is not yet final, but you want to generate API responses from API Gateway anyway to unblock dependent teams for testing. If you choose this option, skip the rest of the instructions in this topic and see [Configure Mock Integration for a Method \(p. 65\)](#).
3. If you chose **Lambda Function**, do the following:
 1. For **Lambda Region**, choose the region identifier that corresponds to the region where you created the Lambda function. For example, if you created the Lambda function in the US East (N. Virginia) region, you would choose `us-east-1`. For a list of region names and identifiers, see [AWS Lambda](#) in the *Amazon Web Services General Reference*.
 2. For **Lambda Function**, type the name of the Lambda function, and then choose the function's corresponding ARN.
 3. Choose **Save**.
4. If you chose **HTTP Proxy**, do the following:
 1. For **HTTP method**, choose the HTTP method type that most closely matches the method in the HTTP proxy.
 2. For **Endpoint URL**, type the URL of the HTTP proxy you want this method to use.
 3. Choose **Save**.
5. If you chose **Mock Integration**, do the following:
 - Choose **Save**.
6. If you chose **Show advanced, AWS Service Proxy**, do the following:
 1. For **AWS Region**, choose the AWS region you want this method to use to call the action.
 2. For **AWS Service**, choose the AWS service you want this method to call.

3. For **HTTP method**, choose the HTTP method type that corresponds to the action. For HTTP method type, see the API reference documentation for the AWS service you chose for **AWS Service**.
4. For **Action**, type the action you want to use. For a list of available actions, see the API reference documentation for the AWS service you chose for **AWS Service**.
5. For **Execution Role**, type the ARN of the IAM role the method will use to call the action.

To create the IAM role, you can adapt the instructions in "To create the Lambda invocation role and its policies" and "To create the Lambda execution role and its policy" in the [Create Lambda Functions \(p. 21\)](#) section of the [Invoke Lambda Functions Synchronously \(p. 20\)](#); and specify an access policy of the following format, with the desired number of action and resource statements:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "action-statement"
      ],
      "Resource": [
        "resource-statement"
      ]
    },
    ...
  ]
}
```

For the action and resource statement syntax, see the documentation for the AWS service you chose for **AWS Service**.

For the IAM role's trust relationship, specify the following, which enables API Gateway to take actions on behalf of your AWS account:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
        "Service": "apigateway.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

6. If the action you typed for **Action** provides a custom resource path you want this method to use, for **Path Override**, type this custom resource path. For the custom resource path, see the API reference documentation for the AWS service you chose for **AWS Service**.
7. Choose **Save**.

7. Do both of the following:

- Specify how the method will receive requests from, and send responses to, callers (which API Gateway refers to as the API's method request/response), and how the method will authorize requests by following the instructions in [Configure How a User Calls an API Method \(p. 62\)](#).
- Specify how the method will send requests to, and receive responses from, the Lambda function, HTTP proxy, or AWS service proxy (which API Gateway refers to as the API's integration request/response) by following the instructions in [Configure How Data Is Mapped between Method and Integration \(p. 64\)](#).

Configure How an API User Calls an API Method in Amazon API Gateway

To use the API Gateway console to specify an API's method request/response and the way in which the method will authorize requests, follow these instructions.

Note

These instructions assume you have already completed the steps in [Configure How a Method Is Integrated with a Back End \(p. 60\)](#).

1. With the method selected in the **Resources** pane, in the **Method Execution** pane, choose **Method Request**.
2. To assign custom access permissions to the method, in the **Authorization Settings** area, for **Authorization Type**, choose **Edit**, and then choose **AWS_IAM**. Only IAM roles with the correct IAM policy attached will be allowed to call this method. If you do not want to assign custom access permissions to the method, choose **NONE**.
 - To create the IAM role, specify an access policy with a format like the following:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "execute-api:Invoke"  
            ],  
            "Resource": [  
                "resource-statement"  
            ]  
        }  
    ]  
}
```

In this access policy, *resource-statement* is the value of the **ARN** field in the **Authorization Settings** section.

To create the IAM role, you can adapt the instructions in "To create the Lambda invocation role and its policy" and "To create the Lambda execution role and its policy" in the [Create Lambda Functions \(p. 21\)](#) section of the [Invoke Lambda Functions Synchronously \(p. 20\)](#).

To save your choice, choose **Update**. Otherwise, choose **Cancel**.

Note

You can also enable an API key. For instructions, see [Use an API Key with the API Gateway Console \(p. 105\)](#).

3. To add a query string parameter to the method, do the following:
 1. Choose the arrow next to **URL Query String Parameters**, and then choose **Add query string**.
 2. For **Name**, type the name of the query string parameter.
 3. Choose **Create a new query string**.

Note

To remove the query string parameter, choose **Cancel or Remove**.

To change the name of the query string parameter, you must remove it and create a new one.

4. To add a header parameter to the method, do the following:
 1. Choose the arrow next to **HTTP Request Headers**, and then choose **Add header**.
 2. For **Name**, type the name of the header parameter, and choose **Create**.

Tip

To remove the header parameter, choose **Cancel or Remove**.

To change the name of the header parameter, you must remove the old header parameter and create a new one in its place.

5. For non-GET method types, expand **Request Models**, and for **Content Type** and **Model name**, type the content type and choose the name of a model that will transform caller-supplied data into the expected format.

To create a model, see [Create a Model \(p. 74\)](#).
6. To send a set of custom response headers, a custom response data format, or both, back to callers based on the HTTP status code returned by the method, do the following:

1. If the **Method Execution** page is not displayed, choose **Method Execution**.

Tip

You will use **Method Response**, to specify all the possible response codes for your API and use **Integration Response** to tell API Gateway how backend errors are mapped to an HTTP status code.

2. Choose **Method Response**. By default, 200 response is included in the method responses. You can modify it, e.g., to have the method return 201 instead. In addition, you can add other responses, e.g., 409 for access denial and 500 for uninitialized stage variables used. Either choose the arrow icon next to **200** to specify settings for the 200 response, or choose **Add Response** to specify settings for any other HTTP response status code. If you choose **Add Response**, for **HTTP Status**, choose the response, choose **Create**, and choose the arrow next to the response.
3. For each custom header you want to include in the response, in the **Response Headers** area, choose **Add Header**, type the name of the header, and then choose **Save**. (Choose **Remove** to remove a header from this list.)

To specify a response model to transform the output's data from one format to another, in the **Response Models** area, choose **Add Response Model**. Type the content type (for **Content type**), choose the model's name (for **Models**), and then choose **Save**. Choose **Add Response Model** to specify an additional model, or choose **Create a model** to define a new model. (Choose **Remove** to remove a response model selection from this list.)

Configure How Data Is Mapped between a Method and its Integration in Amazon API Gateway

To use the API Gateway console to define the API's integration request/response, follow these instructions.

Note

These instructions assume you have already completed the steps in [Configure How a Method Is Integrated with a Back End \(p. 60\)](#).

1. With the method selected in the **Resources** pane, in the **Method Execution** pane, choose **Integration Request**.
2. For an HTTP proxy or an AWS service proxy, to associate a path parameter, a query string parameter, or a header parameter defined in the integration request with a corresponding path parameter, query string parameter, or header parameter in the method request of the HTTP proxy or AWS service proxy, do the following:
 1. Choose the arrow next to **URL Path Parameters**, **URL Query String Parameters**, or **HTTP Headers** respectively, and then choose **Add path**, **Add query string**, or **Add header**, respectively.
 2. For **Name**, type the name of the path parameter, query string parameter, or header parameter in the HTTP proxy or AWS service proxy.
 3. For **Mapped from**, type the mapping value for the path parameter, query string parameter, or header parameter. Use one of the following formats:
 - `method.request.path.parameter-name` for a path parameter named *parameter-name* as defined in the **Method Request** page.
 - `method.request.querystring.parameter-name` for a query string parameter named *parameter-name* as defined in the **Method Request** page.
 - `method.request.header.parameter-name` for a header parameter named *parameter-name* as defined in the **Method Request** page.
4. Alternatively, you can set a literal string value (enclosed by a pair of single quotes) to an integration header.
4. Choose **Create**. (To delete a path parameter, query string parameter, or header parameter, choose **Cancel** or **Remove** next to the parameter you want to delete.)
3. In the **Templates** area (for a Lambda function) or the **Mapping Templates** area (for an HTTP proxy or AWS service proxy), next to **Content-Type**, choose **Add**. In the **Content-Type** box, type the content type of the data that will be passed from the method to the Lambda function, HTTP proxy, or AWS service proxy. Choose **Update**.
4. Select **Input passthrough** to send the data from the method without modification to the Lambda function, HTTP proxy, or AWS service proxy.

If **Input passthrough** is cleared, then for **Input mapping**, specify the input mapping template you want the Lambda function, HTTP proxy, or AWS service proxy to use to receive data from the method. You can type the mapping template manually or choose **Generate template from model** or **Create a model**. For more information, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).
5. You can map an integration response from the backend to a method response of the API returned to the calling app. This includes returning to the client selected response headers from the available ones from the backend, transforming the data format of the backend response payload to an API-specified format. You can specify such mapping by configuring **Method Response** and **Integration Response** from the **Method Execution** page.

1. If the **Method Execution** page is not displayed, choose **Method Execution**. From the **Method Execution** page, choose **Integration Response**. Choose either the arrow next to **200** to specify settings for a 200 HTTP response code from the method, or choose **Add integration response** to specify settings for any other HTTP response status code from the method.
2. For **Lambda error regex** (for a Lambda function) or **HTTP status regex** (for an HTTP proxy or AWS service proxy), type a regular expression to specify which Lambda function error strings (for a Lambda function) or HTTP response status codes (for an HTTP proxy or AWS service proxy) map to this output mapping. For example, to map all 2xx HTTP response status codes from an HTTP proxy to this output mapping, type `2\\d{2}` for **HTTP status regex**.

Note

The error patterns are matched against the `errorMessage` property in the Lambda response, which is populated by `context.fail(errorMessage)` in Node.js or by `throw new MyException(errorMessage)` in Java.

Be aware of the fact that the `.\\` pattern will not match any newline (`\n`).

3. If enabled, for **Method response status**, choose the HTTP response status code you defined in the **Method Response** page.
4. For **Header Mappings**, for each header you defined for the HTTP response status code in the **Method Response** page, specify a mapping value by choosing **Edit**. For **Mapping value**, use the format `integration.response.header.header-name` where `header-name` is the name of a response header from the backend. For example, to return the backend response's `Date` header as an API method's response's `Timestamp` header, the **Response header** column will contain an `Timestamp` entry and the associated **Mapping value** should be set to `integration.response.header.Date`.
5. In the **Template Mappings** area, next to **Content type**, choose **Add**. In the **Content type** box, type the content type of the data that will be passed from the Lambda function, HTTP proxy, or AWS service proxy to the method. Choose **Update**.
6. Select **Output passthrough** if you want the method to receive, but not modify, the data from the Lambda function, HTTP proxy, or AWS service proxy.
7. If **Output passthrough** is cleared, for **Output mapping**, specify the output mapping template you want the Lambda function, HTTP proxy, or AWS service proxy to use to send data to the method. You can either type the mapping template manually or choose a model from **Generate template from model**.
8. Choose **Save**.

Configure Mock Integration for a Method in API Gateway

Amazon API Gateway supports mock integrations for API methods. This feature enables developers to generate API responses from API Gateway directly, without the need for an integration back end. This allows developers to unblock other dependent teams who need to work with an API before development is complete.

Use the API Gateway console to create a mock integration for a method.

Topics

- [Prerequisites \(p. 66\)](#)
- [Mock-Integrate a Method \(Console\) \(p. 66\)](#)
- [Example Request Templates \(p. 67\)](#)
- [Example Response Templates \(p. 68\)](#)

Prerequisites

- You must have the method available in API Gateway. Follow the instructions in [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).

Mock-Integrate a Method (Console)

1. Choose the method you want to mock-integrate. Follow the instructions in [View a Methods List \(p. 197\)](#), and then choose the method.
2. Do one of the following:
 - If the **Setup** pane appears, choose **Mock Integration**, and then choose **Save**.
 - If the **Method Integration** pane appears, choose **Integration Request**. For **Integration type**, choose **Mock Integration**, and then choose **Save**.
3. The **Method Execution** pane appears. Choose **Integration Request**.
4. By default, mock integrations return a 200 HTTP status code. To customize this default behavior, do the following:
 1. Expand **Mapping Templates**.
 2. For **Content-Type**, do one of the following:
 - If the desired content type is already visible (for example, `application/json`), then choose it.
 - If the desired content type is not already visible, then choose **Add mapping template**, type the desired content type (for example, `application/json`), and then choose **Create**.
 3. A section appears next to the content type, displaying the content type as a title (for example, `application/json`). Next to the content type, for **Input passthrough**, choose **Edit**.
 4. In the drop-down list next to **Save** and **Cancel**, choose **Mapping template**.
 5. In the **Template** box, type the contents of the template you want API Gateway to use to determine which HTTP status code to use in the integration response. For more information, see [Example Request Templates \(p. 67\)](#).
 6. Next to **Mapping template**, choose **Save**.
5. For each query string parameter or header parameter you want to add to the method, do the following:
 1. Choose **Method Execution**, and then choose **Method Request**.
 2. Choose the arrow next to **URL Query String Parameters** or **HTTP Request Headers**, and then choose **Add query string** or **Add header**, respectively.
 3. For **Name**, type the name of the query string parameter or header parameter, and then choose **Create a new query string** or **Create**, respectively.
6. Choose **Method Execution**, and then choose **Method Response**.
7. Do one of the following:

Note

To remove a query string parameter or header parameter, choose **Cancel** or **Remove**. To change the name of a query string parameter or header parameter, you must remove it and create a new one in its place.

-
6. Choose **Method Execution**, and then choose **Method Response**.
 7. Do one of the following:

- If all of the **HTTP Status** entries you want to use are already visible (for example, **200**), then skip ahead to step 8.
 - If any of the **HTTP Status** entries you want to use are not already visible, then for each missing **HTTP Status** entry, choose **Add Response**, choose the HTTP status code that you want to use, and then choose **Create**.
8. Choose **Method Execution**, and then choose **Integration Response**.
9. Do one of the following:
- If all of the **Method response status** entries you want to use are already visible (for example, **200**), then skip ahead to step 10.
 - If any of the **Method response status** entries you want to use are not already visible, then for each missing **Method response status** entry, choose **Add integration response**, for **Method response status** choose the **HTTP Status** entry you created earlier, and then choose **Save**.
10. For each **Method response status** entry you want to use, do the following:
1. Expand the row that corresponds to the **Method response status** entry you want to use.
 2. For **HTTP status regex**, type the matching **HTTP Status** entry (for example, type **400** for a 400 **HTTP Status** entry or **500** for a 500 **HTTP Status** entry). Or specify a range of matching HTTP status codes (for example, **5/d{2}** matches all 5XX HTTP status codes).
 3. Expand **Mapping Templates**.
 4. For **Content-Type**, do one of the following:
 - If the desired content type is already visible (for example, **application/json**), then choose it.
 - If the desired content type is not already visible, then choose **Add mapping template**, type the desired content type (for example, **application/json**), and then choose **Create**.
 5. Next to the content type, for **Output passthrough**, choose **Edit**.
 6. In the drop-down list next to **Save** and **Cancel**, choose **Mapping template**.
 7. In the **Template** box, type the contents of the template that you want API Gateway to use to respond to the caller. For more information, see [Example Response Templates \(p. 68\)](#).
 8. Next to **Mapping template**, choose **Save**.
11. Do one of the following to test the method:
- Call the method from the API Gateway console. Follow the instructions in [Test a Method \(p. 233\)](#).
 - Call the method from a web browser, a web debugging proxy tool or the cURL command-line tool, or from your own API. Follow the instructions in [Calling an API \(p. 231\)](#).

Example Request Templates

The following example shows a request template that always uses the 200 HTTP status code.

```
{  
    "statusCode": 200  
}
```

The following example shows a request template that uses the 200 HTTP status code if the request specifies the `petType` parameter of `cat`; 400 if the request specifies `dog`; and uses 500 otherwise. This example is based on the one in the [Map API Request Parameter for HTTP Endpoints \(p. 30\)](#).

```
{  
  #if( $input.params('petType') == "cat" )  
  "statusCode": 200  
  #elseif( $input.params('petType') == "dog" )  
  "statusCode": 400  
  #else  
  "statusCode": 500  
  #end  
}
```

Example Response Templates

The following two examples show response templates that respond with the same information every time. These examples are based on the one in the [Map API Request Parameter for HTTP Endpoints \(p. 30\)](#).

```
## Example 400 response.  
{  
  "Message": "Error: petType not valid."  
}
```

```
## Example 500 response.  
{  
  "Message": "Error: petType not valid or not specified."  
}
```

The following example shows a response template that responds with the same information every time, but includes the value the caller specified for the `petType` parameter. This example is based on the one in the [Map API Request Parameter for HTTP Endpoints \(p. 30\)](#).

```
## Example 200 response for ?petType=cat (response will contain "type": "cat").  
{  
  "id": 1,  
  "name": "Kitty",  
  "type": "$input.params('petType')"  
}
```

Set Up Amazon API Gateway API Request and Response Payload Mappings

In API Gateway, an API's method request can take a payload in a different format from the corresponding integration request payload, as required in the back end. Similarly, the back end may return an integration response payload different from the method response payload, as expected by the front end. API Gateway lets you map the payload from a method request to the corresponding integration request and from an integration response to the corresponding method response. You use mapping templates to specify the mapping and can create model to facilitate the template generation. The section explains how to use the map the API request and response payload using models and mapping templates.

Topics

- [Models \(p. 69\)](#)
- [Mapping Templates \(p. 71\)](#)
- [Tasks for Models and Mapping Templates \(p. 74\)](#)
- [Create a Model in API Gateway \(p. 74\)](#)
- [View a List of Models in API Gateway \(p. 74\)](#)
- [Delete a Model in API Gateway \(p. 75\)](#)
- [Photos Example \(API Gateway Models and Mapping Templates\) \(p. 75\)](#)
- [News Article Example \(API Gateway Models and Mapping Templates\) \(p. 79\)](#)
- [Sales Invoice Example \(API Gateway Models and Mapping Templates\) \(p. 82\)](#)
- [Employee Record Example \(API Gateway Models and Mapping Templates\) \(p. 87\)](#)

Models

In API Gateway, a model defines the format, also known as the schema or shape, of some data. Models are most useful for generating strongly typed SDK of your API. They can also be useful in helping generate a mapping template or validate a payload. Because API Gateway is designed to work primarily with JavaScript Object Notation (JSON)-formatted data, API Gateway uses JSON Schema to define the expected schema of the data.

For example, the following expresses some JSON data:

```
{  
    "department": "produce",  
    "categories": [  
        "fruit",  
        "vegetables"  
    ],  
    "bins": [  
        {  
            "category": "fruit",  
            "type": "apples",  
            "price": 1.99,  
            "unit": "pound",  
            "quantity": 232  
        },  
        {  
            "category": "fruit",  
            "type": "bananas",  
            "price": 0.19,  
            "unit": "each",  
            "quantity": 112  
        },  
        {  
            "category": "vegetables",  
            "type": "carrots",  
            "price": 1.29,  
            "unit": "bag",  
            "quantity": 57  
        }  
    ]  
}
```

In the preceding example:

- The top-level or root object contains a `department` string object, a `categories` array, and a `bins` array.
- The `categories` array contains a collection of string values.
- The `bins` array contains a collection of objects. Each object contains a `category` string object, a `type` string object, a `price` number object, a `unit` string object, and a `quantity` number object.

The corresponding model is expressed in JSON Schema notation:

```
{
    "$schema": "http://json-schema.org/draft-04/schema#",
    "title": "GroceryStoreInputModel",
    "type": "object",
    "properties": {
        "department": { "type": "string" },
        "categories": {
            "type": "array",
            "items": { "type": "string" }
        },
        "bins": {
            "type": "array",
            "items": {
                "type": "object",
                "properties": {
                    "category": { "type": "string" },
                    "type": { "type": "string" },
                    "price": { "type": "number" },
                    "unit": { "type": "string" },
                    "quantity": { "type": "integer" }
                }
            }
        }
    }
}
```

In the preceding example:

- The `$schema` object represents a valid JSON Schema version identifier. In this example, it refers to JSON Schema, draft v4.
- The `title` object is a human-readable identifier for the model. In this example, it is `GroceryStoreInputModel`.
- The top-level, or root, construct in the JSON data is an object.
- The root object in the JSON data contains `department`, `categories`, and `bins` properties.
- The `department` property is a string object in the JSON data.
- The `categories` property is an array in the JSON data. The array contains string values in the JSON data.
- The `bins` property is an array in the JSON data. The array contains objects in the JSON data. Each of these objects in the JSON data contains a `category` string, a `type` string, a `price` number, a `unit` string, and a `quantity` integer (a number without a fraction or exponent part).

The example does not use advanced JSON Schema features, such as specifying required items; minimum and maximum allowed string lengths, numeric values, and array item lengths; regular expressions; and more.

For more information, see [Introducing JSON and JSON Schema](#).

For more complex JSON data formats and their models, see the following examples:

- [Input Model \(Photos Example\) \(p. 76\)](#) and [Output Model \(Photos Example\) \(p. 78\)](#) in the [Photos Example \(p. 75\)](#)
- [Input Model \(News Article Example\) \(p. 80\)](#) and [Output Model \(News Article Example\) \(p. 81\)](#) in the [News Article Example \(p. 79\)](#)
- [Input Model \(Sales Invoice Example\) \(p. 83\)](#) and [Output Model \(Sales Invoice Example\) \(p. 86\)](#) in the [Sales Invoice Example \(p. 82\)](#)
- [Input Model \(Employee Record Example\) \(p. 88\)](#) and [Output Model \(Employee Record Example\) \(p. 91\)](#) in the [Employee Record Example \(p. 87\)](#)

To experiment with models in API Gateway, follow the instructions in [Use Modes and Mapping Templates \(p. 39\)](#), specifically [Step 1: Create Models \(p. 41\)](#).

Mapping Templates

In API Gateway, a mapping template is used to transform some data from one format to another. You create and use input mapping templates and output mapping templates when you need to inform API Gateway about the schema of the data being sent from or returned to the caller, respectively. API Gateway uses the Velocity Template Language (VTL) and JSONPath expressions to define mapping templates.

For an example of an input mapping template, consider the example JSON data from the previous section. The following input mapping template makes no transform to the JSON data as API Gateway receives the JSON data from the caller:

```
#set($inputRoot = $input.path('$'))
{
    "department": "$inputRoot.department",
    "categories": [
#foreach($elem in $inputRoot.categories)
        "$elem"#if($foreach.hasNext),#end

#end
    ],
    "bins" : [
#foreach($elem in $inputRoot.bins)
        {
            "category" : "$elem.category",
            "type" : "$elem.type",
            "price" : $elem.price,
            "unit" : "$elem.unit",
            "quantity" : $elem.quantity
        }#if($foreach.hasNext),#end

#end
    ]
}
```

The preceding input mapping template is expressed as follows:

- Let the variable `$inputRoot` in the input mapping template represent the root object in the original JSON data.
- The values of the `department` object and `categories` and `bins` arrays in the input mapping template (represented by `$inputRoot.department`, `$inputRoot.categories`, and `$inputRoot.bins`)

map to the corresponding values of the `department` object and `categories` and `bins` arrays in the root object in the original JSON data.

- In the input mapping template, each of the values in the `categories` array (represented by the first `$elem`), and each of the objects in the `bins` array (represented by the second `$elem`), map to the corresponding values in the `categories` array and objects in the `bins` array, respectively, within the root object in the original JSON data.
- For each of objects in the `bins` object, the values of the `category`, `type`, `price`, `unit`, and `quantity` objects in the input mapping template (represented by `$elem.category`, `$elem.type`, `$elem.price`, `$elem.unit`, and `$elem.quantity`, respectively) map to the corresponding values of the `category`, `type`, `price`, `unit`, and `quantity` objects in the original JSON data, respectively.

For an example of an output mapping template, first consider the following JSON data schema, which is based on the example JSON data from the previous section.

Note

None of the array and object names in this JSON data schema match the JSON data from the previous section:

```
{
  "choices": [
    {
      "kind": "apples",
      "suggestedPrice": "1.99 per pound",
      "available": 232
    },
    {
      "kind": "bananas",
      "suggestedPrice": "0.19 per each",
      "available": 112
    },
    {
      "kind": "carrots",
      "suggestedPrice": "1.29 per bag",
      "available": 57
    }
  ]
}
```

To transform the example JSON data from the previous section into this JSON data schema, you would use the following model:

```
{
  "$schema": "http://json-schema.org/draft-04/schema#",
  "title": "GroceryStoreOutputModel",
  "type": "object",
  "properties": {
    "choices": {
      "type": "array",
      "items": {
        "type": "object",
        "properties": {
          "kind": { "type": "string" },
          "suggestedPrice": { "type": "string" },
          "available": { "type": "integer" }
        }
      }
    }
  }
}
```

```
    }
}
```

In the preceding example, the JSON schema is expressed as follows:

- The `$schema` object represents a valid JSON Schema version identifier. In this example, it refers to JSON Schema, draft v4.
- The `title` object is a human-readable identifier for the model. In this example, it is `GroceryStoreOutputModel`.
- The top-level, or root, construct in the JSON data is an object.
- The root object in the JSON data contains an array of objects.
- Each object in the array of objects contains a `kind` string, a `suggestedPrice` string, and an `available` integer (a number without a fraction or exponent part).

You would then use the following output mapping template, which is based on this model:

```
#set($inputRoot = $input.path('$'))
{
    "choices": [
#foreach($elem in $inputRoot.bins)
    {
        "kind": "$elem.type",
        "suggestedPrice": "$elem.price per $elem.unit",
        "available": $elem.quantity
    }#if($foreach.hasNext),#end
#end
    ]
}
```

The preceding output mapping template is expressed as follows:

- Let the variable `$inputRoot` in the output mapping template represent the root object in the original JSON data from the previous section. Note the variables in the output mapping template map to the original JSON data, not the desired transformed JSON data schema.
- The `choices` array in the output mapping template maps to the `bins` array with the root object in the original JSON data (`$inputRoot.bins`).
- In the output mapping template, each of the objects in the `choices` array (represented by `$elem`) map to the corresponding objects in the `bins` array within the root object in the original JSON data.
- In the output mapping template, for each of objects in the `choices` object, the values of the `kind` and `available` objects (represented by `$elem.type` and `$elem.quantity`) map to the corresponding values of the `type` and `value` objects in each of the objects in the original JSON data's `bins` array, respectively.
- In the output mapping template, for each of objects in the `choices` object, the value of the `suggestedPrice` object is a concatenation of the corresponding value of the `price` and `unit` objects in each of the objects in the original JSON data, respectively, with each value separated by the word `per`.

For more information about the Velocity Template Language, see [Apache Velocity - VTL Reference](#). For more information about JSONPath, see [JSONPath - XPath for JSON](#).

To explore more complex mapping templates, see the following examples:

- Input Mapping Template (Photos Example) (p. 77) and Output Mapping Template (Photos Example) (p. 78) in the Photos Example (p. 75)
- Input Mapping Template (News Article Example) (p. 80) and Output Mapping Template (News Article Example) (p. 82) in the News Article Example (p. 79)
- Input Mapping Template (Sales Invoice Example) (p. 85) and Output Mapping Template (Sales Invoice Example) (p. 86) in the Sales Invoice Example (p. 82)
- Input Mapping Template (Employee Record Example) (p. 89) and Output Mapping Template (Employee Record Example) (p. 92) in the Employee Record Example (p. 87)

To experiment with mapping templates in API Gateway, follow the instructions in [Use Modes and Mapping Templates \(p. 39\)](#), specifically [Step 5: Specify Method Settings and Test the Methods \(p. 45\)](#).

Tasks for Models and Mapping Templates

For additional things you can do with models and mapping templates, see the following:

- [Create a Model \(p. 74\)](#)
- [View a List of Models \(p. 74\)](#)
- [Delete a Model \(p. 75\)](#)

Create a Model in API Gateway

Use the API Gateway console to create a model for an API.

Topics

- [Prerequisites \(p. 74\)](#)
- [Create a Model With the API Gateway Console \(p. 74\)](#)

Prerequisites

- You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).

Create a Model With the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API where you want to create the model, choose **Models**.
3. Choose **Create**.
4. For **Model Name**, type a name for the model.
5. For **Content Type**, type the model's content type (for example, `application/json` for JSON).
6. (Optional) For **Model description**, type a description for the model.
7. For **Model schema**, type the model's schema. For more information about model schemas, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).
8. Choose **Create model**.

View a List of Models in API Gateway

Use the API Gateway console to view a list of models.

Topics

- [Prerequisites \(p. 75\)](#)
- [View a List of Models with the API Gateway Console \(p. 75\)](#)

Prerequisites

- You must have at least one model in API Gateway. Follow the instructions in [Create a Model \(p. 74\)](#).

View a List of Models with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API, choose **Models**.

Delete a Model in API Gateway

Use the API Gateway console to delete a model.

Warning

Deleting a model may cause part or all of the corresponding API to become unusable by API callers. Deleting a model cannot be undone.

Delete a Model with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the model, choose **Models**.
3. In the **Models** pane, choose the model you want to delete, and then choose **Delete Model**.
4. When prompted, choose **Delete**.

Photos Example (API Gateway Models and Mapping Templates)

The following sections provide examples of models and mapping templates that could be used for a sample photo API in API Gateway. For more information about models and mapping templates in API Gateway, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).

Topics

- [Original Data \(Photos Example\) \(p. 75\)](#)
- [Input Model \(Photos Example\) \(p. 76\)](#)
- [Input Mapping Template \(Photos Example\) \(p. 77\)](#)
- [Transformed Data \(Photos Example\) \(p. 77\)](#)
- [Output Model \(Photos Example\) \(p. 78\)](#)
- [Output Mapping Template \(Photos Example\) \(p. 78\)](#)

Original Data (Photos Example)

The following is the original JSON data for the photos example:

```
{
    "photos": {
        "page": 1,
        "pages": "1234",
        "perpage": 100,
        "total": "123398",
        "photo": [
            {
                "id": "12345678901",
                "owner": "23456789@A12",
                "secret": "abc123d456",
                "server": "1234",
                "farm": 1,
                "title": "Sample photo 1",
                "ispublic": 1,
                "isfriend": 0,
                "isfamily": 0
            },
            {
                "id": "23456789012",
                "owner": "34567890@B23",
                "secret": "bcd234e567",
                "server": "2345",
                "farm": 2,
                "title": "Sample photo 2",
                "ispublic": 1,
                "isfriend": 0,
                "isfamily": 0
            }
        ]
    }
}
```

Input Model (Photos Example)

The following is the input model that corresponds to the original JSON data for the photos example:

```
{
    "$schema": "http://json-schema.org/draft-04/schema#",
    "title": "PhotosInputModel",
    "type": "object",
    "properties": {
        "photos": {
            "type": "object",
            "properties": {
                "page": { "type": "integer" },
                "pages": { "type": "string" },
                "perpage": { "type": "integer" },
                "total": { "type": "string" },
                "photo": {
                    "type": "array",
                    "items": {
                        "type": "object",
                        "properties": {
                            "id": { "type": "string" },
                            "owner": { "type": "string" },
                            "secret": { "type": "string" },
                            "server": { "type": "string" },
                            "farm": { "type": "integer" },
                            "title": { "type": "string" },
                            "ispublic": { "type": "boolean" },
                            "isfriend": { "type": "boolean" },
                            "isfamily": { "type": "boolean" }
                        }
                    }
                }
            }
        }
    }
}
```

```

        "secret": { "type": "string" },
        "server": { "type": "string" },
        "farm": { "type": "integer" },
        "title": { "type": "string" },
        "ispublic": { "type": "integer" },
        "isfriend": { "type": "integer" },
        "isfamily": { "type": "integer" }
    }
}
}
}
}
}
}
```

Input Mapping Template (Photos Example)

The following is the input mapping template that corresponds to the original JSON data for the photos example:

```

#set($inputRoot = $input.path('$.'))
{
    "photos": {
        "page": $inputRoot.photos.page,
        "pages": "$inputRoot.photos.pages",
        "perpage": $inputRoot.photos.perpage,
        "total": "$inputRoot.photos.total",
        "photo": [
#foreach($elem in $inputRoot.photos.photo)
            {
                "id": "$elem.id",
                "owner": "$elem.owner",
                "secret": "$elem.secret",
                "server": "$elem.server",
                "farm": $elem.farm,
                "title": "$elem.title",
                "ispublic": $elem.ispublic,
                "isfriend": $elem.isfriend,
                "isfamily": $elem.isfamily
            }#if($foreach.hasNext),#end
        ]
    }
}
```

Transformed Data (Photos Example)

The following is one example of how the original photos example JSON data could be transformed for output:

```
{
    "photos": [
        {
            "id": "12345678901",
            "owner": "12345678901",
            "secret": "12345678901",
            "server": "12345678901",
            "farm": 12345678901,
            "title": "Photo 12345678901",
            "ispublic": 1,
            "isfriend": 1,
            "isfamily": 1
        }
    ]
}
```

```

        "owner": "23456789@A12",
        "title": "Sample photo 1",
        "ispublic": 1,
        "isfriend": 0,
        "isfamily": 0
    },
    {
        "id": "23456789012",
        "owner": "34567890@B23",
        "title": "Sample photo 2",
        "ispublic": 1,
        "isfriend": 0,
        "isfamily": 0
    }
]
}

```

Output Model (Photos Example)

The following is the output model that corresponds to the transformed JSON data format:

```
{
    "$schema": "http://json-schema.org/draft-04/schema#",
    "title": "PhotosOutputModel",
    "type": "object",
    "properties": {
        "photos": {
            "type": "array",
            "items": {
                "type": "object",
                "properties": {
                    "id": { "type": "string" },
                    "owner": { "type": "string" },
                    "title": { "type": "string" },
                    "ispublic": { "type": "integer" },
                    "isfriend": { "type": "integer" },
                    "isfamily": { "type": "integer" }
                }
            }
        }
    }
}
```

Output Mapping Template (Photos Example)

The following is the output mapping template that corresponds to the transformed JSON data format. The template variables here are based on the original, not transformed, JSON data format:

```

#set($inputRoot = $input.path('$'))
{
    "photos": [
#foreach($elem in $inputRoot.photos.photo)
    {
        "id": "$elem.id",
        "owner": "$elem.owner",

```

```
        "title": "$elem.title",
        "ispublic": $elem.ispublic,
        "isfriend": $elem.isfriend,
        "isfamily": $elem.isfamily
    }#if($foreach.hasNext),#end

}#end
]
}
```

News Article Example (API Gateway Models and Mapping Templates)

The following sections provide examples of models and mapping templates that could be used for a sample news article API in API Gateway. For more information about models and mapping templates in API Gateway, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).

Topics

- [Original Data \(News Article Example\) \(p. 79\)](#)
- [Input Model \(News Article Example\) \(p. 80\)](#)
- [Input Mapping Template \(News Article Example\) \(p. 80\)](#)
- [Transformed Data \(News Article Example\) \(p. 81\)](#)
- [Output Model \(News Article Example\) \(p. 81\)](#)
- [Output Mapping Template \(News Article Example\) \(p. 82\)](#)

Original Data (News Article Example)

The following is the original JSON data for the news article example:

```
{
    "count": 1,
    "items": [
        {
            "last_updated_date": "2015-04-24",
            "expire_date": "2016-04-25",
            "author_first_name": "John",
            "description": "Sample Description",
            "creation_date": "2015-04-20",
            "title": "Sample Title",
            "allow_comment": "1",
            "author": {
                "last_name": "Doe",
                "email": "johndoe@example.com",
                "first_name": "John"
            },
            "body": "Sample Body",
            "publish_date": "2015-04-25",
            "version": "1",
            "author_last_name": "Doe",
            "parent_id": 2345678901,
            "article_url": "http://www.example.com/articles/3456789012"
        }
    ]
}
```

```
],
"version": 1
}
```

Input Model (News Article Example)

The following is the input model that corresponds to the original JSON data for the news article example:

```
{
  "$schema": "http://json-schema.org/draft-04/schema#",
  "title": "NewsArticleInputModel",
  "type": "object",
  "properties": {
    "count": { "type": "integer" },
    "items": {
      "type": "array",
      "items": {
        "type": "object",
        "properties": {
          "last_updated_date": { "type": "string" },
          "expire_date": { "type": "string" },
          "author_first_name": { "type": "string" },
          "description": { "type": "string" },
          "creation_date": { "type": "string" },
          "title": { "type": "string" },
          "allow_comment": { "type": "string" },
          "author": {
            "type": "object",
            "properties": {
              "last_name": { "type": "string" },
              "email": { "type": "string" },
              "first_name": { "type": "string" }
            }
          },
          "body": { "type": "string" },
          "publish_date": { "type": "string" },
          "version": { "type": "string" },
          "author_last_name": { "type": "string" },
          "parent_id": { "type": "integer" },
          "article_url": { "type": "string" }
        }
      }
    },
    "version": { "type": "integer" }
  }
}
```

Input Mapping Template (News Article Example)

The following is the input mapping template that corresponds to the original JSON data for the news article example:

```
#set($inputRoot = $input.path('$'))
{
  "count": $inputRoot.count,
```

```

    "items": [
#foreach($elem in $inputRoot.items)
{
    "last_updated_date": "$elem.last_updated_date",
    "expire_date": "$elem.expire_date",
    "author_first_name": "$elem.author_first_name",
    "description": "$elem.description",
    "creation_date": "$elem.creation_date",
    "title": "$elem.title",
    "allow_comment": "$elem.allow_comment",
    "author": {
        "last_name": "$elem.author.last_name",
        "email": "$elem.author.email",
        "first_name": "$elem.author.first_name"
    },
    "body": "$elem.body",
    "publish_date": "$elem.publish_date",
    "version": "$elem.version",
    "author_last_name": "$elem.author_last_name",
    "parent_id": $elem.parent_id,
    "article_url": "$elem.article_url"
}#if($foreach.hasNext),#end
},
],
"version": $inputRoot.version
}

```

Transformed Data (News Article Example)

The following is one example of how the original news article example JSON data could be transformed for output:

```
{
    "count": 1,
    "items": [
        {
            "creation_date": "2015-04-20",
            "title": "Sample Title",
            "author": "John Doe",
            "body": "Sample Body",
            "publish_date": "2015-04-25",
            "article_url": "http://www.example.com/articles/3456789012"
        }
    ],
    "version": 1
}
```

Output Model (News Article Example)

The following is the output model that corresponds to the transformed JSON data format:

```
{
    "$schema": "http://json-schema.org/draft-04/schema#",
    "title": "NewsArticleOutputModel",
```

```
"type": "object",
"properties": {
    "count": { "type": "integer" },
    "items": {
        "type": "array",
        "items": {
            "type": "object",
            "properties": {
                "creation_date": { "type": "string" },
                "title": { "type": "string" },
                "author": { "type": "string" },
                "body": { "type": "string" },
                "publish_date": { "type": "string" },
                "article_url": { "type": "string" }
            }
        }
    },
    "version": { "type": "integer" }
}
```

Output Mapping Template (News Article Example)

The following is the output mapping template that corresponds to the transformed JSON data format. The template variables here are based on the original, not transformed, JSON data format:

```
#set($inputRoot = $input.path('$'))
{
    "count": $inputRoot.count,
    "items": [
#foreach($elem in $inputRoot.items)
    {
        "creation_date": "$elem.creation_date",
        "title": "$elem.title",
        "author": "$elem.author.first_name $elem.author.last_name",
        "body": "$elem.body",
        "publish_date": "$elem.publish_date",
        "article_url": "$elem.article_url"
    }#if($foreach.hasNext),#end
#end
    ],
    "version": $inputRoot.version
}
```

Sales Invoice Example (API Gateway Models and Mapping Templates)

The following sections provide examples of models and mapping templates that could be used for a sample sales invoice API in API Gateway. For more information about models and mapping templates in API Gateway, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).

Topics

- [Original Data \(Sales Invoice Example\) \(p. 83\)](#)

- [Input Model \(Sales Invoice Example\) \(p. 83\)](#)
- [Input Mapping Template \(Sales Invoice Example\) \(p. 85\)](#)
- [Transformed Data \(Sales Invoice Example\) \(p. 85\)](#)
- [Output Model \(Sales Invoice Example\) \(p. 86\)](#)
- [Output Mapping Template \(Sales Invoice Example\) \(p. 86\)](#)

Original Data (Sales Invoice Example)

The following is the original JSON data for the sales invoice example:

```
{  
    "DueDate": "2013-02-15",  
    "Balance": 1990.19,  
    "DocNumber": "SAMP001",  
    "Status": "Payable",  
    "Line": [  
        {  
            "Description": "Sample Expense",  
            "Amount": 500,  
            "DetailType": "ExpenseDetail",  
            "ExpenseDetail": {  
                "Customer": {  
                    "value": "ABC123",  
                    "name": "Sample Customer"  
                },  
                "Ref": {  
                    "value": "DEF234",  
                    "name": "Sample Construction"  
                },  
                "Account": {  
                    "value": "EFG345",  
                    "name": "Fuel"  
                },  
                "LineStatus": "Billable"  
            }  
        }  
    ],  
    "Vendor": {  
        "value": "GHI456",  
        "name": "Sample Bank"  
    },  
    "APRef": {  
        "value": "HIJ567",  
        "name": "Accounts Payable"  
    },  
    "TotalAmt": 1990.19  
}
```

Input Model (Sales Invoice Example)

The following is the input model that corresponds to the original JSON data for the sales invoice example:

```
{  
    "$schema": "http://json-schema.org/draft-04/schema#",  
    "type": "object",  
    "properties": {  
        "DueDate": {  
            "type": "string",  
            "format": "date-time"  
        },  
        "Balance": {  
            "type": "number",  
            "format": "float"  
        },  
        "DocNumber": {  
            "type": "string",  
            "format": "string"  
        },  
        "Status": {  
            "type": "string",  
            "format": "string"  
        },  
        "Line": {  
            "type": "array",  
            "items": {  
                "type": "object",  
                "properties": {  
                    "Description": {  
                        "type": "string",  
                        "format": "string"  
                    },  
                    "Amount": {  
                        "type": "number",  
                        "format": "float"  
                    },  
                    "DetailType": {  
                        "type": "string",  
                        "format": "string"  
                    },  
                    "ExpenseDetail": {  
                        "type": "object",  
                        "properties": {  
                            "Customer": {  
                                "type": "object",  
                                "properties": {  
                                    "value": {  
                                        "type": "string",  
                                        "format": "string"  
                                    },  
                                    "name": {  
                                        "type": "string",  
                                        "format": "string"  
                                    }  
                                }  
                            },  
                            "Ref": {  
                                "type": "object",  
                                "properties": {  
                                    "value": {  
                                        "type": "string",  
                                        "format": "string"  
                                    },  
                                    "name": {  
                                        "type": "string",  
                                        "format": "string"  
                                    }  
                                }  
                            },  
                            "Account": {  
                                "type": "object",  
                                "properties": {  
                                    "value": {  
                                        "type": "string",  
                                        "format": "string"  
                                    },  
                                    "name": {  
                                        "type": "string",  
                                        "format": "string"  
                                    }  
                                }  
                            },  
                            "LineStatus": {  
                                "type": "string",  
                                "format": "string"  
                            }  
                        }  
                    }  
                }  
            }  
        },  
        "Vendor": {  
            "type": "object",  
            "properties": {  
                "value": {  
                    "type": "string",  
                    "format": "string"  
                },  
                "name": {  
                    "type": "string",  
                    "format": "string"  
                }  
            }  
        },  
        "APRef": {  
            "type": "object",  
            "properties": {  
                "value": {  
                    "type": "string",  
                    "format": "string"  
                },  
                "name": {  
                    "type": "string",  
                    "format": "string"  
                }  
            }  
        },  
        "TotalAmt": {  
            "type": "number",  
            "format": "float"  
        }  
    }  
}
```

```

"title": "InvoiceInputModel",
"type": "object",
"properties": {
    "DueDate": { "type": "string" },
    "Balance": { "type": "number" },
    "DocNumber": { "type": "string" },
    "Status": { "type": "string" },
    "Line": {
        "type": "array",
        "items": {
            "type": "object",
            "properties": {
                "Description": { "type": "string" },
                "Amount": { "type": "integer" },
                "DetailType": { "type": "string" },
                "ExpenseDetail": {
                    "type": "object",
                    "properties": {
                        "Customer": {
                            "type": "object",
                            "properties": {
                                "value": { "type": "string" },
                                "name": { "type": "string" }
                            }
                        },
                        "Ref": {
                            "type": "object",
                            "properties": {
                                "value": { "type": "string" },
                                "name": { "type": "string" }
                            }
                        }
                    }
                },
                "Account": {
                    "type": "object",
                    "properties": {
                        "value": { "type": "string" },
                        "name": { "type": "string" }
                    }
                },
                "LineStatus": { "type": "string" }
            }
        }
    },
    "Vendor": {
        "type": "object",
        "properties": {
            "value": { "type": "string" },
            "name": { "type": "string" }
        }
    }
},
"APRef": {
    "type": "object",
    "properties": {
        "value": { "type": "string" },
        "name": { "type": "string" }
    }
}
}

```

```

        },
        "TotalAmt": { "type": "number" }
    }
}
```

Input Mapping Template (Sales Invoice Example)

The following is the input mapping template that corresponds to the original JSON data for the sales invoice example:

```
#set($inputRoot = $input.path('$'))
{
    "DueDate": "$inputRoot.DueDate",
    "Balance": $inputRoot.Balance,
    "DocNumber": "$inputRoot.DocNumber",
    "Status": "$inputRoot.Status",
    "Line": [
#foreach($elem in $inputRoot.Line)
    {
        "Description": "$elem.Description",
        "Amount": $elem.Amount,
        "DetailType": "$elem.DetailType",
        "ExpenseDetail": {
            "Customer": {
                "value": "$elem.ExpenseDetail.Customer.value",
                "name": "$elem.ExpenseDetail.Customer.name"
            },
            "Ref": {
                "value": "$elem.ExpenseDetail.Ref.value",
                "name": "$elem.ExpenseDetail.Ref.name"
            },
            "Account": {
                "value": "$elem.ExpenseDetail.Account.value",
                "name": "$elem.ExpenseDetail.Account.name"
            },
            "LineStatus": "$elem.ExpenseDetail.LineStatus"
        }
    }#if($foreach.hasNext),#end
}
#end
],
"Vendor": {
    "value": "$inputRoot.Vendor.value",
    "name": "$inputRoot.Vendor.name"
},
"APRef": {
    "value": "$inputRoot.APRef.value",
    "name": "$inputRoot.APRef.name"
},
"TotalAmt": $inputRoot.TotalAmt
}
```

Transformed Data (Sales Invoice Example)

The following is one example of how the original sales invoice example JSON data could be transformed for output:

```
{  
    "DueDate": "2013-02-15",  
    "Balance": 1990.19,  
    "DocNumber": "SAMP001",  
    "Status": "Payable",  
    "Line": [  
        {  
            "Description": "Sample Expense",  
            "Amount": 500,  
            "DetailType": "ExpenseDetail",  
            "Customer": "ABC123 (Sample Customer)",  
            "Ref": "DEF234 (Sample Construction)",  
            "Account": "EFG345 (Fuel)",  
            "LineStatus": "Billable"  
        }  
    ],  
    "TotalAmt": 1990.19  
}
```

Output Model (Sales Invoice Example)

The following is the output model that corresponds to the transformed JSON data format:

```
{  
    "$schema": "http://json-schema.org/draft-04/schema#",  
    "title": "InvoiceOutputModel",  
    "type": "object",  
    "properties": {  
        "DueDate": { "type": "string" },  
        "Balance": { "type": "number" },  
        "DocNumber": { "type": "string" },  
        "Status": { "type": "string" },  
        "Line": {  
            "type": "array",  
            "items": {  
                "type": "object",  
                "properties": {  
                    "Description": { "type": "string" },  
                    "Amount": { "type": "integer" },  
                    "DetailType": { "type": "string" },  
                    "Customer": { "type": "string" },  
                    "Ref": { "type": "string" },  
                    "Account": { "type": "string" },  
                    "LineStatus": { "type": "string" }  
                }  
            }  
        },  
        "TotalAmt": { "type": "number" }  
    }  
}
```

Output Mapping Template (Sales Invoice Example)

The following is the output mapping template that corresponds to the transformed JSON data format. The template variables here are based on the original, not transformed, JSON data format:

```
#set($inputRoot = $input.path('$'))
{
    "DueDate": "$inputRoot.DueDate",
    "Balance": $inputRoot.Balance,
    "DocNumber": "$inputRoot.DocNumber",
    "Status": "$inputRoot.Status",
    "Line": [
        #foreach($elem in $inputRoot.Line)
        {
            "Description": "$elem.Description",
            "Amount": $elem.Amount,
            "DetailType": "$elem.DetailType",
            "Customer": "$elem.ExpenseDetail.Customer.value ($elem.ExpenseDetail.Customer.name)",
            "Ref": "$elem.ExpenseDetail.Ref.value ($elem.ExpenseDetail.Ref.name)",
            "Account": "$elem.ExpenseDetail.Account.value ($elem.ExpenseDetail.Account.name)",
            "LineStatus": "$elem.ExpenseDetail.LineStatus"
        }#if($foreach.hasNext),#end
    ]
    "TotalAmt": $inputRoot.TotalAmt
}
```

Employee Record Example (API Gateway Models and Mapping Templates)

The following sections provide examples of models and mapping templates that can be used for a sample employee record API in API Gateway. For more information about models and mapping templates in API Gateway, see [Set Up Request and Response Payload Mappings \(p. 68\)](#).

Topics

- [Original Data \(Employee Record Example\) \(p. 87\)](#)
- [Input Model \(Employee Record Example\) \(p. 88\)](#)
- [Input Mapping Template \(Employee Record Example\) \(p. 89\)](#)
- [Transformed Data \(Employee Record Example\) \(p. 90\)](#)
- [Output Model \(Employee Record Example\) \(p. 91\)](#)
- [Output Mapping Template \(Employee Record Example\) \(p. 92\)](#)

Original Data (Employee Record Example)

The following is the original JSON data for the employee record example:

```
{
    "QueryResponse": {
        "maxResults": "1",
        "startPosition": "1",
        "Employee": {
            "Organization": "false",
            "Title": "Mrs.",
            "GivenName": "Jane",
```

```
"MiddleName": "Lane",
"FamilyName": "Doe",
"DisplayName": "Jane Lane Doe",
"PrintOnCheckName": "Jane Lane Doe",
"Active": "true",
"PrimaryPhone": { "FreeFormNumber": "505.555.9999" },
"PrimaryEmailAddr": { "Address": "janedoe@example.com" },
"EmployeeType": "Regular",
"status": "Synchronized",
"Id": "ABC123",
"SyncToken": "1",
"MetaData": {
    "CreateTime": "2015-04-26T19:45:03Z",
    "LastUpdatedTime": "2015-04-27T21:48:23Z"
},
"PrimaryAddr": {
    "Line1": "123 Any Street",
    "City": "Any City",
    "CountrySubDivisionCode": "WA",
    "PostalCode": "01234"
}
},
"time": "2015-04-27T22:12:32.012Z"
}
```

Input Model (Employee Record Example)

The following is the input model that corresponds to the original JSON data for the employee record example:

```
{
    "$schema": "http://json-schema.org/draft-04/schema#",
    "title": "EmployeeInputModel",
    "type": "object",
    "properties": {
        "QueryResponse": {
            "type": "object",
            "properties": {
                "maxResults": { "type": "string" },
                "startPosition": { "type": "string" },
                "Employee": {
                    "type": "object",
                    "properties": {
                        "Organization": { "type": "string" },
                        "Title": { "type": "string" },
                        "GivenName": { "type": "string" },
                        "MiddleName": { "type": "string" },
                        "FamilyName": { "type": "string" },
                        "DisplayName": { "type": "string" },
                        "PrintOnCheckName": { "type": "string" },
                        "Active": { "type": "string" },
                        "PrimaryPhone": {
                            "type": "object",
                            "properties": {
                                "FreeFormNumber": { "type": "string" }
                            }
                        }
                    }
                }
            }
        }
    }
}
```

```
        }
    },
    "PrimaryEmailAddr": {
        "type": "object",
        "properties": {
            "Address": { "type": "string" }
        }
    },
    "EmployeeType": { "type": "string" },
    "status": { "type": "string" },
    "Id": { "type": "string" },
    "SyncToken": { "type": "string" },
    "MetaData": {
        "type": "object",
        "properties": {
            "CreateTime": { "type": "string" },
            "LastUpdatedTime": { "type": "string" }
        }
    },
    "PrimaryAddr": {
        "type": "object",
        "properties": {
            "Line1": { "type": "string" },
            "City": { "type": "string" },
            "CountrySubDivisionCode": { "type": "string" },
            "PostalCode": { "type": "string" }
        }
    }
},
"time": { "type": "string" }
}
```

Input Mapping Template (Employee Record Example)

The following is the input mapping template that corresponds to the original JSON data for the employee record example:

```
#set($inputRoot = $input.path('$'))
{
    "QueryResponse": {
        "maxResults": "$inputRoot.QueryResponse.maxResults",
        "startPosition": "$inputRoot.QueryResponse.startPosition",
        "Employee": {
            "Organization": "$inputRoot.QueryResponse.Employee.Organization",
            "Title": "$inputRoot.QueryResponse.Employee.Title",
            "GivenName": "$inputRoot.QueryResponse.Employee.GivenName",
            "MiddleName": "$inputRoot.QueryResponse.Employee.MiddleName",
            "FamilyName": "$inputRoot.QueryResponse.Employee.FamilyName",
            "DisplayName": "$inputRoot.QueryResponse.Employee.DisplayName",
            "PrintOnCheckName": "$inputRoot.QueryResponse.Employee.PrintOnCheckName",
            "Active": "$inputRoot.QueryResponse.Employee.Active",
            "EmployeeType": "$inputRoot.QueryResponse.Employee.EmployeeType"
        }
    }
}
```

```
    "PrimaryPhone": { "FreeFormNumber": "$inputRoot.QueryResponse.Employee.PrimaryPhone.FreeFormNumber" },
    "PrimaryEmailAddr": { "Address": "$inputRoot.QueryResponse.Employee.PrimaryEmailAddr.Address" },
    "EmployeeType": "$inputRoot.QueryResponse.Employee.EmployeeType",
    "status": "$inputRoot.QueryResponse.Employee.status",
    "Id": "$inputRoot.QueryResponse.Employee.Id",
    "SyncToken": "$inputRoot.QueryResponse.Employee.SyncToken",
    "MetaData": {
        "CreateTime": "$inputRoot.QueryResponse.Employee.MetaData.CreateTime",
        "LastUpdatedTime": "$inputRoot.QueryResponse.Employee.MetaData.LastUpdateTime"
    },
    "PrimaryAddr" : {
        "Line1": "$inputRoot.QueryResponse.Employee.PrimaryAddr.Line1",
        "City": "$inputRoot.QueryResponse.Employee.PrimaryAddr.City",
        "CountrySubDivisionCode": "$inputRoot.QueryResponse.Employee.PrimaryAddr.CountrySubDivisionCode",
        "PostalCode": "$inputRoot.QueryResponse.Employee.PrimaryAddr.PostalCode"
    }
},
"time": "$inputRoot.time"
}
```

Transformed Data (Employee Record Example)

The following is one example of how the original employee record example JSON data could be transformed for output:

```
{
    "QueryResponse": {
        "maxResults": "1",
        "startPosition": "1",
        "Employees": [
            {
                "Title": "Mrs.",
                "GivenName": "Jane",
                "MiddleName": "Lane",
                "FamilyName": "Doe",
                "DisplayName": "Jane Lane Doe",
                "PrintOnCheckName": "Jane Lane Doe",
                "Active": "true",
                "PrimaryPhone": "505.555.9999",
                "Email": [
                    {
                        "type": "primary",
                        "Address": "janedoe@example.com"
                    }
                ],
                "EmployeeType": "Regular",
                "PrimaryAddr": {
                    "Line1": "123 Any Street",
                    "City": "Any City",

```

```
        "CountrySubDivisionCode": "WA",
        "PostalCode": "01234"
    }
]
},
"time": "2015-04-27T22:12:32.012Z"
}
```

Output Model (Employee Record Example)

The following is the output model that corresponds to the transformed JSON data format:

```
{
  "$schema": "http://json-schema.org/draft-04/schema#",
  "title": "EmployeeOutputModel",
  "type": "object",
  "properties": {
    "QueryResponse": {
      "type": "object",
      "properties": {
        "maxResults": { "type": "string" },
        "startPosition": { "type": "string" },
        "Employees": {
          "type": "array",
          "items": {
            "type": "object",
            "properties": {
              "Title": { "type": "string" },
              "GivenName": { "type": "string" },
              "MiddleName": { "type": "string" },
              "FamilyName": { "type": "string" },
              "DisplayName": { "type": "string" },
              "PrintOnCheckName": { "type": "string" },
              "Active": { "type": "string" },
              "PrimaryPhone": { "type": "string" },
              "Email": {
                "type": "array",
                "items": {
                  "type": "object",
                  "properties": {
                    "type": { "type": "string" },
                    "Address": { "type": "string" }
                  }
                }
              }
            },
            "EmployeeType": { "type": "string" },
            "PrimaryAddr": {
              "type": "object",
              "properties": {
                "Line1": { "type": "string" },
                "City": { "type": "string" },
                "CountrySubDivisionCode": { "type": "string" },
                "PostalCode": { "type": "string" }
              }
            }
          }
        }
      }
    }
  }
}
```

```
        }
    }
},
"time": { "type": "string" }
}
```

Output Mapping Template (Employee Record Example)

The following is the output mapping template that corresponds to the transformed JSON data format. The template variables here are based on the original, not transformed, JSON data format:

```
#set($inputRoot = $input.path('$'))
{
    "QueryResponse": {
        "maxResults": "$inputRoot.QueryResponse.maxResults",
        "startPosition": "$inputRoot.QueryResponse.startPosition",
        "Employees": [
            {
                "Title": "$inputRoot.QueryResponse.Employee.Title",
                "GivenName": "$inputRoot.QueryResponse.Employee.GivenName",
                "MiddleName": "$inputRoot.QueryResponse.Employee.MiddleName",
                "FamilyName": "$inputRoot.QueryResponse.Employee.FamilyName",
                "DisplayName": "$inputRoot.QueryResponse.Employee.DisplayName",
                "PrintOnCheckName": "$inputRoot.QueryResponse.Employee.PrintOnCheckName",

                "Active": "$inputRoot.QueryResponse.Employee.Active",
                "PrimaryPhone": "$inputRoot.QueryResponse.Employee.PrimaryPhone.FreeFormNumber",
                "Email" : [
                    {
                        "type": "primary",
                        "Address": "$inputRoot.QueryResponse.Employee.PrimaryEmailAddr.Address"
                    }
                ],
                "EmployeeType": "$inputRoot.QueryResponse.Employee.EmployeeType",
                "PrimaryAddr": {
                    "Line1": "$inputRoot.QueryResponse.Employee.PrimaryAddr.Line1",
                    "City": "$inputRoot.QueryResponse.Employee.PrimaryAddr.City",
                    "CountrySubDivisionCode": "$inputRoot.QueryResponse.Employee.PrimaryAddr.CountrySubDivisionCode",
                    "PostalCode": "$inputRoot.QueryResponse.Employee.PrimaryAddr.PostalCode"
                }
            ]
        },
        "time": "$inputRoot.time"
    }
}
```

Amazon API Gateway API Request and Response Parameter-Mapping Reference

This section explains how to set up data mappings from an API's method request data, including other data stored in [context \(p. 96\)](#), [stage \(p. 100\)](#) or [util \(p. 100\)](#) variables, to the corresponding integration request parameters and from an integration response data, including the other data, to the method response parameters. The method request data includes request parameters (path, query string, headers) and the body. The integration response data includes response parameters (headers), and the body. For more information about using the stage variables, see [Amazon API Gateway Stage Variables Reference \(p. 220\)](#).

Topics

- [Map Data to Integration Request Parameters \(p. 93\)](#)
- [Map Data to Method Response Headers \(p. 94\)](#)
- [Transform Request and Response Bodies \(p. 95\)](#)

Map Data to Integration Request Parameters

Integration request parameters, in the form of path variables, query strings or headers, can be mapped from any defined method request parameters and the payload.

Integration request data mapping expressions

Mapped data source	Mapping expression
Method request path	method.request.path. <i>PARAM_NAME</i>
Method request query string	method.request.querystring. <i>PARAM_NAME</i>
Method request header	method.request.header. <i>PARAM_NAME</i>
Method request body	method.request.body
Method request body (JsonPath)	method.request.body. <i>JSONPath_EXPRESSION</i> .
Stage variables	stageVariables. <i>VARIABLE_NAME</i>
Context variables	context. <i>VARIABLE_NAME</i> that must be one of the supported context variables (p. 96) .
Static value	' <i>STATIC_VALUE</i> '. The <i>STATIC_VALUE</i> is a string literal and must be enclosed within a pair of single quotes.

Here, *PARAM_NAME* is the name of a method request parameter of the given parameter type. It must have been defined before it can be referenced. *JSONPath_EXPRESSION* is a JSONPath expression for a JSON field of the body of a request or response. However, the "\$." prefix is omitted in this syntax.

Example mappings from method request parameter in Swagger

The following example shows a Swagger snippet that maps 1) the method request's header, named `methodRequestHeaderParam`, into the integration request path parameter, named `integrationPathParam`; 2) the method request query string, named `methodRequestQueryParam`, into the integration request query string, named `integrationQueryParam`.

```

...
"requestParameters" : {

    "integration.request.path.integrationPathParam" : "method.request.header.methodRequestHeaderParam",
    "integration.request.querystring.integrationQueryParam" : "method.request.querystring.methodRequestQueryParam"

}
...

```

Integration request parameters can also be mapped from fields in the JSON request body using a [JSONPath expression](#). The following table shows the mapping expressions for a method request body and its JSON fields.

Example mapping from method request body in Swagger

The following example shows a Swagger snippet that maps 1) the method request body to the integration request header, named `body-header`, and 2) a JSON field of the body, as expressed by a JSON expression (`petstore.pets[0].name`, without the `$.` prefix).

```

...
"requestParameters" : {

    "integration.request.header.body-header" : "method.request.body",
    "integration.request.path.pet-name" : "method.request.body.pet
store.pets[0].name",

}
...

```

Map Data to Method Response Headers

Method response header parameters can be mapped from any integration response header, from the integration response body, or from the method request body.

Method response header mapping expressions

Mapped Data Source	Mapping expression
Integration response header	<code>integration.response.header.PARAM_NAME</code>
Integration response body	<code>integration.response.body</code>

Mapped Data Source	Mapping expression
Integration response body (JsonPath)	integration.response.body. <i>JSONPath_EXPRESSION</i>
Method request body	method.request.body
Method request body (JsonPath)	method.request.body. <i>JSONPath_EXPRESSION</i>
Stage variable	stageVariables. <i>VARIABLE_NAME</i>
Context variable	context. <i>VARIABLE_NAME</i> that must be one of the supported context variables (p. 96) .
Static value	' <i>STATIC_VALUE</i> '. The <i>STATIC_VALUE</i> is a string literal and must be enclosed within a pair of single quotes.

Example data mapping from integration response in Swagger

The following example shows a Swagger snippet that maps 1) the integration response's `redirect.url`, `JSONPath` field into the request response's `location` header; and 2) the integration response's `x-app-id` header to the method response's `id` header.

```

...
"responseParameters" : {
    "method.response.header.location" : "integration.response.body.redirect.url",
    "method.response.header.id" : "integration.response.header.x-app-id",
}
...

```

Transform Request and Response Bodies

Integration request and method response bodies can be transformed from the method request and integration response bodies, respectively, by using [Mapping Templates \(p. 71\)](#) written in [Velocity Template Language \(VTL\)](#). JSON data can be manipulated using VTL logic and JSONPath expressions, and additional data can be included from HTTP parameters, the calling context, and stage variables.

Select Mapping Templates

The request mapping template used to transform the method request body into the integration request body is selected by the value of the "Content-Type" header sent in the client request.

The response mapping template used to transform the integration response body into the method response body is selected by the value of the "Accept" header sent in the client request.

For example, if the client sends headers of "Content-Type : application/xml", and "Accept : application/json", the request template with the `application/xml` key will be used for the integration request, and the response template with the `application/json` key will be used for the method response.

Only the MIME type is used from the `Accept` and `Content-Type` headers when selecting a mapping template. For example, a header of `"Content-Type: application/json; charset=UTF-8"` will have a request template with the `application/json` key selected.

API Gateway API Request and Response Payload-Mapping Template Reference

Amazon API Gateway defines a set of variables for working with models and mapping templates. This document describes those functions and provides examples for working with input payloads.

Topics

- [Accessing the \\$context Variable \(p. 96\)](#)
- [Accessing the \\$input Variable \(p. 98\)](#)
- [Accessing the \\$stageVariables Variable \(p. 100\)](#)
- [Accessing the \\$util Variable \(p. 100\)](#)

Accessing the `$context` Variable

The `$context` variable holds all the contextual information of your API call.

`$context` Variable Reference

Parameter	Description
<code>\$context.apiId</code>	The identifier API Gateway assigns to your API.
<code>\$context.authorizer.principalId</code>	The principal user identification associated with the token sent by the client.
<code>\$context.httpMethod</code>	The HTTP method used. Valid values include: <code>DELETE</code> , <code>GET</code> , <code>HEAD</code> , <code>OPTIONS</code> , <code>PATCH</code> , <code>POST</code> , and <code>PUT</code> .
<code>\$context.identity.accountId</code>	The AWS account ID associated with the request.
<code>\$context.identity.apiKey</code>	The API owner key associated with your API.
<code>\$context.identity.caller</code>	The principal identifier of the caller making the request.
<code>\$context.identity.cognitoAuthenticationProvider</code>	The Amazon Cognito authentication provider used by the caller making the request. Available only if the request was signed with Amazon Cognito credentials. For information related to this and the other Amazon Cognito <code>\$context</code> variables, see Amazon Cognito Identity .
<code>\$context.identity.cognitoAuthenticationType</code>	The Amazon Cognito authentication type of the caller making the request. Available only if the request was signed with Amazon Cognito credentials.

Parameter	Description
\$context.identity.cognitoIdentityId	The Amazon Cognito identity ID of the caller making the request. Available only if the request was signed with Amazon Cognito credentials.
\$context.identity.cognitoIdentityPoolId	The Amazon Cognito identity pool ID of the caller making the request. Available only if the request was signed with Amazon Cognito credentials.
\$context.identity.sourceIp	The IP address of the API caller as determined by the X-Forwarded-For header. Note that this should not be used for security purposes. To get the full value of X-Forwarded-For addresses you can define the X-Forwarded-For header in your API and then map it to your integration.
\$context.identity.user	The principal identifier of the user making the request.
\$context.identity.userAgent	The User Agent of the API caller.
\$context.identity.userArn	The Amazon Resource Name (ARN) of the effective user identified after authentication.
\$context.requestId	An automatically generated ID for the API call.
\$context.resourceId	The identifier API Gateway assigns to your resource.
\$context.resourcePath	The path to your resource. For more information, see Build and Deploy an API Gateway API Step by Step (p. 13) .
\$context.stage	The deployment stage of the API call (for example, Beta or Prod).

Example

You may want to use the \$context variable if you're using AWS Lambda as the target backend that the API method calls. For example, you may want to perform two different actions depending on whether the stage is in Beta or in Prod.

Context Variables Template Example

The following example shows how to get context variables:

```
{
    "stage" : "$context.stage",
    "request_id" : "$context.requestId",
    "api_id" : "$context.apiId",
    "resource_path" : "$context.resourcePath",
    "resource_id" : "$context.resourceId",
    "http_method" : "$context.httpMethod",
    "source_ip" : "$context.identity.sourceIp",
    "user-agent" : "$context.identity.userAgent",
    "account_id" : "$context.identity.accountId",
```

```

"api_key" : "$context.identity.apiKey",
"caller" : "$context.identity.caller",
"user" : "$context.identity.user",
"user_arn" : "$context.identity.userArn"
}

```

Accessing the \$input Variable

The `$input` variable represents the input payload and parameters to be processed by your template. It provides four functions:

Function Reference

Variable and Function	Description
<code>\$input.body</code>	Returns the raw payload as a string.
<code>\$input.json(x)</code>	This function evaluates a JSONPath expression and returns the results as a JSON string. For example, <code>\$input.json('\$.pets')</code> will return a JSON string representing the pets structure. For more information about JSONPath, see JSONPath or JSONPath for Java .
<code>\$input.params()</code>	Returns a map of all the request parameters of your API call.
<code>\$input.params(x)</code>	Returns the value of a method request parameter from the path, query string, or header value (in that order) given a parameter name string x.
<code>\$input.path(x)</code>	Takes a JSONPath expression string (x) and returns an object representation of the result. This allows you to access and manipulate elements of the payload natively in Apache Velocity Template Language (VTL) . For example, <code>\$input.path('\$.pets').size()</code> For more information about JSONPath, see JSONPath or JSONPath for Java .

Examples

You may want to use the `$input` variable to get query strings and the request body with or without using models. You may also want to get the parameter and the payload, or a subsection of the payload, into your AWS Lambda function. The examples below show how to do this.

Example JSON Mapping Template

The following example shows how to use a mapping to read a name from the query string and then include the entire POST body in an element:

```
{  
    "name" : "$input.params('name')",  
    "body" : $input.json('$')  
}
```

If the JSON input contains unescaped characters that cannot be parsed by JavaScript, a 400 response may be returned. Applying `$util.escapeJavaScript($input.json('$'))` above will ensure that the JSON input can be parsed properly.

Example Inputs Mapping Template

The following example shows how to pass a JSONPath expression to the `json()` method. You could also read a specific property of your request body object by using a period (`.`), followed by your property name:

```
{  
    "name" : "$input.params('name')",  
    "body" : $input.json('$..mykey')  
}
```

If a method request payload contains unescaped characters that cannot be parsed by JavaScript, you may get 400 response. In this case, you need to call `$util.escapeJavaScript()` function in the mapping template, as shown as follows:

```
{  
    "name" : "$input.params('name')",  
    "body" : $util.escapeJavaScript($input.json('$..mykey'))  
}
```

Param Map Template Example

The following example shows how to map parameters:

```
{  
    "all-params" : "$input.params()",  
    "path-params" : "$input.params().path",  
    "query-params" : "$input.params().querystring",  
    "header-params" : "$input.params().header",  
    "header-param-names" : "$input.params().header.keySet()",  
    "content-type-value" : "$input.params().header.get('Content-Type')"  
}
```

Example Request and Response

Here's an example that uses all three functions:

Request Template:

```
Resource: /things/{id}  
  
With input template:  
{  
    "id" : "$input.params('id')",  
    "count" : "$input.path('$..things').size()",  
}
```

```

        "things" : $util.escapeJavaScript($input.json('$.things'))
    }

POST /things/abc
{
    "things" : {
        "1" : {},
        "2" : {},
        "3" : {}
    }
}

```

Response:

```
{
    "id": "abc",
    "count": "3",
    "things": {
        "1": {},
        "2": {},
        "3": {}
    }
}
```

For more mapping examples, see [Set Up Request and Response Payload Mappings \(p. 68\)](#)

Accessing the \$stageVariables Variable

The syntax for inserting a stage variable looks like this: `$stageVariables`.

\$stageVariables Reference

Syntax	Description
<code>\$stageVariables.<variable_name></code>	<code><variable_name></code> represents a stage variable name.
<code>\$stageVariables['<variable_name>']</code>	<code><variable_name></code> represents any stage variable name.
<code>\$(stageVariables['<variable_name>'])</code>	<code><variable_name></code> represents any stage variable name.

Accessing the \$util Variable

The `$util` variable contains utility functions for use in mapping templates.

Note

Unless otherwise specified, the default character set is UTF-8.

\$util Variable Reference

Function	Description
<code>\$util.escapeJavaScript()</code>	Escapes the characters in a string using JavaScript string rules.
<code>\$util.parseJson()</code>	<p>Takes "stringified" JSON and returns an object representation of the result. You can use the result from this function to access and manipulate elements of the payload natively in Apache Velocity Template Language (VTL). For example, if you have the following payload:</p> <pre>{ "errorMessage": { "key1": "var1", "key2": { "arr": [1,2,3] } }}</pre> <p>and use the following mapping template</p> <pre>#set (\$errorMessageObj = \$util.parseJson(\$input.path('\$errorMessage'))) { "errorMessageObjKey2ArrVal" : \$errorMessageObj.key2.arr[0] }</pre> <p>You will get the following output:</p> <pre>{ "errorMessageObjKey2ArrVal" : 1 }</pre>
<code>\$util.urlEncode()</code>	Converts a string into "application/x-www-form-urlencoded" format.
<code>\$util.urlDecode()</code>	Decodes an "application/x-www-form-urlencoded" string.
<code>\$util.base64Encode()</code>	Encodes the data into a base64-encoded string.
<code>\$util.base64Decode()</code>	Decodes the data from a base64-encoded string.

Enable CORS for a Resource through a Method in API Gateway

To serve requests of a resource of your API originated from a domain other than the API's own domain, you must enable cross-origin resource sharing (CORS) for selected methods on the resource. This amounts to having your API to respond to the `OPTIONS` preflight request with at least the following CORS-required response headers:

- `Access-Control-Allow-Methods`

- Access-Control-Allow-Headers
- Access-Control-Allow-Origin

In API Gateway you do this by setting up an OPTIONS request with the MOCK integration and then applying header mappings to map the above headers (with to-be-specified static values) to the method response headers.

This section describes how to enable CORS for a method in API Gateway using the API Gateway console or the API Gateway [Import API](#) (p. 117).

Topics

- [Prerequisites](#) (p. 102)
- [Enable CORS on a Resource Using the API Gateway Console](#) (p. 102)
- [Enable CORS for a Resource Using the API Gateway Import API](#) (p. 103)

Prerequisites

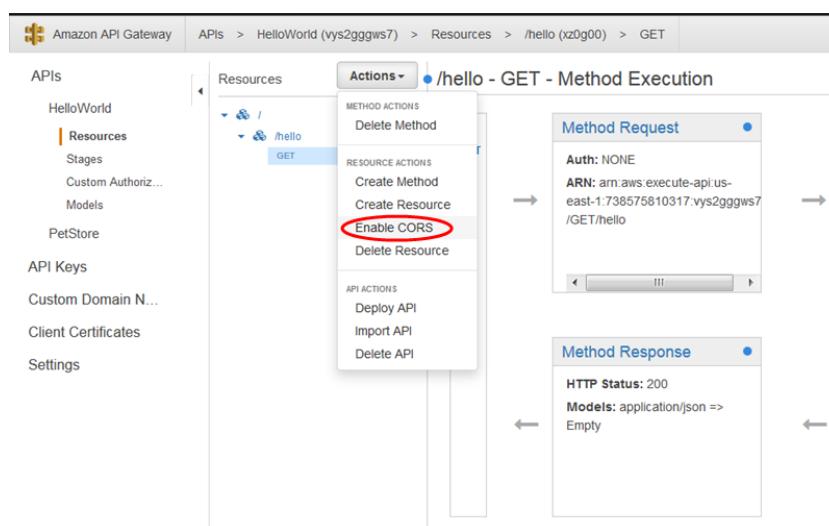
- You must have the method available in API Gateway. For instructions on how to create and configure a method, see [Build and Deploy an API Gateway API Step by Step](#) (p. 13).

Enable CORS on a Resource Using the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the API Gateway console, choose an API under **APIs**.
3. Choose a resource under **Resources**. This will enable CORS for all the methods on the resource.

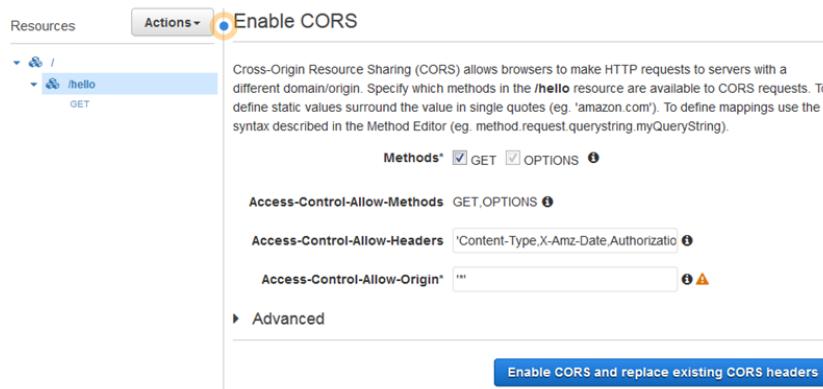
Alternatively, you could choose a method under the resource to enable CORS for just this method.

4. Choose **Enable CORS** from the **Actions** drop-down menu.

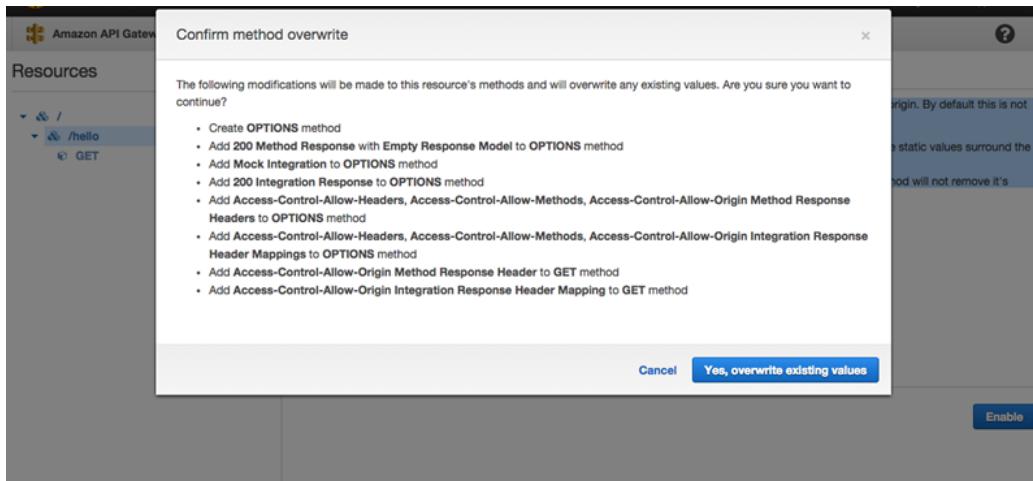


5. In the **Enable CORS** form, do the following:

- a. In the **Access-Control-Allow-Headers** input field, type a static string of a comma-separated list of headers that the client must submit in the actual request of the resource. Use the console-provided header list of
 'Content-Type,X-Amz-Date,Authorization,X-Api-Key,X-Amz-Security-Token' or specify your own headers.
- b. Use the console-provided value of '*' as the **Access-Control-Allow-Origin** header value to allow access requests from all domains, or specify a named domain to all access requests from the specified domain.
- c. Choose **Enable CORS and replace existing CORS headers**.



6. In **Confirm method changes**, choose **Yes, overwrite existing values** to confirm the new CORS settings.



7. The API Gateway console makes all the required changes and displays the summary screen.

Enable CORS for a Resource Using the API Gateway Import API

If you are using the [API Gateway Import API \(p. 117\)](#), you can set up CORS support using a Swagger file. You must first define an OPTIONS method in your resource that returns the required headers.

Note

Web browsers expect Access-Control-Allow-Headers, and Access-Control-Allow-Origin headers to be set up in each API method that accepts CORS requests. In addition, some browsers first make an HTTP request to an OPTIONS method in the same resource, and then expect to receive the same headers.

The following example creates an OPTIONS method and specifies mock integration. For more information, see [Configure Mock Integration for a Method \(p. 65\)](#).

```
/users
  options:
    summary: CORS support
    description: |
      Enable CORS by returning correct headers
    consumes:
      - application/json
    produces:
      - application/json
    tags:
      - CORS
    x-amazon-apigateway-integration:
      type: mock
      requestTemplates:
        application/json: |
          {
            "statusCode" : 200
          }
      responses:
        "default":
          statusCode: "200"
          responseParameters:
            method.response.header.Access-Control-Allow-Headers : "'Content-Type,X-Amz-Date,Authorization,X-Api-Key'"
            method.response.header.Access-Control-Allow-Methods : "'*'"
            method.response.header.Access-Control-Allow-Origin : "'*'"

        responseTemplates:
          application/json: |
            {}

      responses:
        200:
          description: Default response for CORS method
          headers:
            Access-Control-Allow-Headers:
              type: "string"
            Access-Control-Allow-Methods:
              type: "string"
            Access-Control-Allow-Origin:
              type: "string"
```

Once you have configured the OPTIONS method for your resource, you can add the required headers to the other methods in the same resource that need to accept CORS requests.

1. Declare the **Access-Control-Allow-Origin** and **Headers** to the response types.

```
responses:
  200:
    description: Default response for CORS method
```

```
headers:  
  Access-Control-Allow-Headers:  
    type: "string"  
  Access-Control-Allow-Methods:  
    type: "string"  
  Access-Control-Allow-Origin:  
    type: "string"
```

2. In the `x-amazon-apigateway-integration` tag, set up the mapping for those headers to your static values:

```
responses:  
  "default":  
    statusCode: "200"  
    responseParameters:  
      method.response.header.Access-Control-Allow-Headers : "'Content-Type,X-Amz-Date,Authorization,X-Api-Key'"  
      method.response.header.Access-Control-Allow-Methods : "'*'"  
      method.response.header.Access-Control-Allow-Origin : " '*'"
```

Use an API Key in API Gateway

You can use an API key in API Gateway to enable an API's methods. When you enable an API key, callers must supply it as part of the call.

Topics

- [Prerequisites \(p. 105\)](#)
- [Use an API Key with the API Gateway Console \(p. 105\)](#)

Prerequisites

1. You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).
2. You must have deployed the API in API Gateway at least once. Follow the instructions in [Deploying an API \(p. 199\)](#).

Use an API Key with the API Gateway Console

To enable an API key with the API Gateway console, follow these instructions:

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. Choose the **GET** method under a resource of your choosing.
3. Choose the **Method Request** box
4. If **API Key Required** is set to **false**, choose the pencil icon next to it. From the drop-down menu list, choose **true**. Finally, choose the check-mark icon to save the setting.

Note

The steps above configures the API Gateway to enforce using API key. Otherwise, the API key created following the instructions below will not be used.

5. In the secondary navigation bar, in the first list next to the console home button, choose **API Keys**.

6. Choose **Create API Key**.
7. For **Name**, type a name for the API key entry.
8. (Optional) For **Description**, type a description for the API key entry.
9. To enable the API key, select **Enabled**.
10. Choose **Save**. Make a note of the key displayed in **API key**.
11. For **API Stage Association**, for **Select API**, choose the name of the API.
12. For **Select stage**, choose the name of the stage.
13. Choose **Add**, and then choose **Save**.
14. Callers must now add to each call a custom header named `x-api-key`, along with the value of the API key. For example, if the API key value is `bkayZOMvuy8aZOhIgxq94K9Oe7Y70Hw55`, the custom header would be as follows:

```
x-api-key: bkayZOMvuy8aZOhIgxq94K9Oe7Y70Hw55
```

Note

In addition to, or instead of, enabling an API key, you can restrict access to certain IAM users only. For instructions, see [Configure How a User Calls an API Method \(p. 62\)](#).

Enable Client-Side SSL Authentication of an API with the API Gateway Console

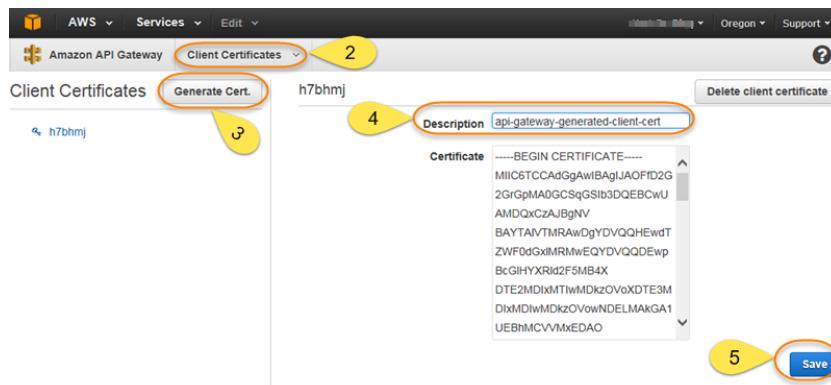
You can use API Gateway to generate a SSL certificate and configure the certificate for an API to be authenticated by the HTTP back end for all such calls, from any client, through the API. SSL Certificates are self-signed and only the public key of the certificate is visible in the API Gateway console and through the APIs.

Topics

- [Generate a Client Certificate \(p. 106\)](#)
- [Configure an API to Use SSL Certificates \(p. 107\)](#)
- [Test Invoke \(p. 107\)](#)
- [Configure Back End to Authenticate API \(p. 108\)](#)

Generate a Client Certificate

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the secondary navigation bar, choose **Client Certificates**.
3. From **Client Certificates** pane, choose **Generate Cert**.
4. Optionally, For **Description**, enter a short descriptive title for the generated certificate. API Gateway generates a new certificate and returns the new certificate GUID, along with the PEM-encoded public key.
5. Choose the **Save** button to save the certificate to API Gateway.

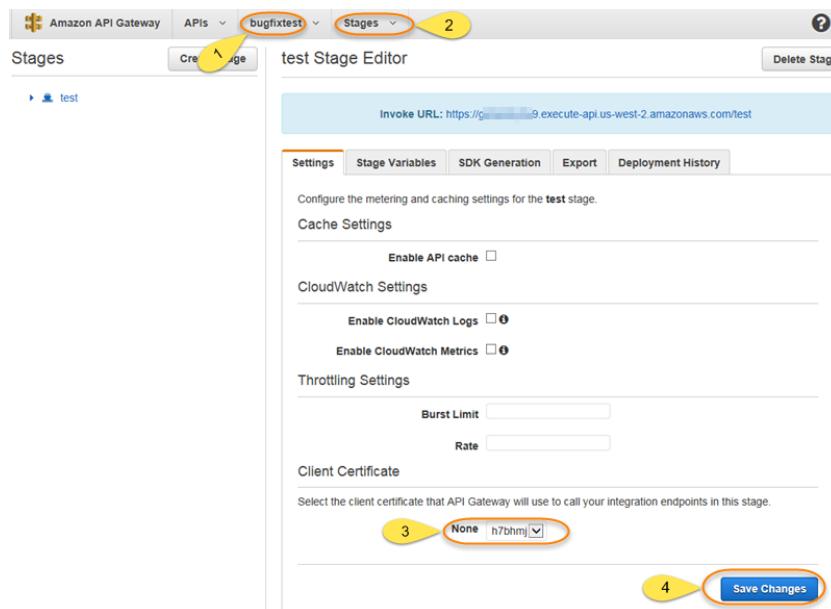


You are now ready to configure an API to use the certificate.

Configure an API to Use SSL Certificates

These instructions assume you have already completed [Generate a Client Certificate \(p. 106\)](#).

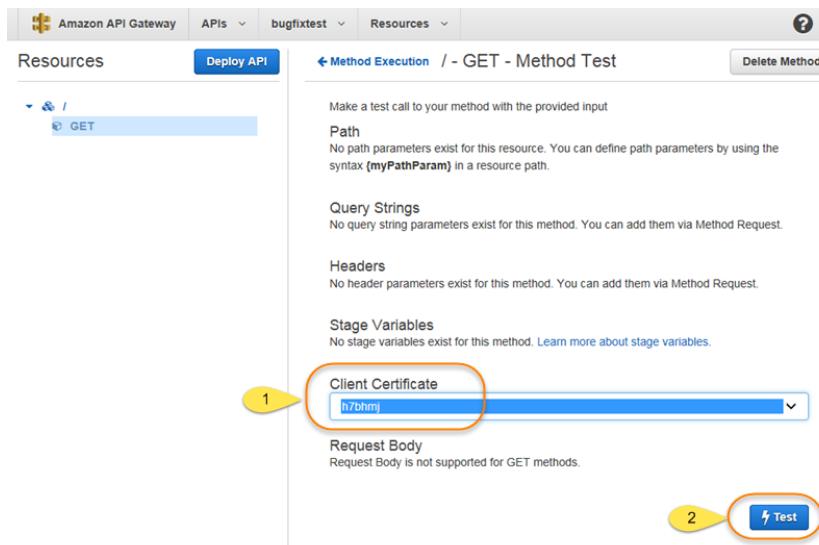
1. In the API Gateway console, create or open an API for which you want to use the client certificate. Make sure the API has been deployed to a stage.
2. In the secondary navigation bar, choose **Stages**.
3. In the **Stage Editor** panel, select a certificate under the **Client Certificate** section.
4. Choose the **Save Changes** button to save the settings.



Once a certificate is selected and saved, API Gateway will use the certificate for all calls to HTTP integrations in your API.

Test Invoke

1. From the **Method Execution**, of an API method, select a certificate
2. Choose **Test** to invoke the method request.



API Gateway will present the chosen SSL certificate for the HTTP back end to authenticate the API.

Configure Back End to Authenticate API

These instructions assume you have already completed [Generate a Client Certificate \(p. 106\)](#) and [Configure an API to Use SSL Certificates \(p. 107\)](#).

When receiving HTTPS requests from API Gateway, your back end can authenticate your API using the PEM-encoded certificate generated by API Gateway, provided that the back end is properly configured. Most Web servers can be easily configured to do so.

For example, in Node.js you can use the [https](#) module to create a HTTPS back end and use the [client-certificate-auth](#) modules to authenticate client requests with PEM-encoded certificates.

Enable Amazon API Gateway Custom Authorization

Topics

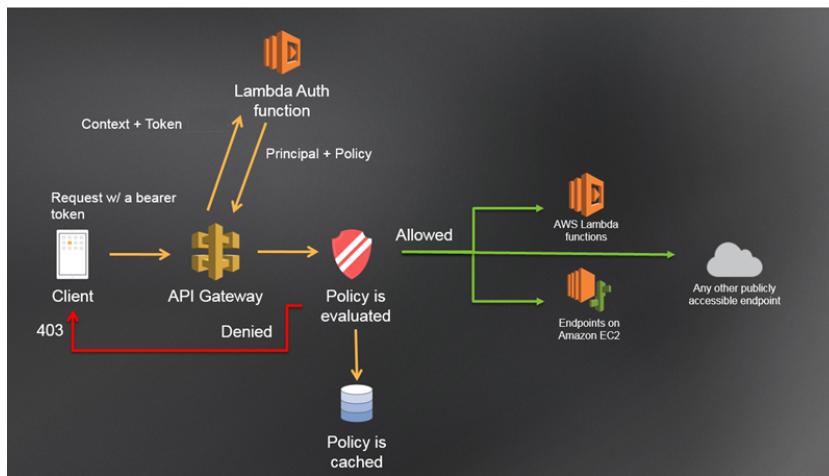
- [Amazon API Gateway Custom Authorization Overview \(p. 108\)](#)
- [Create the API Gateway Custom Authorizer Lambda Function \(p. 109\)](#)
- [Input to an Amazon API Gateway Custom Authorizer \(p. 111\)](#)
- [Output from an Amazon API Gateway Custom Authorizer \(p. 111\)](#)
- [Configure Custom Authorization Using the API Gateway Console \(p. 112\)](#)
- [Call an API Using API Gateway Custom Authorization \(p. 114\)](#)

Amazon API Gateway Custom Authorization Overview

With Amazon API Gateway custom authorization, you can control access to your APIs using bearer token authentication strategies, such as OAuth or SAML. To do so, you provide and configure a custom

authorizer, a Lambda function you own, for API Gateway to use to authorize the client requests for the configured APIs.

When an API request is made, API Gateway verifies whether a custom authorizer is configured for the API. If so, API Gateway calls the Lambda function, supplying the authorization token extracted from a custom request header. You use this Lambda function to implement various authorization strategies, such as JSON Web Token (JWT) verification and OAuth provider callout, to return IAM policies that authorize the request. If the returned policy is invalid or the permissions are denied, the API call will not succeed. For a valid policy, API Gateway caches the returned policy, associated with the incoming token and used for the current and subsequent requests, over a pre-configured time-to-live (TTL) period of up to 3600 seconds. You can set the TTL period to zero seconds to disable the policy caching. The default TTL value is 300 seconds. Currently, the maximum TTL value of 3600 seconds cannot be increased.



Create the API Gateway Custom Authorizer Lambda Function

Before creating an API Gateway custom authorizer, you must first create the AWS Lambda function that implements the logic to authenticate and authorize the caller. You can do so in the Lambda console, using the code template available from the API Gateway Custom Authorizer blueprint. Or you can create one from scratch. For illustration purposes, we will explain here the creation of the Lambda function without using the blueprint.

Note

The custom authorizer Lambda function presented here is for illustration purposes. In production code, you should follow the API Gateway Custom Authorizer blueprint to implement your authorizer Lambda function.

When creating the Lambda function for your API Gateway custom authorizer, you will be asked to assign an execution role for the Lambda function if it calls other AWS services. For the following example, the basic `AWSLambdaRole` will suffice. For more involved use cases, follow the [instructions](#) to grant permissions in an execution role for the Lambda function.

In the code editor of the Lambda console, enter the following Node.js code.

```
console.log('Loading function');

exports.handler = function(event, context) {
    var token = event.authorizationToken;
```

```

// Call oauth provider, crack jwt token, etc.
// In this example, the token is treated as the status for simplicity.

switch (token) {
    case 'allow':
        context.succeed(generatePolicy('user', 'Allow', event.methodArn));

        break;
    case 'deny':
        context.succeed(generatePolicy('user', 'Deny', event.methodArn));
        break;
    case 'unauthorized':
        context.fail("Unauthorized");
        break;
    default:
        context.fail("error");
}
};

var generatePolicy = function(principalId, effect, resource) {
    var authResponse = {};
    authResponse.principalId = principalId;
    if (effect && resource) {
        var policyDocument = {};
        policyDocument.Version = '2012-10-17'; // default version
        policyDocument.Statement = [];
        var statementOne = {};
        statementOne.Action = 'execute-api:Invoke'; // default action
        statementOne.Effect = effect;
        statementOne.Resource = resource;
        policyDocument.Statement[0] = statementOne;
        authResponse.policyDocument = policyDocument;
    }
    return authResponse;
}

```

The preceding Lambda function returns an Allow IAM policy on a specified method if the request's authorization token contains an 'allow' value, thereby permitting a caller to invoke the specified method. The function returns a Deny policy against the specified method if the authorization token has a 'deny' value, thus blocking the caller from calling the method. If the token is 'unauthorized', the client will receive a 401 Forbidden response. If the value that returns is 'fail' or anything else, the function returns a 500 Internal Server Error error message. In both of the last two cases, the calls will not succeed.

Note

In production code, you may need to authenticate the user before granting authorizations. If so, you can add authentication logic in the Lambda function as well. Consult the provider-specific documentation for instructions on how to call such an authentication provider.

Before going further, you may want to test the Lambda function from within the Lambda Console. To do this, configure the sample event to provide the input and verify the result by examining the output. The next two sections explain the [Input to a Custom Authorizer \(p. 111\)](#) and [Output from a Custom Authorizer \(p. 111\)](#).

Input to an Amazon API Gateway Custom Authorizer

When a custom authorizer is enabled on an API method, you must specify a custom header for the method caller to pass the required authorization token in the initial client request. Upon receiving the request, API Gateway extracts the token from the custom header as the input `authorizationToken` parameter value into the Lambda function and calls the custom authorizer with the following request payload.

```
{  
    "type": "TOKEN",  
    "authorizationToken": "<caller-supplied-token>",  
    "methodArn": "arn:aws:execute-api:<regionId>:<accountId>:<apiId>/<stage>/<methodArn>/<resourcePath>"  
}
```

In this example, the `type` property specifies the payload type. Currently, the only valid value is the `TOKEN` literal. The `<caller-supplied-token>` originates from the custom authorization header in a client request. The `methodArn` is the ARN of the submitted method request and is populated by API Gateway in accordance with the custom authorizer configuration.

For the custom authorizer shown in the preceding section, the `<caller-supplied-token>` string is `allow`, `deny`, `unauthorized`, or any other string value. An empty string value is the same as `unauthorized`. The following shows an example of such an input to obtain an `Allow` policy on the `GET` method of an API (`ymy8tbxw7b`) of the AWS account (123456789012) in any stage (*).

```
{  
    "type": "TOKEN",  
    "authorizationToken": "allow",  
    "methodArn": "arn:aws:execute-api:us-west-2:123456789012:ymy8tbxw7b/*/GET/"  
}
```

Output from an Amazon API Gateway Custom Authorizer

The custom authorizer's Lambda function must return a response that includes the principal identifier (`principalId`) and a policy document containing a list of policy statements. The following shows an example of a response.

```
{  
    "principalId": "xxxxxxxx", // The principal user identification associated  
    // with the token send by the client.  
    "policyDocument": {  
        "Version": "2012-10-17",  
        "Statement": [  
            {  
                "Action": "execute-api:Invoke",  
                "Effect": "Allow|Deny",  
                "Resource": "arn:aws:execute-api:us-west-2:123456789012:ymy8tbxw7b/*/  
            }  
        ]  
    }  
}
```

```
        "Resource": "arn:aws:execute-api:<regionId>:<accountId>:<ap  
pId>/<stage>/<httpVerb>/[<resource>/<httpVerb>/[...]]"  
    }  
]  
}  
}
```

Here, a policy statement stipulates whether to allow or deny (`Effect`) the API Gateway execution service to invoke (`Action`) the specified API method (`Resource`). You can use a wild card (*) to specify a resource type (`method`).

You can access the `principalId` value in a mapping template using the `$context.authorizer.principalId` variable. This is useful if you want to pass the value to the back end. For more information, see [Accessing the \\$context Variable \(p. 96\)](#).

For information about valid policies for calling an API, see [Permissions to call an API in API Gateway \(p. 240\)](#).

The following shows example output from the example custom authorizer. The example output contains a policy statement to block (`Deny`) calls to the `GET` method in an API (`ymy8tbxw7b`) of an AWS account (123456789012) in any stage (*).

```
{  
  "principalId": "user",  
  "policyDocument": {  
    "Version": "2012-10-17",  
    "Statement": [  
      {  
        "Action": "execute-api:Invoke",  
        "Effect": "Deny",  
        "Resource": "arn:aws:execute-api:us-west-  
2:123456789012:ymy8tbxw7b/*/GET/"  
      }  
    ]  
  }  
}
```

Configure Custom Authorization Using the API Gateway Console

After you create the Lambda function and verify that it works, you can configure the API Gateway Custom Authorizer in the API Gateway console.

Create a Custom Authorization for API Methods

1. Sign in to the API Gateway console.
2. Create a new or select an existing API and choose **Custom Authorizers**.
3. Under **New Custom Authorizer**, do the following:
 - In **Name**, choose a name for your new custom authorization.
 - In **Lambda region**, select the region where you upload your custom authorizer's Lambda function.

- In **Lambda function**, select the Lambda function for your custom authorizer.

Note

You must first create a custom authorizer Lambda function in the region for it to be available in the drop-down list.

- In **Identity token source**, type the mapping expression for your authorizer's custom header.

Note

The custom header mapping expression is of the `method.request.header.<name>` format, where `<name>` is the name of a custom authorization header submitted as part of the client request. In the following example, this custom header name is `Auth`.

- In **Token validation expression**, you can optionally provide a RegEx statement for API Gateway to validate the input token before calling the custom authorizer Lambda function. This helps you avoid or reduce the chances of being charged for processing invalid tokens.
- In **Result TTL in seconds**, you can change or use the default (300) value to enable caching (>0) or disable caching (=0) of the policy returned from the Lambda function.

Note

The policy caching uses a cache key generated from the supplied token for the targeted API and custom authorizer in a specified stage. To enable caching, your authorizer must return a policy that is applicable to all methods across an API. To enforce method-specific policy, you can set the TTL value to zero to disable policy caching for the API.

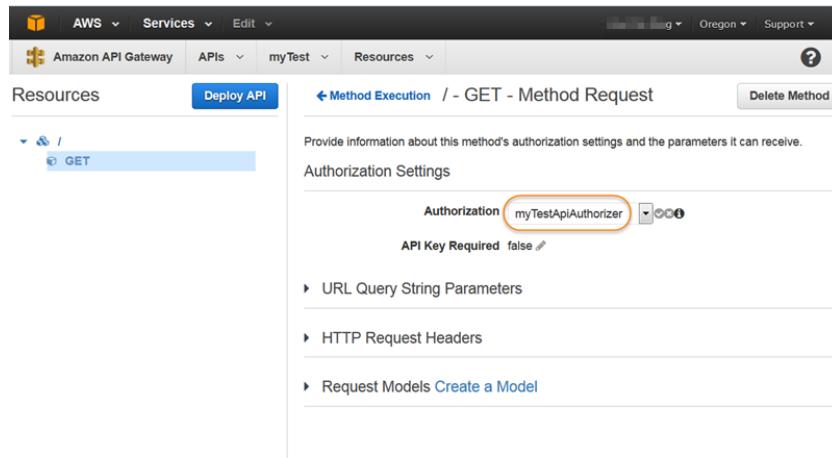
The screenshot shows the 'Create New Custom Authorizer' dialog in the AWS Management Console. The 'Name' field is set to 'myTestApiAuthorizer'. The 'Lambda region' is 'us-west-2'. The 'Lambda function' is 'myCustomAuthorizer'. The 'Identity token source' is 'method.request.header.Auth'. The 'Result TTL in seconds' is '300'. The 'Create custom authorization' button is highlighted with an orange border.

4. In **Add Permission to Lambda Function**, choose **OK**. After the custom authorization is created, you can test it with appropriate authorization token values to verify that it works as expected.

This completes the procedure to create a custom authorization. The next procedure shows how to configure an API method to use the custom authorizer.

Configure an API Method to Use the Custom Authorizer

1. Go back to the API. Create a new method or choose an existing method. If necessary, create a new resource.
2. In **Method Execution**, choose the **Method Request** link.
3. In **Authorization Settings**, choose the pencil icon to edit the **Authorization** field. Select the custom authorization you just created (**myTestApiAuthorizer**), and then choose the checkmark icon to save the choice.



4. Optionally, while still on the **Method Request** page, choose **Add header** if you also want to pass the custom authorization header to the back end. In **Name**, type a custom header name that matches the header mapping expression you used when you created the custom authorization, and then choose the checkmark icon to save the settings.
5. Choose **Deploy API** to deploy the API to a stage. Make a note of the **Invoke URL** value. You will need it when calling the API.

Call an API Using API Gateway Custom Authorization

After you configure your API to use the custom authorizer, you or your customers can call the API using the custom authorizer. Because it involves submitting a custom authorization token header in the requests, you need a REST client that supports this. In the following examples, API calls are made using the [Postman Chrome App](#).

Calling an API with Custom Authorization Tokens

1. Open the [Postman Chrome App](#), choose the **GET** method and paste the API's **Invoke URL** into the adjacent URL field.

Add the custom authorization token header and set the value to `allow`. Choose **Send**.

Amazon API Gateway Developer Guide

Call an API with Custom authorization

The screenshot shows the Postman interface with a GET request to `https://[REDACTED].execute-api.us-west-2.amazonaws.com/test`. The Authorization tab is selected, showing a header named "Auth" with the value "allow". The response status is 200 OK, and the JSON body contains a successful response message.

```
1 {  
2   "args": {},  
3   "headers": {  
4     "Accept": "application/json",  
5     "Host": "httpbin.org",  
6     "User-Agent": "AmazonAPIGateway_ymy8ttxw7b",  
7     "X-Amzn-Apigateway-Api-Id": "ymy8ttxw7b"  
8   },  
9   "origin": "54.186.57.107",  
10  "url": "http://httpbin.org/get"  
11 }
```

The response shows that the API Gateway custom authorizer returns a **200 OK** response and successfully authorizes the call to access the HTTP endpoint (`http://httpbin.org/get`) integrated with the method.

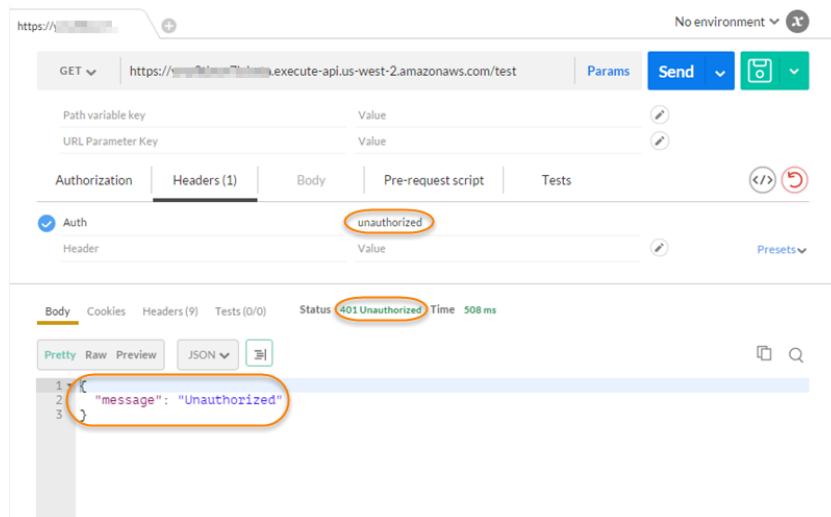
2. Still in Postman, change the custom authorization token header value to deny. Choose **Send**.

The screenshot shows the Postman interface with the same GET request and Authorization setup as the previous screenshot, but with the "Auth" value changed to "deny". The response status is 403 Forbidden, and the JSON body contains a denial message.

```
1 {  
2   "Message": "User is not authorized to access this resource"  
3 }
```

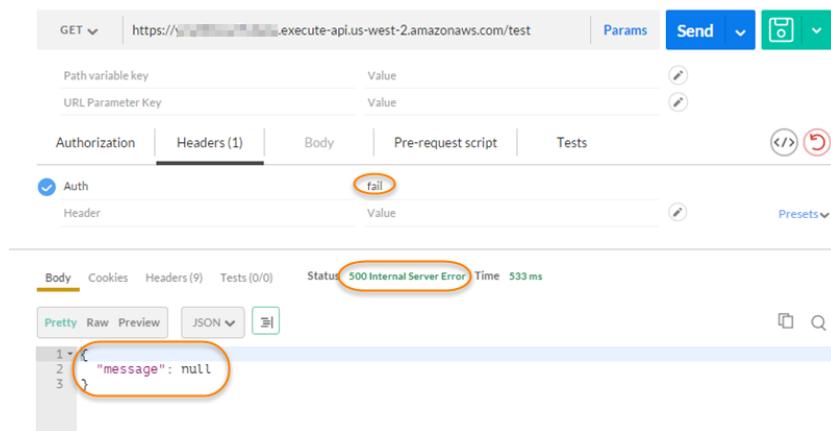
The response shows that the API Gateway custom authorizer returns a **403 Forbidden** response without authorizing the call to access the HTTP endpoint.

3. In Postman, change the custom authorization token header value to `unauthorized` and choose **Send**.



The response shows that API Gateway returns a **401 Unauthorized** response without authorizing the call to access the HTTP endpoint.

4. Now, change the custom authorization token header value to `fail`. Choose **Send**.



The response shows that API Gateway returns a **500 Internal Server Error** response without authorizing the call to access the HTTP endpoint.

Import and Export API Gateway API with Swagger Definition Files

As an alternative to using the Amazon API Gateway console to create and update your API, you can use the API Gateway Import API feature to upload API definitions into API Gateway from external API definition files, such as those using the [Swagger specification](#) with the [API Gateway extensions \(p. 122\)](#).

After an API is created and configured in API Gateway, you can download it as a Swagger definition file using the Amazon API Gateway Export API. The API Gateway console has enabled this feature for you to export an API using intuitive visual interfaces.

Topics

- Import an API into API Gateway (p. 117)
- Export an API from API Gateway (p. 120)
- API Gateway Extensions to Swagger (p. 122)

Import an API into API Gateway

You can use the API Gateway Import API feature to import an API from an external definition file into API Gateway. Currently, the Import API feature supports [Swagger v2.0](#) definition files.

With the Import API, you can either create a new API by submitting a [POST request](#) that includes a definition as the payload, or you can update an existing API by using a [PUT request](#) that contains a definition in the payload. You can update an API by overwriting it with a new definition, or merge a definition with an existing API. You specify the options in the request URL using a `mode` query parameter.

Note

For RAML API definitions, you can continue to use [API Gateway Importer](#).

Besides making explicit calls to the REST API, as described below, you can also use the Import API feature in the API Gateway console. The option is available as an item in the **Actions** drop-down menu. For an example of using the Import API feature from the API Gateway console, see [Build and Test an API Gateway API from an Example \(p. 5\)](#).

Use the Import API to Create a New API

To use the Import API feature to create a new API, POST your API definition file to <https://apigateway.<region>.amazonaws.com/restapis?mode=import>. This request results in a new `RestApi`, along with `Resources`, `Models`, and other items defined in the definition file.

The following code snippet shows an example of the POST request with the payload of a JSON-formatted Swagger definition:

```
POST /restapis?mode=import
Host:apigateway.<region>.amazonaws.com
Content-Type: application/json
Content-Length: ...


```

Swagger API definition in JSON (p. 162)

Use the Import API to Update an Existing API

You can use the Import API feature to update an existing API when there are aspects of that API you would like to preserve, such as stages and stage variables, or references to the API from API Keys.

An API update can occur in two modes: merge or overwrite. Merging an API is useful when you have decomposed your external API definitions into multiple, smaller parts and only want to apply changes from one of those parts at a time. For example, this might occur if multiple teams are responsible for different parts of an API and have changes available at different rates. In this mode, items from the existing API that are not specifically defined in the imported definition will be left alone.

Overwriting an API is useful when an external API definition contains the complete definition of an API. In this mode, items from an existing API that are not specifically defined in the imported definition will be deleted.

To merge an API, submit a PUT request to https://apigateway.<region>.amazonaws.com/restapis/<restapi_id>?mode=merge. The `restapi_id` path parameter value specifies the API to which the supplied API definition will be merged.

The following code snippet shows an example of the PUT request to merge a Swagger API definition in JSON, as the payload, with the specified API already in API Gateway.

```
PUT /restapis/<restapi_id>?mode=merge
Host:apigateway.<region>.amazonaws.com
Content-Type: application/json
Content-Length: ...
```

A Swagger API definition in JSON (p. 162)

The merging update operation takes two complete API definitions and merges them together. For a small and incremental change, you can use the [resource update](#) operation.

To overwrite an API, submit a PUT request to

https://apigateway.<region>.amazonaws.com/restapis/<restapi_id>?mode=overwrite. The `restapi_id` path parameter specifies the API that will be overwritten with the supplied API definitions.

The following code snippet shows an example of an overwriting request with the payload of a JSON-formatted Swagger definition:

```
PUT /restapis/<restapi_id>?mode=overwrite
Host:apigateway.<region>.amazonaws.com
Content-Type: application/json
Content-Length: ...
```

A Swagger API definition in JSON (p. 162)

When the `mode` query parameter is not specified, merge is assumed.

Note

The PUT operations are idempotent, but not atomic. That means if a system error occurs part way through processing, the API can end up in a bad state. However, repeating the operation will put the API into the same final state as if the first operation had succeeded.

Swagger `basePath`

In Swagger, you can use the `basePath` property to provide one or more path parts that precede each path defined in the `paths` property. Because API Gateway has several ways to express a resource's path, the Import API feature provides three options for interpreting the `basePath` property during an import:

ignore

If the Swagger file has a `basePath` value of `"/a/b/c"` and the `paths` property contains `"/e"` and `"/f"`, the following POST or PUT request:

```
POST /restapis?mode=import&basepath=ignore
```

```
PUT /restapis/api_id?basepath=ignore
```

will result in the following resources in the API:

- /
- /e
- /f

The effect is to treat the `basePath` as if it was not present, and all of the declared API resources are served relative to the host. This can be used, for example, when you have a custom domain name with an API mapping that does not include a *Base Path* and a *Stage* value that refers to your production stage.

Note

API Gateway will automatically create a root resource for you, even if it is not explicitly declared in your definition file.

prepend

If the Swagger file has a `basePath` value of `"/a/b/c"` and the `paths` property contains `"/e"` and `"/f"`, the following POST or PUT request:

```
POST /restapis?mode=import&basepath=prepend
```

```
PUT /restapis/api_id?basepath=prepend
```

will result in the following resources in the API:

- /
- /a
- /a/b
- /a/b/c
- /a/b/c/e
- /a/b/c/f

The effect is to treat the `basePath` as specifying additional resources (without methods) and to add them to the declared resource set. This can be used, for example, when different teams are responsible for different parts of an API and the `basePath` could reference the path location for each team's API part.

Note

API Gateway will automatically create intermediate resources for you, even if they are not explicitly declared in your definition.

split

If the Swagger file has a `basePath` value of `"/a/b/c"` and the `paths` property contains `"/e"` and `"/f"`, the following POST or PUT request:

```
POST /restapis?mode=import&basepath=split
```

```
PUT /restapis/api_id?basepath=split
```

will result in the following resources in the API:

- /
- /b
- /b/c
- /b/c/e
- /b/c/f

The effect is to treat top-most path part, " /a ", as the beginning of each resource's path, and to create additional (no method) resources within the API itself. This could, for example, be used when "a" is a stage name that you want to expose as part of your API.

Errors during Import

During the import, errors can be generated for major issues like an invalid Swagger document. Errors are returned as exceptions (e.g., `BadRequestException`) in an unsuccessful response. When an error occurs, the new API definition is discarded and no change is made to the existing API.

Warnings during Import

During the import, warnings can be generated for minor issues like a missing model reference. If a warning occurs, the operation will continue if the `failonwarnings=false` query expression is appended to the request URL. Otherwise, the updates will be rolled back. By default, `failonwarnings` is set to `false`. In such cases, warnings are returned as a field in the resulting [RestAPI resource](#). Otherwise, warnings are returned as a message in the exception.

Export an API from API Gateway

Once you created and configured an API in API Gateway, using the API Gateway console or otherwise, you can export it to a Swagger file using the API Gateway Export API, which is part of the Amazon API Gateway Control Service. You have options to include the API Gateway integration extensions, as well as the [Postman](#) extensions, in the exported Swagger definition file.

Request to Export an API

With the Export API, you export an existing API by submitting a GET request, specifying the to-be-exported API as part of URL paths. The request URL is of the following format:

```
https://<host>/restapis/<restapi_id>/stages/<stage_name>/exports/swagger
```

You can append the `extensions` query string to specify whether to include API Gateway extensions (with the `integration` value) or Postman extensions (with the `postman` value).

In addition, you can set the `Accept` header to `application/json` or `application/yaml` to receive the API definition output in JSON or YAML format, respectively.

For more information about submitting GET requests using the API Gateway Export API, see [Making HTTP Requests](#).

Download API Swagger Definition in JSON

To export and download an API in Swagger definitions in JSON format:

```
GET /restapis/<restapi_id>/stages/<stage_name>/exports/swagger
Host: apigateway.<region>.amazonaws.com
Accept: application/json
```

Here, `<region>` could be, for example, `use-east-1`. For all the regions where API Gateway is available, see [Regions and Endpoints](#)

Download API Swagger Definition in YAML

To export and download an API in Swagger definitions in YAML format:

```
GET /restapis/<restapi_id>/stages/<stage_name>/exports/swagger
Host: apigateway.<region>.amazonaws.com
Accept: application/yaml
```

Download API Swagger Definition with Postman Extensions in JSON

To export and download an API in Swagger definitions with the Postman extension in JSON format:

```
GET /restapis/<restapi_id>/stages/<stage_name>/exports/swagger?extensions=postman
Host: apigateway.<region>.amazonaws.com
Accept: application/json
```

Download API Swagger Definition with API Gateway Integration in YAML

To export and download an API in Swagger definitions with API Gateway integration in YAML format:

```
GET /restapis/<restapi_id>/stages/<stage_name>/exports/swagger?extensions=integration
Host: apigateway.<region>.amazonaws.com
Accept: application/yaml
```

Export API Using the API Gateway Console

From the **stage configuration** page in the API Gateway console, choose the **Export** tab and then one of the available options (**Export as Swagger**, **Export as Swagger + API Gateway Integrations** and **Export as Postman**) to download your API's Swagger definition.



API Gateway Extensions to Swagger

The API Gateway extensions support the AWS-specific authorization and API Gateway-specific API integrations. In this section, we will describe the API Gateway extensions to the Swagger specification.

Topics

- [x-amazon-apigateway-auth Object \(p. 122\)](#)
- [x-amazon-apigateway-authorizer Object \(p. 123\)](#)
- [x-amazon-apigateway-authtype Property \(p. 124\)](#)
- [x-amazon-apigateway-integration Object \(p. 125\)](#)
- [x-amazon-apigateway-integration.requestTemplates Object \(p. 127\)](#)
- [x-amazon-apigateway-integration.requestParameters Object \(p. 128\)](#)
- [x-amazon-apigateway-integration.responses Object \(p. 129\)](#)
- [x-amazon-apigateway-integration.response Object \(p. 130\)](#)
- [x-amazon-apigateway-integration.responseTemplates Object \(p. 131\)](#)
- [x-amazon-apigateway-integration.responseParameters Object \(p. 131\)](#)

x-amazon-apigateway-auth Object

Enables AWS authorization for calling the specified API. This object functions as an extended property of the [Swagger Operation](#) object.

Properties

Property Name	Type	Description
type	string	Authorization type. Set it to <code>aws_iam</code> to use AWS role- or resource-based authorization. Otherwise, set it to <code>none</code> .

x-amazon-apigateway-auth Example

The following example sets the API to use AWS [Signature Version 4](#) authentication.

```
"x-amazon-apigateway-auth" : {
    "type" : "aws_iam"
}
```

x-amazon-apigateway-authorizer Object

Defines a custom authorizer to be applied for authorization of method invocations in API Gateway. This object is an extended property of the [Swagger Security Definitions Operation](#) object.

Properties

Property Name	Type	Description
type	string	The type of the authorizer. This is a required property and the value must be "token".
authorizerUri	string	<p>The Uniform Resource Identifier (URI) of the authorizer (a Lambda function). For example,</p> <div style="border: 1px solid black; padding: 5px; margin-top: 10px;"> <code>"arn:aws:apigateway:us-east-1:lambda:path/2015-03-31/functions/arn:aws:lambda:us-east-1:account-id:function:auth_function_anem/invocations"</code> </div>
authorizerCredentials	string	Credentials required for the authorizer, if any, in the form of an ARN of an IAM execution role. For example, "arn:aws:iam:: account-id : IAM_role ".
identityValidationExpression	string	A regular expression for validating the incoming identity. For example, " <code>^x-[a-z]+</code> ".
authorizerResultTtlInSeconds	string	The number of seconds during which the resulting IAM policy is cached.

x-amazon-apigateway-authorizer Example

The following Swagger security definitions example specifies a custom authorizer named `test-authorizer`.

```

"securityDefinitions" : {
    "test-authorizer" : {
        "type" : "apiKey", // Required and the value must
        be "apiKey" for an API Gateway API.
        "name" : "Authorization", // The source header name
        identifying this authorizer.
        "in" : "header", // Required and the value must
        be "header" for an API Gateway API.
        "x-amazon-apigateway-authtype" : "oauth2", // For human readability only,
        ignored by API Gateway.
        "x-amazon-apigateway-authorizer" : { // An API Gateway custom au
        thorizer definition
            "type" : "token", // Required property and the
            value must "token"
            "authorizerUri" : "arn:aws:apigateway:us-east-1:lambda:path/2015-03-
            31/functions/arn:aws:lambda:us-east-1:account-id:function:function-name/invoca
            tions",
            "authorizerCredentials" : "arn:aws:iam::account-id:role",
            "identityValidationExpression" : "^x-[a-z]+",
            "authorizerResultTtlInSeconds" : 60
        }
    }
}

```

The following Swagger operation object snippet sets the GET /http to use the custom authorizer specified above.

```

"/http" : {
    "get" : {
        "responses" : { },
        "security" : [ {
            "test-authorizer" : [ ]
        }],
        "x-amazon-apigateway-integration" : {
            "type" : "http",
            "responses" : {
                "default" : {
                    "statusCode" : "200"
                }
            },
            "httpMethod" : "GET",
            "uri" : "http://api.example.com"
        }
    }
}

```

x-amazon-apigateway-authtype Property

Specify the type of a custom authorizer. It is parsed by the API Gateway API import and export features and for human readability only.

x-amazon-apigateway-authtype Example

The following example sets the type of a custom authorizer using OAuth 2.

This property is an extended property of the [Swagger Security Definitions Operation object](#).

```
"cust-authorizer" : {
    "type" : "...", // required
    "name" : "...", // name of the identity source header
    "in" : "header", // must be header
    "x-amazon-apigateway-authtype" : "oauth2", // a customer-supplied string,
    ignored by API Gateway.
    "x-amazon-apigateway-authorizer" : {
        ...
    }
}
```

x-amazon-apigateway-integration Object

Specifies details of the back-end integration used for this method. This extension is an extended property of the [Swagger Operation object](#). The result is an [API Gateway integration object](#).

Properties

Property Name	Type	Description
type	string	The type of integration with the specified back end. The valid value is <code>http</code> (for integration with an HTTP back end) or <code>aws</code> (for integration with AWS Lambda functions or other AWS services, such as DynamoDB, SNS or SQS).
uri	string	The endpoint URI of the back end. For integrations of the <code>aws</code> type, this is an ARN value. For the HTTP integration, this is the URL of the HTTP endpoint including the <code>https</code> or <code>http</code> scheme.
httpMethod	string	The HTTP method used in the integration request. For Lambda function invocations, the value must be <code>POST</code> .

Property Name	Type	Description
credentials	string	For AWS IAM role-based credentials, specify the ARN of an appropriate IAM role. If unspecified, credentials will default to resource-based permissions that must be added manually to allow the API to access the resource. For more information, see Granting Permissions Using a Resource Policy . Note: when using IAM credentials, please ensure that AWS STS regional endpoints are enabled for the region where this API is deployed for best performance.
requestTemplates	x-amazon-apigateway-integration.requestTemplates (p. 127)	Mapping templates for a request payload of specified MIME types.
requestParameters	x-amazon-apigateway-integration.requestParameters (p. 128)	Specifies mappings from method request parameters to integration request parameters. Supported request parameters are <code>querystring</code> , <code>path</code> , <code>header</code> , and <code>body</code> .
cacheNamespace	string	An API-specific tag group of related cached parameters.
cacheKeyParameters	An array of string	A list of request parameters whose values are to be cached.
responses	x-amazon-apigateway-integration.responses (p. 129)	Defines the method's responses and specifies desired parameter mappings or payload mappings from integration responses to method responses.

x-amazon-apigateway-integration Example

The following example integrates an API's `POST` method with a Lambda function in the back end. For demonstration purposes, the sample mapping templates shown in `requestTemplates` and `responseTemplates` of the examples below are assumed to apply to the following JSON-formatted payload: `{ "name": "value_1", "key": "value_2", "redirect": { "url": "..." } }` to generate a JSON output of `{ "stage": "value_1", "user-id": "value_2" }` or an XML output of `<stage>value_1</stage>`.

```
"x-amazon-apigateway-integration" : {
    "type" : "aws",
    "uri" : "arn:aws:apigateway:us-east-1:lambda:path/2015-03-31/functions/arn:aws:lambda:us-east-1:012345678901:function:HelloWorld/invocations",
    "httpMethod" : "POST",
    "credentials" : "arn:aws:iam::012345678901:role/apigateway-invoke-lambda-exec-role",
    "requestTemplates" : {
```

```

        "application/json" : "#set ($root=$input.path('$')) { \"stage\":
\"$root.name\", \"user-id\": \"$root.key\" }",
        "application/xml" : "#set ($root=$input.path('$'))"
<stage>$root.name</stage> "
},
"requestParameters" : {
    "integration.request.path.stage" : "method.request.querystring.version",

    "integration.request.querystring.provider" : "method.request.querystring.vendor"
},
"cacheNamespace" : "cache namespace",
"cacheKeyParameters" : [],
"responses" : {
    "2\\d{2}" : {
        "statusCode" : "200",
        "responseParameters" : {
            "method.response.header.requestId" : "integration.response.header.cid"
        },
        "responseTemplates" : {
            "application/json" : "#set ($root=$input.path('$')) { \"stage\":
\"$root.name\", \"user-id\": \"$root.key\" }",
            "application/xml" : "#set ($root=$input.path('$'))"
<stage>$root.name</stage> "
        }
    },
    "302" : {
        "statusCode" : "302",
        "responseParameters" : {
            "method.response.header.Location" : "integration.response.body.redirect.url"
        }
    },
    "default" : {
        "statusCode" : "400",
        "responseParameters" : {
            "method.response.header.test-method-response-header" : "'static
value'"
        }
    }
}
}

```

Note that double quotes ("") of the JSON string in the mapping templates must be string-escaped (\").

x-amazon-apigateway-integration.requestTemplates Object

Specifies mapping templates for a request payload of the specified MIME types.

Properties

Property Name	Type	Description
<i>MIME type</i>	string	An example of the MIME type is application/json. For information about creating a mapping template, see Mapping Templates (p. 71) .

x-amazon-apigateway-integration.requestTemplates Example

The following example sets mapping templates for a request payload of the application/json and application/xml MIME types.

```
"requestTemplates" : {
    "application/json" : "#set ($root=$input.path('$')) { \"stage\":
\$root.name\", \"user-id\": \"$root.key\" }",
    "application/xml" : "#set ($root=$input.path('$')) <stage>$root.name</stage>
"
}
```

x-amazon-apigateway-integration.requestParameters Object

Specifies mappings from named method request parameters to integration request parameters. The method request parameters must be defined before being referenced.

Properties

Property Name	Type	Description
integration.request.<param-type>.<param-name>	string	The value must be a predefined method request parameter of the method.request.<param-type>.<param-name> format, where <param-type> can be querystring, path, header, or body. For the body parameter, the <param-name> is a JSON path expression without the \$ prefix.

x-amazon-apigateway-integration.requestParameters Example

The following request parameter mappings example translates a method request's query (version), header (x-user-id) and path (service) parameters to the integration request's query (stage), header (x-userid), and path parameters (op), respectively.

```
"requestParameters" : {
    "integration.request.querystring.stage" : "method.request.querystring.version",
```

```

        "integration.request.header.x-userid" : "method.request.header.x-user-id",
        "integration.request.path.op" : "method.request.path.service"
    },
}

```

x-amazon-apigateway-integration.responses Object

Defines the method's responses and specifies parameter mappings or payload mappings from integration responses to method responses.

Properties

Property Name	Type	Description
<i>Response Id</i>	x-amazon-apigateway-integration.response (p. 130)	Selection regular expression used to match the integration response to the method response. For HTTP integrations, this regex applies to the integration response status code. For Lambda invocations, the regex applies to the <code>errorMessage</code> field of the error information object returned by AWS Lambda as a failure response body when the Lambda function execution throws an exception .

x-amazon-apigateway-integration.responses Example

The following example shows a list of responses from 2xx and 302 responses. For the 2xx response, the method response is mapped from the integration response's payload of the `application/json` or `application/xml` MIME type. This response uses the supplied mapping templates. For the 302 response, the method response returns a `Location` header whose value is derived from the `redirect.url` property on the integration response's payload.

```

"responses" : {
    "2\\d{2}" : {
        "statusCode" : "200",
        "responseTemplates" : {
            "application/json" : "#set ($root=$input.path('$')) { \\\"stage\\\": \\\"$root.name\\\", \\\"user-id\\\": \\\"$root.key\\\" }",
            "application/xml" : "#set ($root=$input.path('$'))
<stage>$root.name</stage> "
        }
    },
    "302" : {
        "statusCode" : "302",
        "responseParameters" : {
            "method.response.header.Location" : "integration.response.body.redirect.url"
        }
    }
}

```

}

x-amazon-apigateway-integration.response Object

Defines a response and specifies parameter mappings or payload mappings from the integration response to the method response.

Properties

Property Name	Type	Description
statusCode	string	HTTP status code for the method response; for example, "200". This must correspond to a matching response in the Swagger Operation responses field.
responseTemplates	x-amazon-apigateway-integration.responseParameters (p. 131)	Specifies MIME type-specific mapping templates for the response's payload.
responseParameters	x-amazon-apigateway-integration.responseTemplates (p. 131)	Specifies parameter mappings for the response. Only the <code>header</code> and <code>body</code> parameters of the integration response can be mapped to the <code>header</code> parameters of the method.

x-amazon-apigateway-integration.response Example

The following example defines a 302 response for the method that derives a payload of the `application/json` or `application/xml` MIME type from the back end. The response uses the supplied mapping templates and returns the redirect URL from the integration response in the method's `Location` header.

```
{
    "statusCode" : "302",
    "responseTemplates" : {
        "application/json" : "#set ($root=$input.path('$')) { \"stage\": \"$root.name\", \"user-id\": \"$root.key\" }",
        "application/xml" : "#set ($root=$input.path('$'))
<stage>$root.name</stage> "
    },
    "responseParameters" : {
        "method.response.header.Location": "integration.response.body.redirect.url"
    }
}
```

x-amazon-apigateway-integration.responseTemplates Object

Specifies mapping templates for a response payload of the specified MIME types.

Properties

Property Name	Type	Description
<i>MIME type</i>	string	Specifies a mapping template to transform the integration response body to the method response body for a given MIME type. For information about creating a mapping template, see Mapping Templates (p. 71) . An example of the <i>MIME type</i> is application/json.

x-amazon-apigateway-integration.responseTemplate Example

The following example sets mapping templates for a request payload of the application/json and application/xml MIME types.

```
"responseTemplates" : {
    "application/json" : "#set ($root=$input.path('$')) { \"stage\": \"$root.name\", \"user-id\": \"$root.key\" }",
    "application/xml" : "#set ($root=$input.path('$')) <stage>$root.name</stage>
"
}
```

x-amazon-apigateway-integration.responseParameters Object

Specifies mappings from integration method response parameters to method response parameters. Only the header and body types of the integration response parameters can be mapped to the header type of the method response.

Properties

Property Name	Type	Description
method.response.header.< <i>param-name</i> >	string	The named parameter value can be derived from the header and body types of the integration response parameters only.

x-amazon-apigateway-integration.responseParameters Example

The following example maps body and header parameters of the integration response to two header parameters of the method response.

```
"responseParameters" : {
    "method.request.header.Location" : "integration.request.body.redirect.url",
    "method.request.header.x-user-id" : "integration.request.header.x-userid"
}
```

Create an API as an Amazon S3 Proxy

As an example to showcase using an API in API Gateway to proxy Amazon S3, this section describes how to create and configure an API to expose the following Amazon S3 operations:

- Expose GET on the API's root resource to [list all of the Amazon S3 buckets of a caller](#).
- Expose GET on a Folder resource to [view a list of all of the objects in an Amazon S3 bucket](#).
- Expose PUT on a Folder resource to [add a bucket to Amazon S3](#).
- Expose DELETE on a Folder resource to [remove a bucket from Amazon S3](#).
- Expose GET on a Folder/Item resource to [view or download an object from an Amazon S3 bucket](#).
- Expose PUT on a Folder/Item resource to [upload an object to an Amazon S3 bucket](#).
- Expose HEAD on a Folder/Item resource to [get object metadata in an Amazon S3 bucket](#).
- Expose DELETE on a Folder/Item resource to [remove an object from an Amazon S3 bucket](#).

You may want to import the sample API as an Amazon S3 proxy, as shown in [A Sample Amazon S3 Proxy API in Swagger with API Gateway Extensions \(p. 141\)](#). To do so, copy the Swagger definition and paste it into a file. Use the [API Gateway Swagger Importer](#). For more information, see [Getting Started with the API Gateway Swagger Importer](#).

To use the API Gateway console to create the API, you must first sign up for an AWS account.

If you do not have an AWS account, use the following procedure to create one.

To sign up for AWS

1. Open <http://aws.amazon.com/> and choose **Create an AWS Account**.
2. Follow the online instructions.

To allow the API to invoke the Amazon S3 actions, you must have appropriate IAM policies attached to an IAM role. The next section describes how to verify and to create, if necessary, the required IAM role and policies.

Create an IAM Role and Policy for the API to Access Amazon S3

For read-only operations, including `Get*` and `List*` actions in Amazon S3, you can use the `AmazonS3ReadOnlyAccess` policy provided by the IAM , whose ARN is `arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess`:

The AmazonS3ReadOnlyAccess Policy

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "s3:Get*",  
                "s3>List*"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

This policy document states that any of the Amazon S3 `Get*` and `List*` actions can be invoked on any of the Amazon S3 resources.

To allow Amazon S3 buckets and objects to be updated, you can use a custom policy for any of the Amazon S3 `Put*` actions.

An Amazon S3 Put-only Policy

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": "s3:Put*",  
            "Resource": "*"  
        }  
    ]  
}
```

You can attach a read-only and a put-only policy to an IAM role if your API works with Amazon S3 `Get*`, `List*` and `Put*` actions only.

To invoke the Amazon S3 `Post` actions, you must include `s3:Post*` action in an Allow policy document. For a complete list of Amazon S3 actions, see [Specifying Amazon S3 Permissions in a Policy](#).

For an API to create, view, update, and delete buckets and objects in Amazon S3, you can attach a single full-access policy. For this, you can use the `AmazonS3FullAccess` policy, which is provided by the IAM console and whose ARN is `arn:aws:iam::aws:policy/AmazonS3FullAccess`.

The AmazonS3FullAccess Policy

```
{  
    "Version": "2012-10-17",  
}
```

```
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "s3:*",
            "Resource": "*"
        }
    ]
}
```

This policy covers all actions on any resources in Amazon S3. Using the IAM role and policies ensures that API Gateway can call the specifically allowed Amazon S3 actions on the specified Amazon S3 resources.

After you have decided which IAM policy documents to use, create an IAM role. Attach the policies to the role. The resulting IAM role must contain the following trust policy for the attached policies to be applied on API Gateway.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "",
            "Effect": "Allow",
            "Principal": {
                "Service": "apigateway.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

When using the IAM console to create the role, choose the **Amazon API Gateway** role type to ensure that this trust policy is automatically included.

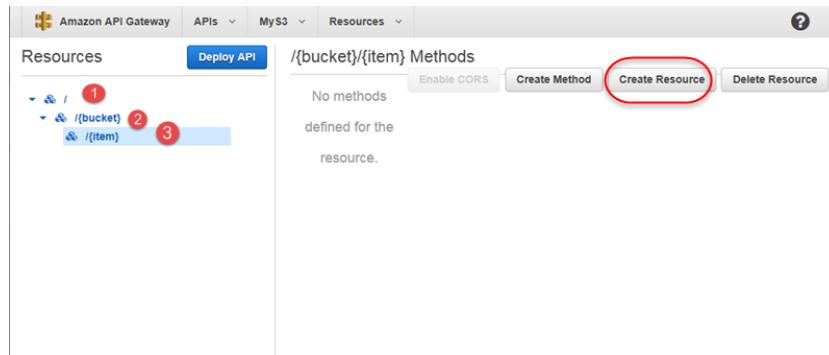
Create API Resources for Amazon S3 Features

The following procedure describes how to use the API Gateway console to create an API that exposes the Amazon S3 service features.

To create an API resource that exposes the Amazon S3 service features

1. In the API Gateway console, create an API named **MyS3**. This API's root resource (**/**) represents the Amazon S3 service. Later, we will expose the GET method to list the Amazon S3 buckets of the caller.
2. For the API's root resource, create a child resource named **Folder**, setting the required **Resource Path** as **/{folder}**. The **folder** path parameter enables the client to specify a bucket name in the URL when the client calls the API to work with the bucket. Later, we will expose the DELETE, GET, and PUT methods on this **Folder** resource to work with an Amazon S3 bucket in the back end. We will also declare a **bucket** path parameter for the back-end URL and specify a mapping expression to translate **folder** to **bucket**.
3. For the API's **Folder** resource, identified by the **/{folder}** resource path, create an **Item** child resource. Set the required **Resource Path** as **/{item}**. The **folder** and **item** path parameters enable the client

to specify, in the request's URL, an object name in a given **folder** when the client calls the API to work with the object. Later, we will expose the DELETE, HEAD, GET and PUT methods on this **Item** resource. We will also declare **bucket** and **object** path parameters for the back-end URL and specify mapping expressions to translate **folder** and **item** to **bucket** and **object**, respectively.



Expose a GET Method on an API Root as Get Service Action in Amazon S3

Use **Create Method** in the API Gateway console to create a GET method for the API's root resource, `(/)`. In the **Set up** pane for the method, configure the GET method to integrate with the GET Service action in Amazon S3, as follows.

To set up the newly created GET method on the API root

1. For the **Integration type**, choose **AWS Service Proxy**.
2. From the list, choose an **AWS Region**.
3. From **AWS Service**, choose **S3**.
4. From **HTTP method**, choose **GET**.
5. For **Action Type**, choose **Use path override**.
6. (Optional) In **Path override** type `/`.
7. Copy the previously created IAM role's ARN (from the IAM console) and paste it into **Execution role**.
8. Choose **Save** to finish setting up this method.

Amazon API Gateway Developer Guide
Expose a GET Method on an API Root as Get Service
Action in Amazon S3

Choose the integration point for your new method. ⓘ

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy 1

AWS Region 2

AWS Service 3

AWS Subdomain

HTTP method 4

Action Type Use action name
 Use path override 5

Path override (optional) 6

Execution role 7 ⓘ

8 Save

Note

After the method is set up, you can modify these settings in the method's **Integration Request** page.

By default, API Gateway assumes the request and response payload to be of the "application/json" type. However, Amazon S3 returns results in an XML-formatted response payload. This means that you must override the default Content-Type header value of the method response with the Content-Type header value from the integration response. Otherwise, the client will see "application/json" in the response Content-Type header when the response body is an XML string.

To translate the integration response header to the method response header, use response header mappings. The process involves declaring the Content-Type header explicitly for the method response and specifying a header mapping expression for the integration response to pass the back-end Content-Type header value to the front-end Content-Type header value.

To set up header mappings to return an integration response Content-Type header

1. In the API Gateway console, choose **Method Response**. Add the **Content-Type** header.

Amazon API Gateway Developer Guide
Expose a GET Method on an API Root as Get Service
Action in Amazon S3

Provide information about this method's response types, their headers and content types.

HTTP Status	
▼ 200	

Response Headers for 200 Response Models for 200 Create a model

Name	Content type	Models
Timestamp	application/json	Empty
Content-Length		
Content-Type		

Add Response Model

Add Header

Add Response

2. In **Integration Response**, for **Content-Type**, type `integration.response.header.Content-Type` for the method response.

First, declare response types using [Method Response](#). Then, map the possible responses from the backend to this method's response types.

HTTP status regex	Method response status	Output model	Default mapping
▼ \d{3}	200		No

Map the output from your HTTP endpoint to the headers and output model of the 200 method response.

HTTP status regex `\d{3}`

Method response status 200

Cancel Save

▼ Header Mappings

Response header	Mapping value ⓘ
Timestamp	integration.response.header.Date
Content-Length	integration.response.header.Content-Length
Content-Type	integration.response.header.Content-Type

► Mapping Templates

Add integration response

Test the GET method on the API root resource

- In **Method Execution**, choose **Test**. An example result is shown in the following figure.

Amazon API Gateway Developer Guide

Expose Methods on an API Folder Resource as Bucket Actions in Amazon S3

[Method Execution](#) / - GET - Method Test [Delete Method](#)

Make a test call to your method with the provided input

Path
No path parameters exist for this resource. You can define path parameters by using the syntax `{myPathParam}` in a resource path.

Query Strings
No query string parameters exist for this method. You can add them via Method Request.

Headers
No header parameters exist for this method. You can add them via Method Request.

Stage Variables
No stage variables exist for this method.

Client Certificate

Request Body
Request Body is not supported for GET methods.

[Test](#)

```
<?xml version="1.0" encoding="UTF-8"?>
<ListAllMyBucketsResult xmlns="http://s3.amazonaws.com/doc/2006-03-01/"><Owner><ID>0e611111111111111111111111111111</ID><DisplayName>aws-lambda</DisplayName><Bucket><Name>aws-lambda</Name><CreationDate>2016-02-15T22:05:55.000Z</CreationDate><Bucket><Name>pig-demo1</Name><CreationDate>2016-02-15T23:38:35.000Z</CreationDate><Bucket><Name>aws-dev-od-test-kd</Name><CreationDate>2015-10-27T22:59:29.000Z</CreationDate><Bucket><Name>aws-lambda-fileproc-test-kd-eventarchive-17</Name><CreationDate>2015-10-21T22:39:57.000Z</CreationDate><Bucket><Name>aws-lambda-fileproc-test-kd-eventarchive-18</Name><CreationDate>2015-10-21T23:14:14.000Z</CreationDate><Bucket><Name>aws-lab-triggers-kd</Name><CreationDate>2015-10-22T23:48:01.000Z</CreationDate><Bucket><Name>f-templates-10</Name><CreationDate>2015-10-21T20:59:45.000Z</CreationDate><Bucket><Name>elasticbeanstalk-us-east-1-70000000007</Name><CreationDate>2015-10-16T23:13:15.000Z</CreationDate><Bucket><Name>jaws.dev.ubuntu1404.20150219141414</Name><CreationDate>2015-12-01T19:34:56.000Z</CreationDate><Bucket><Name>jaws.dev.useast1.myapp-evolvun.com</Name><CreationDate>2015-12-01T20:22:47.000Z</CreationDate><Bucket><Name>my-pictures-02-16-2016</Name><CreationDate>2016-02-18T08:14:40.000Z</CreationDate><Bucket><Name>myvr-contentdelivery-mobilehub-1161765204</Name><CreationDate>2015-10-26T05:07:51.000Z</CreationDate><Bucket><Name>myvr-userfiles-mobilehub-1161765204</Name><CreationDate>2015-10-26T05:09:38.000Z</CreationDate><Bucket><Name>node-sdk-sample-33358871-858c-40fd-8f90-90be0096ee</Name><CreationDate>2015-10-21T23:49:15.000Z</CreationDate><Bucket></Buckets></ListAllMyBucketsResult>
```

Response Headers

Logs

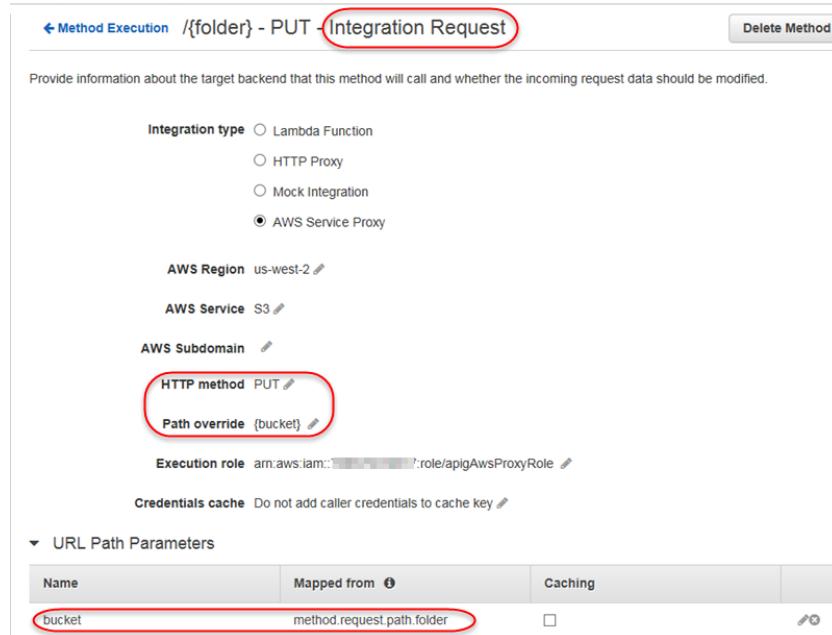
Expose Methods on an API Folder Resource as Bucket Actions in Amazon S3

You can apply a similar procedure to expose methods on the Folder (`/{folder}`) resources, with minor modifications that include using different HTTP verbs and setting up path parameter mappings between the method request and integration request. To illustrate, we expose the GET, PUT, and DELETE methods for listing objects in a bucket, creating a new bucket, and deleting an existing bucket, respectively. As an example, we will cover the PUT method setup in detail.

To expose methods on a folder resource

1. On the `/{folder}` resource under the root, choose **Create Method** to create a PUT method.
2. Follow the setup steps for the GET method on the root resource explained in the previous procedure, except that you will specify **PUT** for **HTTP method** and `/ {bucket}` for **Path override**.
3. Set up the mappings for the Content-Type and possibly other headers from method request to integration request and from integration response to method response. The instructions are the same or similar to the header mapping from the integration response to the method response for the GET method on the API's root resource. You must add the mapping for the Content-Type header from the method request to integration because you must supply an XML payload in the PUT request; this mapping overrides the default Content-Type header value (i.e., application/json) to reflect the actual payload content type.
4. In **Integration Request**, expand the **URL Path Parameters** section. Add the path parameter name, for example, **bucket**, as specified in **Path override**. Type a path-mapping expression, namely **method.request.path.folder**, to map the front-end path parameter (**folder**) for the method request to the back-end path parameter (**bucket**) for the integration request.

Amazon API Gateway Developer Guide
Expose Methods on an API Folder Resource as Bucket
Actions in Amazon S3



Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-west-2

AWS Service S3

AWS Subdomain

HTTP method PUT

Path override {bucket}

Execution role arn:aws:iam::██████████:role/apigAwsProxyRole

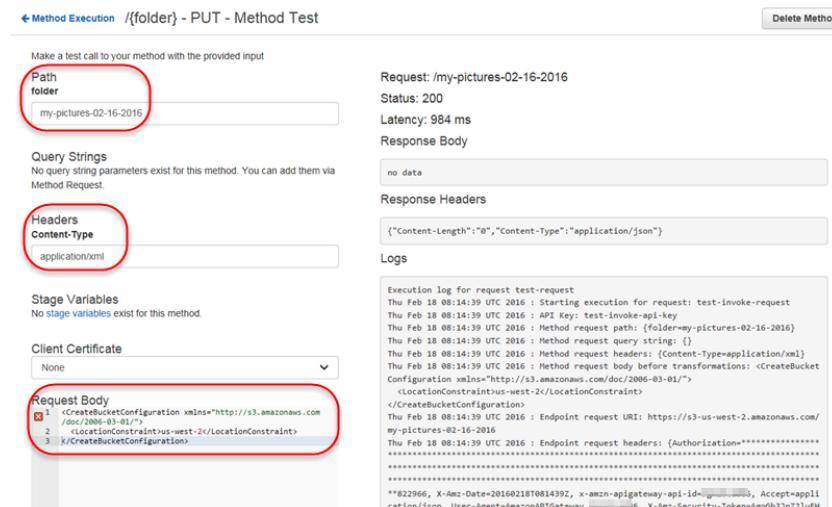
Credentials cache Do not add caller credentials to cache key

▼ URL Path Parameters

Name	Mapped from	Caching
bucket	method.request.path.folder	<input type="checkbox"/>

- Choose **Test** from the **Method Execution** pane to test this PUT method. In **folder**, type a bucket name, and then in **Content-Type**, type **application/xml**. In **Request Body**, provide the bucket region as the location constraint; it is declared in an XML fragment as the payload of the request.

```
<CreateBucketConfiguration>
<LocationConstraint>us-west-2</LocationConstraint>
</CreateBucketConfiguration>
```



Make a test call to your method with the provided input

Path
my-pictures-02-16-2016

Query Strings
No query string parameters exist for this method. You can add them via Method Request.

Headers
Content-Type
application/xml

Stage Variables
No stage variables exist for this method.

Client Certificate
None

Request Body
<CreateBucketConfiguration xmlns="http://s3.amazonaws.com/>
1 <LocationConstraint>us-west-2</LocationConstraint>
2 </CreateBucketConfiguration>

Request: /my-pictures-02-16-2016
Status: 200
Latency: 984 ms
Response Body
no data

Response Headers
{"Content-Length": "0", "Content-Type": "application/json"}

Logs

```
Execution log for request test-request
Thu Feb 18 08:14:39 UTC 2016 : Starting execution for request: test-invoke-request
Thu Feb 18 08:14:39 UTC 2016 : API Key: test-invoke-api-key
Thu Feb 18 08:14:39 UTC 2016 : Method request path: /{folder}/my-pictures-02-16-2016
Thu Feb 18 08:14:39 UTC 2016 : Method request query string: {}
Thu Feb 18 08:14:39 UTC 2016 : Method request headers: (Content-Type=application/xml)
Thu Feb 18 08:14:39 UTC 2016 : Method request body before transformations: <CreateBucketConfiguration xmlns="http://s3.amazonaws.com/doc/2006-03-01">
<LocationConstraint>us-west-2</LocationConstraint>
</CreateBucketConfiguration>
Thu Feb 18 08:14:39 UTC 2016 : Endpoint request URI: https://s3-us-west-2.amazonaws.com/my-pictures-02-16-2016
Thu Feb 18 08:14:39 UTC 2016 : Endpoint request headers: (Authorization=*****)
*****
```

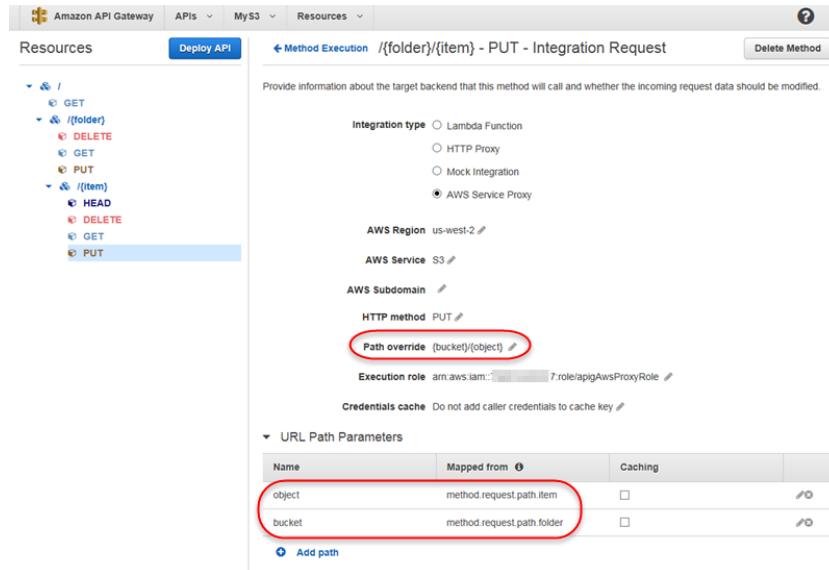
- Repeat the preceding steps to create and configure the GET and DELETE method on the API's **{/folder}** resource.

Expose Methods on an API Item in a Folder as Actions on an Amazon S3 Object in a Bucket

The following procedure describes how to expose methods on an API item in a folder as an action on an Amazon S3 object in a bucket.

To expose methods on an item resource

1. Repeat the preceding steps to create and configure the DELETE, HEAD, GET and PUT methods on the `/{folder}/{item}` resource. The following image shows integration settings for the PUT method. The settings for the other methods are similar.



2. To test the GET method on a Folder/Item resource using the API Gateway console, choose **Test** in the **Method Execution** page. The following image shows the result of an example test, where the README.txt file in the apig-demo bucket in Amazon S3 contains a string of plain text ("Welcome to README.txt").

Amazon API Gateway Developer Guide
Swagger Definitions of a Sample API as an Amazon S3 Proxy

Method Execution /{folder}/{item} - GET - Method Test Delete Method

Make a test call to your method with the provided input

Path
folder
apig-demo

item
README.txt

Request: /apig-demo/README.txt
Status: 200
Latency: 486 ms

Response Body
Welcome to README.txt

Response Headers
{"Content-Type": "text/plain"}

Query Strings
No query string parameters exist for this method.
You can add them via Method Request.

Headers
No header parameters exist for this method. You can add them via Method Request.

Stage Variables
No stage variables exist for this method.

Client Certificate
None

Request Body
Request Body is not supported for GET methods.

Logs

Execution log for request test-request
Fri Feb 19 03:35:20 UTC 2016 : Starting execution for request: test-invoke-request
Fri Feb 19 03:35:20 UTC 2016 : API Key: test-invoke-api-key
Fri Feb 19 03:35:20 UTC 2016 : Method request path: {item=README.txt, folder=apig-demo}
Fri Feb 19 03:35:20 UTC 2016 : Method request query string: {}
Fri Feb 19 03:35:20 UTC 2016 : Method request headers: {}
Fri Feb 19 03:35:20 UTC 2016 : Method request body before transformations: null
Fri Feb 19 03:35:20 UTC 2016 : Endpoint request URI: http://s3-us-west-2.amazonaws.com/apig-demo/README.txt
Fri Feb 19 03:35:20 UTC 2016 : Endpoint request headers: {Authorization=*****

*****a99006, X-Amz

Test

A Sample Amazon S3 Proxy API in Swagger with API Gateway Extensions

```
{  
  "swagger": "2.0",  
  "info": {  
    "version": "2016-02-19T04:30:12Z",  
    "title": "MyS3"  
  },  
  "host": "1234567890.execute-api.us-east-1.amazonaws.com",  
  "basePath": "/S3",  
  "schemes": [  
    "https"  
  ],  
  "paths": {  
    "/": {  
      "get": {  
        "produces": [  
          "application/json"  
        ],  
        "parameters": [],  
        "responses": {  
          "200": {  
            "description": "200 response",  
            "schema": {  
              "$ref": "#/definitions/Empty"  
            }  
          }  
        }  
      }  
    }  
  }  
}
```

```
        },
        "headers": {
            "Content-Length": {
                "type": "string"
            },
            "Timestamp": {
                "type": "string"
            },
            "Content-Type": {
                "type": "string"
            }
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
        "responses": {
            "\\\d{3}": {
                "statusCode": "200",
                "responseParameters": {
                    "method.response.header.Content-Type": "integration.response.header.Content-Type",
                    "method.response.header.Content-Length": "integration.response.header.Content-Length",
                    "method.response.header.Timestamp": "integration.response.header.Timestamp",
                    "method.response.header.Date": ""
                },
                "responseTemplates": {
                    "application/json": "__passthrough__"
                }
            }
        },
        "uri": "arn:aws:apigateway:us-west-2:s3:path//",
        "httpMethod": "GET",
        "type": "aws"
    }
},
"/{folder)": {
    "get": {
        "produces": [
            "application/json"
        ],
        "parameters": [
            {
                "name": "folder",
                "in": "path",
                "required": true,
                "type": "string"
            }
        ],
        "responses": {
            "200": {
                "description": "200 response",
                "schema": {
                    "$ref": "#/definitions/Empty"
                },
                "headers": {

```

```
        "Content-Length": {
            "type": "string"
        },
        "Date": {
            "type": "string"
        },
        "Content-Type": {
            "type": "string"
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
        "\d{3}": {
            "statusCode": "200",
            "responseParameters": {
                "method.response.header.Content-Type": "integration.response.header.Content-Type",
                "method.response.header.Date": "integration.response.header.Date",
                "method.response.header.Content-Length": "integration.response.header.content-length"
            },
            "responseTemplates": {
                "application/json": "__passthrough__"
            }
        }
    },
    "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}",
    "httpMethod": "GET",
    "requestParameters": {
        "integration.request.path.bucket": "method.request.path.folder"
    },
    "type": "aws"
},
"put": {
    "produces": [
        "application/json"
    ],
    "parameters": [
        {
            "name": "Content-Type",
            "in": "header",
            "required": false,
            "type": "string"
        },
        {
            "name": "folder",
            "in": "path",
            "required": true,
            "type": "string"
        }
    ],
    "responses": {
        "200": {

```

Amazon API Gateway Developer Guide
Swagger Definitions of a Sample API as an Amazon S3
Proxy

```
        "description": "200 response",
        "schema": {
            "$ref": "#/definitions/Empty"
        },
        "headers": {
            "Content-Length": {
                "type": "string"
            },
            "Content-Type": {
                "type": "string"
            }
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
        "responses": {
            "\\\d{3}": {
                "statusCode": "200",
                "responseParameters": {
                    "method.response.header.Content-Type": "integration.response.header.Content-Type",
                    "method.response.header.Content-Length": "integration.response.header.Content-Length"
                },
                "responseTemplates": {
                    "application/json": "__passthrough__"
                }
            }
        },
        "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}",
        "httpMethod": "PUT",
        "requestParameters": {
            "integration.request.path.bucket": "method.request.path.folder",
            "integration.request.header.Content-Type": "method.request.header.Content-Type"
        },
        "type": "aws"
    },
    "delete": {
        "produces": [
            "application/json"
        ],
        "parameters": [
            {
                "name": "folder",
                "in": "path",
                "required": true,
                "type": "string"
            }
        ],
        "responses": {
            "200": {
                "description": "200 response",
                "schema": {
                    "$ref": "#/definitions/Empty"
                },
                "headers": {
                    "Content-Type": {
                        "type": "string"
                    }
                }
            }
        }
    }
}
```

```
    "headers": {
        "Date": {
            "type": "string"
        },
        "Content-Type": {
            "type": "string"
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
        "\\\d{3}": {
            "statusCode": "200",
            "responseParameters": {
                "method.response.header.Content-Type": "integration.response.header.Content-Type",
                "method.response.header.Date": "integration.response.header.Date"
            },
            "responseTemplates": {
                "application/json": "__passthrough__"
            }
        }
    },
    "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}",
    "httpMethod": "DELETE",
    "requestParameters": {
        "integration.request.path.bucket": "method.request.path.folder"
    },
    "type": "aws"
}
},
"/{folder}/{item)": {
    "get": {
        "produces": [
            "application/json"
        ],
        "parameters": [
            {
                "name": "item",
                "in": "path",
                "required": true,
                "type": "string"
            },
            {
                "name": "folder",
                "in": "path",
                "required": true,
                "type": "string"
            }
        ],
        "responses": {
            "200": {
                "description": "200 response",
                "schema": {

```

```
        "$ref": "#/definitions/Empty"
    },
    "headers": {
        "content-type": {
            "type": "string"
        },
        "Content-Type": {
            "type": "string"
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
        "\d{3}": {
            "statusCode": "200",
            "responseParameters": {
                "method.response.header.content-type": "integration.response.header.content-type",
                "method.response.header.Content-Type": "integration.response.header.Content-Type"
            },
            "responseTemplates": {
                "application/json": "__passthrough__"
            }
        }
    },
    "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}/{object}",
    "httpMethod": "GET",
    "requestParameters": {
        "integration.request.path.object": "method.request.path.item",
        "integration.request.path.bucket": "method.request.path.folder"
    },
    "type": "aws"
},
"head": {
    "produces": [
        "application/json"
    ],
    "parameters": [
        {
            "name": "item",
            "in": "path",
            "required": true,
            "type": "string"
        },
        {
            "name": "folder",
            "in": "path",
            "required": true,
            "type": "string"
        }
    ],
    "responses": {
        "200": {
            "description": "200 response",

```

```
    "schema": {
      "$ref": "#/definitions/Empty"
    },
    "headers": {
      "Content-Length": {
        "type": "string"
      },
      "Content-Type": {
        "type": "string"
      }
    }
  },
  "x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
      "\\\d{3}": {
        "statusCode": "200",
        "responseParameters": {
          "method.response.header.Content-Type": "integration.response.header.Content-Type",
          "method.response.header.Content-Length": "integration.response.header.Content-Length"
        },
        "responseTemplates": {
          "application/json": "__passthrough__"
        }
      }
    },
    "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}/{object}",
    "httpMethod": "HEAD",
    "requestParameters": {
      "integration.request.path.object": "method.request.path.item",
      "integration.request.path.bucket": "method.request.path.folder"
    },
    "type": "aws"
  },
  "put": {
    "produces": [
      "application/json"
    ],
    "parameters": [
      {
        "name": "Content-Type",
        "in": "header",
        "required": false,
        "type": "string"
      },
      {
        "name": "item",
        "in": "path",
        "required": true,
        "type": "string"
      },
      {
        "name": "folder",
        "in": "path",
        "type": "string"
      }
    ]
  }
}
```

```
        "required": true,
        "type": "string"
    },
],
"responses": {
    "200": {
        "description": "200 response",
        "schema": {
            "$ref": "#/definitions/Empty"
        },
        "headers": {
            "Content-Length": {
                "type": "string"
            },
            "Content-Type": {
                "type": "string"
            }
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
        "\\\d{3}": {
            "statusCode": "200",
            "responseParameters": {
                "method.response.header.Content-Type": "integration.response.header.Content-Type",
                "method.response.header.Content-Length": "integration.response.header.Content-Length"
            },
            "responseTemplates": {
                "application/json": "__passthrough__"
            }
        }
    },
    "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}/{object}",
    "httpMethod": "PUT",
    "requestParameters": {
        "integration.request.path.object": "method.request.path.item",
        "integration.request.header.content-type": "method.request.header.Content-Type",
        "integration.request.path.bucket": "method.request.path.folder",
        "integration.request.header.Content-Type": "method.request.header.Content-Type"
    },
    "type": "aws"
},
"delete": {
    "produces": [
        "application/json"
    ],
    "parameters": [
        {
            "name": "item",
            "in": "path",
            "required": true,

```

```
        "type": "string"
    },
{
    "name": "folder",
    "in": "path",
    "required": true,
    "type": "string"
}
],
"responses": {
    "200": {
        "description": "200 response",
        "schema": {
            "$ref": "#/definitions/Empty"
        },
        "headers": {
            "Content-Length": {
                "type": "string"
            },
            "Content-Type": {
                "type": "string"
            }
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
        "d\\{3}": {
            "statusCode": "200",
            "responseParameters": {
                "method.response.header.Content-Type": "integration.response.header.Content-Type",
                "method.response.header.Content-Length": "integration.response.header.Content-Length"
            },
            "responseTemplates": {
                "application/json": "__passthrough__"
            }
        }
    },
    "uri": "arn:aws:apigateway:us-west-2:s3:path/{bucket}/{object}",
    "httpMethod": "DELETE",
    "requestParameters": {
        "integration.request.path.object": "method.request.path.item",
        "integration.request.path.bucket": "method.request.path.folder"
    },
    "type": "aws"
}
}
},
"definitions": {
    "Empty": {
        "type": "object"
    }
}
```

}

Create an API Gateway API as an AWS Lambda Proxy

If your API makes only synchronous calls to Lambda functions in the back end, you should use the **Lambda Function** integration type. For instructions, see [Invoke Lambda Functions Synchronously \(p. 20\)](#).

If your API makes asynchronous calls to Lambda functions, you must use the **AWS Service Proxy** integration type described in this section. The instructions apply to requests for synchronous Lambda function invocations as well. For the asynchronous invocation, you must explicitly add the `X-Amz-Invocation-Type: Event` header to the integration request. For the synchronous invocation, you can add the `X-Amz-Invocation-Type: RequestResponse` header to the integration request or leave it unspecified. The following example shows the integration request of an asynchronous Lambda function invocation:

```
POST /2015-03-31/functions/FunctionArn/invocations?Qualifier=Qualifier HTTP/1.1
X-Amz-Invocation-Type: Event
...
Content-Type: application/json
Content-Length: PayloadSize

Payload
```

In this example, `FunctionArn` is the ARN of the Lambda function to be invoked. For more information, see the [Invoke](#) action in the *AWS Lambda Developer Guide*.

To illustrate how to create and configure an API as an AWS service proxy for Lambda, we will create a Lambda function (`Calc`) that performs addition (+), subtraction (-), multiplication (*), and division (/). When a client submits a method request to perform any of these operations, API Gateway will post the corresponding integration request to call the specified Lambda function, passing the required input (two operands and one operator) as a JSON payload. A synchronous call will return the result, if any, as the JSON payload. An asynchronous call will return no data.

The API can expose a GET or POST method on the `/calc` resource to invoke the Lambda function. With the GET method, a client supplies the input to the back-end Lambda function through three query string parameters (`operand1`, `operand2`, and `operator`). These are mapped to the JSON payload of the integration request. With the POST method, a client provides the input to the Lambda function as a JSON payload of the method request, which is then passed through to the integration request. Alternatively, the API can expose a GET method on the `/calc/{operand1}/{operand2}/{operator}` resource. With this method, the client specifies the Lambda function input as the values of the path parameters. Parameter mappings and mapping templates are used to translate the method request data into the Lambda function input and to translate the output from the integration responses to the method response.

This section provides more detailed discussions for the following tasks:

- Create the `Calc` Lambda function to implement the arithmetic operations, accepting and returning JSON-formatted input and output.
- Expose GET on the `/calc` resource to invoke the Lambda function, supplying the input as query strings.

- Expose POST on the /calc resource to invoke the Lambda function, supplying the input in the payload.
- Expose GET on the /calc/{operand1}/{operand2}/{operator} resource to invoke the Lambda function, specifying the input in the path parameters.

You can import the sample API as a Lambda proxy from the [Swagger Definitions of a Sample API as Lambda Proxy \(p. 162\)](#). To do so, copy the Swagger definition, paste it into a file, and use the [API Gateway Swagger Importer](#). For more information, see [Getting Started with the API Gateway Swagger Importer](#).

To use the API Gateway console to create the API, you must first sign up for an AWS account.

If you do not have an AWS account, use the following procedure to create one.

To sign up for AWS

1. Open <http://aws.amazon.com/> and choose **Create an AWS Account**.
2. Follow the online instructions.

To allow the API to invoke Lambda functions, you must have an IAM role that has appropriate IAM policies attached to it. The next section describes how to verify and to create, if necessary, the required IAM role and policies.

Topics

- [Set Up an IAM Role and Policy for an API to Invoke Lambda Functions \(p. 151\)](#)
- [Create a Lambda Function in the Back End \(p. 152\)](#)
- [Create API Resources for the Lambda Function \(p. 153\)](#)
- [Create a GET Method with Query Strings to Call the Lambda Function \(p. 154\)](#)
- [Create a POST Method with a JSON Payload to Call the Lambda Function \(p. 156\)](#)
- [Create a GET Method with Path Parameters to Call the Lambda Function \(p. 158\)](#)
- [A Sample API as a Lambda Proxy in Swagger with API Gateway Extensions \(p. 162\)](#)

Set Up an IAM Role and Policy for an API to Invoke Lambda Functions

The API will use the [InvokeFunction](#) action to call a Lambda function. At minimum, you must attach the following IAM policy to an IAM role for API Gateway to assume the policy.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": "lambda:InvokeFunction",  
            "Resource": "*"  
        }  
    ]  
}
```

If you do not enact this policy, the API caller will receive a 500 Internal Server Error response. The response contains the "Invalid permissions on Lambda function" error message. For a complete list of error messages returned by Lambda, see the [Invoke](#) topic.

An API Gateway assumable role is an IAM role with the following trusted relationship:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "",  
            "Effect": "Allow",  
            "Principal": {  
                "Service": "apigateway.amazonaws.com"  
            },  
            "Action": "sts:AssumeRole"  
        }  
    ]  
}
```

Create a Lambda Function in the Back End

Copy the following Lambda function and paste it into the code editor in the Lambda console.

```
exports.handler = function(event, context) {  
    //console.log('Received event:', JSON.stringify(event, null, 2));  
    var res = {};  
    res.a = event.a;  
    res.b = event.b;  
    res.op = event.op;  
  
    switch(event.op)  
    {  
        case "+":  
            res.c = Number(event.a) + Number(event.b);  
            break;  
        case "-":  
            res.c = Number(event.a) - Number(event.b);  
            break;  
        case "*":  
            res.c = Number(event.a) * Number(event.b);  
            break;  
        case "/":  
            res.c = Number(event.b) === 0 ? NaN : Number(event.a) / Number(event.b);  
            break;  
        default:  
            res.c = "Invalid op";  
    }  
    context.succeed(res);  
};
```

This function requires two operands (`a` and `b`) and an operator (`op`) from the `event` input parameter. The input is a JSON object of the following format:

```
{
  "a": "Number" | "String",
  "b": "Number" | "String",
  "op": "String"
}
```

This function returns the calculated result (*c*) and the input. For an invalid input, the function returns either the null value or the "Invalid op" string as the result. The output is of the following JSON format:

```
{
  "a": "Number",
  "b": "Number",
  "op": "String",
  "c": "Number" | "String"
}
```

You should test the function in the Lambda console before integrating it with the API, which is created next.

Create API Resources for the Lambda Function

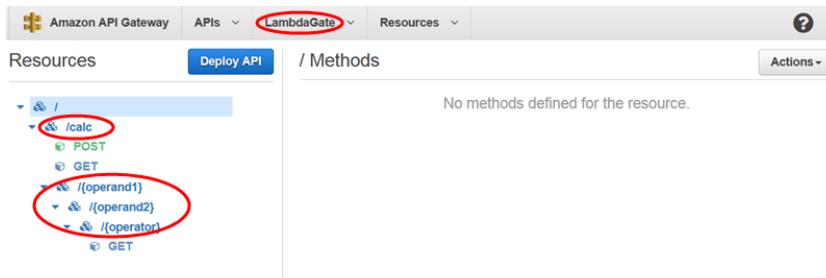
The following procedure describes how to create API resources for the Lambda function.

To create API resources for Lambda functions

1. In the API Gateway console, create an API named **LambdaGate**. You can create child resources to represent different Lambda functions; in the following, you will create a single child resource of the API root.
2. For the simple calculator function you created, create the **/calc** resource off the API's root. You will expose the GET and POST methods on this resource for the client to invoke the back-end Lambda function, supplying the required input as query string parameters (to be declared as `?operand1=...&operand2=...&operator=...`) in the GET request and as a JSON payload in the POST request, respectively.

You will also create the **/calc/{operand1}/{operand2}/{operator}** to expose the GET method to invoke the Lambda function and to supply the required input as the three path parameters (**operand1**, **operand2**, and **operator**).

We will show how to apply API Gateway request and response data mapping to normalize the input to the back end Lambda function.



Create a GET Method with Query Strings to Call the Lambda Function

Use the following steps to expose a GET method with query strings to call a Lambda function.

To set up the GET method with query strings to invoke the Lambda function

1. Choose **Create Method** in the API Gateway console to create a GET method for the API's `/calc` resource.

In the method's **Set up** pane, configure the method with the following settings.

Method Execution /calc - GET - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-west-2

AWS Service Lambda

AWS Subdomain

HTTP method POST

Path override /2015-03-31/functions/arn:aws:lambda:us-west-2:...:function:Calc/invocations

Execution role arn:aws:iam::...:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

URL Path Parameters

URL Query String Parameters

HTTP Headers

Mapping Templates

You must use the POST method for the integration request when calling a Lambda function, although you can use any other HTTP verbs for the method request.

The **Path override** value must the URL path of the Lambda [Invoke](#) action. The path is of the following format:

```
/2015-03-31/functions/FunctionName/invocations?Qualifier=version
```

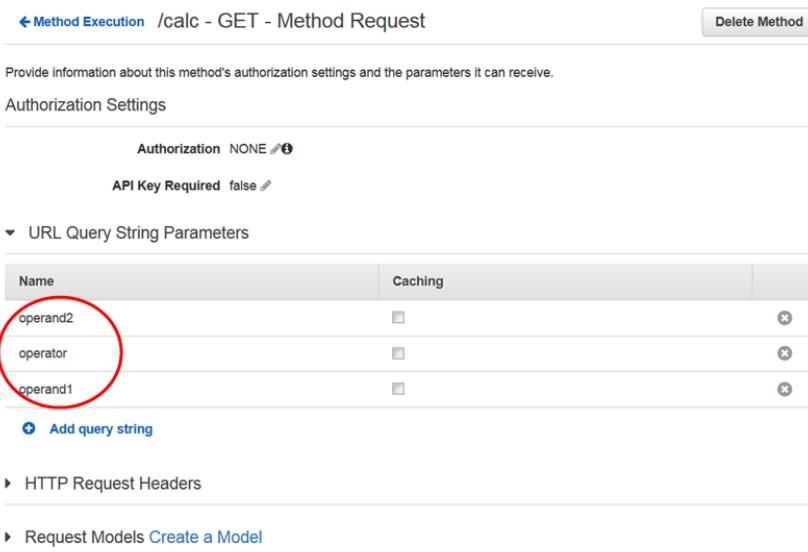
where *FunctionName* is the ARN of the Lambda function to be invoked. The optional *Qualifier* query string can be used to select a version of the function. If it not specified, the \$LATEST version will be used.

You can also add the `X-Amz-Invocation-Type: Event | RequestResponse | DryRun` header to have the action invoked asynchronously, as request and response, or as a test run, respectively. If the header is not specified, the action will be invoked as request and response. For the example shown here, this header has the default value.

Amazon API Gateway Developer Guide
Create a GET Method with Query Strings to Call the Lambda Function

We will come back to setting up **Mapping Templates** after setting up the query string parameters to hold the input data for the Lambda function.

2. In **Method Request** for the **GET** method on **/calc**, expand the **URL Query String Parameters** section. Choose **Add query string** to add the **operand1**, **operand2**, and **operator** query string parameters.

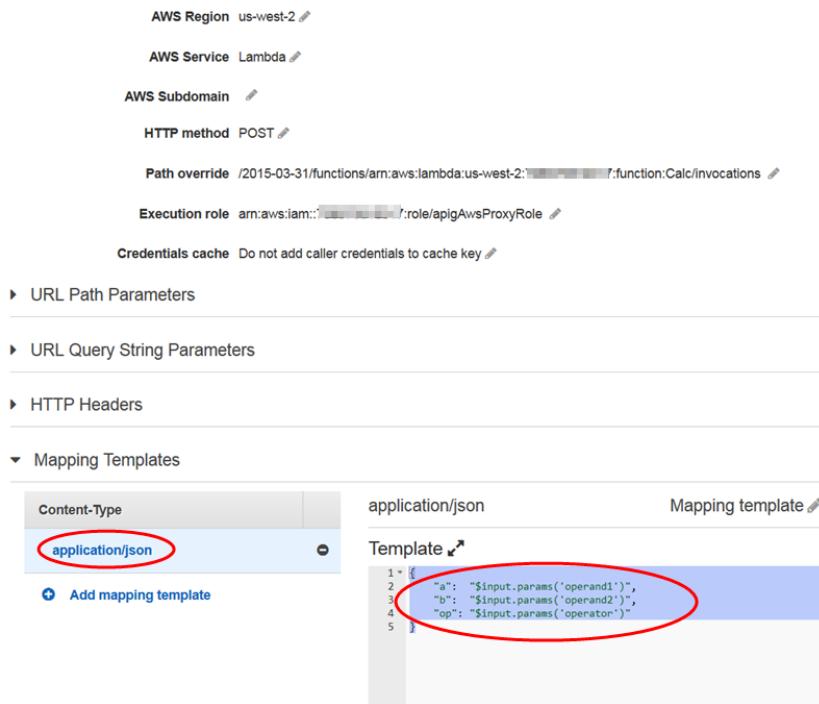


The screenshot shows the AWS API Gateway interface for creating a method request. At the top, it says 'Method Execution /calc - GET - Method Request' and 'Delete Method'. Below that, there's a note about authorization settings: 'Provide information about this method's authorization settings and the parameters it can receive.' Under 'Authorization Settings', it shows 'Authorization NONE' and 'API Key Required false'. The main focus is the 'URL Query String Parameters' section, which is expanded. It contains three rows: 'operand2' (with a red circle around it), 'operator', and 'operand1'. Each row has a 'Name' column, a 'Caching' column with a checkbox, and a delete icon. At the bottom of this section is a blue 'Add query string' button. Below this section are collapsed sections for 'HTTP Request Headers' and 'Request Models'.

3. Go back to **Integration Request**. Expand the **Mapping Templates** section. If necessary, in **Content-Type**, under **application/json**, choose **Add mapping template**. Type the following in the **Mapping template** editor:

```
{  
    "a": "$input.params('operand1')",  
    "b": "$input.params('operand2')",  
    "op": "$input.params('operator')"  
}
```

Amazon API Gateway Developer Guide
Create a POST Method with a JSON Payload to Call the Lambda Function



AWS Region us-west-2

AWS Service Lambda

AWS Subdomain

HTTP method POST

Path override /2015-03-31/functions/arn:aws:lambda:us-west-2:
Execution role arn:aws:iam::
Credentials cache Do not add caller credentials to cache key

▶ URL Path Parameters

▶ URL Query String Parameters

▶ HTTP Headers

▼ Mapping Templates

Content-Type	application/json	Mapping template
application/json		

Add mapping template

Template

```
1 var {  
2   "a": "$input.params('operand1')",  
3   "b": "$input.params('operand2')",  
4   "op": "$input.params('operator')"  
5 }
```

This template maps the three query string parameters declared in **Method Request** into designated property values of the JSON object as the input to the back-end Lambda function. The transformed JSON object will be included as the integration request payload.

4. You can now choose **Test** to verify that the GET method on the **/calc** resource has been properly set up to invoke the Lambda function.

Create a POST Method with a JSON Payload to Call the Lambda Function

The following steps describe how to expose a POST method with a JSON payload.

To set up the POST method with a JSON payload to invoke a Lambda function

1. Choose **Create Method** in the API Gateway console to create a POST method for the **LambdaGate** API's **/calc** resource.

In the method's **Set Up** panel, configure the POST method with the following settings.

Amazon API Gateway Developer Guide
Create a POST Method with a JSON Payload to Call the Lambda Function

◀ Method Execution /calc - POST - Integration Request Delete Method

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-west-2

AWS Service Lambda

AWS Subdomain

HTTP method POST

Path override /2015-03-31/functions/arn:aws:lambda:us-west-2:REDACTED:function:Calc/invocations

Execution role arn:aws:iam:REDACTED:role/apiGwProxyRole

Credentials cache Do not add caller credentials to cache key

▶ URL Path Parameters

▶ URL Query String Parameters

▶ HTTP Headers

▶ Mapping Templates

Using a POST request with a JSON payload is the simplest way to invoke a Lambda function, because no mappings are needed.

2. You can now choose **Test** to verify the POST method works as expected. The following input:

```
{  
    "a": 1,  
    "b": 2,  
    "op": "+"  
}
```

should produce the following output:

```
{  
    "a": 1,  
    "b": 2,  
    "op": "+",  
    "c": 3  
}
```

If you would like to implement POST as an asynchronous call, you can add an `InvocationType:Event` header in the method request and map it to the `X-Amz-Invocation-Type` header in the integration request, using the header mapping expression of `method.request.header.InvocationType`. You must inform the clients to include the `InvocationType:Event` header in the method request. Alternatively, you can set the `X-Amz-Invocation-Type` header with the 'Event' string literal in the integration

request, without requiring the client to include the header. The asynchronous call will return an empty response, instead.

Create a GET Method with Path Parameters to Call the Lambda Function

The following steps describe how to set up the GET method with path parameters to call the Lambda function.

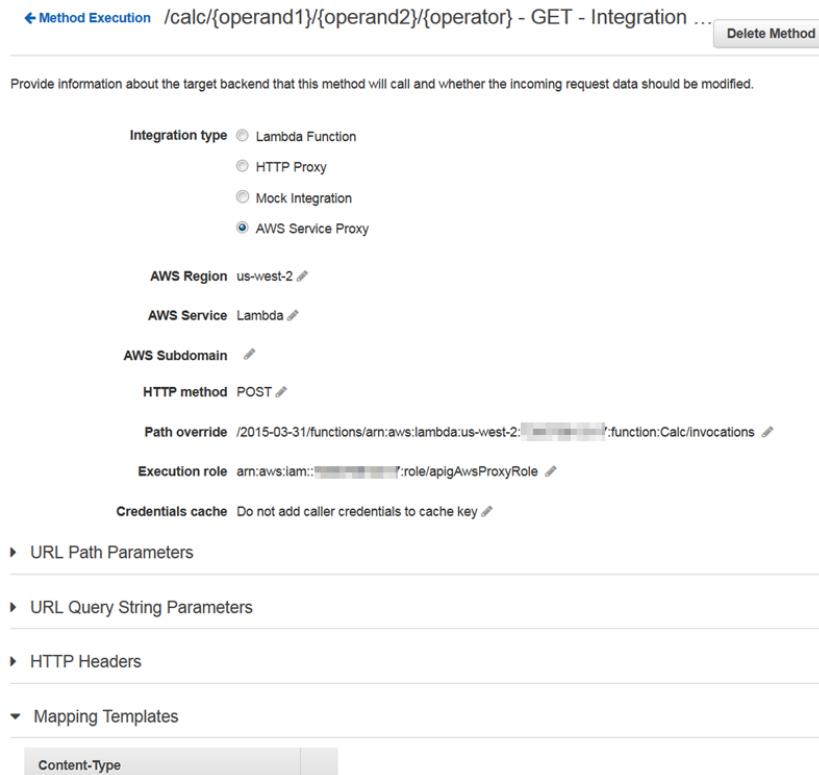
To set up the GET method with URL path parameters to invoke the Lambda function

Next, expand the **Mapping Templates** section, highlight the **application/json** entry under the **Content-Type** header (of integration response), open the **Mapping template** editor, enter and save the following mapping script:

```
$input.path('$a') $input.path('$op') $input.path('$b') = $input.path('$c')
```

1. Choose **Create Method** in the API Gateway console to create a GET method for the API's **/calc/{operand1}/{operand2}/{operator}** resource.

In the method's **Set up** pane, configure this GET method with the following settings.



Next, we will set up **Mapping Templates** to translate the URL path parameters into the integration request JSON payload as the input to the Lambda function.

Amazon API Gateway Developer Guide
Create a GET Method with Path Parameters to Call the Lambda Function

2. In **Method Request** for the **GET** method on **/calc**, expand the **URL Query String Parameters** section. Choose **Add query string** to add the **operand1**, **operand2**, and **operator** query string parameters.

[Method Execution](#) /calc - GET - Method Request Delete Method

Provide information about this method's authorization settings and the parameters it can receive.

Authorization Settings

Authorization	NONE
API Key Required	false

URL Query String Parameters

Name	Caching
operand2	<input type="checkbox"/>
operator	<input type="checkbox"/>
operand1	<input type="checkbox"/>

[Add query string](#)

HTTP Request Headers

Request Models [Create a Model](#)

3. Go back to **Integration Request**. Expand the **Mapping Templates** section. If necessary, in **Content-Type**, under **application/json**, choose **Add mapping template**.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-west-2

AWS Service Lambda

AWS Subdomain

HTTP method POST

Path override /2015-03-31/functions/arm:aws:lambda:us-west-2:REDACTED:function:Calc/invocations

Execution role arn:aws:iam::REDACTED:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

URL Path Parameters

URL Query String Parameters

HTTP Headers

Mapping Templates

Content-Type	application/json	Mapping template
--------------	------------------	------------------

[application/json](#)

Template Select a model to generate a template

```
1 <!
2   "a": "$input.params('operand1')",
3   "b": "$input.params('operand2')",
4   "op": #if($input.params('operator')=='%2F')#${else}"$input
5   params('operator')#end
6 }
```

Amazon API Gateway Developer Guide
Create a GET Method with Path Parameters to Call the Lambda Function

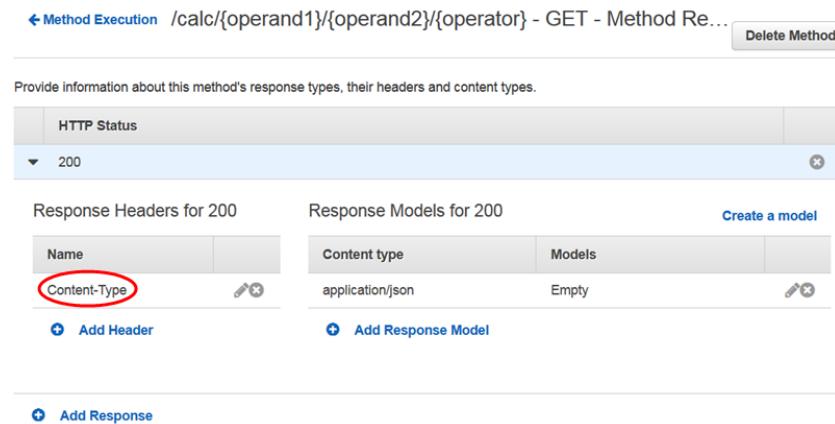
Type the following in the **Mapping Template** editor:

```
{  
    "a": "$input.params('operand1')",  
    "b": "$input.params('operand2')",  
    "op": "#if($input.params('operator')=='%2F')/${input.params('operator')}#end"  
}
```

This template maps the three URL path parameters, declared when the **/calc/{operand1}/{operand2}/{operator}** resource was created, into designated property values of the JSON object. Because URL paths must be URL-encoded, the division operator must be specified as %2F instead of /. This template maps these translations as well. The transformed JSON object will be included as the integration request payload.

4. As another exercise, we demonstrate how to translate the JSON returned from the Lambda function to show the output as a plain text string to the caller. This involves resetting the method request's Content-Type header to "text/plain" and providing a mapping template to translate the JSON output into a plain string.

First, make sure that **Content-Type** header is included in the **Response Headers for 200** section in **Method Response**.



HTTP Status	
▼ 200	X
Response Headers for 200	
Name	Content-type
Content-Type	application/json
+ Add Header	+ Add Response Model
Response Models for 200	
Content type	Models
application/json	Empty
+ Create a model	+ Add Response Model

In **Integration Response**, expand the 200 method response entry. Expand the **Header Mappings** section. In **Mapping value** for **Content-Type**, type '**'text/plain'**'. This header mapping expression overrides the Content-Type header with a literal string, which must be enclosed within a pair of single quotes.

Amazon API Gateway Developer Guide

Create a GET Method with Path Parameters to Call the Lambda Function

[Method Execution](#) /calc/{operand1}/{operand2}/{operator} - GET - Integration ... [Delete Method](#)

First, declare response types using [Method Response](#). Then, map the possible responses from the backend to this method's response types.

HTTP status regex	Method response status	Output model	Default mapping
200	200		Yes

Map the output from your HTTP endpoint to the headers and output model of the 200 method response.

HTTP status regex [default](#)

Method response status200

[Cancel](#) [Save](#)

[Header Mappings](#)

Response header	Mapping value
Content-Type	'text/plain'

[Mapping Templates](#)

Content-Type	application/json	Mapping template
application/json		Edit

[Template](#)

```
1 $input.path('$a') $input.path('$op') $input.path('$b') +  
$input.path('$c')
```

5. Choose **Test** to verify the GET method on the **/calc/{operand1}/{operand2}/{operator}** works as expected. The following input:

```
{  
    "a": 1,  
    "b": 2,  
    "op": "2%F"  
}
```

should produce the following plain text output:

```
1 / 2 = 0.5
```

6. After testing the API using the Test Invoke in the API Gateway console, you must deploy the API to make it public available. If you update the API, such as adding, modifying or deleting a resource or method, updating any data mapping, you must redeploy the API to make the new features or updates available.

A Sample API as a Lambda Proxy in Swagger with API Gateway Extensions

```
{  
    "swagger": "2.0",  
    "info": {  
        "version": "2016-02-23T05:35:54Z",  
        "title": "LambdaGate"  
    },  
    "host": "a123456789.execute-api.us-east-1.amazonaws.com",  
    "basePath": "/test",  
    "schemes": [  
        "https"  
    ],  
    "paths": {  
        "/calc": {  
            "get": {  
                "produces": [  
                    "application/json"  
                ],  
                "parameters": [  
                    {  
                        "name": "operand2",  
                        "in": "query",  
                        "required": false,  
                        "type": "string"  
                    },  
                    {  
                        "name": "operator",  
                        "in": "query",  
                        "required": false,  
                        "type": "string"  
                    },  
                    {  
                        "name": "operand1",  
                        "in": "query",  
                        "required": false,  
                        "type": "string"  
                    }  
                ],  
                "responses": {  
                    "200": {  
                        "description": "200 response",  
                        "schema": {  
                            "$ref": "#/definitions/Empty"  
                        },  
                        "headers": {  
                            "operand_1": {  
                                "type": "string"  
                            },  
                            "operand_2": {  
                                "type": "string"  
                            },  
                            "operator": {  
                                "type": "string"  
                            }  
                        }  
                    }  
                }  
            }  
        }  
    }  
}
```

```

        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
    "responses": {
        "default": {
            "statusCode": "200",
            "responseParameters": {
                "method.response.header.operator": "integration.re
sponse.body.op",
                "method.response.header.operand_2": "integration.re
sponse.body.b",
                "method.response.header.operand_1": "integration.response.body.a"
            },
            "responseTemplates": {
                "application/json": "#set($res= $input.path('$'))\n{\n\"result\": \"$res.a, $res.b, $res.op => $res.c\"\n}"
            }
        },
        "requestTemplates": {
            "application/json": "{\n    \"a\": \"$input.params('operand1')\", \n    \"b\": \"$input.params('operand2')\", \n    \"op\": \"$input.params('oper
ator')\"\n}"
        },
        "uri": "arn:aws:apigateway:us-west-2:lambda:path//2015-03-31/func
tions/arn:aws:lambda:us-west-2:123456789012:function:Calc/invocations",
        "httpMethod": "POST",
        "type": "aws"
    }
},
"post": {
    "produces": [
        "application/json"
    ],
    "parameters": [],
    "responses": {
        "200": {
            "description": "200 response",
            "schema": {
                "$ref": "#/definitions/Empty"
            },
            "headers": {}
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
        "responses": {
            "default": {
                "statusCode": "200",
                "responseTemplates": {
                    "application/json": "__passthrough__"
                }
            }
        }
    }
}

```

```

        "uri": "arn:aws:apigateway:us-west-2:lambda:path//2015-03-31/functions/arn:aws:lambda:us-west-2:123456789012:function:Calc/invocations",
        "httpMethod": "POST",
        "type": "aws"
    }
},
"/calc/{operand1}/{operand2}/{operator}": {
    "get": {
        "produces": [
            "application/json"
        ],
        "parameters": [
            {
                "name": "operand2",
                "in": "path",
                "required": true,
                "type": "string"
            },
            {
                "name": "operator",
                "in": "path",
                "required": true,
                "type": "string"
            },
            {
                "name": "operand1",
                "in": "path",
                "required": true,
                "type": "string"
            }
        ],
        "responses": {
            "200": {
                "description": "200 response",
                "schema": {
                    "$ref": "#/definitions/Empty"
                },
                "headers": {
                    "Content-Type": {
                        "type": "string"
                    }
                }
            }
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::123456789012:role/apigAwsProxyRole",
        "responses": {
            "default": {
                "statusCode": "200",
                "responseParameters": {
                    "method.response.header.Content-Type": "'text/plain'"
                },
                "responseTemplates": {
                    "application/json": "\"$input.path('$a') $input.path('$op')\n$input.path('$b') = $input.path('$c')\""
                }
            }
        }
    }
}

```

```

        },
        "requestTemplates": {
            "application/json": "\n{n\n  \"a\": \"$input.params('operand1')\",\\n\n\"b\": \"$input.params('operand2')\",\\n  \"op\": #if($input.params('operator')=='%2F')/#else{$input.params('operator')}#end\\n  \\n}"
        },
        "uri": "arn:aws:apigateway:us-west-2:lambda:path//2015-03-31/functions/arn:aws:lambda:us-west-2:123456789012:function:Calc/invocations",
        "httpMethod": "POST",
        "type": "aws"
    }
}
},
"definitions": {
    "Empty": {
        "type": "object"
    }
}
}

```

Create an API Gateway API as an Amazon Kinesis Proxy

This section describes how to create and configure an API Gateway API as an AWS proxy to access Amazon Kinesis.

For the purpose of illustration, we will create an example API to enable a client to do the following:

1. List the user's available streams in Amazon Kinesis
2. Create, describe, or delete a specified stream
3. Read data records from or write data records into the specified stream

To accomplish the preceding tasks, the API exposes methods on various resources to invoke the following, respectively:

1. The [ListStreams](#) action in Amazon Kinesis
2. The [CreateStream](#), [DescribeStream](#), or [DeleteStream](#) action
3. The [GetRecords](#) or [PutRecords](#) (including [PutRecord](#)) action in Amazon Kinesis

Specifically, we will build the API as follows:

- Expose an HTTP GET method on the API's `/streams` resource and integrate the method with the [ListStreams](#) action in Amazon Kinesis to list the streams in the caller's account.
- Expose an HTTP POST method on the API's `/streams/{stream-name}` resource and integrate the method with the [CreateStream](#) action in Amazon Kinesis to create a named stream in the caller's account.
- Expose an HTTP GET method on the API's `/streams/{stream-name}` resource and integrate the method with the [DescribeStream](#) action in Amazon Kinesis to describe a named stream in the caller's account.

- Expose an HTTP DELETE method on the API's `/streams/{stream-name}` resource and integrate the method with the [DeleteStream](#) action in Amazon Kinesis to delete a stream in the caller's account.
- Expose an HTTP PUT method on the API's `/streams/{stream-name}/record` resource and integrate the method with the [PutRecord](#) action in Amazon Kinesis. This enables the client to add a single data record to the named stream.
- Expose an HTTP PUT method on the API's `/streams/{stream-name}/records` resource and integrate the method with the [PutRecords](#) action in Amazon Kinesis. This enables the client to add a list of data records to the named stream.
- Expose an HTTP GET method on the API's `/streams/{stream-name}/records` resource and integrate the method with the [GetRecords](#) action in Amazon Kinesis. This enables the client to list data records in the named stream, with a specified shard iterator. A shard iterator specifies the shard position from which to start reading data records sequentially.
- Expose an HTTP GET method on the API's `/streams/{stream-name}/sharditerator` resource and integrate the method with the [GetShardIterator](#) action in Amazon Kinesis. This helper method must be supplied to the [ListStreams](#) action in Amazon Kinesis.

You can apply the instructions presented here to other Amazon Kinesis actions. For the complete list of the Amazon Kinesis actions, see [Amazon Kinesis API Reference](#).

Instead of using the API Gateway console to create the sample API, you can import the sample API into API Gateway, using either the API Gateway [Import API](#) or the [API Gateway Swagger Importer](#). For information on how to use the Import API, see [Import an API \(p. 117\)](#). For information on how to use the API Gateway Swagger Importer, see [Getting Started with the API Gateway Swagger Importer](#).

If you do not have an AWS account, use the following procedure to create one.

To sign up for AWS

1. Open <http://aws.amazon.com/> and choose **Create an AWS Account**.
2. Follow the online instructions.

Create an IAM Role and Policy for the API to Access Amazon Kinesis

To allow the API to invoke Amazon Kinesis actions, you must have appropriate IAM policies attached to an IAM role. This section explains how to verify and to create, if necessary, the required IAM role and policies.

To enable read-only access to Amazon Kinesis, you can use the `AmazonKinesisReadOnlyAccess` policy that allows the `Get*`, `List*`, and `Describe*` actions in Amazon Kinesis to be invoked.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "kinesis:Get*",  
                "kinesis>List*",  
                "kinesis:Describe*"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

```
        ]
    }
```

This policy is available from the IAM console and its ARN is
arn:aws:iam::aws:policy/AmazonKinesisReadOnlyAccess.

To enable read-write actions in Amazon Kinesis, you can use the AmazonKinesisFullAccess policy.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "kinesis:*",
            "Resource": "*"
        }
    ]
}
```

This policy is also available from the IAM console. Its ARN is
arn:aws:iam::aws:policy/AmazonKinesisFullAccess.

After you decide which IAM policy to use, attach it to a new or existing IAM role. Make sure that the API Gateway control service (`apigateway.amazonaws.com`) is a trusted entity of the role and is allowed to assume the execution role (`sts:AssumeRole`).

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "",
            "Effect": "Allow",
            "Principal": {
                "Service": "apigateway.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

If you create the execution role in the IAM console and choose the **Amazon API Gateway** role type, this trust policy is automatically attached.

Note the ARN of the execution role. You will need it when creating an API method and setting up its integration request.

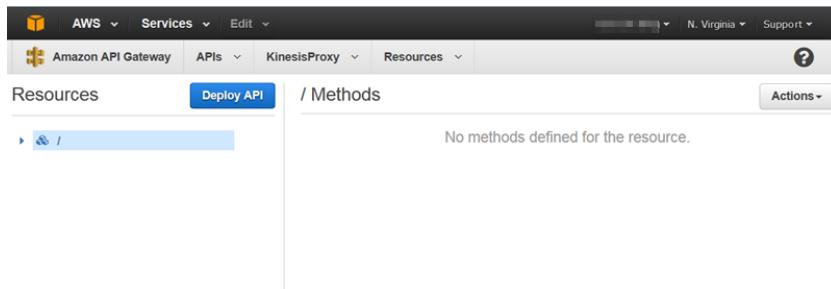
Start to Create an API as an Amazon Kinesis Proxy

Use the following steps to create the API in the API Gateway console.

To create an API as an AWS service proxy for Amazon Kinesis

1. In the API Gateway console, choose **Create API**.
2. In **API name**, type **KinesisProxy**. Leave the default values in the other fields.
3. For **Clone from API**, choose **Do not clone from existing API**.
4. Type a description in **Description**.
5. Choose **Create API**.

After the API is created, the API Gateway console displays the **Resources** page, which contains only the API's root (/) resource.



List Streams in Amazon Kinesis

To list streams in Amazon Kinesis, add a /streams resource to the API's root, expose a GET method on the resource, and integrate the method to the `ListStreams` action of Amazon Kinesis.

The following procedure describes how to list Amazon Kinesis streams by using the API Gateway console.

To list Amazon Kinesis streams by using the API Gateway console

1. Select the API root resource. In **Actions**, choose **Create Resource**.
In **Resource Name**, type **Streams**, leave **Resource Path** as the default, and choose **Create Resource**.
2. Select the /Streams resource. From **Actions**, choose **Create Method**, choose **GET** from the list, and then choose the checkmark icon to finish creating the method.

Note

You can choose any of the available HTTP verbs for a method request. We use `GET` here, because listing streams is a READ operation.

3. In the method's **Setup** pane, choose **Show Advanced** and then choose **AWS Service Proxy**.
 - a. For **AWS Region**, choose a region (e.g., **us-east-1**).
 - b. For **AWS Service**, choose **Kinesis**.
 - c. For **HTTP method**, choose **POST**.

Note

For the integration request with Amazon Kinesis, you must choose the `POST` HTTP verb to invoke the action, although you can use any of the available HTTP verbs for the API's method request.

- d. For **Action Type**, choose **Use action name**.
- e. For **Action**, type `ListStreams`.
- f. For **Execution role**, type the ARN for your execution role.
- g. Choose **Save** to finish the initial setup of the method.

[Method Execution](#) /streams - GET - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1

AWS Service Kinesis

AWS Subdomain

HTTP method POST

Action ListStreams

Execution role arn:aws:iam::██████████:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

The initial setup of the integration request will suffice if there is no need to map data between the method and integration requests and/or between the method and integration responses. Examples discussed in this topic require data mapping, which is covered in the second half of the **Integration Request** pane.

4. In the **Integration Request** pane, expand the **HTTP Headers** section:
 - a. Choose **Add header**.
 - b. In the **Name** column, type **Content-Type**.
 - c. In the **Mapped from** column, type '**application/x-amz-json-1.1**'.
 - d. Choose the checkmark icon to save the setting.
5. Expand the **Body Mapping Templates** section:
 - a. Choose **Add mapping template**.
 - b. For **Content-Type**, type **application/json**.
 - c. Choose the checkmark icon to save the setting.
 - d. Choose the pencil icon to the right of **Mapping template**.
 - e. Choose **Mapping template** from the drop-down list to open the **Template** editor.
 - f. Type **{}** in the template editor.
 - g. Choose the checkmark icon to save the mapping template.

The [ListStreams](#) request takes a payload of the following JSON format:

```
{  
    "ExclusiveStartStreamName": "string",  
    "Limit": number  
}
```

However, the properties are optional. To use the default values, we opted for an empty JSON payload here.

The screenshot shows the AWS Lambda function configuration interface. Under the 'HTTP Headers' section, there is a table with one row: 'Content-Type' mapped from 'application/x-amz-json-1.1'. Under the 'Body Mapping Templates' section, there is a table with one entry for 'Content-Type: application/json', which maps to the template '{ }'.

6. Test the GET method on the Streams resource to invoke the `ListStreams` action in Amazon Kinesis:

From the API Gateway console, select the `/streams/GET` entry from the **Resources** pane, choose the **Test** invocation option, and then choose **Test**.

If you have already created two streams named "myStream" and "yourStream" in Amazon Kinesis, the successful test will return a 200 OK response containing the following payload:

```
{
    "HasMoreStreams": false,
    "StreamNames": [
        "myStream",
        "yourStream"
    ]
}
```

Create, Describe, and Delete a Stream in Amazon Kinesis

Creating, describing, and deleting a stream in Amazon Kinesis involves making the following Amazon Kinesis REST API requests, respectively:

```
POST /?Action=CreateStream HTTP/1.1
Host: kinesis.region.domain
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "ShardCount": number,
```

```
        "StreamName": "string"
    }
```

```
POST /?Action=DescribeStream HTTP/1.1
Host: kinesis.region.domain
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "ExclusiveStartShardId": "string",
    "Limit": number,
    "StreamName": "string"
}
```

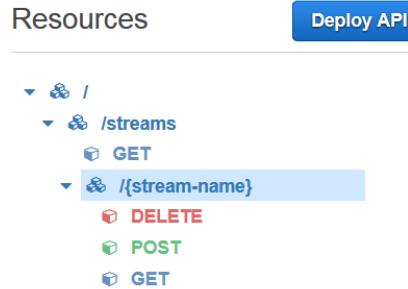
```
POST /?Action=DeleteStream HTTP/1.1
Host: kinesis.region.domain
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "StreamName": "string"
}
```

We can build our API to accept the required input as a JSON payload of the method request and pass the payload through to the integration request. However, to provide more examples of data mapping between method and integration requests, and method and integration responses, we will create our API slightly differently.

We will expose the `GET`, `POST`, and `Delete` HTTP methods on a to-be-named `Stream` resource. We will use the `{stream-name}` path variable to hold the to-be-named stream resource and integrate these API methods with the Amazon Kinesis' `DescribeStream`, `CreateStream`, and `DeleteStream` actions, respectively. We require that the client pass other input data as headers, query parameters, or the payload of a method request, and we provide mapping templates to transform the data to the required integration request payload.

After the methods are created on a to-be-named stream resource, the structure of the API looks like the following:



To configure and test the GET method on a stream resource

1. Set up the GET method to describe a named stream in Amazon Kinesis, as shown in the following.

[Method Execution](#) /streams/{stream-name} - GET - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1

AWS Service Kinesis

AWS Subdomain

HTTP method POST

Action [DescribeStream](#)

Execution role arn:aws:iam::7...:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

2. Map data from the GET method request to the integration request, as shown in the following:

HTTP method **POST**

Action **DescribeStream**

Execution role **arn:aws:iam::717171717171:role/apigAwsProxyRole**

Credentials cache **Do not add caller credentials to cache key**

- ▶ URL Path Parameters
- ▶ URL Query String Parameters
- ▶ HTTP Headers
- ▼ Body Mapping Templates

Content-Type	application/json	Mapping template
application/json	✖	Template ↗ <pre>1 ▼ { 2 "StreamName": "\$input.params('stream-name')" 3 }</pre>
+ Add mapping template		

3. Test the GET method to invoke the `DescribeStream` action in Amazon Kinesis:

From the API Gateway console, select `/streams/{stream-name}/GET` in the **Resources** pane, choose **Test** to start testing, type the name of an existing Amazon Kinesis stream in the **Path** field for `stream-name`, and choose **Test**. If the test is successful, a 200 OK response is returned with a payload similar to the following:

```
{
  "StreamDescription": {
    "HasMoreShards": false,
    "RetentionPeriodHours": 24,
    "Shards": [
      {
        "HashKeyRange": {
          "EndingHashKey": "68056473384187692692674921486353642290",
          "StartingHashKey": "0"
        },
        "SequenceNumberRange": {
          "StartingSequenceNumber": "49559266461454070523309915164834022007924120923395850242",
          "EndingSequenceNumber": "49559266461543273504104037657400164881014714369419771970"
        },
        "ShardId": "shardId-000000000000"
      },
      ...
    ],
    "StreamARN": "arn:aws:kinesis:us-east-1:123456789012:stream/test-stream"
  }
}
```

```
        },
        "ShardId": "shardId-000000000004"
    },
],
"StreamARN": "arn:aws:kinesis:us-east-1:12345678901:stream/myStream",
"StreamName": "myStream",
"StreamStatus": "ACTIVE"
}
}
```

After you deploy the API, you can make a REST request against this API method:

```
GET https://your-api-id.execute-api.region.amazonaws.com/stage/streams/myStream
HTTP/1.1
Host: your-api-id.execute-api.region.amazonaws.com
Content-Type: application/json
Authorization: ...
X-Amz-Date: 20160323T194451Z
```

To configure and test the POST method on a stream resource

1. Set up the POST method on a stream resource to create the stream in Amazon Kinesis, as shown in the following:

[!\[\]\(1763d178402133739f19689796b980d4_img.jpg\) Method Execution](#)
/streams/{stream-name} - POST - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1 

AWS Service Kinesis 

AWS Subdomain 

HTTP method POST 

Action CreateStream 

Execution role arn:aws:iam::7...:role/apigAwsProxyRole 

Credentials cache Do not add caller credentials to cache key 

2. Map data from the POST method request to the integration request, as shown in the following:

HTTP method **POST**

Action **CreateStream**

Execution role **arn:aws:iam::7...:role/apigAwsProxyRole**

Credentials cache **Do not add caller credentials to cache key**

- ▶ URL Path Parameters
- ▶ URL Query String Parameters
- ▼ HTTP Headers

Name	Mapped from	Caching	
Content-Type	'application/x-amz-json-1.1'		

+ Add header

- ▼ Body Mapping Templates

Content-Type	application/json	Mapping template
application/json		Template ↗ <pre> 1 { 2 "ShardCount": 5, 3 "StreamName": "\$input.params('stream-name')" 4 }</pre>

+ Add mapping template

In this example, we set the `ShardCount` to a fixed value in the provided mapping template. We could also require the client to supply a JSON payload of `{ "ShardCount" : 5 }` to the method request. In this case, you would use the following mapping template for the integration request.

```
{
    "ShardCount": $input.path('$ShardCount'),
    "StreamName": "$input.params('stream-name')"
}
```

3. Test the POST method to create a named stream in Amazon Kinesis:

From the API Gateway console, select **/streams/{stream-name}/POST** in the **Resources** pane, choose **Test** to start testing, type the name of an existing Amazon Kinesis stream in **Path** for `stream-name`, and choose **Test**. If the test is successful, a 200 OK response is returned with no data.

After you deploy the API, you can also make a REST API request against the POST method on a Stream resource to invoke the `CreateStream` action in Amazon Kinesis:

```

POST https://your-api-id.execute-api.region.amazonaws.com/stage/streams/yourStream HTTP/1.1
Host: your-api-id.execute-api.region.amazonaws.com
Content-Type: application/json
Authorization: ...
X-Amz-Date: 20160323T194451Z

```

```
{  
    "ShardCount": 5  
}
```

Configure and test the DELETE method on a stream resource

1. Set up the DELETE method to invoke the `DeleteStream` action in Amazon Kinesis, as shown in the following.

[Method Execution](#)
`/streams/{stream-name}` - DELETE - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1

AWS Service Kinesis

AWS Subdomain

HTTP method POST

Action DeleteStream

Execution role arn:aws:iam::7...7:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

2. Map data from the DELETE method request to the integration request, as shown in the following:

HTTP method **POST**

Action **DeleteStream**

Execution role **arn:aws:iam::738575810317:role/apigAwsProxyRole**

Credentials cache **Do not add caller credentials to cache key**

- ▶ URL Path Parameters
- ▶ URL Query String Parameters
- ▼ HTTP Headers

Name	Mapped from	Caching
Content-Type	'application/x-amz-json-1.1'	

+ Add header

- ▼ Body Mapping Templates

Content-Type	application/json	Mapping template
application/json		Template ↗ <pre>1 var { 2 "StreamName": "\$input.params('stream-name')" 3 }</pre>

+ Add mapping template

3. Test the DELETE method to delete a named stream in Amazon Kinesis:

From the API Gateway console, select the **/streams/{stream-name}/DELETE** method node in the **Resources** pane, choose **Test** to start testing, type the name of an existing Amazon Kinesis stream in **Path** for **stream-name**, and choose **Test**. If the test is successful, a 200 OK response is returned with no data.

After you deploy the API, you can also make the following REST API request against the DELETE method on the Stream resource to call the `DeleteStream` action in Amazon Kinesis:

```
DELETE https://your-api-id.execute-api.region.amazonaws.com/stage/streams/yourStream HTTP/1.1
Host: your-api-id.execute-api.region.amazonaws.com
Content-Type: application/json
Authorization: ...
X-Amz-Date: 20160323T194451Z

{ }
```

Get Records from and Add Records to a Stream in Amazon Kinesis

After you create a stream in Amazon Kinesis, you can add data records to the stream and read the data from the stream. Adding data records involves calling the [PutRecords](#) or [PutRecord](#) action in Amazon Kinesis. The former adds multiple records whereas the latter adds a single record to the stream.

```
POST /?Action=PutRecords HTTP/1.1
Host: kinesis.region.domain
Authorization: AWS4-HMAC-SHA256 Credential=..., ...
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "Records": [
        {
            "Data": blob,
            "ExplicitHashKey": "string",
            "PartitionKey": "string"
        }
    ],
    "StreamName": "string"
}
```

or

```
POST /?Action=PutRecord HTTP/1.1
Host: kinesis.region.domain
Authorization: AWS4-HMAC-SHA256 Credential=..., ...
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "Data": blob,
    "ExplicitHashKey": "string",
    "PartitionKey": "string",
    "SequenceNumberForOrdering": "string",
    "StreamName": "string"
}
```

Here, `StreamName` identifies the target stream to add records. `StreamName`, `Data`, and `PartitionKey` are required input data. In our example, we use the default values for all of the optional input data and will not explicitly specify values for them in the input to the method request.

Reading data in Amazon Kinesis amounts to calling the [GetRecords](#) action:

```
POST /?Action=GetRecords HTTP/1.1
Host: kinesis.region.domain
```

```
Authorization: AWS4-HMAC-SHA256 Credential=..., ...
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "ShardIterator": "string",
    "Limit": number
}
```

Here, the source stream from which we are getting records is specified in the required `ShardIterator` value, as is shown in the following Amazon Kinesis action to obtain a shard iterator:

```
POST /?Action=GetShardIterator HTTP/1.1
Host: kinesis.region.domain
Authorization: AWS4-HMAC-SHA256 Credential=..., ...
...
Content-Type: application/x-amz-json-1.1
Content-Length: PayloadSizeBytes

{
    "ShardId": "string",
    "ShardIteratorType": "string",
    "StartingSequenceNumber": "string",
    "StreamName": "string"
}
```

For the `GetRecords` and `PutRecords` actions, we expose the `GET` and `PUT` methods, respectively, on a `/records` resource that is appended to a named stream resource (`/{{stream-name}}`). Similarly, we expose the `PutRecord` action as a `PUT` method on a `/record` resource.

Because the `GetRecords` action takes as input a `ShardIterator` value, which is obtained by calling the `GetShardIterator` helper action, we expose a `GET` helper method on a `ShardIterator` resource (`/sharditerator`).

The following figure shows the API structure of resources after the methods are created:

Resources Deploy API

```

    - /
      - /streams
        - GET
        - DELETE
        - POST
        - GET
      - /records
        - GET
        - PUT
      - /record
        - PUT
      - /sharditerator
        - GET
  
```

The following four procedures describe how to set up each of the methods, how to map data from the method requests to the integration requests, and how to test the methods.

To configure and test the PUT method on the record resource in the API to invoke the PutRecord action in Amazon Kinesis:

1. Set up the PUT method, as shown in the following:

[Method Execution](#) /streams/{stream-name}/record - PUT - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1

AWS Service **Kinesis**

AWS Subdomain

HTTP method **POST**

Action **PutRecord**

Execution role arn:aws:iam::7...:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

2. Configure data mapping for the PUT-on-Record method, as shown in the following:

Amazon API Gateway Developer Guide

Get Records from and Add Records to a Stream in Amazon Kinesis

Action **PutRecord** 

Execution role arn:aws:iam::7...7:role/apigAwsProxyRole 

Credentials cache Do not add caller credentials to cache key 

▶ URL Path Parameters

▶ URL Query String Parameters

▼ HTTP Headers

Name	Mapped from 	Caching
Content-Type	'application/x-amz-json-1.1'	 

 [Add header](#)

▼ Body Mapping Templates

Content-Type	application/json	Mapping template 
application/json	 	  1 « { 2 "StreamName": "\$input.params('stream-name')", 3 "Data": "\$util.base64Encode(\$input.path('\$Data'))", 4 "PartitionKey": "\$input.path('\$PartitionKey')" 5 }

 [Add mapping template](#)

The preceding mapping template assumes that the method request payload is of the following format:

```
{  
  "Data": "some data",  
  "PartitionKey": "some key"  
}
```

This data can be modeled by the following JSON schema:

```
{  
  "$schema": "http://json-schema.org/draft-04/schema#",  
  "title": "PutRecord proxy single-record payload",  
  "type": "object",  
  "properties": {  
    "Data": { "type": "string" },  
    "PartitionKey": { "type": "string" }  
  }  
}
```

You can create a model to include this schema and use the model to facilitate generating the mapping template. However, you can generate a mapping template without using any model.

3. To test the PUT method, set the `stream-name` path variable to an existing stream, supply a payload of the preceding format, and then submit the method request. The successful result is a 200 OK response with a payload of the following format:

```
{  
  "SequenceNumber":  
  "49559409944537880850133345460169886593573102115167928386",  
  "ShardId": "shardId-000000000004"
```

}

To configure and test the PUT method on the records resource in the API to invoke the PutRecords action in Amazon Kinesis

1. Set up the PUT method, as shown in the following:

[◀ Method Execution /streams/{stream-name}/records - PUT - Integration Request](#)

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1

AWS Service **Kinesis**

AWS Subdomain

HTTP method **POST**

Action **PutRecords**

Execution role arn:aws:iam::7...:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

2. Configure data mapping for the PUT method, as shown in the following:

Action PutRecords

Execution role arn:aws:iam::7...:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

- ▶ URL Path Parameters
- ▶ URL Query String Parameters
- ▶ HTTP Headers
- ▼ Body Mapping Templates

Content-Type	application/json	Mapping template
application/json		<p>Template </p> <pre> 1 { 2 "StreamName": "\$input.params('stream-name')", 3 "Records": [4 #foreach(\$elem in \$input.path('\$..records')) 5 { 6 "Data": "\$util.base64Encode(\$elem.data)", 7 "PartitionKey": "\$elem.partition-key" 8 }#if(\$foreach.hasNext),#end 9] 10 } 11 </pre>

Add mapping template

The preceding mapping template assumes the method request payload can be modeled by following JSON schema:

```
{  
  "$schema": "http://json-schema.org/draft-04/schema#",  
  "title": "PutRecords proxy payload data",  
  "type": "object",  
  "properties": {  
    "records": {  
      "type": "array",  
      "items": {  
        "type": "object",  
        "properties": {  
          "data": { "type": "string" },  
          "partition-key": { "type": "string" }  
        }  
      }  
    }  
  }  
}
```

3. To test the PUT method, set the `stream-name` path variable to an existing stream, supply a payload as previously shown, and submit the method request. The successful result is a 200 OK response with a payload of the following format:

```
{  
  "records": [  
    {  
      "data": "some data",  
      "partition-key": "some key"  
    },  
    {  
      "data": "some other data",  
      "partition-key": "some key"  
    }  
  ]  
}
```

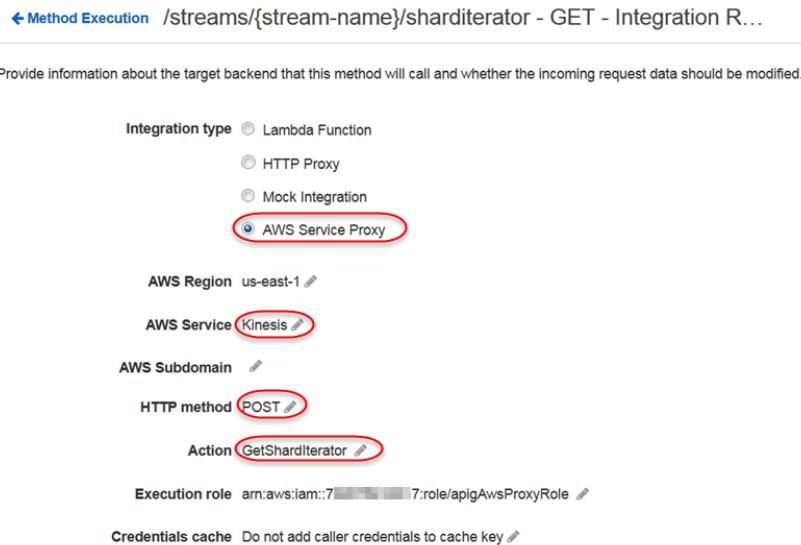
The response payload will be similar to the following output:

```
{  
  "FailedRecordCount": 0,  
  "Records": [  
    {  
      "SequenceNumber":  
"49559409944537880850133345460167468741933742152373764162",  
      "ShardId": "shardId-000000000004"  
    },  
    {  
      "SequenceNumber":  
"49559409944537880850133345460168677667753356781548470338",  
      "ShardId": "shardId-000000000004"  
    }  
  ]  
}
```

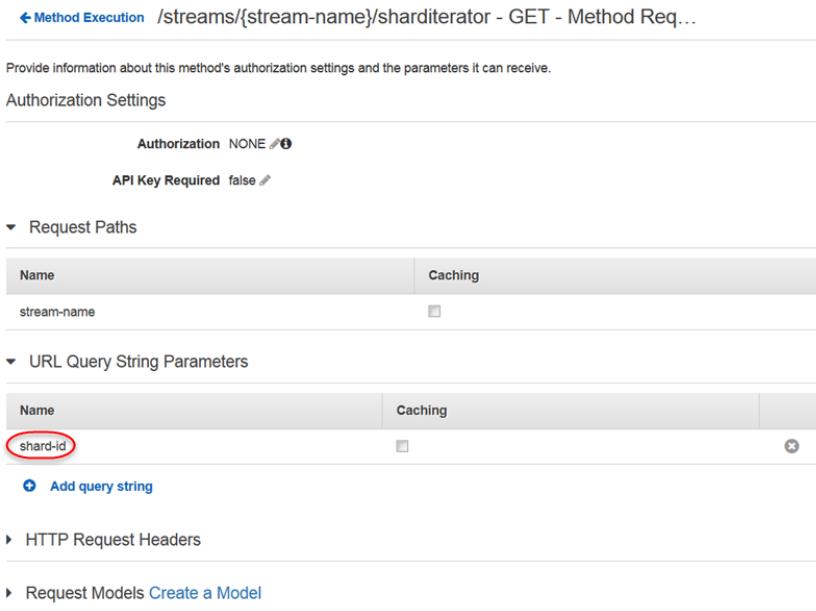
To configure and test the GET method on the ShardIterator resource in the API to invoke the GetShardIterator action in Amazon Kinesis

The GET-on-ShardIterator method is a helper method to acquire a required shard iterator before calling the GET-on-Records method.

1. Set up the GET-on-ShardIterator method, as shown in the following:



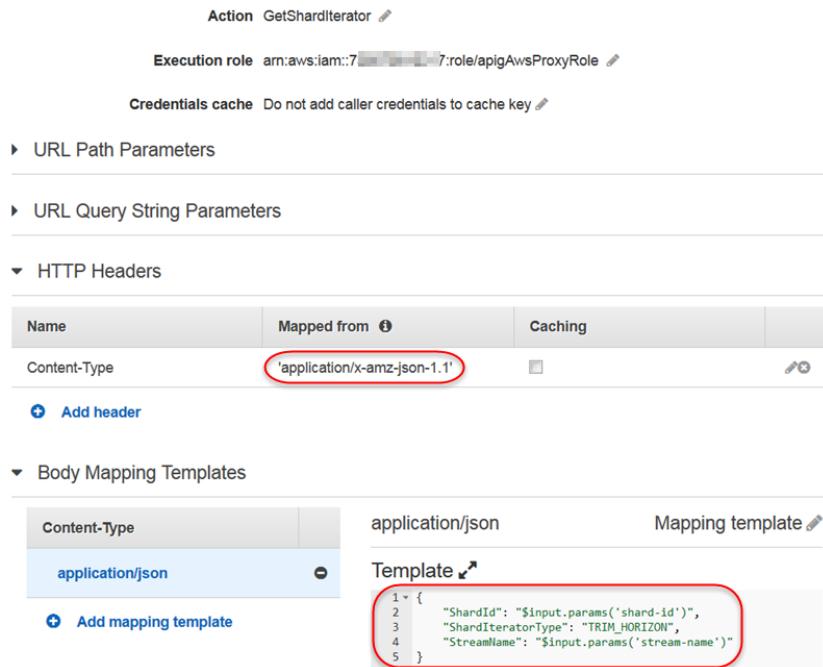
2. The `GetShardIterator` action requires an input of a `ShardId` value. To pass a client-supplied `ShardId` value, we add a `shard-id` query parameter to the method request, as shown in the following:



In the following mapping template, we add the translation of the `shard-id` query parameter value to the `ShardId` property value of the JSON payload for the `GetShardIterator` action in Amazon Kinesis.

Amazon API Gateway Developer Guide
Get Records from and Add Records to a Stream in
Amazon Kinesis

3. Configure data mapping for the GET-on-ShardIterator method:



The screenshot shows the 'Action' dropdown set to 'GetShardIterator'. Under 'HTTP Headers', the 'Content-Type' header is listed with a value of 'application/x-amz-json-1.1'. Below this, under 'Body Mapping Templates', there is a table with two rows. The first row has 'Content-Type' in the 'Content-Type' column and 'application/json' in the 'Mapping template' column. The second row has 'application/json' in the 'Content-Type' column and a 'Template' input field containing the following JSON template:

```
1+ {  
2 "ShardId": "$input.params('shard-id')",  
3 "ShardIteratorType": "TRIM_HORIZON",  
4 "StreamName": "$input.params('stream-name')"  
5 }
```

4. Using the **Test** option in the API Gateway console, enter an existing stream name as the **stream-name** **Path** variable value, set the **shard-id** **Query string** to an existing ShardID value (e.g., shard-000000000004), and choose **Test**.

The successful response payload will be similar to the following output:

```
{  
  "ShardIterator": "AAAAAAAAAAFYVN3V1Fy..."}
```

Make note of the `ShardIterator` value. You will need it to get records from a stream.

To configure and test the GET Method on the records resource in the API to invoke the GetRecords action in Amazon Kinesis

1. Set up the GET method, as shown in the following:

Amazon API Gateway Developer Guide

Get Records from and Add Records to a Stream in Amazon Kinesis

[Method Execution](#) /streams/{stream-name}/records - GET - Integration Request

Provide information about the target backend that this method will call and whether the incoming request data should be modified.

Integration type Lambda Function
 HTTP Proxy
 Mock Integration
 AWS Service Proxy

AWS Region us-east-1

AWS Service **Kinesis**

AWS Subdomain

HTTP method **POST**

Action **GetRecords**

Execution role arn:aws:iam::**7**:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

2. The GetRecords action requires an input of a ShardIterator value. To pass a client-supplied ShardIterator value, we add a Shard-Iterator header parameter to the method request, as shown in the following:

[Method Execution](#) /streams/{stream-name}/records - GET - Method Request

Provide information about this method's authorization settings and the parameters it can receive.

Authorization Settings

Authorization **NONE**

API Key Required **false**

Request Paths

Name	Caching
stream-name	

URL Query String Parameters

HTTP Request Headers

Name	Caching
Shard-Iterator	

[Add header](#)

Request Models [Create a Model](#)

In the following mapping template, we add the mapping from the Shard-Iterator header value to the shardIterator property value of the JSON payload for the GetRecords action in Amazon Kinesis.

3. Configure data mapping for the GET-on-Records method:

Action GetRecords

Execution role arn:aws:iam::7...:role/apigAwsProxyRole

Credentials cache Do not add caller credentials to cache key

- ▶ URL Path Parameters
- ▶ URL Query String Parameters
- ▼ HTTP Headers

Name	Mapped from	Caching
Content-Type	application/x-amz-json-1.1	<input type="checkbox"/>

Add header

- ▼ Body Mapping Templates

Content-Type	application/json	Mapping template
application/json		<p>Template </p> <pre>1+ { 2 "ShardIterator": "\$input.params('Shard-Iterator')" 3 }</pre>

Add mapping template

4. Using the **Test** option in the API Gateway console, type an existing stream name as the `stream-name` **Path** variable value, set the `Shard-Iterator` **Header** to the `ShardIterator` value obtained from the test run of the GET-on-ShardIterator method (above), and choose **Test**.

The successful response payload will be similar to the following output:

```
{
  "MillisBehindLatest": 0,
  "NextShardIterator": "AAAAAAAAAAF...",
  "Records": [ ... ]
}
```

Swagger Definitions of a Sample API as an Amazon Kinesis Proxy

```
{
  "swagger": "2.0",
  "info": {
    "version": "2016-03-31T18:25:32Z",
    "title": "KinesisProxy"
  },
  "host": "wd4zclrobb.execute-api.us-east-1.amazonaws.com",
  "basePath": "/test",
  "schemes": [
    "https"
  ],
  "paths": {
    "/streams": {
      "get": {
        "consumes": [
          "application/json"
        ]
      }
    }
  }
}
```

```

],
"produces": [
    "application/json"
],
"responses": {
    "200": {
        "description": "200 response",
        "schema": {
            "$ref": "#/definitions/Empty"
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::738575810317:role/apigAwsProxyRole",
    "responses": {
        "default": {
            "statusCode": "200"
        }
    },
    "requestTemplates": {
        "application/json": "{\n}"
    },
    "uri": "arn:aws:apigateway:us-east-1:kinesis:action/ListStreams",
    "httpMethod": "POST",
    "requestParameters": {
        "integration.request.header.Content-Type": "'application/x-amz-json-1.1'"
    },
    "type": "aws"
}
},
"/streams/{stream-name)": {
    "get": {
        "consumes": [
            "application/json"
        ],
        "produces": [
            "application/json"
        ],
        "parameters": [
            {
                "name": "stream-name",
                "in": "path",
                "required": true,
                "type": "string"
            }
        ],
        "responses": {
            "200": {
                "description": "200 response",
                "schema": {
                    "$ref": "#/definitions/Empty"
                }
            }
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::738575810317:role/apigAwsProxyRole",

```

```

    "responses": {
        "default": {
            "statusCode": "200"
        }
    },
    "requestTemplates": {
        "application/json": "{\n            \"StreamName\": \"$input.params('stream-\nname')\"\n        }\n    "
    },
    "uri": "arn:aws:apigateway:us-east-1:kinesis:action/DescribeStream",
    "httpMethod": "POST",
    "type": "aws"
},
"post": {
    "consumes": [
        "application/json"
    ],
    "produces": [
        "application/json"
    ],
    "parameters": [
        {
            "name": "stream-name",
            "in": "path",
            "required": true,
            "type": "string"
        }
    ],
    "responses": {
        "200": {
            "description": "200 response",
            "schema": {
                "$ref": "#/definitions/Empty"
            }
        }
    }
},
"x-amazon-apigateway-integration": {
    "credentials": "arn:aws:iam::738575810317:role/apigAwsProxyRole",
    "responses": {
        "default": {
            "statusCode": "200"
        }
    },
    "requestTemplates": {
        "application/json": "{\n            \"ShardCount\": 5,\n            \"StreamName\": \"$input.params('stream-name')\"\n        }\n    "
    },
    "uri": "arn:aws:apigateway:us-east-1:kinesis:action/CreateStream",
    "httpMethod": "POST",
    "requestParameters": {
        "integration.request.header.Content-Type": "'application/x-amz-json-\n1.1'"
    },
    "type": "aws"
}
},

```

```
  "delete": {
    "consumes": [
      "application/json"
    ],
    "produces": [
      "application/json"
    ],
    "parameters": [
      {
        "name": "stream-name",
        "in": "path",
        "required": true,
        "type": "string"
      }
    ],
    "responses": {
      "200": {
        "description": "200 response",
        "schema": {
          "$ref": "#/definitions/Empty"
        },
        "headers": {
          "Content-Type": {
            "type": "string"
          }
        }
      },
      "400": {
        "description": "400 response",
        "headers": {
          "Content-Type": {
            "type": "string"
          }
        }
      },
      "500": {
        "description": "500 response",
        "headers": {
          "Content-Type": {
            "type": "string"
          }
        }
      }
    },
    "x-amazon-apigateway-integration": {
      "credentials": "arn:aws:iam::738575810317:role/apigAwsProxyRole",
      "responses": {
        "4\d{2}": {
          "statusCode": "400",
          "responseParameters": {
            "method.response.header.Content-Type": "integration.response.header.Content-Type"
          }
        },
        "default": {
          "statusCode": "200",
          "responseParameters": {
            "method.response.header.Content-Type": "integration.response.head
```

```

        "er.Content-Type"
            }
        },
        "5\\d{2}": {
            "statusCode": "500",
            "responseParameters": {
                "method.response.header.Content-Type": "integration.response.header.Content-Type"
            }
        },
        "requestTemplates": {
            "application/json": "{\n                \"StreamName\": \"$input.params('stream-name')\"\n            }"
        },
        "uri": "arn:aws:apigateway:us-east-1:kinesis:action/DeleteStream",
        "httpMethod": "POST",
        "requestParameters": {
            "integration.request.header.Content-Type": "'application/x-amz-json-1.1'"
        },
        "type": "aws"
    }
},
"/streams/{stream-name}/record": {
    "put": {
        "consumes": [
            "application/json"
        ],
        "produces": [
            "application/json"
        ],
        "parameters": [
            {
                "name": "stream-name",
                "in": "path",
                "required": true,
                "type": "string"
            }
        ],
        "responses": {
            "200": {
                "description": "200 response",
                "schema": {
                    "$ref": "#/definitions/Empty"
                }
            }
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::738575810317:role/apigAwsProxyRole",
        "responses": {
            "default": {
                "statusCode": "200"
            }
        },
        "requestTemplates": {
            "application/json": "{\n                \"StreamName\": \"$input.params('stream-name')\"\n            }"
        }
    }
}

```

```

name')\",\n      \"Data\": \"$util.base64Encode($input.path('$.Data'))\",\\n
\"PartitionKey\": \"$input.path('$.PartitionKey')\"\n    },
    \"uri\": \"arn:aws:apigateway:us-east-1:kinesis:action/PutRecord\",
    \"httpMethod\": \"POST\",
    \"requestParameters\": {
        \"integration.request.header.Content-Type\": \"application/x-amz-json-
1.1\""
    },
    \"type\": \"aws"
}
},
\"/streams/{stream-name}/records\": {
    \"get\": {
        \"consumes\": [
            \"application/json\"
        ],
        \"produces\": [
            \"application/json\"
        ],
        \"parameters\": [
            {
                \"name\": \"stream-name\",
                \"in\": \"path\",
                \"required\": true,
                \"type\": \"string\"
            },
            {
                \"name\": \"Shard-Iterator\",
                \"in\": \"header\",
                \"required\": false,
                \"type\": \"string\"
            }
        ],
        \"responses\": {
            \"200\": {
                \"description\": \"200 response\",
                \"schema\": {
                    \"$ref\": \"#/definitions/Empty\"
                }
            }
        }
    },
    \"x-amazon-apigateway-integration\": {
        \"credentials\": \"arn:aws:iam::738575810317:role/apigAwsProxyRole\",
        \"responses\": {
            \"default\": {
                \"statusCode\": \"200\"
            }
        },
        \"requestTemplates\": {
            \"application/json\": \"{\\n      \\\"ShardIterator\\\": \\\"$in
put.params('Shard-Iterator')\\\"\\n}\\"
        },
        \"uri\": \"arn:aws:apigateway:us-east-1:kinesis:action/GetRecords\",
        \"httpMethod\": \"POST\",
        \"requestParameters\": {
            \"integration.request.header.Content-Type\": \"application/x-amz-json-

```

```

1.1' "
        },
        "type": "aws"
    },
},
"put": {
    "consumes": [
        "application/json",
        "application/x-amz-json-1.1"
    ],
    "produces": [
        "application/json"
    ],
    "parameters": [
        {
            "name": "Content-Type",
            "in": "header",
            "required": false,
            "type": "string"
        },
        {
            "name": "stream-name",
            "in": "path",
            "required": true,
            "type": "string"
        },
        {
            "in": "body",
            "name": "PutRecordsMethodRequestPayload",
            "required": true,
            "schema": {
                "$ref": "#/definitions/PutRecordsMethodRequestPayload"
            }
        }
    ],
    "responses": {
        "200": {
            "description": "200 response",
            "schema": {
                "$ref": "#/definitions/Empty"
            }
        }
    },
    "x-amazon-apigateway-integration": {
        "credentials": "arn:aws:iam::738575810317:role/apigAwsProxyRole",
        "responses": {
            "default": {
                "statusCode": "200"
            }
        },
        "requestTemplates": {
            "application/json": "\n      \"StreamName\": \"$input.params('stream-\nname')\", \n      \"Records\": [\n          #foreach($elem in $input.path('$.\nrecords'))\n          {\n              \"Data\": \"$util.base64En-\ncode($elem.data)\",\n              \"PartitionKey\": \"$elem.partition-key\"\n          }#if($foreach.hasNext),#end\n          #end\n      ]\n    ",
            "application/x-amz-json-1.1": "#set($inputRoot = $in-\nput.path('$'))\n{\n    \"StreamName\": \"$input.params('stream-name')\""
        }
    }
}

```

```

\"records\" : [\n      #foreach($elem in $inputRoot.records)\n          {\n            \"Data\": \"$elem.data\", \n            \"partition-key\": \"$elem.partition-key\"\n            }#if($foreach.hasNext),#end\n          #end\n      ]\n    },\n    \"uri\": \"arn:aws:apigateway:us-east-1:kinesis:action/PutRecords\" ,\n    \"httpMethod\": \"POST\" ,\n    \"requestParameters\": {\n      \"integration.request.header.Content-Type\": \"application/x-amz-json-\n      1.1\" \n    },\n    \"type\": \"aws\"\n  }\n},\n\"/streams/{stream-name}/sharditerator\": {\n  \"get\": {\n    \"consumes\": [\n      \"application/json\"\n    ],\n    \"produces\": [\n      \"application/json\"\n    ],\n    \"parameters\": [\n      {\n        \"name\": \"stream-name\" ,\n        \"in\": \"path\" ,\n        \"required\": true,\n        \"type\": \"string\"\n      },\n      {\n        \"name\": \"shard-id\" ,\n        \"in\": \"query\" ,\n        \"required\": false,\n        \"type\": \"string\"\n      }\n    ],\n    \"responses\": {\n      \"200\": {\n        \"description\": \"200 response\" ,\n        \"schema\": {\n          \"$ref\": \"#/definitions/Empty\"\n        }\n      }\n    },\n    \"x-amazon-apigateway-integration\": {\n      \"credentials\": \"arn:aws:iam::738575810317:role/apigAwsProxyRole\" ,\n      \"responses\": {\n        \"default\": {\n          \"statusCode\": \"200\"\n        }\n      },\n      \"requestTemplates\": {\n        \"application/json\": \"{\n          \"ShardId\": \"$input.params('shard-\n          id')\" ,\n          \"ShardIteratorType\": \"TRIM_HORIZON\" ,\n          \"StreamName\": \"$input.params('stream-name')\"\n        }\" ,\n      },\n      \"uri\": \"arn:aws:apigateway:us-east-1:kinesis:action/GetShardIterator\" ,\n    }\n  }\n}
```

```
        "httpMethod": "POST",
        "requestParameters": {
            "integration.request.header.Content-Type": "'application/x-amz-json-1.1'"
        },
        "type": "aws"
    }
},
"definitions": {
    "PutRecordsMethodRequestPayload": {
        "type": "object",
        "properties": {
            "records": {
                "type": "array",
                "items": {
                    "type": "object",
                    "properties": {
                        "data": {
                            "type": "string"
                        },
                        "partition-key": {
                            "type": "string"
                        }
                    }
                }
            }
        }
    },
    "Empty": {
        "type": "object"
    }
}
}
```

Maintaining an API in Amazon API Gateway

Topics

- [View a List of APIs in API Gateway \(p. 196\)](#)
- [Delete an API in API Gateway \(p. 196\)](#)
- [Delete a Resource in API Gateway \(p. 197\)](#)
- [View a Methods List in API Gateway \(p. 197\)](#)
- [Delete a Method in API Gateway \(p. 198\)](#)

View a List of APIs in API Gateway

Use the API Gateway console to view a list of APIs.

Topics

- [Prerequisites \(p. 196\)](#)
- [View a List of APIs with the API Gateway Console \(p. 196\)](#)

Prerequisites

- You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).

View a List of APIs with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. The list of APIs is displayed.

Delete an API in API Gateway

Use the API Gateway console to delete an API.

Warning

Deleting an API means that you can no longer call it. This action cannot be undone.

Topics

- [Prerequisites \(p. 197\)](#)
- [Delete an API with the API Gateway Console \(p. 197\)](#)

Prerequisites

- You must have deployed the API at least once. Follow the instructions in [Deploying an API \(p. 199\)](#).

Delete an API with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API you want to delete, choose **Resources**.
3. Choose **Delete API**.
4. When prompted to delete the API, choose **Ok**.

Delete a Resource in API Gateway

Use the API Gateway console to delete a resource.

Warning

When you delete a resource, you also delete its child resources and methods. Deleting a resource may cause part of the corresponding API to be unusable. Deleting a resource cannot be undone.

Delete a Resource with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the resource you want to delete, choose **Resources**.
3. In the **Resources** pane, choose the resource, and then choose **Delete Resource**.
4. When prompted, choose **Delete**.

View a Methods List in API Gateway

Use the API Gateway console to view a list of methods for a resource.

Topics

- [Prerequisites \(p. 197\)](#)
- [View a Methods List with the API Gateway Console \(p. 198\)](#)

Prerequisites

- You must have methods available in API Gateway. Follow the instructions in [Build and Deploy an API Gateway API Step by Step \(p. 13\)](#).

View a Methods List with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API, choose **Resources**.
3. The list of methods is displayed in the **Resources** pane.

Tip

You may need to choose the arrow next to one or more resources to display all of the available methods.

Delete a Method in API Gateway

Use the API Gateway console to delete a method.

Warning

Deleting a method may cause part of the corresponding API to become unusable. Deleting a method cannot be undone.

Delete a Method with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the method, choose **Resources**.
3. In the **Resources** pane, choose the arrow next to the resource for the method.
4. Choose the method, and then choose **Delete Method**.
5. When prompted, choose **Delete**.

Deploying an API in Amazon API Gateway

Topics

- [Deploy an API with the Amazon API Gateway Console \(p. 199\)](#)
- [Manage API Request Throttling \(p. 201\)](#)
- [Deploy an API in Stages in Amazon API Gateway \(p. 201\)](#)
- [Enable Amazon API Gateway Caching in a Stage to Enhance API Performance \(p. 205\)](#)
- [Deploy an API with Stage Variables in Amazon API Gateway \(p. 210\)](#)
- [Generate an SDK for an API in API Gateway \(p. 221\)](#)
- [Use a Custom Domain Name in API Gateway \(p. 227\)](#)

Deploy an API with the Amazon API Gateway Console

Prerequisites

- You must specify settings for all of the methods in the API you want to deploy. Follow the instructions in [Set up Method and Integration \(p. 59\)](#).

Deploy an API with the API Gateway Console

Note

If you want to change a stage in API Gateway to use a different deployment, see [Change a Stage to Use a Different Deployment with the API Gateway Console \(p. 200\)](#) instead.

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API you want to deploy, choose **Resources**.
3. In the **Resources** pane, choose **Deploy API**.
4. For **Deployment stage**, do one of the following:

- To deploy the API to an existing stage, choose the name of the stage.
- To deploy the API to a new stage, choose **New Stage**. For **Stage name**, type the name of the stage you want to use for the deployment.

Tip

The stage name should be meaningful, but short enough to be easy and fast to type. Your users will specify this name as part of the URL they will use to invoke the API.

5. (Optional) For **Stage description**, type a description for the stage.
6. (Optional) For **Deployment description**, type a description for the deployment.
7. Choose **Deploy**.

Update deployment configuration with the API Gateway Console

After an API is deployed to a stage, you can, optionally, modify the deployment by updating the stage settings or stage variables. After making any changes, you must redeploy the API. The following procedure demonstrates how to accomplish with the API Gateway Console.

1. If needed, choose the **Settings** tab in the **Stage Editor** pane of the API Gateway Console.
You can then choose to use or not use API cache, to enable or disable CloudWatch logs, to change throttling settings, or to select or deselect a client certificate.
2. If needed, choose the **Stage Variables** tab in the **Stage Editor** pane of the API Gateway Console.
You can then choose to update the values of selected stage variables.
3. If you made any change, choose the **Save Changes** button; go back to the **Resources** window; and then choose **Deploy API** again.

Note

If the updated settings, such as enabling logging, requires a new IAM role, you can add the required IAM role without redeploying the API. However, it can take a few minutes before the new IAM role takes effect. Before that happens, traces of your API calls will not be logged even if you have enabled the logging option.

Change a Stage to Use a Different Deployment with the API Gateway Console

Once you have deployed an API more than once, you can choose a specific deployment for a given stage. The following procedure shows how to do this.

1. You must have deployed to the stage at least twice. Follow the instructions in [Deploy an API with the API Gateway Console \(p. 199\)](#).
2. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
3. In the box that contains the name of the API with the stage you want to change, choose **Stages**.
4. Choose the stage you want to update the deployment.
5. On the **Deployment History** tab, choose the option button next to the deployment you want the stage to use.
6. Choose **Change Deployment**.

Manage API Request Throttling

Topics

- [Account-Level Throttling \(p. 201\)](#)
- [Stage-Level and Method-Level Throttling \(p. 201\)](#)

Amazon API Gateway throttles API requests to your API using the token bucket algorithm. For more information, see [token bucket algorithm](#).

Account-Level Throttling

By default, API Gateway limits the steady-state request rates to 500 requests per second (rps) and allows bursts of up to 1000 rps across all APIs, stages, and methods within an AWS account. If necessary, you can request an increase to your account-level limits. For more information, see [API Gateway Limits \(p. 243\)](#).

You can view account-level throttling limits in the API Gateway console. The console displays the default account-level settings before these settings are overridden by any customization. You can also read the account-level throttling limits by using the [Control Service API \(p. 245\)](#).

Stage-Level and Method-Level Throttling

As an API owner, you can override the account-level request throttling limits for a specific stage or for individual methods in an API. Observed stage-level and method-level throttling limits are bounded by the account-level rate limits, even if you set the stage-level or method-level throttling limits greater than the account-level limits.

You can set the stage-level or method-level throttling limits by using the API Gateway console or by calling the [Control Service API \(p. 245\)](#). For instructions using the console, see [Specify a Stage's Settings in API Gateway \(p. 202\)](#).

Deploy an API in Stages in Amazon API Gateway

In API Gateway, a stage defines the path through which an API deployment is accessible.

Use the API Gateway console to deploy an API in stages.

- [Create a Stage \(p. 201\)](#)
- [View a List of Stages \(p. 202\)](#)
- [Specify a Stage's Settings in API Gateway \(p. 202\)](#)
- [Generate an SDK for an API \(p. 221\)](#)
- [Delete a Stage \(p. 205\)](#)

Create a Stage in API Gateway

Use the API Gateway console to create a stage for an API.

Topics

- [Prerequisites \(p. 202\)](#)
- [Create a Stage with the API Gateway Console \(p. 202\)](#)

Prerequisites

1. You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).
2. You must have deployed the API at least once. Follow the instructions in [Deploying an API \(p. 199\)](#).

Create a Stage with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API, choose **Stages**.
3. Choose **Create Stage**.
4. For **Stage name**, type a name for the stage.
5. (Optional) For **Stage description**, type a description for the stage.
6. For **Deployment**, choose the date and time of the existing API deployment you want to associate with this stage.
7. Choose **Create**.

View a List of Stages in API Gateway

Use the API Gateway console to view a list of stages in API Gateway.

Topics

- [Prerequisites \(p. 202\)](#)
- [View a List of Stages with the API Gateway Console \(p. 202\)](#)

Prerequisites

1. You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).
2. You must have deployed the API in API Gateway at least once. Follow the instructions in [Deploying an API \(p. 199\)](#).

View a List of Stages with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API, choose **Stages**.

Specify a Stage's Settings in API Gateway

Use the API Gateway console to specify stage settings.

Topics

- [Prerequisites \(p. 202\)](#)
- [Specify a Stage's Settings with the API Gateway Console \(p. 203\)](#)

Prerequisites

- You must have a stage available in API Gateway. Follow the instructions in [Create a Stage \(p. 201\)](#).

Specify a Stage's Settings with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the stage where you want to specify settings, choose **Stages**.
3. In the **Stages** pane, choose the name of the stage.
4. To enable a cache for the API, on the **Settings** tab, in the **Cache Settings** area, select **Enable API cache**. Then, for **Cache capacity**, choose a cache size. Finally, choose **Save Changes**.
5. To generate code to call the API from Android, iOS, or JavaScript, you use the **SDK Generation** tab. For more information, see [Generate an SDK for an API \(p. 221\)](#).
6. To enable Amazon CloudWatch Logs for all of the methods associated with this API in API Gateway, do the following:

1. On the **Settings** tab, in the **CloudWatch Settings** area, select **Enable CloudWatch Logs**.

Important

By selecting this box, your AWS account may be charged for using CloudWatch.

Tip

To enable CloudWatch Logs for some methods only, clear **Enable CloudWatch Logs**.

In the **Stages** pane, choose each of the methods for which you want to enable CloudWatch Logs. On the **Settings** tab for each method, choose **Override for this method**, and in the **CloudWatch Settings** area, select **Log to CloudWatch Logs**. Finally, choose **Save Changes**.

2. For **Log level**, choose **ERROR** to write only error-level entries to CloudWatch Logs, or choose **INFO** to write only informational-level entries to CloudWatch Logs.
3. To write entries to CloudWatch Logs that contain full API call request and response information, select **Log full requests/responses**.
4. Choose **Save Changes**. The new settings will take effect after a new deployment.

Important

Whether you enable CloudWatch Logs for all or only some of the methods, you must also specify the ARN of an IAM role that enables API Gateway to write information to CloudWatch Logs on behalf of your IAM user. To do this, in the secondary navigation bar, in the first list next to the console home button, choose **Settings**. Then type the ARN of the IAM role in the **CloudWatch Logging role ARN** box. For common application scenarios, the IAM role could attach the managed policy of `AmazonAPIGatewayPushToCloudWatchLogs`, which contains the following access policy statement:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "logs:CreateLogGroup",  
                "logs:CreateLogStream",  
                "logs:DescribeLogGroups",  
                "logs:DescribeLogStreams",  
                "logs:PutLogEvents",  
                "logs:GetLogEvents",  
                "logs:FilterLogEvents"  
            ],  
            "Resource": "*"  
        }  
    ]  
}
```

```
    ]  
}
```

The IAM role must also contain the following trust relationship statement:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "",  
            "Effect": "Allow",  
            "Principal": {  
                "Service": "apigateway.amazonaws.com"  
            },  
            "Action": "sts:AssumeRole"  
        }  
    ]  
}
```

To create the IAM role, you can adapt the instructions in "To create the Lambda invocation role and its policy" and "To create the Lambda execution role and its policy" in the [Create Lambda Functions \(p. 21\)](#) section of the [Invoke Lambda Functions Synchronously \(p. 20\)](#).

For more information about CloudWatch, see the [Amazon CloudWatch Developer Guide](#).

7. To enable Amazon CloudWatch metrics for all of the methods associated with this API in API Gateway, in the **Stage Editor** pane, on the **Settings** tab, in the **CloudWatch Settings** area, select **Enable CloudWatch metrics**, and then choose **Save Changes**. The new settings will take effect after a new deployment.

Important

By selecting this box, your AWS account may be charged for using CloudWatch.

Tip

To enable CloudWatch metrics for only some methods, clear **Enable CloudWatch metrics**. In the **Stages** pane, choose each of the methods for which you want to enable CloudWatch metrics. For each method you choose, on the **Settings** tab for the method, choose **Override for this method**, and in the **CloudWatch Settings** area, select **Enable CloudWatch metrics**. Finally, choose **Save Changes**.

For more information about CloudWatch, see the [Amazon CloudWatch Developer Guide](#).

8. To set a default throttle limit for all of the methods associated with this API in API Gateway, in the **Stage Editor** pane, on the **Settings** tab, in the **Throttle Settings** area, do the following, and then choose **Save Changes**:

- For **Burst Limit**, type the absolute maximum number of times API Gateway will allow this method to be called per second. (The value of **Burst Limit** must be equal to or greater than the value of **Rate**.) The default setting is 1000 request per second.
- For **Rate**, type the number of times API Gateway will allow this method to be called per second on average. (The value of **Rate** must be equal to or less than the value of **Burst Limit**.) The default setting is 500 request per second.

Note

- When creating a stage, if not supplied, API Gateway will enforce the default values of 1000 for **Burst Limit** and 500 for **Rate** in the stage settings.
 - In addition, API Gateway enforces overall account level throttling at the default values of 1000 for **Burst Limit** and 500 for **Rate**. If you require a higher level of throttling on your account, contact the [AWS Support Center](#) to request an increase.
 - API Gateway uses the [token bucket](#) algorithm, including [average rate](#) and [burst size](#), for both account and method throttling.
9. To change the stage to use a different deployment, in the **Stage Editor** pane, on the **Change Deployment** tab, choose the option button next to the deployment you want the stage to use, and then choose **Change Deployment**.

Delete a Stage in API Gateway

Use the API Gateway console to delete a stage in API Gateway.

Warning

Deleting a stage may cause part or all of the corresponding API to be unusable by API callers. Deleting a stage cannot be undone.

Delete a Stage with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the stage, choose **Stages**.
3. In the **Stages** pane, choose the stage you want to delete, and then choose **Delete Stage**.
4. When prompted, choose **Delete**.

Enable Amazon API Gateway Caching in a Stage to Enhance API Performance

Topics

- [Amazon API Gateway Caching Overview \(p. 205\)](#)
- [Enable Amazon API Gateway Caching \(p. 206\)](#)
- [Override API Gateway Stage-Level Caching for Method Caching \(p. 207\)](#)
- [Use Method or Integration Parameters as Cache Keys to Index Cached Responses \(p. 207\)](#)
- [Flush the API Stage Cache in API Gateway \(p. 208\)](#)
- [Invalidate an API Gateway Cache Entry \(p. 209\)](#)

Amazon API Gateway Caching Overview

You can enable API caching in Amazon API Gateway to cache your endpoint's response. With caching, you can reduce the number of calls made to your endpoint and also improve the latency of the requests to your API. When you enable caching for a stage, API Gateway caches responses from your endpoint for a specified time-to-live (TTL) period, in seconds. API Gateway then responds to the request by looking up the endpoint response from the cache instead of making a request to your endpoint. The default TTL

value for API caching is 300 seconds. The maximum TTL value is 3600 seconds. TTL=0 means caching is disabled.

Note

Caching is charged by the hour and is not eligible for the AWS free tier.

Enable Amazon API Gateway Caching

In API Gateway, you can enable caching for all methods for a specified stage. When you enable caching, you must choose a cache capacity. In general, a larger capacity gives a better performance, but also costs more.

API Gateway enables caching by creating a dedicated cache instance. This process can take up to 4 minutes.

API Gateway changes caching capacity by removing the existing cache instance and recreating a new one with a modified capacity. All existing cached data is deleted.

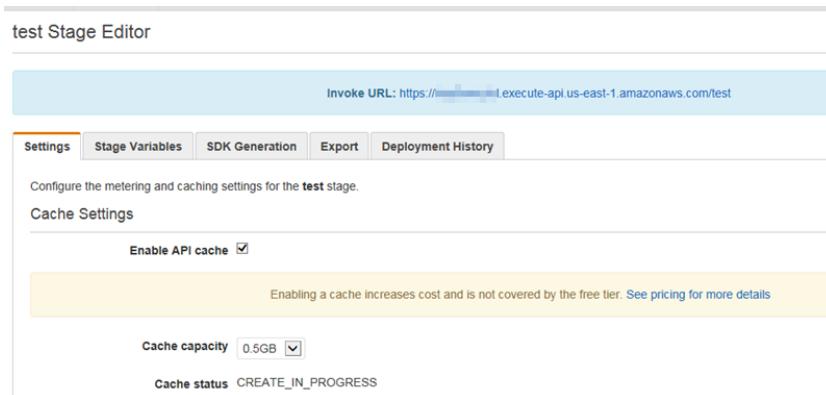
In the API Gateway console, you configure caching in the **Settings** tab of a named **Stage Editor**.

1. Go to the API Gateway console.
2. Navigate to the **Stage Editor** for the stage for which you want to enable caching.
3. Choose **Settings**.
4. Select **Enable API cache**.
5. Wait for the cache creation to complete.

Note

Creating or deleting a cache takes about 4 minutes for API Gateway to complete. When cache is created, the **Cache status** value changes from CREATE_IN_PROGRESS to AVAILABLE. When cache deletion is completed, the **Cache status** value changes from DELETE_IN_PROGRESS to an empty string.

When you enable caching within a stage's **Cache Settings**, you enable caching for all methods in that stage.



If you would like to verify if caching is functioning as expected, you have two general options:

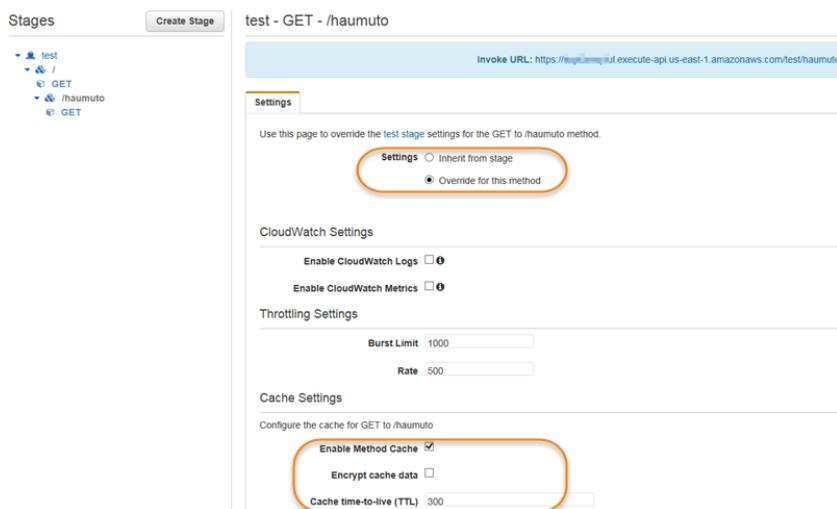
- Inspect the CloudWatch metrics of CacheHits and CacheMisses counts for your API and stage.
- Put a timestamp in the response.

Note

You should not use the X-Cache header from the CloudFront response to determine if your API is being served from your API Gateway cache instance.

Override API Gateway Stage-Level Caching for Method Caching

If you want more granularity in your caching settings, you can override the stage-level caching for individual methods. This includes disabling caching for a specific method, increasing or decreasing its TTL period, and turning on or off encryption of the cached response. If you anticipate that a method will receive sensitive data in its responses, in **Cache Settings**, choose **Encrypt cache data**.



Use Method or Integration Parameters as Cache Keys to Index Cached Responses

When a cached method or integration has parameters, which can take the form of custom headers, URL paths, or query strings, you can use some or all of the parameters to form cache keys. API Gateway can cache the method's responses, depending on the parameter values used.

For example, suppose you have a request of the following format:

```
GET /users?type=... HTTP/1.1
host: example.com
...
```

In this request, `type` can take a value of `admin` or `regular`. If you include the `type` parameter as part of the cache key, the responses from `GET /users?type=admin` will be cached separately from those from `GET /users?type=regular`.

When a method or integration request takes more than one parameter, you can choose to include some or all of the parameters to create the cache key. For example, you can include only the `type` parameter in the cache key for the following request, made in the listed order within a TTL period:

```
GET /users?type=admin&department=A HTTP/1.1
host: example.com
...
```

The response from this request will be cached and will be used to serve the following request:

```
GET /users?type=admin&department=B HTTP/1.1
host: example.com
...
```

To include a method or integration request parameter as part of a cache key in the API Gateway console, select **Caching** after you add the parameter.

[Method Execution](#) /haumuto - GET - Method Request

Provide information about this method's authorization settings and the parameters it can receive.

Authorization Settings

Authorization type: NONE 

ARN: arn:aws:execute-api:us-east-1:/7:h//GET/haumuto

API Key Required: false 

URL Query String Parameters

Name	Caching
query	<input checked="" type="checkbox"/>

 Add query string

HTTP Request Headers

Name	Caching
No headers	

 Add header

Request Models [Create a Model](#)

Flush the API Stage Cache in API Gateway

When API caching is enabled, you can flush your API stage's entire cache to ensure your API's clients get the most recent responses from your integration endpoints.

To flush the API stage cache, you can choose the **Flush Cache** button under the **Stage** tab in the API Gateway console. Notice that flushing the cache will cause the responses to ensuing requests to be serviced from the back end until the cache is build up again. During this period, the number of requests sent to the integration endpoint may increase. That may affect the overall latency of your API.

Invalidate an API Gateway Cache Entry

A client of your API can invalidate an existing cache entry and reloads it from the integration endpoint for individual requests. The client must send a request that contains the `Cache-Control: max-age=0` header. The client receives the response directly from the integration endpoint instead of the cache, provided that the user is authorized to do so. This replaces the existing cache entry with the new response, which is fetched from the integration endpoint.

To grant permission for a caller, attach a policy of the following format to an IAM execution role for the user.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "execute-api:InvalidateCache"  
            ],  
            "Resource": [  
                "arn:aws:execute-api:region:account-id:api-id/stage-name/HTTP-VERB/resource-path-specifier"  
            ]  
        }  
    ]  
}
```

This policy allows the API Gateway execution service to invalidate cache for requests on the specified resource (or resources). To specify a group of targeted resources, use a wildcard (*) character for `account-id`, `api-id`, and other entries in the ARN value of `Resource`. For more information on how to set permissions for the API Gateway execution service, see [Access Permissions \(p. 236\)](#)

If you do not impose an `InvalidateCache` policy, any client can invalidate the API cache. If all or most of the clients invalidate the API cache, there could be significant latency impact on your API.

When the policy is in place, caching is enabled, and authorization is required, you can control how unauthorized requests are handled by choosing an option from **Handle unauthorized requests** in the API Gateway console.

prod Stage Editor Delete Stage

Invoke URL: <https://V...n.execute-api.us-east-1.amazonaws.com/prod>

Settings Stage Variables SDK Generation Export Deployment History

Configure the metering and caching settings for the prod stage.

Cache Settings

Enable API cache

Enabling a cache increases cost and is not covered by the free tier. [See pricing for more details](#)

Cache capacity 0.5GB

Encrypt cache data

Cache time-to-live (TTL) 3600

Invalidation

Require authorization

Handle unauthorized requests Ignore cache control header; Add a warning in response header

CloudWatch Settings

Ignore cache control header
Ignore cache control header; Add a warning in response header
Fail the request with 403 status code

The three options result in the following behaviors:

- **Fail the request with 403 status code:** returns a 403 Unauthorized response.
To set this option using the API, use `FAIL_WITH_403`.
- **Ignore cache control header; Add a warning in response header:** process the request and add a warning header in the response.
To set this option using the API, use `SUCCEED_WITH_RESPONSE_HEADER`.
- **Ignore cache control header:** process the request and do not add a warning header in the response.
To set this option using the API, use `SUCCEED_WITHOUT_RESPONSE_HEADER`.

Deploy an API with Stage Variables in Amazon API Gateway

Stage variables are name-value pairs that you can define as configuration attributes associated with a deployment stage of an API. They act like environment variables and can be used in your API setup and mapping templates.

For example, you can define a stage variable in a stage configuration, and then set its value as the URL string of an HTTP integration for a method in your API. Later, you can reference the URL string using the associated stage variable name from the API setup. This way, you can use the same API setup with a different endpoint at each stage by resetting the stage variable value to the corresponding URLs. You can also access stage variables in the mapping templates, or pass configuration parameters to your AWS Lambda or HTTP back end.

For more information about mapping templates, see [Request and Response Payload-Mapping Reference \(p. 96\)](#).

Use Cases

With deployment stages in API Gateway, you can manage multiple release stages for each API, such as alpha, beta, and production. Using stage variables you can configure an API deployment stage to interact with different back-end endpoints. For example, your API can pass a GET request as an HTTP proxy to the back-end web host (for example, `http://example.com`). In this case, the back-end web host is configured in a stage variable so that when developers call your production endpoint, API Gateway calls `example.com`. When you call your beta endpoint, API Gateway uses the value configured in the stage variable for the beta stage, and calls a different web host (for example, `beta.example.com`). Similarly, stage variables can be used to specify a different AWS Lambda function name for each stage in your API.

You can also use stage variables to pass configuration parameters to a Lambda function through your mapping templates. For example, you may want to re-use the same Lambda function for multiple stages in your API, but the function should read data from a different Amazon DynamoDB table depending on which stage is being called. In the mapping templates that generate the request for the Lambda function, you can use stage variables to pass the table name to Lambda.

Examples

To use a stage variable to customize the HTTP integration endpoint, you must first configure a stage variable of a specified name, e.g., `url`, and then assign it a value, e.g., `example.com`. Next, from your method configuration, set up an HTTP proxy integration, and instead of entering the endpoint's URL, you can tell API Gateway to use the stage variable value, `http://${stageVariables.url}`. This value tells API Gateway to substitute your stage variable `${}` at runtime, depending on which stage your API is running. You can reference stage variables in a similar way to specify a Lambda function name, an AWS Service Proxy path, or an AWS role ARN in the credentials field.

When specifying a Lambda function name as a stage variable value, you must configure the permissions on the Lambda function manually. You can use the AWS Command Line Interface to do this.

```
aws lambda add-permission --function-name arn:aws:lambda:XXXXXX:your-lambda-function-name --source-arn arn:aws:execute-api:us-east-1:YOUR_AC COUNT_ID:api_id/*/HTTP_METHOD/resource --principal apigateway.amazonaws.com --statement-id apigateway-access --action lambda:InvokeFunction
```

The following example assigns API Gateway permission to invoke a Lambda function named `helloWorld` hosted in the US West (Oregon) region of an AWS account on behalf of the API method.

```
arn arn:aws:execute-api:us-west-2:123123123123:bmmuvptwze/*/GET/hello
```

Here is the same command using the AWS CLI.

```
aws lambda add-permission --function-name arn:aws:lambda:us-east-1:123123123123:function:helloWorld --source-arn arn:aws:execute-api:us-west-2:123123123123:bmmuvptwze/*/GET/hello --principal apigateway.amazonaws.com --statement-id apigateway-access --action lambda:InvokeFunction
```

Set Stage Variables Using the Amazon API Gateway Console

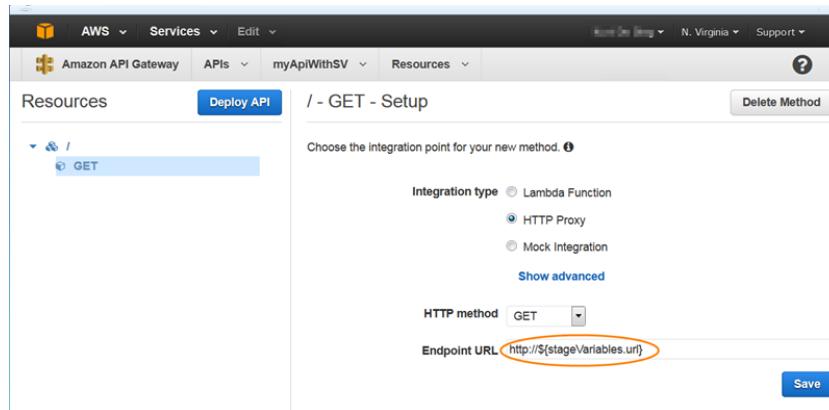
In this tutorial, you will learn how to set stage variables for two deployment stages of a sample API, using the Amazon API Gateway console.

Prerequisites

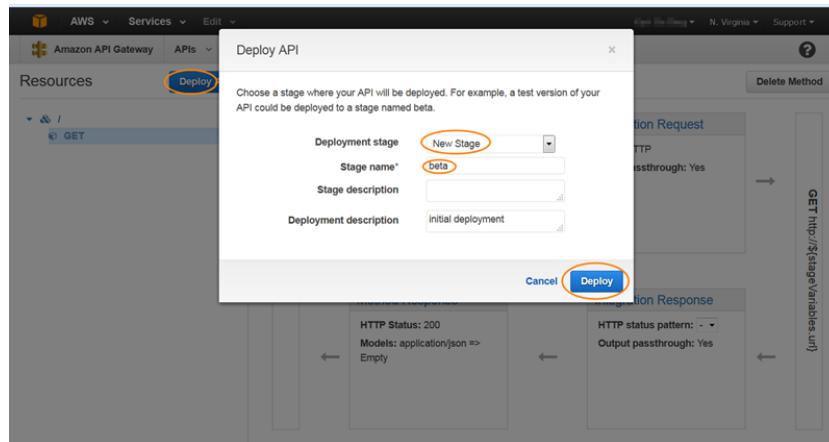
1. You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).
2. You must have deployed the API at least once. Follow the instructions in [Deploying an API \(p. 199\)](#).
3. You must have created the first stage for a deployed API. Follow the instructions in [Create a Stage \(p. 201\)](#).

To Declare Stage Variables Using the API Gateway Console

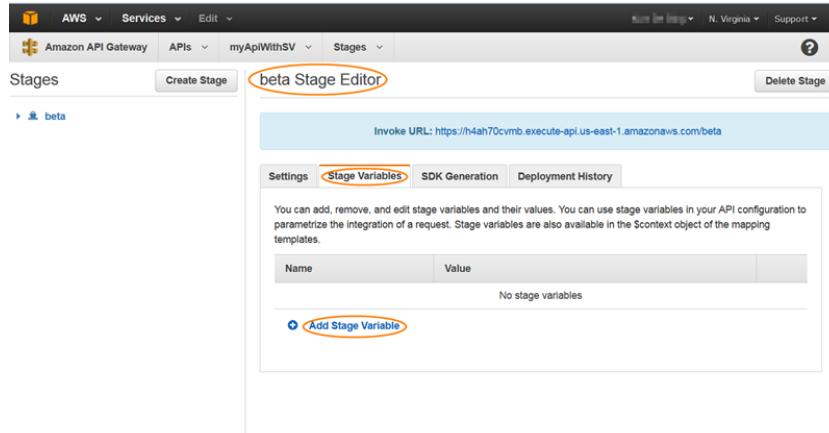
1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. Create an API, create a GET method on the API's root resource, if you have not already done so. Set the HTTP **Endpoint URL** value as "http://\${stageVariables.url}", and then choose **Save**.



3. Choose **Deploy API**. Choose **New Stage** and enter "beta" for **Stage name**. Choose **Deploy**.

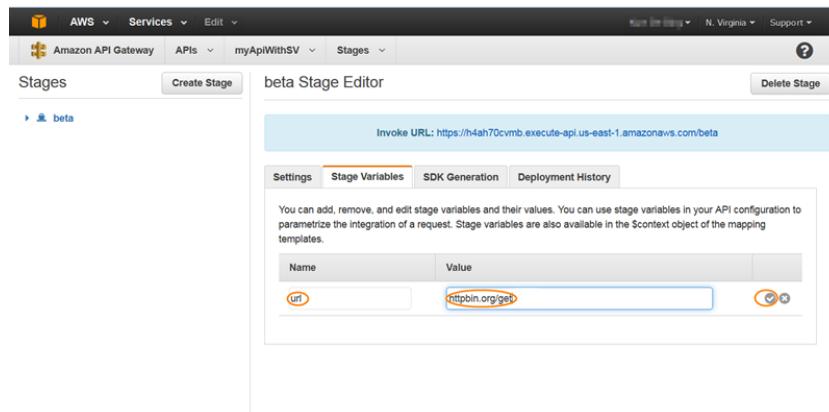


4. In the **beta Stage Editor** panel; choose the **Stage Variables** tab; and then choose **Add Stage Variable**.



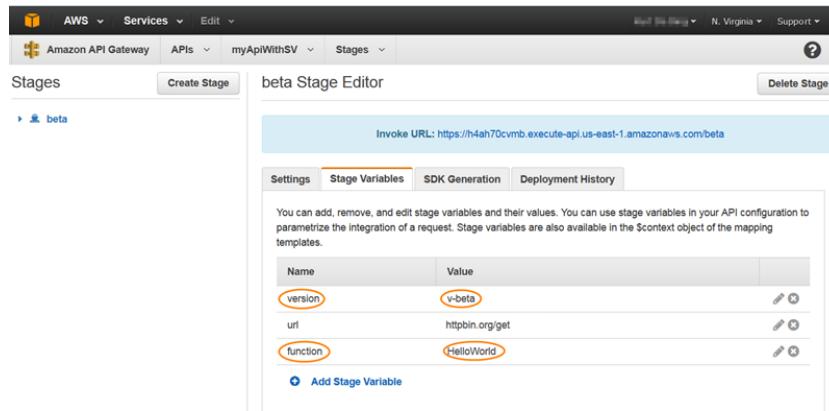
The screenshot shows the AWS Lambda console with the 'beta Stage Editor' open. The 'Stage Variables' tab is selected. A red circle highlights the 'Add Stage Variable' button at the bottom left of the table.

5. Enter the "url" string in the **Name** field and the "httpbin.org/get" in the **Value** field. Choose the checkmark icon to save the setting for the stage variable.



The screenshot shows the AWS Lambda console with the 'beta Stage Editor' open. The 'Stage Variables' tab is selected. A red circle highlights the 'url' entry in the Name column and another red circle highlights the 'httpbin.org/get' entry in the Value column.

6. Repeat the above step to add two more stage variables: `version` and `function`. Set their values as "`v-beta`" and "`HelloWorld`", respectively.



The screenshot shows the AWS Lambda console with the 'beta Stage Editor' open. The 'Stage Variables' tab is selected. Three stage variables are listed in the table:

Name	Value
version	v-beta
url	httpbin.org/get
function	HelloWorld

A red circle highlights the 'version' entry in the Name column and another red circle highlights the 'HelloWorld' entry in the Value column.

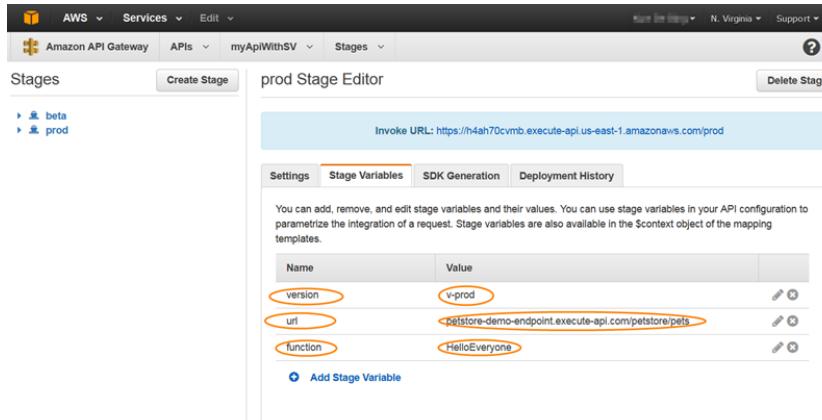
Note

When setting a Lambda function as the value of a stage variable, use the function's local name, possibly including its alias or version specification, as in `HelloWorld`, `HelloWorld:1` or `HelloWorld:alpha`. Do not use the function's ARN (for example, `arn:aws:lambda:us-east-1:123456789012:function:HelloWorld`). The API Gateway console assumes the stage variable value for a Lambda function as the unqualified function name and will expand the given stage variable into an ARN.

7. Choose **Create Stage**. Enter the "prod" string for the **Stage name**. Select a (recent) deployment from the **Deployment** dropdown menu. Then choose **Create**.



8. As with the **beta** stage, set the same three stage variables (**url**, **version**, and **function**) to different values ("petstore-demo-endpoint.execute-api.com/petstore/pets", "v-prod", and "HelloEveryone"), respectively.



Use Amazon API Gateway Stage Variables

You can use API Gateway stage variables to access the HTTP and Lambda back ends for different API deployment stages and to pass stage-specific configuration metadata into an HTTP back end as a query parameter and into a Lambda function as a payload generated in an input mapping template.

Prerequisites

You must create two stages with a `url` variable set to two different HTTP endpoints: a `function` stage variable assigned to two different Lambda functions, and a `version` stage variable containing

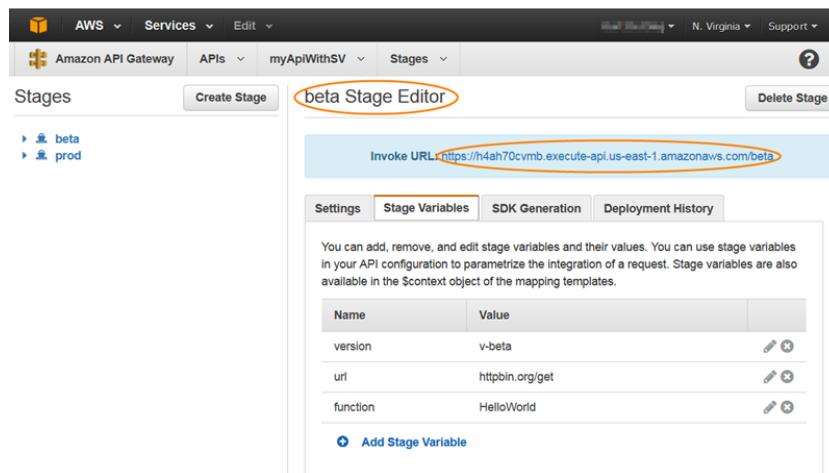
stage-specific metadata. Follow the instructions in [Set Stage Variables Using the Amazon API Gateway Console \(p. 212\)](#).

Access an HTTP endpoint through API with a stage variable

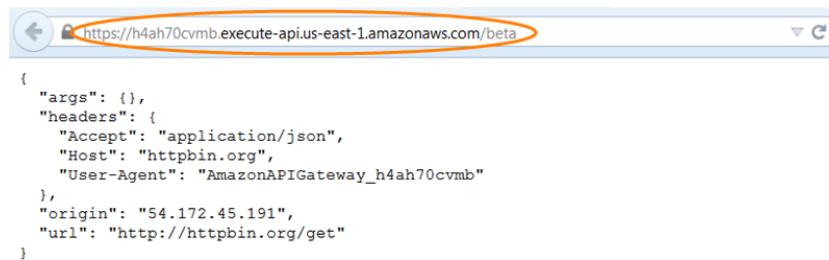
1. In the **Stages** navigation pane, choose **beta**. In **beta Stage Editor**, choose the **Invoke URL** link. This starts the **beta** stage GET request on the root resource of the API.

Note

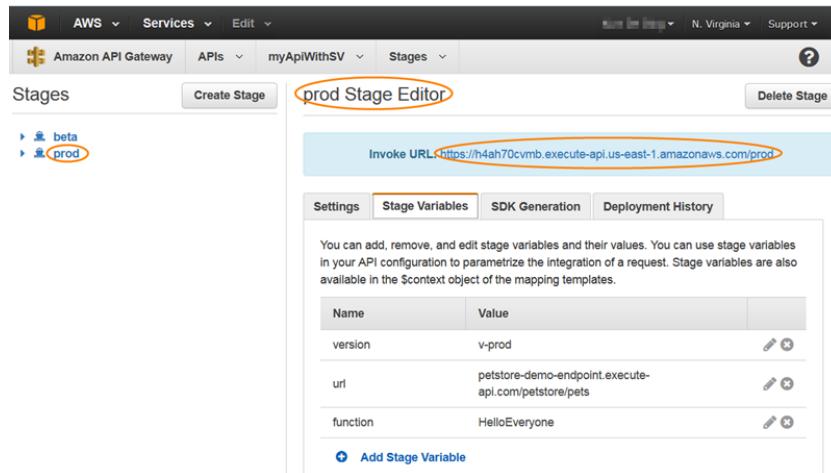
The **Invoke URL** link points to the root resource of the API in its **beta** stage. Navigating to the URL by choosing the link calls the **beta** stage GET method on the root resource. If methods are defined on child resources and not on the root resource itself, choosing the **Invoke URL** link will return a `{"message": "Missing Authentication Token"}` error response. In this case, you must append the name of a specific child resource to the **Invoke URL** link.



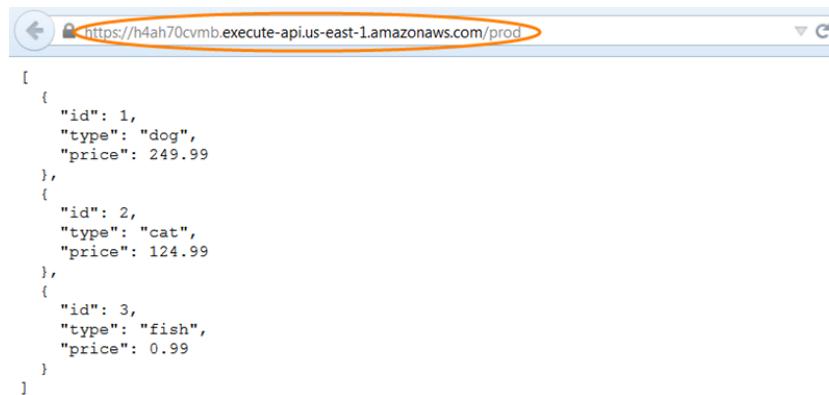
2. The response you get from the **beta** stage GET request is shown next. You can also verify the result by using a browser to navigate to <http://httpbin.org/get>. This value was assigned to the `url` variable in the **beta** stage. The two responses are identical.



3. In the **Stages** navigation pane, choose the **prod** stage. From **prod Stage Editor**, choose the **Invoke URL** link. This starts the **prod** stage GET request on the root resource of the API.



4. The response you get from the **prod** stage GET request is shown next. You can verify the result by using a browser to navigate to <http://petstore-demo-endpoint-execute-api.com/petstore/pets>. This value was assigned to the `url` variable in the **prod** stage. The two responses are identical.



Pass stage-specific metadata to an HTTP back end via a stage variable in a query parameter expression

This procedure describes how to use a stage variable value in a query parameter expression to pass stage-specific metadata into an HTTP back end. We will use the `version` stage variable declared in [Set Stage Variables Using the Amazon API Gateway Console \(p. 212\)](#).

1. In the **Resource** navigation pane, choose the **GET** method. To add a query string parameter to the method's URL, in **Method Execution**, choose **Method Request**. Type `version` for the parameter name.

Amazon API Gateway Developer Guide

Use Stage Variables

The screenshot shows the AWS Lambda console interface. At the top, the navigation bar includes 'AWS', 'Services', 'Edit', 'Amazon API Gateway', 'APIs', 'myApiWithSV', and 'Resources'. Below the navigation, the main area has tabs for 'Resources' and 'Deploy API'. A blue button labeled 'Deploy API' is highlighted. Underneath, there's a section for a 'GET' method. The URL path is shown as '/'. The 'Authorization Settings' section indicates 'Authorization type: NONE' and 'ARN: arn:aws:execute-api:us-east-1:7...:7h4ah70cvmb/*:GET/'. The 'API Key Required' field is set to 'false'. The 'URL Query String Parameters' section contains a table with one row, where 'version' is circled in red. There are also sections for 'HTTP Request Headers' and 'Request Models'.

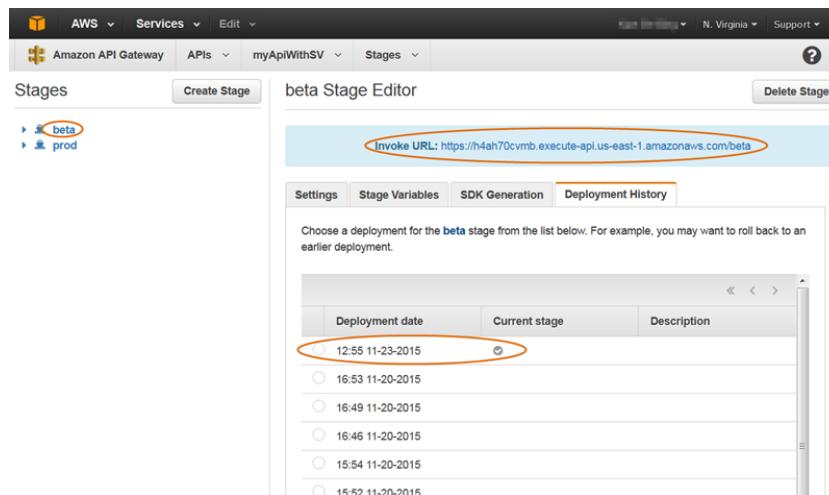
2. In **Method Execution** choose **Integration Request**. Edit the **Endpoint URL** value to append `?version=${stageVariables.version}` to the previously defined URL value, which, in this case, is also expressed with the `url` stage variable. Choose **Deploy API** to deploy these changes.

This screenshot shows the 'Integration Request' configuration for the 'GET' method. The 'Deploy API' button is circled in red. The 'Integration type' section has 'HTTP Proxy' selected. The 'HTTP method' is set to 'GET'. The 'Endpoint URL' field contains the value `http://${stageVariables.url}?version=${stageVariables.version}`, with the entire URL circled in red. There are sections for 'URL Path Parameters', 'URL Query String Parameters', and 'HTTP Headers'.

3. In the **Stages** navigation pane, choose the **beta** stage. From **beta Stage Editor**, verify that the current stage is in the most recent deployment, and then choose the **Invoke URL** link.

Note

We use the beta stage here because the HTTP endpoint, as specified by the `url` variable, "http://httpbin.org/get", accepts query parameter expressions and returns them as the `args` object in its response.



The screenshot shows the 'beta Stage Editor' in the AWS Management Console. At the top, there's a navigation bar with 'AWS', 'Services', 'Edit', 'N. Virginia', and 'Support'. Below it, the 'Amazon API Gateway' logo, 'APIs', and 'myApiWithSV' are visible. A 'Stages' tab is selected, showing two stages: 'beta' (highlighted with a red circle) and 'prod'. A 'Create Stage' button is also present. The main area is titled 'beta Stage Editor' with a 'Delete Stage' button. It includes tabs for 'Settings', 'Stage Variables', 'SDK Generation', and 'Deployment History' (which is currently selected). A message says: 'Choose a deployment for the beta stage from the list below. For example, you may want to roll back to an earlier deployment.' Below this is a table with columns 'Deployment date', 'Current stage', and 'Description'. The first row, '12:55 11-23-2015', is highlighted with a red circle.

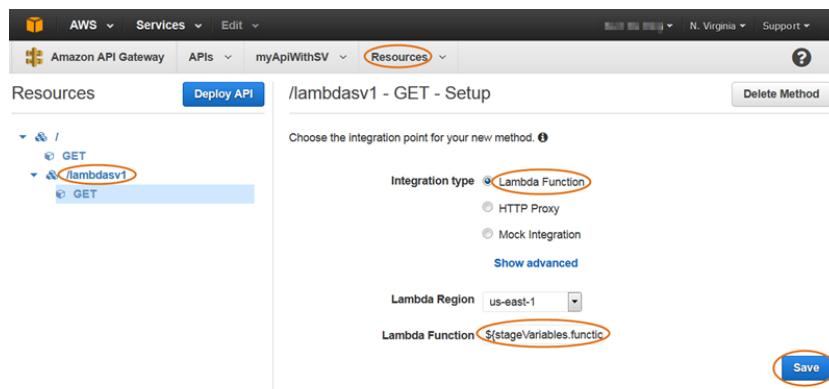
4. The response is shown next. Notice that v-beta, assigned to the `version` stage variable, is passed in the back end as the `version` argument.

```
{
  "args": {
    "version": "v-beta"
  },
  "headers": {
    "Accept": "application/json",
    "Host": "httpbin.org",
    "User-Agent": "AmazonAPIGateway_h4ah70cvmb"
  },
  "origin": "52.91.42.97",
  "url": "http://httpbin.org/get?version=v-beta"
}
```

Call Lambda function through API with a stage variable

This procedure describes how to use a stage variable to call a Lambda function as a back end of your API. We will use the `function` stage variable declared earlier. For more information, see [Set Stage Variables Using the Amazon API Gateway Console \(p. 212\)](#).

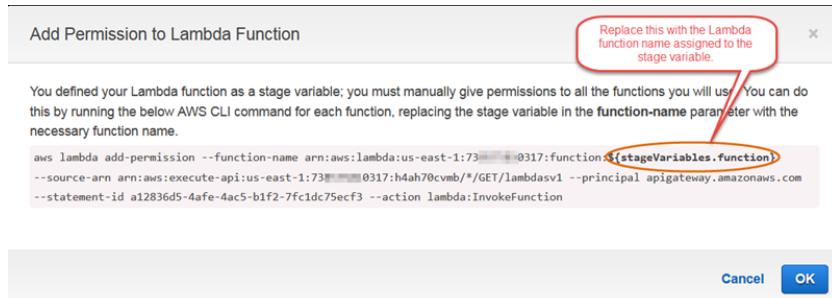
1. In the **Resources** pane, create a `/lambda` child resource under the root directory, and then create a GET method on the child resource. Set the **Integration type** to **Lambda Function**, and in **Lambda Function**, type `$(stageVariables.function)`. Choose **Save**.



The screenshot shows the 'Resources' pane in the AWS Management Console. At the top, there's a navigation bar with 'AWS', 'Services', 'Edit', 'N. Virginia', and 'Support'. Below it, the 'Amazon API Gateway' logo, 'APIs', and 'myApiWithSV' are visible. A 'Resources' tab is selected, showing a tree structure with a root node '/' containing a 'GET' method. A child node '`/lambda`' is expanded, showing another 'GET' method. The right-hand panel is titled '`/lambda` - GET - Setup' with a 'Deploy API' button. It asks 'Choose the integration point for your new method.' Below that, the 'Integration type' is set to 'Lambda Function' (radio button selected), with other options like 'HTTP Proxy' and 'Mock Integration' available. Under 'Show advanced', 'Lambda Region' is set to 'us-east-1' and 'Lambda Function' is set to `$(stageVariables.function)`. A large 'Save' button is at the bottom right.

Tip

When prompted with **Add Permission to Lambda Function**, make a note of the AWS CLI command before choosing **OK**. You must run the command on each Lambda function that is or will be assigned to the function stage variable for each of the newly created API methods. Failing to do so results in a **500 Internal Server Error** response when invoking the method. Make sure to replace `${stageVariables.function}` with the Lambda function name that is assigned to the stage variable.



2. Deploy the API to available stages.
3. In the **Stages** navigation pane, choose the **beta** stage. Verify that your most recent deployment is in **beta Stage Editor**. Copy the **Invoke URL** link, paste it into the address bar of your browser, and append `/lambdasv1` to that URL. This calls the underlying Lambda function through the `GET` method on the **LambdaSv1** child resource of the API.

Note

Your `HelloWorld` Lambda function implements the following code.

```
exports.handler = function(event, context) {
  if (event.version)
    context.succeed('Hello, World! (' + event.version + ')');
  else
    context.succeed("Hello, world! (v-unknown)");
};
```

This implementation results in the following response.

```
"Hello, world! (v-unknown)"
```

Pass stage-specific metadata to a Lambda function via a stage variable

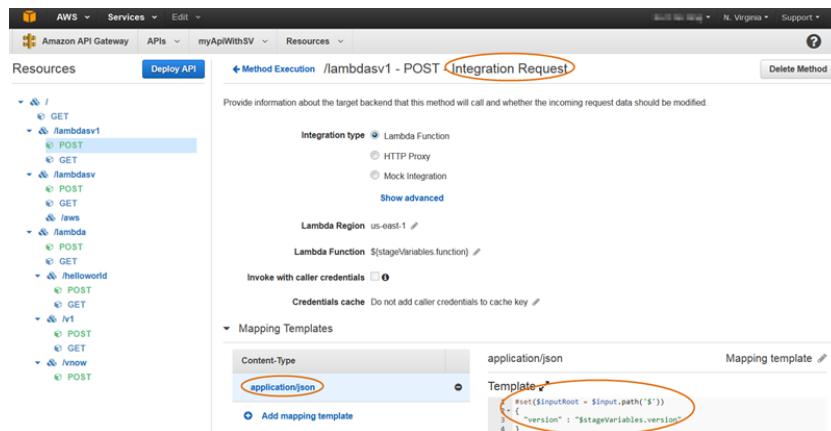
This procedure describes how to use a stage variable to pass stage-specific configuration metadata into a Lambda function. We will use a `POST` method and an input mapping template to generate payload using the `version` stage variable declared earlier.

1. In the **Resources** pane, choose the `/lambdasv1` child resource. Create a `POST` method on the child resource, set the **Integration type** to **Lambda Function**, and type `${stageVariables.function}` in **Lambda Function**. Choose **Save**.

Tip

This step is similar to the step we used to create the `GET` method. For more information, see [Call Lambda function through API with a stage variable \(p. 218\)](#).

2. From the **/Method Execution** pane, choose **Integration Request**. In the **Integration Request** pane, expand **Mapping Templates**, and then choose **Add mapping template** to add a template for the `application/json` content-type, as shown in the following.



Note

In a mapping template, a stage variable must be referenced within quotes (as in `\"$stageVariables.version"` or `"${stageVariables.version}"`), whereas elsewhere it must be referenced without quotes (as in `${stageVariables.function}`).

3. Deploy the API to available stages.
4. In the **Stages** navigation pane, choose **beta**. In **beta Stage Editor**, verify that the current stage has the most recent deployment. Copy the **Invoke URL** link, paste it into the URL input field of a REST API client, append `/lambdasv1` to that URL, and then submit a `POST` request to the underlying Lambda function.

Note

You will get the following response.

```
"Hello, world! (v-beta)"
```

To summarize, we have demonstrated how to use API Gateway stage variables to target different HTTP and Lambda back ends for different stages of API deployment. In addition, we also showed how to use the stage variables to pass stage-specific configuration data into HTTP and Lambda back ends. Together, these procedures demonstrate the versatility of the API Gateway stage variables in managing API development.

Amazon API Gateway Stage Variables Reference

You can use API Gateway stage variables in the following cases.

Parameter Mapping Expressions

A stage variable can be used in a parameter mapping expression for an API method's request or response header parameter, without any partial substitution. In the following example, the stage variable is referenced without the `$` and the enclosing `{ ... }`.

- `stageVariables.<variable_name>`

Mapping Templates

A stage variable can be used anywhere in a mapping template, as shown in the following examples.

- `{ "name" : "$stageVariables.<variable_name>" }`
- `{ "name" : "${stageVariables.<variable_name>}" }`

HTTP Integration URIs

A stage variable can be used as part of an HTTP integration URL, as shown in the following examples.

- A full URI without protocol, e.g., `http://${stageVariables.<variable_name>}`
- A full domain: e.g., `http://${stageVariables.<variable_name>}/resource/operation`
- A subdomain: e.g.,
`http://${stageVariables.<variable_name>}.example.com/resource/operation`
- A path, e.g., `http://example.com/${stageVariables.<variable_name>}/bar`
- A query string, e.g., `http://example.com/foo?q=${stageVariables.<variable_name>}`

AWS Integration URIs

A stage variable can be used as part of AWS URI action or path components, as shown in the following example.

- `arn:aws:apigateway:<region>:<service>:${stageVariables.<variable_name>}`

AWS Integration URIs (Lambda Functions)

A stage variable can be used in place of a Lambda function name, or version/alias, as shown in the following examples.

- `arn:aws:apigateway:<region>:lambda:path/2015-03-31/functions/<alias>:<account_id>:function:${stageVariables.function_variable_name}/invocations`
- `arn:aws:apigateway:<region>:lambda:path/2015-03-31/functions/<alias>:<account_id>:function:<function_name>:${stageVariables.version_variable_name}/invocations`

AWS Integration Credentials

A stage variable can be used as part of AWS user/role credential ARN, as shown in the following example.

- `arn:aws:iam::<account_id>:${stageVariables.<variable_name>}`

Generate an SDK for an API in API Gateway

You can generate an SDK for a specific stage of an API in API Gateway. The SDK contains code you can use to call the API from Android, iOS, or JavaScript. Use the API Gateway console to generate the SDK.

Topics

- [Prerequisites \(p. 222\)](#)

- Generate an SDK for an API with the API Gateway Console (p. 222)
- Integrate an API Gateway-Generated Android SDK into Your Android Project (p. 223)
- Integrate an API Gateway-Generated iOS SDK into Your iOS Project (p. 224)
- Integrate an API Gateway-Generated JavaScript SDK into Your JavaScript Code (p. 225)

Prerequisites

- You must have deployed the API at least once in API Gateway. Follow the instructions in [Deploying an API \(p. 199\)](#).

Generate an SDK for an API with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the stage, choose **Stages**.
3. In the **Stages** pane, choose the name of the stage.
4. On the **SDK Generation** tab, for **Platform**, choose the platform.
5. If you chose **Android**, specify the following:
 - For **Group ID**, type the unique identifier for the corresponding project. This is used in the `pom.xml` file (for example, `com.mycompany`).
 - For **Invoker package**, type the namespace for the generated client classes (for example, `com.mycompany.clientsdk`).
 - For **Artifact ID**, type the name of the compiled .jar file without the version. This is used in the `pom.xml` file (for example, `mycompany-client`).
 - For **Artifact version**, type the artifact version number for the generated client. This is used in the `pom.xml` file and should follow a `major.minor.patch` pattern (for example, `1.0.0`).
6. If you chose **iOS**, in the **Prefix** box, type the unique prefix for the generated classes. (For example, typing `CLI` will result in classes named `CLIRequestModel.h` and `CLIRequestModel.m`.)
7. Choose **Generate SDK**, and then follow the on-screen directions to download the API Gateway-generated SDK.
8. Do one of the following:
 - If you chose **Android** for **Platform**, follow the instructions in [Integrate an API Gateway-Generated Android SDK into Your Android Project \(p. 223\)](#).
 - If you chose **iOS** for **Platform**, follow the instructions in [Integrate an API Gateway-Generated iOS SDK into Your iOS Project \(p. 224\)](#).
 - If you chose **JavaScript** for **Platform**, follow the instructions in [Integrate an API Gateway-Generated JavaScript SDK into Your JavaScript Code \(p. 225\)](#).

Integrate an API Gateway-Generated Android SDK into Your Android Project

Note

These instructions assume you have already completed the steps in [Generate an SDK for an API with the API Gateway Console \(p. 222\)](#).

1. Extract the contents of the API Gateway-generated .zip file you downloaded earlier.
2. Download and install [Apache Maven](#) (preferably version 3.x).
3. Download and install the [JDK](#) (preferably version 1.7 or later).
4. Set the JAVA_HOME environment variable.
5. Run the command **mvn install** to install the compiled artifact files to your local Maven repository.
6. Copy the aws-backplane-client-1.0.0.jar file from the target folder, along with all of the other libraries from the target/lib folder, into your project's lib folder.
7. Use the `ApiClientFactory` class to initialize the API Gateway-generated SDK. For example:

```
ApiClientFactory factory = new ApiClientFactory()  
    .endpoint("https://localhost/");  
// Create a client.  
final IntTestApiClient client = factory.build(IntTestApiClient.class);  
  
// Invoke your parentPath1Get method.  
OriginalModel output = client.parentPath1Get(param1,body);  
  
// You also have access to your API's models.  
OriginalModel myModel = new OriginalModel();  
myModel.setStreetAddress(streetAddress);  
myModel.setCity(city);  
myModel.setState(state);  
myModel.setStreetNumber(streetNumber);  
myModel.setNested(nested);  
myModel.setPoBox(poBox);
```

8. To use a Amazon Cognito credentials provider to authorize calls to your API, use the `ApiClientFactory` class to pass a set of AWS credentials by using the API Gateway-generated SDK. For example:

```
// Use CognitoCachingCredentialsProvider to provide AWS credentials  
// for the ApiClientFactory  
AWSCredentialsProvider credentialsProvider = new CognitoCachingCredentialsProvider(  
    context, // activity context  
    "identityPoolId", // Cognito identity pool id  
    Regions.US_EAST_1 // region of Cognito identity pool  
);  
  
ApiClientFactory factory = new ApiClientFactory()  
    .credentialsProvider(credentialsProvider);
```

9. To set an API key by using the API Gateway-generated SDK, use code similar to the following:

```
ApiClientFactory factory = new ApiClientFactory()  
    .apiKey("YOUR_API_KEY");
```

Integrate an API Gateway-Generated iOS SDK into Your iOS Project

Note

These instructions assume you have already completed the steps in [Generate an SDK for an API with the API Gateway Console \(p. 222\)](#).

1. Extract the contents of the API Gateway-generated .zip file you downloaded earlier.
2. Import the AWS Mobile SDK for iOS into your project by using [CocoaPods](#) or Frameworks.

To import the AWS Mobile SDK for iOS into your project by using CocoaPods, do the following:

1. Install CocoaPods by running the command **sudo gem install cocoapods**.
2. Copy the `Podfile` file from the extracted .zip file into the same directory as your Xcode project file. If your Xcode project file already contains a file named `Podfile`, you can simply add the following line of code to it:

```
pod 'AWSAPIGateway', '~> 2.2.1'
```

3. Run the **pod install** command .
4. Use Xcode to open the `*.xcworkspace` file.
5. Copy all of the `.h` and `.m` files from the extracted .zip file's `generated-src` directory into your Xcode project.

To import the AWS Mobile SDK for iOS into your project by using Frameworks, do the following:

1. Download the [AWS Mobile SDK for iOS](#), version 2.2.1 or later.
2. With your project already open in Xcode, press and hold the Ctrl key while choosing **Frameworks**, and then choose **Add files to "<project name>"...**
3. In Finder, browse to and select both the `AWSCore.framework` and `AWSAPIGateway.framework` files, and then choose **Add**.
4. Open a target for your project, choose **Build Phases**, expand **Link Binary With Libraries**, choose the **+** button, and then add the following: `sqlite3.dylib`, `libz.dylib`, and `SystemConfiguration.framework`.
3. Import the `.h` file from the API Gateway-generated SDK. For example:

```
#import "<generated header file name>"
```

4. Get the `defaultClient` from your code. For example:

```
APIGIntTestApiClient *client = [APIGIntTestApiClient defaultClient];
```

5. Use the API Gateway-generated SDK to call your API's method. For example:

```
[[client parentPath:childPath1Get:@"/test" body:APIGOrginalModel] continue
WithBlock:^id(AWSTask *task) {
    if (task.error) {
        NSLog(@"Error: %@", task.error);
        return nil;
    }
    if (task.result) {
        APIGOrginalModel * output = task.result;
        //Do something with the output.
    }
    return nil;
}];
```

6. To use a Amazon Cognito credentials provider to authorize calls to your API, create an `AWSCognitoCredentialsProvider` object as the default provider for the API Gateway-generated SDK. For example:

```
AWSCognitoCredentialsProvider *creds = [[AWSCognitoCredentialsProvider alloc]
initWithRegionType:AWSRegionUSEast1
identityPoolId:CognitoPoolID];
AWSServiceConfiguration *configuration = [[AWSServiceConfiguration alloc]
initWithRegion:AWSRegionUSEast1 credentialsProvider:creds];
AWSServiceManager.defaultServiceManager.defaultServiceConfiguration = configuration;
```

7. To send an API key in your requests, set the `apiKey` property of the API Gateway-generated SDK. For example:

```
client.apiKey = @"Your API key";
```

Integrate an API Gateway-Generated JavaScript SDK into Your JavaScript Code

Note

These instructions assume you have already completed the instructions in [Generate an SDK for an API with the API Gateway Console \(p. 222\)](#).

1. Extract the contents of the API Gateway-generated .zip file you downloaded earlier.
2. Enable cross-origin resource sharing (CORS) for all of the methods the API Gateway-generated SDK will call. For instructions, see [Enable CORS for a Resource \(p. 101\)](#).
3. In your web page, include references to the following scripts:

```
<script type="text/javascript" src="lib/axios/dist/axios.standalone.js"></script>
<script type="text/javascript" src="lib/CryptoJS/rollups/hmac-sha256.js"></script>
<script type="text/javascript" src="lib/CryptoJS/rollups/sha256.js"></script>
<script type="text/javascript" src="lib/CryptoJS/components/hmac.js"></script>
<script type="text/javascript" src="lib/CryptoJS/components/enc-base64.js"></script>
<script type="text/javascript" src="lib/url-template/url-template.js"></script>
<script type="text/javascript" src="lib/apiGatewayCore/sigV4Client.js"></script>
<script type="text/javascript" src="lib/apiGatewayCore/apiGatewayClient.js"></script>
<script type="text/javascript" src="lib/apiGatewayCore/simpleHttpClient.js"></script>
<script type="text/javascript" src="lib/apiGatewayCore/utils.js"></script>
<script type="text/javascript" src="apigClient.js"></script>
```

4. In your code, initialize the API Gateway-generated SDK by using code similar to the following:

```
var apigClient = apigClientFactory.newClient();
```

5. Call the API in API Gateway by using code similar to the following. Each call returns a promise with a success and failure callbacks:

```
var params = {
    // This is where any modeled request parameters should be added.
    // The key is the parameter name, as it is defined in the API in API
    // Gateway.
    param0: '',
    param1: ''
};

var body = {
    // This is where you define the body of the request,
};

var additionalParams = {
    // If there are any unmodeled query parameters or headers that must be
    // sent with the request, add them here.
    headers: {
        param0: '',
        param1: ''
    },
    queryParams: {
        param0: '',
        param1: ''
    }
};

apigClient.methodName(params, body, additionalParams)
    .then(function(result){
        // Add success callback code here.
    })
    .catch( function(result){
```

```
// Add error callback code here.  
});
```

6. To initialize the API Gateway-generated SDK with AWS credentials, use code similar to the following. If you use AWS credentials, all requests to the API will be signed. This means you must set the appropriate CORS Accept headers for each request:

```
var apigClient = apigClientFactory.newClient({  
    accessKey: 'ACCESS_KEY',  
    secretKey: 'SECRET_KEY',  
});
```

7. To use an API key with the API Gateway-generated SDK, you can pass the API key as a parameter to the Factory object by using code similar to the following. If you use an API key, it is specified as part of the `x-api-key` header and all requests to the API will be signed. This means you must set the appropriate CORS Accept headers for each request:

```
var apigClient = apigClientFactory.newClient({  
    apiKey: 'API_KEY'  
});
```

Use a Custom Domain Name in API Gateway

You can use API Gateway to let users call your API with a simpler, more intuitive URL that contains a custom domain name (for example, `https://api.example.com`) instead of the default URL (for example, `https://my-api-id.execute-api.region-id.amazonaws.com`).

API Gateway supports custom domain names by enforcing Server Name Indication (SNI) on Amazon CloudFront distributions. For more information on using custom domain names, including the required certificate format and the maximum size of a certificate key length, see [Using Alternate Domain Names and HTTPS](#).

Note

API Gateway is not currently integrated with domain name registrars or managed DNS providers. You must register your custom domain name with a domain name registrar before setting it up in API Gateway. After you've set up your domain name and base path mapping(s) in the Amazon API Gateway console, you must update your domain name with your authoritative DNS name server provider to direct traffic to the CloudFront distribution API Gateway created for your domain name. The domain name of this CloudFront distribution is displayed as the "Distribution Domain Name" in the Amazon API Gateway console.

Topics

- [Prerequisites \(p. 227\)](#)
- [Specify a Custom Domain Name with the API Gateway Console \(p. 228\)](#)
- [Specify API Mappings for a Custom Domain Name with the API Gateway Console \(p. 229\)](#)
- [Call Your API with Custom Domain Names \(p. 229\)](#)

Prerequisites

1. Register your custom domain name. See the [Accredited Registrar Directory](#).

2. Get an SSL certificate for your custom domain name from a certificate authority. API Gateway is not currently integrated with certificate authorities. For a partial list, see [Third-Party Certificate Authorities](#).
3. Generate a private key for your SSL certificate.
4. Get the certificate chain from the same certificate authority or construct your own certificate chain.

Specify a Custom Domain Name with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the secondary navigation bar, in the first list next to the console home button, choose **Custom Domain Names**.
3. Choose **Create**.
4. Specify the following:
 - For **Domain name**, type your domain name (for example, `api.example.com`).
 - For **Certificate name**, type a name for future reference (for example, `my-example-certificate`).
 - For **Certificate body**, type or paste the body of the server certificate from your certificate authority.
 - For **Certificate private key**, type or paste your certificate's private key.
 - For **Certificate chain**, type or paste the intermediate certificates and, optionally, the root certificate, one after the other without any blank lines. If you include the root certificate, your certificate chain must start with intermediate certificates and end with the root certificate. Use the intermediate certificates provided by your certificate authority. Do not include any intermediaries that are not in the chain of trust path.
5. Choose **Save**.

Note

To delete a custom domain name, choose the domain name, and then choose **Delete Custom Domain Name**. When prompted, choose **Delete**.

If you need to edit the custom domain name entry, you must delete it and then start this procedure over again.

6. The console displays text similar to the following:

Create a CNAME resource record with your DNS provider to map api.example.com to d111111abcdef8.cloudfront.net

Certificate name my-example-certificate

Distribution domain name d111111abcdef8.cloudfront.net

These settings configure DNS to route DNS queries for your custom domain name or subdomain name (such as `api.example.com`) to the domain name for the Amazon CloudFront distribution. For most DNS services, including Amazon Route 53, your custom domain name is added to the hosted zone as a CNAME resource record set. The CNAME record name specifies the custom domain name you typed earlier for **Domain Name** (for example, `api.example.com`). The CNAME record value specifies the domain name for the CloudFront distribution (for example, `d111111abcdef8.cloudfront.net`). However, use of a CNAME record will not work if your custom domain is a zone apex (i.e., `example.com` instead of `api.example.com`). A zone apex is also commonly known as the root domain of your organization.

With Amazon Route 53 as the DNS service for your domain, you can also create an alias resource record (or A record) set for your custom domain name and specify the CloudFront distribution as the alias target. This means that Amazon Route 53 can route your custom domain name even if it is a

zone apex. For more information, see [Choosing Between Alias and Non-Alias Resource Record Sets in the Amazon Route 53 Developer Guide](#).

Specify API Mappings for a Custom Domain Name with the API Gateway Console

Specify the variations of the URLs that can be used to call your APIs.

Note

These instructions assume you have already completed the steps in [Specify a Custom Domain Name with the API Gateway Console \(p. 228\)](#).

1. For each URL variation you want to enable, choose **Create API mapping**.
2. (Optional) For **Base path**, type the base path name API callers must provide as part of the URL. This value must be unique for all of the mappings across a single API. Leave this blank if you do not want callers to specify a base path name after the domain name.

Note

Base path enables hosting multiple APIs behind a single domain name.

3. For **API**, choose the name of the API to which to apply this mapping.
4. (Optional) For **Stage**, choose the name of the API's stage you want to use for this mapping. Leave this blank if you want callers to explicitly specify the stage name after any base path name.
5. Choose **Save**.

Note

To delete a mapping after you create it, next to the mapping that you want to delete, choose **Remove**.

For example, assuming a custom domain name of `api.example.com`:

- Leave **Base Path** blank, specify an **API** of `MyDemoAPI`, and specify a **Stage** value of `prod` to enable calls to `https://api.example.com` to be forwarded to `https://my-api-id.execute-api.region-id.amazonaws.com/prod` (where `my-api-id` is the identifier API Gateway assigns to the API named `MyDemoAPI`).
- Leave **Base Path** blank, specify an **API** of `MyDemoAPI`, and leave **Stage** blank to enable calls to `https://api.example.com/prod` to be forwarded to `https://my-api-id.execute-api.region-id.amazonaws.com/prod` (where `my-api-id` is the identifier API Gateway assigns to the API named `MyDemoAPI`).
- Specify a **Base Path** value of `billing`, specify an **API** of `MyDemoAPI`, and leave **Stage** blank to enable calls to `https://api.example.com/billing/beta` to be forwarded to `https://my-api-id.execute-api.region-id.amazonaws.com/beta` (where `my-api-id` is the identifier API Gateway assigns to the API named `MyDemoAPI`).
- Specify a **Base Path** value of `scheduling`, specify an **API** of `MyDemoAPI`, and specify a **Stage** value of `gamma` to enable calls to `https://api.example.com/scheduling` to be forwarded to `https://my-api-id.execute-api.region-id.amazonaws.com/gamma` (where `my-api-id` is the identifier API Gateway assigns to the API named `MyDemoAPI`).

Call Your API with Custom Domain Names

API Gateway supports custom domain names for an API by using [Server Name Indication \(SNI\)](#). After a custom domain name is configured with the API, you can call the API with the custom domain name by using a browser or a client library that supports Server Name Indication (SNI).

API Gateway enforces SNI on the [CloudFront distributions](#). For information on how CloudFront uses custom domain names, see [Amazon CloudFront Custom SSL](#).

Calling an API in Amazon API Gateway

Use the API Gateway console to call an API.

Topics

- [Prerequisites \(p. 231\)](#)
- [Call an API with the API Gateway Console \(p. 231\)](#)
- [Use the API Dashboard in the API Gateway Console \(p. 232\)](#)
- [Test a Method in API Gateway \(p. 233\)](#)
- [Log Amazon API Gateway API Calls by Using AWS CloudTrail \(p. 234\)](#)

Prerequisites

- You must have already deployed the API in API Gateway. Follow the instructions in [Deploying an API \(p. 199\)](#).

Call an API with the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API you want to call, choose **Stages**.
3. In the **Stages** pane, choose the name of the stage.
4. The URL displayed next to **Invoke URL** should look something like this, where *my-api-id* is the identifier API Gateway assigns to your API, *region-id* is the AWS region identifier (for example, us-east-1) where you deployed your API, and *stage-name* is the name of the stage for the API you want to call:

```
https://my-api-id.execute-api.region-id.amazonaws.com/stage-name/{resource Path}
```

5. Depending on the method type you want to call and the tool you want to use, copy this URL to your clipboard, and then paste and modify it to call the API from a web browser, a web debugging proxy tool or the cURL command-line tool, or from your own API.

If you are not familiar with which method to call or the format you must use to call it, browse the list of available methods by following the instructions in [View a Methods List \(p. 197\)](#).

To call the method directly from the API Gateway console, see [Test a Method \(p. 233\)](#).

For more options, contact the API owner.

Use the API Dashboard in the API Gateway Console

The API dashboard in the API Gateway Console displays a summary of API activity over time. You can use it to get the URL for invoking an API.

Topics

- [Prerequisites \(p. 232\)](#)
- [Use the API Dashboard in the API Gateway Console \(p. 232\)](#)

Prerequisites

1. You must have an API available in API Gateway. Follow the instructions in [Creating an API \(p. 58\)](#).
2. You must have deployed the API at least once. Follow the instructions in [Deploying an API \(p. 199\)](#).
3. Enable Amazon CloudWatch metrics if you want to get metrics for the API's methods. You do this by specifying settings for the stage associated with the methods. Follow the instructions in [Specify a Stage's Settings in API Gateway \(p. 202\)](#).

Use the API Dashboard in the API Gateway Console

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. Choose the name of the API.
3. In the secondary navigation bar, in the third list next to the console home button, choose **Dashboard**.
4. Next to **Invoke this API at** is a URL you can use to invoke the API. If you are unsure about how to invoke the API, follow the instructions in [Calling an API \(p. 231\)](#).
5. To display a summary of API activity over time, for **Stage**, choose the desired stage.
6. Choose **ALL** to display activity for all of the API's methods, or choose a single method to display activity for that method only.
7. Use **From** and **To** to enter the date range.
8. The following activity information is displayed:
 - The **API Calls** graph displays the number of calls that were made to the methods in API Gateway.
 - The **Latency** graph displays the time, in milliseconds, between the receipt of requests from the caller and returned responses.
 - The **4xx Error** graph displays the number of responses with 4xx HTTP status codes.
 - The **5xx Error** graph displays the number of responses with 5xx HTTP status codes.

Test a Method in API Gateway

Use the API Gateway console to test a method.

Topics

- [Prerequisites \(p. 233\)](#)
- [Test a Method with the API Gateway Console \(p. 233\)](#)

Prerequisites

- You must specify the settings for the methods you want to test. Follow the instructions in [Set up Method and Integration \(p. 59\)](#).

Test a Method with the API Gateway Console

Important

Testing methods with the API Gateway console may result in changes to resources that cannot be undone. Testing a method with the API Gateway console is the same as calling the method outside of the API Gateway console. For example, if you use the API Gateway console to call a method that deletes an API's resources, if the method call is successful, the API's resources will be deleted.

1. Sign in to the API Gateway console at <https://console.aws.amazon.com/apigateway>.
2. In the box that contains the name of the API for the method, choose **Resources**.
3. In the **Resources** pane, choose the method you want to test.
4. In the **Method Execution** pane, in the **Client** box, choose **TEST**. Type values in any of the displayed boxes (such as **Query Strings**, **Headers**, and **Request Body**).

For additional options you may need to specify, contact the API owner.

5. Choose **Test**. The following information will be displayed:
 - **Request** is the resource's path that was called for the method.
 - **Status** is the response's HTTP status code.
 - **Latency** is the time between the receipt of the request from the caller and the returned response.
 - **Response Body** is the HTTP response body.
 - **Response Headers** are the HTTP response headers.

Tip

Depending on the mapping, the HTTP status code, response body, and response headers may be different from those sent from the Lambda function, HTTP proxy, or AWS service proxy.

- **Logs** are the simulated Amazon CloudWatch Logs entries that would have been written if this method were called outside of the API Gateway console.

Note

Although the CloudWatch Logs entries are simulated, the results of the method call are real.

Log Amazon API Gateway API Calls by Using AWS CloudTrail

API Gateway is integrated with CloudTrail, a service that captures API calls made by or on behalf of API Gateway in your AWS account and delivers the log files to an Amazon S3 bucket you specify. Examples of these API calls include creating a new API, resource, or method in API Gateway. CloudTrail captures API calls from the API Gateway console or from the API Gateway APIs directly. Using the information collected by CloudTrail, you can determine which request was made to API Gateway, the source IP address from which the request was made, who made the request, when it was made, and so on. To learn more about CloudTrail, including how to configure and enable it, see the [AWS CloudTrail User Guide](#).

API Gateway Information in CloudTrail

When CloudTrail logging is enabled in your AWS account, API calls made to API Gateway actions are tracked in log files. API Gateway records are written together with other AWS service records in a log file. CloudTrail determines when to create and write to a new file based on a time period and file size.

All of the API Gateway actions are logged and documented in the [Control Service API \(p. 245\)](#). For example, calls to create a new API, resource, or method in API Gateway generate entries in CloudTrail log files.

Every log entry contains information about who generated the request. The user identity information in the log helps you determine whether the request was made with root or IAM user credentials, with temporary security credentials for a role or federated user, or by another AWS service. For more information, see the **userIdentity** field in the [CloudTrail Event Reference](#).

You can store your log files in your bucket for as long as you want, but you can also define Amazon S3 lifecycle rules to archive or delete log files automatically. By default, your log files are encrypted by using Amazon S3 server-side encryption (SSE).

You can choose to have CloudTrail publish Amazon SNS notifications when new log files are delivered so you can take action quickly. For more information, see [Configuring Amazon SNS Notifications](#).

You can also aggregate API Gateway log files from multiple AWS regions and multiple AWS accounts into a single Amazon S3 bucket. For more information, see [Aggregating CloudTrail Log Files to a Single Amazon S3 Bucket](#).

Understanding API Gateway Log File Entries

CloudTrail log files can contain one or more log entries where each entry is made up of multiple JSON-formatted events. A log entry represents a single request from any source and includes information about the requested action, any parameters, the date and time of the action, and so on. The log entries are not guaranteed to be in any particular order. That is, they are not an ordered stack trace of the public API calls.

The following example shows a CloudTrail log entry that demonstrates the API Gateway get resource action:

```
{  
  Records: [  
    {  
      eventVersion: "1.03",  
      userIdentity: {  
        type: "Root",
```

```
principalId: "AKIAI44QH8DHBEXAMPLE",
arn: "arn:aws:iam::123456789012:root",
accountId: "123456789012",
accessKeyId: "AKIAIOSFODNN7EXAMPLE",
sessionContext: {
    attributes: {
        mfaAuthenticated: "false",
        creationDate: "2015-06-16T23:37:58Z"
    }
},
eventTime: "2015-06-17T00:47:28Z",
eventSource: "apigateway.amazonaws.com",
eventName: "GetResource",
awsRegion: "us-east-1",
sourceIPAddress: "203.0.113.11",
userAgent: "example-user-agent-string",
requestParameters: {
    restApiId: "3rbEXAMPLE",
    resourceId: "5tfEXAMPLE",
    template: false
},
responseElements: null,
requestID: "6d9c4bfc-148a-11e5-81b6-7577cEXAMPLE",
eventID: "4d293154-a15b-4c33-9e0a-ff5eeEXAMPLE",
readOnly: true,
eventType: "AwsApiCall",
recipientAccountId: "123456789012"
},
... additional entries ...
]
```

Amazon API Gateway Access Permissions

When working with Amazon API Gateway, you will most likely deal with two services. You use one to create and configure, manage, your API and the other to actually execute your deployed API. When setting access permissions, you reference the API-managing service as `apigateway` and the API-executing service as `execute-api`. While `apigateway` supports the actions of GET, POST, PUT, PATCH, DELETE, OPTIONS, etc., `execute-api` supports only the `Invoke` action.

You can use IAM to allow IAM users and roles in your AWS account to manage only certain API Gateway entities (for example, APIs, resources, methods, models, and stages) and perform only certain actions against those entities. You may want to do this, for example, if you have IAM users you want to allow to list, but not create, resources and methods for selected APIs. You may have other IAM users you want to allow to list and create new resources and methods for any API they have access to in API Gateway.

In the [Get Ready to Use API Gateway \(p. 3\)](#) instructions, you attached an access policy to an IAM user in your AWS account that contains a policy statement similar to this:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "apigateway:*"  
            ],  
            "Resource": [  
                "*"  
            ]  
        }  
    ]  
}
```

This statement allows the IAM user in your AWS account to perform all available actions and access all available resources in API Gateway to which your AWS account has access. In practice, you may not want to give the IAM users in your AWS account this much access.

You can also use IAM to enable users inside your organization to interact with only certain API methods in API Gateway.

In the [Configure How a User Calls an API Method \(p. 62\)](#) instructions, the API Gateway console may have displayed a resource ARN you used to create a policy statement similar to this:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "execute-api:Invoke"  
            ],  
            "Resource": [  
                "arn:aws:execute-api:us-east-1:my-aws-account-id:my-api-id/my-stage/GET/my-resource-path"  
            ]  
        }  
    ]  
}
```

This statement allows the IAM user to call the GET method for the resource path associated with the specified resource ARN in API Gateway. In practice, you may want to give IAM users access to more methods.

Note

IAM policies are effective only if IAM authentication is enabled. If you, as the API owner, has enabled AWS identity and access management on a specific resource, users from other AWS accounts cannot access your API. If you do not enable IAM authentication on the resource, that resource is effectively public accessible.

Create and Attach a Policy to an IAM User

To create and attach an access policy to an IAM user that restricts the API Gateway entities the IAM user can manage or the API methods the IAM user can call, do the following:

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. Choose **Policies**, and then choose **Create Policy**. (If a **Get Started** button appears, choose it, and then choose **Create Policy**.)
3. Next to **Create Your Own Policy**, choose **Select**.
4. For **Policy Name**, type any value that will be easy for you to refer to later, if needed.
5. For **Policy Document**, type a policy statement with the following format, and then choose **Create Policy**:

```
{  
    "Version": "2012-10-17",  
    "Statement" : [  
        {  
            "Effect" : "Allow",  
            "Action" : [  
                "action-statement"  
            ],  
            "Resource": [  
                "arn:aws:execute-api:us-east-1:my-aws-account-id:my-api-id/my-stage/GET/my-resource-path"  
            ]  
        }  
    ]  
}
```

```
    "Resource" : [
        "resource-statement"
    ]
},
{
    "Effect" : "Allow",
    "Action" : [
        "action-statement"
    ],
    "Resource" : [
        "resource-statement"
    ]
}
]
```

In this statement, substitute *action-statement* and *resource-statement* as needed, and add additional statements as needed, to specify the API Gateway entities you want to allow the IAM user to manage, the API methods the IAM user can call, or both. (By default, the IAM user will not have permissions unless a corresponding `Allow` statement is explicitly stated.)

6. Choose **Users**.
7. Choose the IAM user to whom you want to attach the policy.
8. For **Permissions**, for **Managed Policies**, choose **Attach Policy**.
9. Select the policy you just created, and then choose **Attach Policy**.

Action and Resource Syntax for Managing Entities in API Gateway

The following information describes the format for specifying actions and resources for API Gateway entities (such as APIs, resources, methods, models, and stages) that an IAM user can manage.

Actions follow this general format:

```
apigateway:action
```

Where *action* is an available API Gateway action:

- *, which represents all of the following actions.
- **GET**, which is used to get information about resources.
- **POST**, which is primarily used to create child resources.
- **PUT**, which is primarily used to update resources (and, although not recommended, can be used to create child resources).
- **DELETE**, which is used to delete resources.
- **PATCH**, which can be used to update resources.
- **HEAD**, which is the same as GET but does not return the resource representation. HEAD is used primarily in testing scenarios.
- **OPTIONS**, which can be used by callers to get information about available communication options for the target service.

Resources follow this general format:

```
arn:aws:apigateway:region::resource-path-specifier
```

Where *region* is a target AWS region (such as `us-east-1` or `*` for all supported AWS regions), and *resource-path-specifier* is the path to the target resources.

Some example actions statements include:

- `apigateway:*` for all API Gateway actions.
- `apigateway:GET` for just the GET action in API Gateway.

Some example resource statements include:

- `arn:aws:apigateway:region::/restapis/*` for all resources, methods, models, and stages in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/*` for all resources, methods, models, and stages in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/resources/resource-id/*` for all resources and methods in the resource with the identifier *resource-id*, which is in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/resources/resource-id/methods/*` for all of the methods in the resource with the identifier *resource-id*, which is in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/resources/resource-id/methods/GET` for just the GET method in the resource with the identifier *resource-id*, which is in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/models/*` for all of the models in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/models/model-name` for the model with the name of *model-name*, which is in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/stages/*` for all of the stages in the API with the identifier of *api-id* in the AWS region of *region*.
- `arn:aws:apigateway:region::/restapis/api-id/stages/stage-name` for just the stage with the name of *stage-name* in the API with the identifier of *api-id* in the AWS region of *region*.

The following policy statement gives the user permission to get information about all of the resources, methods, models, and stages in the API with the identifier of `a123456789` in the AWS region of `us-east-1`:

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "apigateway:GET"  
            ],  
            "Resource": [  
                "arn:aws:apigateway:us-east-1::/restapis/a123456789/*"  
            ]  
        }  
    ]  
}
```

The following example policy statement gives the IAM user permission to list information for all resources, methods, models, and stages in any region. The user also has permission to perform all available API Gateway actions for the API with the identifier of a123456789 in the AWS region of us-east-1:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "apigateway:GET"
            ],
            "Resource": [
                "arn:aws:apigateway:*:::/restapis/*"
            ]
        },
        {
            "Effect": "Allow",
            "Action": [
                "apigateway:*
            ],
            "Resource": [
                "arn:aws:apigateway:us-east-1:::/restapis/a123456789/*"
            ]
        }
    ]
}
```

Action and Resource Syntax for Calling an API's Methods in API Gateway

The following information describes the format for specifying actions and resources for an API's methods.

Actions follow this general format:

```
execute-api:Invoke
```

Resources follow this general format:

```
arn:aws:execute-api:region:account-id:api-id/stage-name/HTTP-VERB/resource-path-specifier
```

Where:

- *region* is the AWS region (such as **us-east-1** or ***** for all AWS regions) that corresponds to the deployed API for the method.
- *account-id* is the 12-digit AWS account Id of the REST API owner.
- *api-id* is the identifier API Gateway has assigned to the API for the method. (***** can be used for all APIs, regardless of the API's identifier.)
- *stage-name* is the name of the stage associated with the method (***** can be used for all stages, regardless of the stage's name.)

- *HTTP-VERB* is the HTTP verb for the method. It can be one of the following: GET, POST, PUT, DELETE, PATCH, HEAD, OPTIONS.
- *resource-path-specifier* is the path to the desired method. (* can be used for all paths).

Some example resource statements include:

- `arn:aws:execute-api:*:*:*` for any resource path in any stage, for any API in any AWS region. (This is equivalent to *).
- `arn:aws:execute-api:us-east-1:*:*` for any resource path in any stage, for any API in the AWS region of us-east-1.
- `arn:aws:execute-api:us-east-1:*:api-id/*` for any resource path in any stage, for the API with the identifier of `api-id` in the AWS region of us-east-1.
- `arn:aws:execute-api:us-east-1:*:api-id/test/*` for resource path in the stage of test, for the API with the identifier of `api-id` in the AWS region of us-east-1.
- `arn:aws:execute-api:us-east-1:*:api-id/test/*/mydemoresource/*` for any resource path along the path of mydemoresource, for any HTTP method in the stage of test, for the API with the identifier of `api-id` in the AWS region of us-east-1.
- `arn:aws:execute-api:us-east-1:*:api-id/test/GET/mydemoresource/*` for GET methods under any resource path along the path of mydemoresource, in the stage of test, for the API with the identifier of `api-id` in the AWS region of us-east-1.

The following policy statement gives the user permission to call any POST method along the path of mydemoresource, in the stage of test, for the API with the identifier of a123456789, assuming the corresponding API has been deployed to the AWS region of us-east-1:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "execute-api:Invoke"
      ],
      "Resource": [
        "arn:aws:execute-api:us-east-1:*:a123456789/test/POST/mydemoresource/*"
      ]
    }
  ]
}
```

The following example policy statement gives the user permission to call any method on the resource path of petstorewalkthrough/pets, in any stage, for the API with the identifier of a123456789, in any AWS region where the corresponding API has been deployed:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "execute-api:Invoke"
      ],
      "Resource": [
        "arn:aws:execute-api:*:a123456789/*"
      ]
    }
  ]
}
```

```
        "Resource": [
            "arn:aws:execute-api:*:::a123456789/test/*/petstorewalkthrough/pets"
        ]
    }
}
```

Amazon API Gateway Limits and Pricing

Topics

- [API Gateway Pricing \(p. 243\)](#)
- [API Gateway Limits \(p. 243\)](#)
- [Known Issues \(p. 244\)](#)

API Gateway Pricing

For API Gateway region-specific pricing information, see [Amazon API Gateway Pricing](#).

Note

API caching in Amazon API Gateway is not eligible for the AWS Free Tier.

API Gateway Limits

These are the current API Gateway limits that can be increased upon request, unless noted otherwise:

- For API caching, the default TTL value is 300 seconds (5 minutes) and the maximum TTL value is 3600 seconds (an hour), which cannot be increased currently.
- 60 APIs maximum per AWS account.
- 10,000 API keys per AWS account by default.
- 60 client certificates maximum per AWS account.
- 300 resources maximum per API.
- 10 stages maximum per API.
- 10-second timeout for both AWS Lambda and HTTP back-end integrations. This limit cannot be changed currently.
- 500 requests per second (sustained) per AWS account, across all account APIs. Bursts of up to 1,000 requests per second are allowed, with request throttling enforced using the token bucket algorithm.
- 10MB payload size, which cannot be changed currently.

- 1000 maximum number of iterations in a `#foreach ... #end` loop in mapping templates. Currently, this limit cannot be increased.
- When authorization is enabled on a method, the maximum length of the method's ARN (e.g., `arn:aws:execute-api:{region-id}:{account-id}:{api-id}/{stage-id}/{method}/{resource}/{path}`) is 1600 bytes. When the ARN includes URL path parameters, the size of path parameter values can cause the ARN length to exceed the limit. If the limit is exceeded, the API client will receive a `414 Request URI too long` response. The limit cannot be increased.

If you need to increase any of the limits, you must request and obtain an approval. To request a limit increase, contact the [AWS Support Center](#).

Known Issues

- Cross-account authentication is not currently supported in API Gateway. An API caller must be an IAM user of the same AWS account of the API owner.

API Gateway Control Service API

When you use the Amazon API Gateway console to create, configure, update, and deploy an API, the console submits appropriate RESTful requests to the API Gateway Control Service API to make things happen behind the scenes.

You can also use the AWS CLI to create, configure, update, and deploy an API. When you submit an AWS CLI command, you actually submit a REST request to the API Gateway Control Service API, and you receive the corresponding response, including NULL, from the API.

In addition, you can create your own app to programmatically create, configure, update, and deploy an API in API Gateway by calling the Control Service API directly or using an AWS SDK.

For more information on how use the API Gateway Control Service API, see [Amazon API Gateway REST API Reference](#).

Document History

The following table describes the important changes to the documentation since the last release of this guide.

- **Latest documentation update:** April 5, 2016

Change	Description	Date Changed
Documenting changes for the updated Amazon API Gateway console.	Learn how to create and set up an API using the updated API Gateway console. For more information, see Build and Test an API Gateway API from an Example (p. 5) and Build and Deploy an API Gateway API Step by Step (p. 13) .	April 5, 2016
Enabling the Import API feature to create a new or update an existing API from external API definitions.	With the Import API features, you can create a new API or update an existing one by uploading an external API definition expressed in Swagger 2.0 with the API Gateway extensions. For more information about the Import API, see Import an API (p. 117) .	April 5, 2016
Exposing the <code>\$input.body</code> variable to access the raw payload as string and the <code>\$util.parseJson()</code> function to turn a JSON string into a JSON object in a mapping template.	For more information about <code>\$input.body</code> and <code>\$util.parseJson()</code> , see Request and Response Payload-Mapping Reference (p. 96) .	April 5, 2016
Enabling client requests with method-level cache invalidation, and improving request throttling management.	Flush API stage-level cache and invalidate individual cache entry. For more information, see Flush the API Stage Cache in API Gateway (p. 208) and Invalidate an API Gateway Cache Entry (p. 209) . Improve the console experience for managing API request throttling. For more information, see Manage API Request Throttling (p. 201) .	March 25, 2016

Change	Description	Date Changed
Enabling and calling API Gateway API using custom authorization	Create and configure an AWS Lambda function to implement custom authorization. The function returns an IAM policy document that grants the Allow or Deny permissions to client requests of an API Gateway API. For more information, see Enable Custom authorization (p. 108) .	February 11, 2016
Importing and exporting API Gateway API using a Swagger definition file and extensions	Create and update your API Gateway API using the Swagger specification with the API Gateway extensions. Import the Swagger definitions using the API Gateway Importer. Export an API Gateway API to a Swagger definition file using the API Gateway console or API Gateway Export API. For more information, see Import and Export API (p. 116) .	December 18, 2015
Mapping request or response body or body's JSON fields to request or response parameters.	Map method request body or its JSON fields into integration request's path, query string, or headers. Map integration response body or its JSON fields into request response's headers. For more information, see Request and Response Parameter-Mapping Reference (p. 93) .	December 18, 2015
Working with Stage Variables in Amazon API Gateway	Learn how to associate configuration attributes with a deployment stage of an API in Amazon API Gateway. For more information, see Deploy an API with Stage Variables in Amazon API Gateway (p. 210) .	November 5, 2015
How to: Enable CORS for a Method	It is now easier to enable cross-origin resource sharing (CORS) for methods in Amazon API Gateway. For more information, see Enable CORS for a Resource (p. 101) .	November 3, 2015
How to: Use Client Side SSL Authentication	Use Amazon API Gateway to generate SSL certificates that you can use to authenticate calls to your HTTP backend. For more information, see Enable Client-Side SSL Authentication of an API (p. 106) .	September 22, 2015
Mock integration of methods	Learn how to mock-integrate an API with Amazon API Gateway (p. 65) . This feature enables developers to generate API responses from API Gateway directly without the need for a final integration back end beforehand.	September 1, 2015
Amazon Cognito Identity support	Amazon API Gateway has expanded the scope of the \$context variable so that it now returns information about Amazon Cognito Identity when requests are signed with Amazon Cognito credentials. In addition, we have added a \$util variable for escaping characters in JavaScript and encoding URLs and strings. For more information, see Request and Response Payload-Mapping Reference (p. 96) .	August 28, 2015
Swagger integration	Use the Swagger import tool on GitHub to import Swagger API definitions into Amazon API Gateway. Learn more about Import and Export API (p. 116) to create and deploy APIs and methods using the import tool. With the Swagger importer tool you can also update existing APIs.	July 21, 2015
Mapping Template Reference	Read about the \$input parameter and its functions in the Request and Response Payload-Mapping Reference (p. 96) .	July 18, 2015
Initial public release	This is the initial public release of the <i>Amazon API Gateway Developer Guide</i> .	July 9, 2015

AWS Glossary

For the latest AWS terminology, see the [AWS Glossary](#) in the *AWS General Reference*.